

Life and Mind

Did Consciousness Precede Life or Emerge From It?

Noushin Nabavi

2026

Before life could ask what consciousness is, consciousness may already have been asking what form life could take.

Perhaps consciousness did not appear at the end of life's story, but quietly shaped the conditions under which the story could begin. Life is matter that begins to matter to itself.

Contents

Preface	13
0.1 Why This Book?	13
0.2 Who This Book Is For	13
0.3 How to Read This Book	13
0.4 Conventions Used	14
0.5 Acknowledgments	14
 I Part I: Framing the Problem	 15
 1 The Fundamental Question	 16
1.1 Chapter Overview	16
1.2 The Paradox	16
1.3 Why the Question Has Been Avoided	17
1.4 The Three Possible Answers	18
1.5 Scope and Boundaries	19
1.6 A Map of the Book	19
1.7 How This Chapter Changes the Central Question	20
1.8 Chapter Summary	20
 2 What Is Life?	 22
2.1 Chapter Overview	22
2.2 Why Defining Life Is Hard	22
2.3 Biological Definitions	23
2.4 Thermodynamic and Physical Definitions	24
2.5 Informational Definitions	24
2.6 Autopoietic Definitions	25
2.7 Functional vs Ontological Definitions	26
2.8 Implications for the Central Question	26
2.9 How This Chapter Changes the Central Question	27
2.10 Chapter Summary	27

<i>CONTENTS</i>	3
3 What Is Consciousness?	28
3.1 Chapter Overview	28
3.2 The Problem of Definition	28
3.3 Phenomenal Consciousness	29
3.4 Access Consciousness	30
3.5 Minimal Consciousness and Proto-Consciousness	30
3.6 Self-Consciousness and Higher-Order Consciousness	31
3.7 Key Distinctions	31
3.8 The Explanatory Gap	32
3.9 Implications for the Central Question	33
3.10 How This Chapter Changes the Central Question	33
3.11 Chapter Summary	34
 II Part II: Historical and Philosophical Roots	 35
4 Ancient and Classical Perspectives	36
4.1 Chapter Overview	36
4.2 Pre-Socratic Thought	36
4.3 The Atomists	37
4.4 Plato and the Soul	38
4.5 Aristotle and the Scale of Nature	38
4.6 Stoicism and Pneuma	39
4.7 Neoplatonism	40
4.8 Implications for the Central Question	40
4.9 How This Chapter Changes the Central Question	41
4.10 Chapter Summary	41
 5 Eastern Philosophical Traditions	 42
5.1 Chapter Overview	42
5.2 Hindu Philosophy and Consciousness	42
5.3 Buddhist Perspectives	43
5.4 Daoist Thought	44
5.5 Jain Philosophy	45
5.6 Indigenous and Non-Western Perspectives	45
5.7 Comparative Analysis	46
5.8 Implications for the Central Question	47
5.9 How This Chapter Changes the Central Question	47
5.10 Chapter Summary	48

6	Modern Philosophy of Mind	49
6.1	Chapter Overview	49
6.2	Descartes and the Mind-Body Split	49
6.3	Spinoza and Neutral Monism	50
6.4	Leibniz and the Monadology	51
6.5	Kant and the Limits of Knowledge	51
6.6	Phenomenology	52
6.7	Physicalism and Its Variants	53
6.8	Property Dualism and Emergence	53
6.9	Contemporary Panpsychism	54
6.10	Implications for the Central Question	55
6.11	How This Chapter Changes the Central Question	55
6.12	Chapter Summary	56
III	Part III: The Origin of Life	57
7	Origin of Life Models	58
7.1	Chapter Overview	58
7.2	The Problem of Abiogenesis	58
7.3	RNA World Hypothesis	59
7.4	Metabolism-First Models	60
7.5	Protocells and Compartmentalization	61
7.6	Autocatalytic Sets and Chemical Self-Organization	61
7.7	Thermodynamic and Dissipative Approaches	62
7.8	Assembly Theory	63
7.9	Implications for the Central Question	63
7.10	How This Chapter Changes the Central Question	64
7.11	Chapter Summary	64
8	Information, Complexity, and Self-Organization	66
8.1	Chapter Overview	66
8.2	Information Theory Foundations	66
8.3	Information and the Origin of Life	67
8.4	Complexity and Emergence	68
8.5	What Is Self-Organization?	69
8.6	Self-Organization in the Origin of Life	69
8.7	Autopoiesis: Life as Self-Making	70

8.8	Recursion and Self-Reference	71
8.9	From Self-Organization to Self-Awareness	71
8.10	Feedback Loops as a Bridging Concept	72
8.11	Information Integration	73
8.12	Implications for the Central Question	73
8.13	How This Chapter Changes the Central Question	74
8.14	Chapter Summary	74

IV Part IV: Scientific Theories of Consciousness 76

9 Emergentist Models of Consciousness 77

9.1	Chapter Overview	77
9.2	What Is Emergence?	77
9.3	Neural Correlates of Consciousness	78
9.4	Biological Naturalism	79
9.5	Neurobiological Theories of Consciousness	80
9.6	Higher-Order Theories	80
9.7	Strengths and Limitations of Emergentism	81
9.8	Implications for the Central Question	82
9.9	How This Chapter Changes the Central Question	83
9.10	Chapter Summary	83

10 Contemporary Theories of Consciousness 84

10.1	Chapter Overview	84
10.2	Global Workspace Theory	84
10.3	Integrated Information Theory	85
10.4	Predictive Processing and the Free Energy Principle	86
10.5	Attention Schema Theory	87
10.6	Recurrent Processing Theory	88
10.7	Orchestrated Objective Reduction — Brief Preview	88
10.8	Theory Comparison	89
10.9	The Adversarial Collaboration	90
10.10	Implications for the Central Question	91
10.11	How This Chapter Changes the Central Question	91
10.12	Chapter Summary	91

V Part V: Consciousness First? 93

11 Consciousness-First Theories 94

11.1 Chapter Overview	94
11.2 The Core Claim	94
11.3 Panpsychism	95
11.4 Cosmopsychism	96
11.5 Idealism	96
11.6 Taheri's T-Consciousness Framework	97
11.7 Russellian Monism	98
11.8 Dual-Aspect Monism	99
11.9 Process Philosophy	99
11.10 How This Chapter Changes the Central Question	100
11.11 Chapter Summary	100

12 How Consciousness Could Have Enabled Life 102

12.1 Chapter Overview	102
12.2 The Explanatory Challenge	102
12.3 Consciousness as an Organizing Principle	103
12.4 Taheri's View of the Origin of Life	103
12.5 Proto-Experience and Chemical Selection	104
12.6 Information and Meaning	105
12.7 Quantum Approaches to Consciousness-Enabled Life	106
12.8 The Biogenic Universe Hypothesis	106
12.9 Critiques and Limitations	107
12.10 Implications for the Central Question	108
12.11 How This Chapter Changes the Central Question	108
12.12 Chapter Summary	108

13 Quantum and Speculative Frameworks 110

13.1 Chapter Overview	110
13.2 Why Quantum Mechanics Enters the Consciousness Debate	110
13.3 Orchestrated Objective Reduction	111
13.4 Quantum Coherence in Biology	112
13.5 The Observer Problem and Consciousness	113
13.6 Quantum Mind Theories Beyond Orch-OR	114
13.7 Non-Quantum Consciousness Field Theories	114
13.8 Speculative Extensions	115

13.9 Evaluating the Speculative Landscape	116
13.10 Implications for the Central Question	117
13.11 How This Chapter Changes the Central Question	117
13.12 Chapter Summary	117
13.13 Further Reading	118

VI Part VI: Life Produced Consciousness? 119

14 The Evolution of Consciousness 120

14.1 Chapter Overview	120
14.2 Minimal Cognition: Defining the Starting Point	120
14.3 Bacterial Intelligence	121
14.4 Protists and Single-Celled Decision-Making	122
14.5 Plant Cognition and Signaling	122
14.6 Fungal Networks and Distributed Intelligence	123
14.7 Sensing vs Feeling: The Critical Divide	124
14.8 Memory and Learning Without Brains	125
14.9 The Emergence of Nervous Systems	125
14.10 Key Evolutionary Candidates	126
14.11 Theoretical Frameworks for Basal Cognition	127
14.12 Implications for the Central Question	127
14.13 How This Chapter Changes the Central Question	128
14.14 Chapter Summary	128

15 The Hard Transition 130

15.1 Chapter Overview	130
15.2 Stating the Problem Precisely	130
15.3 Gradualism: Consciousness as a Continuum	131
15.4 Threshold Models: A Discrete Transition	132
15.5 Developmental Perspectives	132
15.6 The Phenomenal vs the Functional	133
15.7 How Major Theories Locate the Transition	134
15.8 The Explanatory Gap and Its Implications	135
15.9 Implications for the Central Question	136
15.10 How This Chapter Changes the Central Question	136
15.11 Chapter Summary	137

VII Part VII: Artificial Intelligence and Consciousness 138

16 Artificial Consciousness 139

16.1 Chapter Overview	139
16.2 The Question AI Poses	139
16.3 Functionalism and Machine Consciousness	140
16.4 What Current AI Systems Do and Don't Have	141
16.5 Distinguishing Intelligence, Language, Agency, and Sentience	141
16.6 Theoretical Perspectives on AI Consciousness	142
16.7 The Substrate Question	143
16.8 Artificial Life	144
16.9 Risks and Responsibilities	144
16.10 Implications for the Central Question	145
16.11 How This Chapter Changes the Central Question	146
16.12 Chapter Summary	146

VIII Part VIII: Open Questions and Synthesis 147

17 Limits and Testability 148

17.1 Chapter Overview	148
17.2 The Measurement Problem for Consciousness	148
17.3 First-Person vs Third-Person Methods	149
17.4 Falsifiability and Consciousness Theories	150
17.5 The Problem of Other Minds Applied to Life	151
17.6 Mathematical and Formal Approaches	151
17.7 Interdisciplinary Challenges	152
17.8 What Would Count as Progress?	153
17.9 Implications for the Central Question	153
17.10 How This Chapter Changes the Central Question	154
17.11 Chapter Summary	154

18 Ethics and Moral Status 156

18.1 Chapter Overview	156
18.2 Why the Consciousness Question Is an Ethics Question	156
18.3 Clarifying Key Ethical Terms	157
18.4 Sentience and Moral Consideration	157
18.5 Moral Status of Simple Organisms	158

18.6 Moral Status of Ecosystems	159
18.7 Moral Status of Artificial Systems	159
18.8 Rights Frameworks and Consciousness	160
18.9 Equity, Accessibility, and Reconciliation	161
18.10 The Precautionary Principle	161
18.11 Policy Implications	162
18.12 Implications for the Central Question	163
18.13 How This Chapter Changes the Central Question	163
18.14 Chapter Summary	163
19 Toward a Unified Perspective	165
19.1 Chapter Overview	165
19.2 Recapping the Three Positions	165
19.3 What Each Position Explains Well	166
19.4 What Each Position Struggles With	167
19.5 The Author's Position	167
19.6 Criteria for a Satisfying Answer	168
19.7 The Role of Future Science	169
19.8 Living With the Question	170
19.9 How This Chapter Changes the Central Question	170
19.10 Chapter Summary	171
A Glossary	172
A.1 Purpose	172
A.2 Glossary of Terms	172
B Timeline of Major Theories	178
B.1 Purpose	178
B.2 Ancient and Classical Period	178
B.3 Medieval and Early Religious-Philosophical Developments	179
B.4 Early Modern Period	179
B.5 Nineteenth Century	180
B.6 Early Twentieth Century	180
B.7 Mid-Twentieth Century	181
B.8 Late Twentieth Century	181
B.9 Early Twenty-First Century	182
B.10 Recent Developments	183
B.11 Thematic Summary	183

C Comparison Tables	185
C.1 Purpose	185
C.2 Consciousness Theories Compared	185
C.3 Origin-of-Life Models and Consciousness Implications	187
C.4 Where Each Theory Places the Transition	189
C.5 Philosophical Frameworks Compared	191
C.6 Life-First, Consciousness-First, and Co-Emergence Compared	193
C.7 Artificial Consciousness Models Compared	195
C.8 Ethical Implications Compared	197
C.9 Methodological Approaches Compared	199
C.10 Theory Strengths and Weaknesses at a Glance	201
C.11 Summary	203
D Key Philosophical Terms	204
D.1 Purpose	204
D.2 Ontology	204
D.3 Epistemology	204
D.4 Metaphysics	205
D.5 Dualism	205
D.6 Monism	205
D.7 Phenomenology	206
D.8 Intentionality	206
D.9 Supervenience	206
D.10 Teleology	207
D.11 Reductionism	207
D.12 Emergence	208
D.13 Qualia	208
D.14 Zombie	208
D.15 The Explanatory Gap	209
D.16 The Hard Problem	209
D.17 The Hard Transition	209
D.18 Intrinsic Nature	209
D.19 Ontological Primitive	210
D.20 Taheri's T-Consciousness Framework	210
D.21 Consciousness Field	210
D.22 T-Consciousness	211
D.23 Consciousness-First	211

D.24 Co-Emergence	211
D.25 Moral Status	211
D.26 Personhood	212
D.27 Reduction vs Explanation	212
D.28 First-Person Perspective	212
D.29 Third-Person Perspective	212
D.30 Second-Person Perspective	213
E Consciousness Taxonomy	214
E.1 Purpose	214
E.2 Types of Consciousness	214
E.3 By Biological Level	216
E.4 By Ontological Level	218
E.5 Levels and Dimensions	220
E.6 States of Consciousness	222
E.7 Theory Classification	224
E.8 By Mechanism	226
E.9 By Relation to Life	228
E.10 By Evidence Type	230
E.11 Summary of Main Taxonomic Distinctions	232
E.12 Taxonomic Caution	234
F References	235
F.1 Core Texts on Consciousness	235
F.2 Philosophy of Mind and Metaphysics	236
F.3 Consciousness-First, Idealist, Panpsychist, and T-Consciousness Sources	236
F.4 Origin of Life, Abiogenesis, and Self-Organization	237
F.5 Autopoiesis, Enactivism, Biosemiotics, and Basal Cognition	238
F.6 Evolution of Consciousness and Animal Minds	238
F.7 Predictive Processing, Free Energy, and Information-Based Frameworks	239
F.8 Quantum, Speculative, and Physical Frameworks	239
F.9 Artificial Intelligence, Artificial Life, and Machine Consciousness	240
F.10 Ethics, Moral Status, and Sentience	240
F.11 Indigenous, Relational, and Ecological Perspectives	240

Copyright © 2026 Noushin Nabavi

All rights reserved.

This book is a scholarly and interdisciplinary exploration of life, consciousness, and the relationship between biological organization and subjective experience.

First draft edition, 2026

Preface

0.1 Why This Book?

The origin of life and the nature of consciousness are two of the deepest unsolved problems in science and philosophy. They are also usually treated as separate questions. Origin-of-life research often asks how chemistry became biology, while consciousness studies often asks how brains generate experience. Yet these questions may be more closely connected than they first appear.

This book asks what happens when we consider them together. Did consciousness come first, or did life give rise to consciousness? Or are life and consciousness inseparable aspects of a deeper process involving organization, information, meaning, and experience?

The aim of this book is not to force a single answer too quickly. Instead, it explores the strongest versions of three broad possibilities: life first, consciousness first, and co-emergence. Each possibility is examined through science, philosophy, history, and speculative theory, while remaining honest about uncertainty and the limits of current knowledge.

0.2 Who This Book Is For

This book is written for researchers, students, and curious readers across philosophy, neuroscience, biology, physics, origin-of-life studies, artificial intelligence, and cognitive science.

No single disciplinary background is assumed. Scientific ideas are explained with philosophical context, and philosophical ideas are introduced in accessible language. The goal is to make a difficult question understandable without oversimplifying it.

Readers may approach this book as a philosophical investigation, a scientific synthesis, a speculative inquiry, or a guide to one of the deepest interdisciplinary questions of our time.

0.3 How to Read This Book

The book is organized in eight parts.

Parts I and II introduce the central problem, define key terms, and place the question in historical and philosophical context.

Parts III and IV examine the origin of life and scientific theories of consciousness separately, showing how each field frames its own questions.

Parts V and VI compare the two main directional hypotheses: consciousness-first and life-first.

Part VII considers artificial intelligence and artificial consciousness, asking whether consciousness requires biological life.

Part VIII brings the discussion together through questions of testability, ethics, moral status, and synthesis. The appendices provide reference materials, including a glossary, timeline, comparison tables, philosophical terms, consciousness taxonomy, and references.

0.4 Conventions Used

Key terms are defined in the Glossary.

Philosophical terminology is collected in the appendix on key philosophical terms.

Comparison tables are provided to help readers compare theories across different assumptions, mechanisms, strengths, and limitations.

Each chapter ends with a summary and an open question. These questions are not decorative; they are meant to keep the inquiry alive and to remind readers that the relationship between life and consciousness remains unsettled.

Speculative theories are treated respectfully but cautiously. Where a claim is philosophical, metaphysical, or speculative rather than empirically established, the text identifies it as such.

0.5 Acknowledgments

This book grew from a long-standing interest in the relationship between life, mind, matter, and meaning. I am grateful to the many scientists, philosophers, theorists, and contemplative traditions whose work has shaped the questions explored here.

I also acknowledge the importance of interdisciplinary thinking. No single field can fully answer the question of whether consciousness preceded life or emerged from it. Biology, neuroscience, philosophy, physics, information theory, artificial intelligence, ethics, and contemplative traditions each contribute part of the picture.

I am especially grateful to the researchers whose work has helped clarify the origin of life, the hard problem of consciousness, autopoiesis, basal cognition, emergence, artificial consciousness, and consciousness-first theories. Their ideas provided the foundation for this synthesis.

Finally, I acknowledge the readers who approach difficult questions with patience, humility, and curiosity. This book is written for those willing to sit with uncertainty and consider that the deepest questions may not belong to one discipline alone.

Part I

Part I: Framing the Problem

Chapter 1

The Fundamental Question

Before asking whether consciousness came before life or emerged from it, we must first ask what we mean by life, consciousness, emergence, and explanation. The question is not only about sequence, but about relationship: whether living systems and conscious experience are separate mysteries, or two expressions of a deeper problem concerning organization, responsiveness, and meaning.

1.1 Chapter Overview

This book begins with a question that is simple to ask but difficult to answer: **which came first, life or consciousness?**

At first glance, the question may seem like a variation of the familiar chicken-and-egg problem. Living systems appear to require some capacity to sense, respond, regulate, and adapt to their surroundings. Consciousness, in turn, is usually assumed to require a living nervous system, or at least some organized biological substrate. If life requires responsiveness, and consciousness requires life, then the relationship between the two becomes circular.

This chapter introduces the central question of the book: what is the relationship between the origin of life and the emergence of consciousness? It explains why this question has often been avoided, why a synthesis is needed, and what is at stake in considering life and consciousness together rather than as separate problems.

The aim is not to resolve the question immediately. Instead, this chapter sets the stage for the inquiry that follows.

1.2 The Paradox

Life is often described in terms of metabolism, reproduction, evolution, and cellular organization. Yet even the simplest living systems do more than exist passively. They maintain boundaries, exchange matter and energy with their environments, regulate internal conditions, and respond to changing circumstances.

A bacterium moves toward nutrients and away from harmful substances. A cell repairs damage, regulates its internal chemistry, and reacts to signals. A plant bends toward light. None of these examples necessarily implies human-like awareness, thought, or subjective experience. However, they do show that life is never merely inert matter. Life involves responsiveness.

Consciousness, on the other hand, is often described as experience, awareness, subjectivity, or the capacity for an inner point of view. In standard scientific accounts, consciousness appears much later than life. It is usually associated with nervous systems, brains, perception, memory, and complex behaviour. From this perspective, life begins first, and consciousness emerges only after biological systems become sufficiently complex.

This creates a deep conceptual tension. If life begins as organized responsiveness, and consciousness is also connected to responsiveness, where should the line be drawn? At what point does biological sensitivity become something closer to awareness? Is consciousness a late product of life, or is some primitive form of proto-experience already implicit in living organization?

The paradox can be stated simply:

Life seems to require some form of responsiveness to the world. Consciousness seems to require some form of living organization. So which came first?

This is not only a chronological question. It is also a structural question. It asks whether life and consciousness are two separate phenomena, one producing the other, or whether they are different expressions of a deeper process involving organization, information, responsiveness, and experience.

If life came first, then consciousness must be explained as something that emerged from living matter. If consciousness came first, then life may be understood as a biological expression of a more fundamental field or principle of awareness. If they co-emerged, then life and consciousness may be inseparable aspects of self-organizing systems.

Each answer changes how we understand biology, mind, matter, and the place of human beings in nature.

1.3 Why the Question Has Been Avoided

The relationship between life and consciousness has rarely been treated as a unified problem. One reason is that modern knowledge is organized into specialized disciplines.

Origin-of-life research is usually located within chemistry, molecular biology, geoscience, and systems biology. It asks how non-living matter could have produced self-replicating molecules, metabolic networks, protocells, and eventually cellular life. Its central questions concern chemical pathways, environmental conditions, molecular stability, energy flows, and evolutionary transitions.

Consciousness studies, by contrast, is usually located within neuroscience, psychology, philosophy of mind, cognitive science, artificial intelligence, and sometimes physics. It asks how experience arises, how awareness relates to brain activity, how information is integrated, and whether machines or non-human organisms can be conscious.

These two research areas often move in parallel without much contact. Origin-of-life research rarely asks whether primitive life had any experiential dimension. Consciousness studies rarely asks whether the emergence of consciousness is related to the earliest features of life itself.

There are good reasons for this separation. Origin-of-life research depends heavily on experimental chemistry, geological evidence, and plausible models of early Earth. Consciousness research faces difficulties of measurement because subjective experience is not directly observable from the outside. Bringing the two together introduces methodological challenges.

There are also institutional barriers. Scientists are often trained to avoid questions that appear too speculative, too philosophical, or too difficult to test. Consciousness is already a controversial subject in many scientific contexts. Linking it to abiogenesis can seem even more uncertain.

Yet avoiding the question does not make it disappear. If life is the condition under which consciousness later appears, then any theory of consciousness must eventually explain why living organization matters. If

consciousness is more fundamental than life, then origin-of-life research may be missing a deeper dimension of organization or information. If life and consciousness co-emerged, then both fields may require a broader conceptual framework.

The question has been avoided not because it is unimportant, but because it is difficult to place within existing academic boundaries.

1.4 The Three Possible Answers

There are three broad ways to answer the question of whether life or consciousness came first.

1.4.1 Life first, consciousness later

The standard scientific view is that life came first and consciousness emerged later. According to this view, the early Earth produced increasingly complex chemical systems. Some of these systems became capable of replication, metabolism, compartmentalization, and evolution. Over time, biological organisms became more complex. Nervous systems evolved. Eventually, brains developed capacities for perception, memory, learning, emotion, and awareness.

In this view, consciousness is an emergent property of living systems. It is not present at the beginning of life. Instead, it appears when biological organization reaches a certain level of complexity.

This view has the advantage of fitting well with evolutionary biology and neuroscience. It does not require consciousness to exist outside or before living systems. However, it faces a major challenge: it must explain how subjective experience arises from physical and biological processes. This is often called the hard problem of consciousness.

1.4.2 Consciousness first, life as its expression

A second view reverses the standard order. It suggests that consciousness, awareness, or mind-like qualities are more fundamental than biological life. In this view, life does not produce consciousness from nothing. Instead, life expresses, channels, organizes, or localizes a more basic field or principle of consciousness.

This position appears in different forms. Panpsychism suggests that mind-like qualities may be present, in some primitive form, throughout nature. Idealism suggests that consciousness is more fundamental than matter. Some spiritual and metaphysical traditions also treat life as an expression of a deeper consciousness or intelligence.

This view has the advantage of taking consciousness seriously as a fundamental feature of reality rather than treating it as a late accident of evolution. However, it faces its own challenge: it must explain why consciousness appears to be so closely tied to biological organization, nervous systems, and embodied life.

1.4.3 Co-emergence of life and consciousness

A third possibility is that life and consciousness co-emerged. On this view, consciousness is not fully present before life, nor does it suddenly appear only after complex brains evolve. Instead, life and consciousness may be two aspects of a single process of self-organization.

In this framework, the earliest living systems may not have consciousness in the human sense. However, they may possess primitive forms of sensitivity, valuation, or world-directed activity. Consciousness, in its richer forms, may develop gradually from these basic features of living systems.

This view does not require bacteria or cells to think, feel, or reflect. Rather, it asks whether the roots of experience may lie in the same processes that make life possible: boundary formation, metabolism, information processing, responsiveness, regulation, and adaptation.

Each of these three answers carries a different worldview. The life-first view emphasizes biological emergence. The consciousness-first view emphasizes mind as fundamental. The co-emergence view emphasizes continuity between life, information, responsiveness, and experience.

This book examines all three possibilities.

1.5 Scope and Boundaries

This book is not a comprehensive textbook on the origin of life. It does not attempt to review every model of abiogenesis in technical detail. Important theories such as the RNA world, metabolism-first models, hydrothermal vent hypotheses, lipid-world models, autocatalytic networks, mineral-surface theories, and protocell models will be introduced, but the focus is not on declaring a single winning theory.

This book is also not a complete textbook on consciousness science. It does not attempt to settle debates about the neural correlates of consciousness, artificial consciousness, integrated information theory, global workspace theory, higher-order theories, or panpsychism. These approaches will be discussed, but the purpose is not to defend one theory as final.

Instead, the contribution of this book is synthetic. It asks what happens when origin-of-life research and consciousness studies are placed in conversation with each other.

The central concern is the relationship between matter, life, information, organization, and experience. The book asks whether the emergence of life and the emergence of consciousness should be understood as separate transitions or as connected dimensions of a broader evolutionary process.

This means the discussion will move across scientific, philosophical, and speculative boundaries. Some chapters will focus on established scientific models. Others will examine philosophical traditions and consciousness-first frameworks. Still others will explore open questions, conceptual tensions, and future research directions.

The goal is not to collapse science into metaphysics or metaphysics into science. The goal is to clarify what can be known, what remains uncertain, and what kinds of questions become visible when life and consciousness are examined together.

1.6 A Map of the Book

The book is organized as a gradual movement from definitions to synthesis.

Part I: Framing introduces the central question and defines the key terms: life, consciousness, emergence, information, and organization. These early chapters establish the conceptual foundation for the rest of the book.

Part II: Historical and Philosophical Roots examines how different traditions have understood the relationship between life and mind. Ancient, classical, Eastern, and modern philosophical perspectives provide a wider background for contemporary debates.

Part III: Origin of Life introduces major scientific theories about how life may have emerged from non-living matter. This part focuses on chemistry, self-organization, molecular evolution, protocells, and early biological systems.

Part IV: Scientific Theories of Consciousness reviews major scientific and philosophical theories of consciousness, including emergentist models, global workspace theory, integrated information theory, higher-order approaches, biological theories, and other contemporary frameworks.

Part V: Consciousness First? explores theories that treat consciousness as fundamental or prior to biological life. This includes panpsychism, idealism, consciousness-field approaches, and other speculative frameworks.

Part VI: Life Produced Consciousness? examines the standard evolutionary view that consciousness arises from life. It considers minimal cognition, biological agency, nervous systems, and the gradual evolution of awareness.

Part VII: Artificial Intelligence and Consciousness asks whether consciousness requires biology or whether artificial systems could become conscious. This part considers artificial life, machine intelligence, simulation, embodiment, and moral status.

Part VIII: Open Questions and Synthesis returns to the central question. It compares the major positions, identifies unresolved problems, and asks whether life and consciousness may be understood through a unified framework.

Readers from different backgrounds may enter the book differently. Those interested in biology may focus first on the chapters on origin-of-life models and minimal cognition. Those interested in philosophy may begin with the chapters on consciousness-first theories and the hard problem. Those interested in artificial intelligence may move toward the later chapters on artificial consciousness and moral status.

The full argument, however, develops across the whole book. It begins with definitions, moves through history and science, and ends with synthesis and implications.

1.7 How This Chapter Changes the Central Question

This chapter changes the central question by clarifying that life is not merely matter arranged in a complex way. Life involves boundary, metabolism, self-maintenance, adaptation, reproduction, and continuity over time. These features make living systems different from ordinary physical systems.

The question therefore becomes sharper: if consciousness emerged from life, which features of life made that emergence possible? Was it metabolism, self-maintenance, responsiveness, information processing, or the organism's need to preserve itself as a living whole?

1.8 Chapter Summary

This chapter introduced the fundamental question of the book: what is the relationship between the origin of life and the emergence of consciousness?

The chapter argued that this is not merely a chicken-and-egg problem. It is a structural question about the relationship between matter, organization, responsiveness, life, and experience.

Three broad answers were introduced. The first holds that life came first and consciousness emerged later through biological evolution. The second holds that consciousness is fundamental and life is one of its expressions. The third proposes that life and consciousness co-emerged as inseparable aspects of self-organizing systems.

The chapter also explained why this question has often been avoided. Origin-of-life research and consciousness studies belong to different disciplines, use different methods, and face different evidentiary challenges. Yet the question remains important because both fields deal with transitions from matter to organization, from organization to responsiveness, and from responsiveness to experience.

The open question for the rest of the book is this:

Is the relationship between life and consciousness a scientific question, a philosophical question, or both?

Chapter 2

What Is Life?

2.1 Chapter Overview

Before asking whether consciousness preceded life or emerged from it, we need a working definition of life. This is harder than it first appears. Life seems obvious when we point to a tree, an animal, a bacterium, or a human being. Yet the boundary becomes less clear when we consider viruses, prions, self-replicating molecules, synthetic cells, artificial life, or possible life on other planets.

This chapter surveys several major ways of defining life: biological, thermodynamic, informational, autopoietic, functional, and ontological. Each definition highlights something important, but each also leaves something unresolved. More importantly, each definition shapes how we understand the relationship between life and consciousness.

If life is mainly chemistry, then consciousness may appear as a later biological achievement. If life is information processing, then the distance between life and mind becomes smaller. If life is self-production and self-maintenance, then something like minimal cognition may be present from the beginning.

The question “What is life?” is therefore not only biological. It is also philosophical, physical, informational, and perhaps even experiential.

2.2 Why Defining Life Is Hard

Life is one of the most familiar realities we encounter, yet it remains surprisingly difficult to define. We recognize living beings around us, but recognition is not the same as definition. A bird, a mushroom, a bacterium, and a human being are all alive, but they differ enormously in structure, behaviour, complexity, and form of existence.

The difficulty becomes clearer at the edges. Viruses reproduce and evolve, but they cannot reproduce independently outside host cells. Prions can transmit biological effects without containing DNA or RNA. Self-replicating molecules may copy themselves but may not metabolize or maintain a boundary. Synthetic biological systems may be engineered to perform life-like functions, raising the question of whether they are alive, artificial, or something in between.

These edge cases matter because they reveal that life is not defined by a single feature. No one property seems sufficient on its own. Reproduction is important, but sterile organisms are still alive. Metabolism is important, but some dormant organisms can suspend metabolism for long periods. Genetic information

is important, but information alone does not make something alive. Cellular structure is important, but viruses challenge the boundary between living and non-living systems.

The problem is not simply that we lack a good definition. The deeper problem is that life may not be a single thing in the ordinary sense. It may be a process, a pattern, a form of organization, or a cluster of related properties. Depending on which feature we emphasize, we arrive at a different answer.

This matters for the central question of this book. If life is defined narrowly as cellular chemistry, then consciousness appears to be a separate and later problem. If life is defined as adaptive information processing, then the roots of consciousness may already be near the roots of life. If life is defined as self-producing organization, then the gap between life and cognition becomes less clear.

The definition of life is therefore not neutral. It shapes the entire discussion that follows.

2.3 Biological Definitions

Biological definitions often begin with the observable features of living systems. Living beings grow, reproduce, metabolize, respond to stimuli, maintain internal organization, and evolve. These features are useful because they describe what living organisms actually do. They also provide practical criteria for distinguishing living systems from non-living matter.

One influential working definition, often associated with NASA astrobiology, describes life as a self-sustaining chemical system capable of Darwinian evolution. This definition is powerful because it combines chemistry, self-maintenance, and evolution. It is especially useful in the search for life beyond Earth because it does not depend on the details of Earth biology alone. It allows scientists to ask whether another system, perhaps based on different molecules, could sustain itself and evolve.

Yet even this definition has limitations. It emphasizes evolution, but not every individual living being evolves. Populations evolve; individuals live. It emphasizes chemistry, but it does not fully explain why some chemical systems become living systems while others do not. It also leaves open the question of whether a system must be conscious, cognitive, or responsive in any meaningful sense to count as alive.

Metabolic definitions place more emphasis on organized chemistry. From this perspective, life is a network of chemical reactions that extracts energy and matter from the environment in order to maintain itself. A living system is not merely a collection of molecules. It is a dynamic process that continually rebuilds itself.

Genetic definitions emphasize replication and heredity. Life depends on information that can be copied, modified, and transmitted across generations. DNA and RNA are central in known life because they store and transmit biological information. This perspective makes evolution possible because variation in hereditary information can be acted upon by selection.

Cellular definitions emphasize bounded organization. Life, as we know it, is cellular. Cells create boundaries between inside and outside. They regulate what enters and leaves. They maintain internal conditions. They coordinate chemical reactions within a protected space. From this view, life begins not only when molecules replicate, but when they become organized within a boundary that can sustain a self-maintaining process.

Each biological definition captures part of the truth. Life is chemical, metabolic, genetic, cellular, and evolutionary. But none of these alone fully explains why living systems seem different from ordinary matter. Life is not only made of molecules; it is made of molecules organized in a particular way.

That organization becomes central when we turn to physics and thermodynamics.

2.4 Thermodynamic and Physical Definitions

From a physical perspective, life is not separate from matter and energy. Living systems obey the laws of physics. They require energy, produce waste, maintain structure, and exist far from thermodynamic equilibrium.

Erwin Schrödinger famously described life in relation to “negative entropy.” He was not suggesting that life violates the second law of thermodynamics. Rather, living systems maintain internal order by drawing energy and matter from their environments. A living organism can remain highly organized only because it exchanges energy with the world around it.

This perspective is important because living systems are not static objects. They are dynamic processes. A living body is not like a stone. It must continually maintain itself. Its order is not permanent; it is actively produced. The organism survives by resisting decay, repairing damage, regulating internal states, and transforming energy into organized activity.

Ilya Prigogine’s work on dissipative structures further developed this idea. A dissipative structure is an ordered pattern that arises in a system far from equilibrium through the flow of energy. Examples include whirlpools, convection cells, and certain chemical oscillations. These systems show that order can emerge spontaneously under the right physical conditions.

Life can be understood as a highly complex form of organized energy flow. It is not merely a structure, but a structure that persists through activity. It survives by dissipating energy, maintaining boundaries, and creating internal order through external exchange.

This view narrows the distance between life and other forms of self-organization in nature. It suggests that life may be part of a broader physical tendency for complex systems to arise under certain conditions. However, life is not just any dissipative structure. A flame also consumes energy and maintains a dynamic form, but we do not usually call it alive.

So something else is needed. Living systems do not merely process energy. They also store, transmit, and use information.

2.5 Informational Definitions

Information is central to life. DNA stores hereditary information. RNA participates in information transfer and regulation. Cells process signals from their environments. Organisms use information to maintain themselves, respond to change, and reproduce.

From an informational perspective, life is not only chemistry, but chemistry organized by information. A living system does not simply react. It interprets molecular signals in context. The same molecule can have different effects depending on where it appears, when it appears, and what state the cell is in. Biological information is therefore not only stored in genes; it is distributed across networks, structures, interactions, and regulatory processes.

This makes information a bridge between physics and biology. Physical systems involve matter and energy. Biological systems also involve matter and energy, but they add layers of coding, regulation, memory, and control. The living cell is not only a chemical factory. It is also an informational system.

Hereditary information is especially important because it allows continuity across generations. Life persists not only by surviving in the present but by transmitting patterns into the future. Mutation, variation, and selection operate on these patterns. Evolution is therefore both a material and informational process.

Some theorists argue that the origin of life should be understood as a transition in informational architecture. In this view, life begins when information gains causal power over matter: when stored patterns do not merely

describe a system, but help organize and maintain it. The genome is not alive by itself, but in a cellular context it participates in the ongoing production of the organism.

This approach brings life closer to cognition and consciousness. If life is fundamentally about information processing, then the difference between life and mind may be a matter of degree, organization, and complexity rather than an absolute divide.

However, information alone is not enough. A book contains information, but it is not alive. A computer processes information, but whether it is alive or conscious remains contested. Biological information is embedded in self-maintaining processes. It is not detached from the organism. It matters because it participates in the ongoing production of life.

This leads to a deeper definition: life as self-production.

2.6 Autopoietic Definitions

The theory of autopoiesis, developed by Humberto Maturana and Francisco Varela, defines living systems as self-producing systems. The term “autopoiesis” literally means self-making or self-production. A living system is one that continuously produces and maintains the components that make the system possible.

A cell is the classic example. It produces molecules that maintain its membrane, enzymes that support its metabolism, and internal structures that allow it to continue operating as a cell. The cell is not assembled once and then left alone. It is constantly remaking itself.

Autopoiesis emphasizes operational closure. This does not mean that the organism is closed off from the environment. Living systems are materially and energetically open; they exchange matter and energy with their surroundings. But they are organizationally closed in the sense that their internal processes form a network that sustains the system’s own identity.

A living system is also structurally coupled to its environment. It does not merely exist in the world; it changes in relation to the world. The organism and environment shape each other through ongoing interaction. The organism responds to environmental changes according to its own structure, needs, and possibilities.

This is where autopoiesis becomes important for the question of consciousness. Maturana and Varela argued that cognition is not something added to life later; rather, living itself is already a form of cognition. To live is to enact a world of relevance. A living system distinguishes, through its own organization, what matters for its continuation.

This does not mean that all living systems are conscious in the human sense. A bacterium does not reflect on its existence. A cell does not have language or self-awareness. But autopoiesis suggests that the roots of cognition may be present wherever there is self-maintaining, world-directed organization.

If this is true, then cognition is not a late addition to life. It is built into the structure of living systems from the beginning. Consciousness may still emerge later, but it emerges from a field already shaped by responsiveness, self-maintenance, and meaning for the organism.

This view will become important in later chapters on self-organization, minimal cognition, and the possible continuity between life and mind.

2.7 Functional vs Ontological Definitions

Definitions of life can be divided into two broad types: functional and ontological.

Functional definitions ask what life does. A system may be considered alive if it metabolizes, grows, reproduces, evolves, responds to stimuli, maintains homeostasis, and processes information. This approach is practical because it gives us a list of observable features. It is especially useful in biology, astrobiology, and synthetic biology.

The weakness of functional definitions is that no single function is decisive. Some non-living systems perform life-like functions. Fire grows and consumes energy. Crystals grow and replicate patterns. Computer viruses reproduce in a digital environment. Artificial systems can process information, adapt, and respond. Yet we hesitate to call all of these alive.

Ontological definitions ask what life is. They treat life not only as a list of functions but as a distinctive mode of being or process. From this perspective, life is not just replication, metabolism, or responsiveness. It is a self-maintaining, self-producing, historically continuous process that creates and preserves its own organization.

The distinction matters because something might appear functionally life-like without being fully alive. A simulation of a cell may reproduce some biological dynamics, but it may not be materially self-producing. A robot may respond to stimuli, but it may not have metabolism or biological self-maintenance. A virus may evolve, but it depends on host cells for reproduction.

This raises a difficult question: can something be functionally alive but not really alive?

The answer depends on what we mean by “really.” If life is defined by functions, then sufficiently life-like systems may count as life. If life is defined by material self-production, then simulations and machines may not qualify. If life is defined more broadly as adaptive organization, then artificial life may occupy a middle category.

This debate foreshadows later questions about artificial consciousness. A machine may behave intelligently, but is it conscious? A system may simulate life, but is it alive? In both cases, we face the same challenge: whether function is enough, or whether there is something deeper about the organization, embodiment, or existence of the system.

2.8 Implications for the Central Question

The way we define life has direct consequences for the central question of this book.

If life is defined primarily as chemistry, then consciousness appears as a separate problem. Life begins with chemical systems capable of metabolism, replication, and evolution. Consciousness emerges much later, perhaps only with nervous systems and brains. This view supports the standard scientific sequence: matter first, life second, consciousness third.

If life is defined as information processing, the distance between life and consciousness becomes smaller. Living systems are not only chemical systems; they are systems that store, interpret, regulate, and respond to information. Since consciousness is also often associated with information processing, integration, attention, and responsiveness, the roots of consciousness may be closer to the roots of life.

If life is defined through autopoiesis, then cognition may be built into life from the beginning. A living system does not merely exist; it maintains itself, distinguishes itself from its environment, and responds to the world in ways that matter for its continuation. This suggests that the earliest forms of life may already contain the foundations of sense-making, even if not consciousness in a rich experiential sense.

If life is defined thermodynamically, then both life and consciousness may be viewed as forms of organized process. Life maintains order through energy flow. Consciousness may depend on patterns of organization within living systems. This view encourages us to ask whether consciousness is related to the same principles of self-organization that make life possible.

None of these definitions proves that consciousness existed before life. None proves that consciousness emerged only after life. But each definition changes the shape of the question.

A narrow chemical definition creates a larger gap between life and consciousness. An informational definition creates a possible bridge. An autopoietic definition places cognition close to the foundation of life. A thermodynamic definition situates life within broader principles of organization and energy flow.

The central question therefore cannot be answered without first deciding what kind of phenomenon life is.

2.9 How This Chapter Changes the Central Question

This chapter changes the central question by showing that consciousness is not a single simple concept. It may refer to subjective experience, access to information, self-awareness, sentience, reflective thought, or a fundamental field-like reality.

The question therefore cannot be answered until we are clear about which meaning of consciousness is being discussed. If consciousness means reportable awareness, it may appear late in evolution. If it means sentience, it may appear earlier. If it means a fundamental reality, it may not emerge from life at all.

2.10 Chapter Summary

This chapter examined several major approaches to defining life.

Biological definitions emphasize chemistry, metabolism, heredity, cellular organization, and evolution. Thermodynamic definitions emphasize energy flow, order, and far-from-equilibrium organization. Informational definitions emphasize coding, regulation, memory, and causal control. Autopoietic definitions emphasize self-production, operational closure, and structural coupling. Functional definitions focus on what life does, while ontological definitions ask what kind of process life is.

No single definition resolves all difficulties. Edge cases such as viruses, prions, self-replicating molecules, synthetic cells, and artificial life show that life exists near conceptual boundaries. These boundaries matter because they shape how we understand the relationship between life and consciousness.

If life is only chemistry, consciousness may seem like a later addition. If life is information processing, consciousness may appear as a more complex form of a process already present in living systems. If life is autopoiesis, then minimal cognition may be inseparable from life itself.

The open question is therefore:

Does every viable definition of life already smuggle in some form of cognition?

Chapter 3

What Is Consciousness?

3.1 Chapter Overview

Consciousness is at once the most intimate and the most difficult phenomenon to explain. Nothing is more familiar than having an experience: seeing colour, feeling pain, remembering a childhood scene, tasting food, hearing music, or sensing one's own presence in the world. Yet when we try to define consciousness precisely, the concept becomes elusive.

This chapter introduces the major definitions and distinctions used in consciousness studies. It examines phenomenal consciousness, access consciousness, minimal consciousness, proto-consciousness, self-consciousness, and higher-order consciousness. It also clarifies related terms such as awareness, sentience, cognition, wakefulness, and experience.

The purpose of this chapter is not to defend one theory of consciousness as final. Instead, it establishes the vocabulary needed for the central question of this book: did consciousness emerge from life, precede life, or co-emerge with living organization?

The answer depends strongly on what we mean by consciousness. If consciousness means reflective self-awareness, then it may be a late evolutionary development. If it means subjective experience, then it may extend further back into animal life. If it means minimal sensitivity or proto-experience, then it may be linked to the earliest forms of life. If it is fundamental, then it may not depend on life at all.

3.2 The Problem of Definition

Consciousness is difficult to define because it is both obvious and mysterious. We do not usually need a definition in order to know that we are experiencing something. We are directly acquainted with our own perceptions, feelings, sensations, and thoughts. Yet this direct familiarity does not easily translate into scientific or philosophical explanation.

The word “consciousness” is also used in different ways. Sometimes it means wakefulness, as when a patient is described as conscious rather than unconscious. Sometimes it means awareness of the environment. Sometimes it means subjective experience. Sometimes it means self-awareness, reflection, or the ability to report one's mental state. These meanings overlap, but they are not identical.

This ambiguity creates confusion. A scientist studying brain activity during anaesthesia, a philosopher discussing the hard problem, a psychologist measuring attention, and an artificial intelligence researcher

testing machine behaviour may all use the word consciousness, but they may not mean exactly the same thing.

For this reason, it is important to specify which sense of consciousness is under discussion. A theory that explains attention may not explain subjective experience. A theory that explains wakefulness may not explain self-awareness. A theory that explains behavioural flexibility may not explain what it feels like to be an organism.

The central question of this book depends on these distinctions. If consciousness means full reflective self-awareness, then it almost certainly appears long after the origin of life. If consciousness means basic sentience or feeling, then the question becomes more open. If consciousness means proto-experience or a fundamental feature of reality, then it may precede biological life.

Some contemporary consciousness-first frameworks define consciousness even more broadly. For example, Taheri's T-Consciousness framework treats consciousness not as a mental state produced by the brain, but as a fundamental, non-material reality through which matter, life, and individual awareness become organized or manifest. In this view, consciousness is not limited to phenomenal experience, cognition, or neural activity. It is a primary field-like order underlying existence itself. This definition will be examined more fully in later chapters on consciousness-first theories and consciousness-enabled life.

Before we can ask when consciousness appeared, we must ask what kind of consciousness we are talking about.

3.3 Phenomenal Consciousness

Phenomenal consciousness refers to the felt, subjective quality of experience. It is often described as “what it is like” to be a subject of experience. There is something it is like to see red, feel warmth, hear a melody, experience fear, taste sweetness, or feel pain.

This sense of consciousness is central to the philosophical problem of mind. A physical description of the brain may describe neurons, electrical activity, neurotransmitters, and information processing. But many philosophers argue that such a description still seems to leave something out: the felt quality of experience itself.

These felt qualities are often called qualia. The redness of red, the bitterness of coffee, the ache of pain, the warmth of sunlight, and the emotional tone of sadness are all examples of qualitative experience. They are not merely behaviours or functions. They are how experience appears from the inside.

Phenomenal consciousness is especially important because it creates the deepest explanatory challenge. We can imagine explaining how light enters the eye, how the retina processes signals, how visual information travels to the brain, and how behaviour is guided by perception. But why should any of this processing be accompanied by experience? Why should there be something it feels like?

This is the problem often called the hard problem of consciousness. The hard problem is not simply how the brain processes information or controls behaviour. It is why physical and biological processes give rise to subjective experience at all.

For the purposes of this book, phenomenal consciousness raises a major difficulty. If consciousness means subjective experience, then we must ask when experience first appeared in the history of life. Did it arise only with nervous systems? With brains? With animals? With primitive forms of sensing? Or is experience, in some minimal form, more deeply connected to life itself?

3.4 Access Consciousness

Access consciousness refers to information that is available for reasoning, reporting, decision-making, and control of behaviour. A mental content is access conscious when it can be used by the system in a flexible way. For example, if a person sees a red light, can report seeing it, stop walking, remember it, and use it to guide action, that information is access conscious.

This concept is different from phenomenal consciousness. Phenomenal consciousness concerns what experience feels like. Access consciousness concerns what information is available to the system for use. The distinction matters because a system might use information without necessarily having subjective experience.

For example, a camera can detect light, a thermostat can respond to temperature, and a computer can process visual data. These systems have access to information in a functional sense, but we do not usually assume that they have inner experience. At the same time, human beings may have experiences that are difficult to report or use clearly, such as vague feelings, background moods, or fleeting sensations.

The distinction between access consciousness and phenomenal consciousness is especially important for debates about animals and artificial intelligence. An animal may not be able to verbally report its experience, but it may still feel pain or fear. A machine may be able to report information, solve problems, and control behaviour, but whether it has subjective experience remains uncertain.

Access consciousness is easier to study scientifically because it can be linked to behaviour, report, attention, memory, and neural activity. Phenomenal consciousness is harder because it concerns first-person experience.

For the central question of this book, access consciousness offers one possible bridge between life and mind. Living systems process information and use it to regulate behaviour. But does information use become consciousness only when it is globally available, reportable, or integrated into flexible action? Or is access consciousness only one outward expression of a deeper experiential process?

3.5 Minimal Consciousness and Proto-Consciousness

If consciousness exists in degrees, then we must ask whether there is a simplest form of consciousness. Minimal consciousness refers to the most basic form of experience that could still count as consciousness. It does not require language, abstract thought, memory of the self, or reflective awareness. It may involve only a simple feeling of presence, sensation, or world-directed experience.

Proto-consciousness is an even more tentative concept. It refers to pre-reflective, pre-verbal, or primitive forms of experience that may not yet be consciousness in the full sense, but may contain the roots from which consciousness develops. Proto-consciousness is often discussed in theories that seek continuity between matter, life, mind, and experience.

These ideas are important because they challenge a sharp boundary between conscious and non-conscious systems. If consciousness is all-or-nothing, then at some point in evolution it suddenly appeared. Before that point, there was no experience; after that point, experience existed. But if consciousness is graded, then it may have developed gradually through simpler forms of sensitivity, feeling, and awareness.

Minimal consciousness may include basic sentience: the capacity to feel or be affected. A simple animal may not have self-reflection, but it may still experience pain, pleasure, attraction, aversion, hunger, or threat. Such forms of experience may not be complex, but they may still matter morally and biologically.

The concept of minimal consciousness also connects to living organization. Even simple organisms regulate themselves in relation to their surroundings. They distinguish favourable from harmful conditions, maintain boundaries, and act in ways that preserve their existence. This does not prove that they are conscious, but it raises the question of whether the earliest roots of consciousness may lie in biological responsiveness.

For the origin question, minimal consciousness and proto-consciousness are crucial. They allow us to ask whether consciousness began as reflective thought, as animal feeling, as cellular sensitivity, or as a more fundamental feature of organized processes.

3.6 Self-Consciousness and Higher-Order Consciousness

Self-consciousness is awareness of oneself as a subject. It is not merely having an experience, but being aware that one is having an experience. A self-conscious being can, in some sense, distinguish itself from the world and recognize itself as the one who perceives, feels, acts, or thinks.

Higher-order consciousness involves awareness of one's own mental states. A being with higher-order consciousness does not only see or feel; it can know that it sees or feels. It may reflect on its own thoughts, monitor its own uncertainty, recognize its own intentions, or consider how its mind differs from the minds of others.

In humans, self-consciousness is closely connected to language, memory, social interaction, and identity. We tell stories about ourselves. We imagine ourselves in the past and future. We reflect on our choices. We ask what kind of person we are. This form of consciousness is highly developed and culturally shaped.

In animals, self-consciousness is more difficult to assess. The mirror test is often used as one measure of self-recognition, although it is not a perfect measure of consciousness. Some animals appear able to recognize themselves in mirrors, while others show complex social understanding, planning, memory, or emotional behaviour without necessarily passing the test.

Theory of mind and metacognition are related capacities. Theory of mind involves understanding that other beings have perspectives, intentions, or beliefs. Metacognition involves monitoring one's own knowledge or uncertainty. These abilities suggest more advanced forms of consciousness, especially in social and cognitively complex animals.

Self-consciousness is likely a later evolutionary development than basic sentience or minimal experience. It may depend on nervous systems, social environments, memory, and learning. This means that if consciousness is defined as self-consciousness, then consciousness almost certainly did not precede life. It emerged within life, and probably within complex animal life.

However, self-consciousness is not the only possible meaning of consciousness. The deeper question is whether self-consciousness is the essence of consciousness or only one highly developed form of it.

3.7 Key Distinctions

Several distinctions are necessary for avoiding confusion.

Consciousness is not identical to cognition. Cognition refers to processes such as perception, memory, learning, problem-solving, and decision-making. These processes may support consciousness, but they may also occur without conscious awareness. Much of human cognition is unconscious. The brain processes information, regulates behaviour, and forms predictions without everything becoming part of conscious experience.

Consciousness is also not identical to awareness. Awareness often refers to responsiveness or sensitivity to something. A person may be aware of a sound, a danger, or an emotion. But awareness can be used broadly, sometimes even for non-conscious systems. Consciousness usually implies a stronger sense of subjective experience.

Sentience is another important term. Sentience refers to the capacity to feel, especially pleasure, pain, or affective states. A sentient being is not necessarily self-conscious or rational, but it can be affected in a way that matters to it. Sentience is often central in moral discussions because the capacity to suffer or feel is ethically significant.

Wakefulness must also be distinguished from experience. A person can be awake but not fully aware of certain stimuli. A person under anaesthesia may lose conscious experience. A patient with locked-in syndrome may be fully conscious but unable to move or communicate normally. These cases show that outward behaviour is not always a reliable guide to inner experience.

Another distinction is between the state of consciousness and the contents of consciousness. The state refers to the overall condition of being conscious, dreaming, anaesthetized, asleep, or minimally conscious. The contents are the specific experiences within consciousness: a sound, a colour, a memory, a pain, a thought, or an emotion.

Finally, consciousness has both first-person and third-person dimensions. From the first-person perspective, consciousness is lived experience. From the third-person perspective, it is studied through behaviour, brain activity, physiology, and report. The difficulty is that no third-person description seems identical to the first-person experience itself.

These distinctions will recur throughout the book. Without them, the question “when did consciousness emerge?” becomes too vague. We must ask: which consciousness, in what sense, and by what evidence?

3.8 The Explanatory Gap

The explanatory gap is the distance between physical descriptions and subjective experience. Even if we describe every neural event associated with seeing red, feeling pain, or hearing music, there seems to remain a question: why does this physical process feel like something from the inside?

This gap is not simply a lack of scientific detail. It may be a deeper conceptual problem. A complete map of brain activity might show where and when information is processed, but it may still not explain why experience exists at all.

Several philosophical arguments illustrate the problem. The zombie argument asks us to imagine a being physically and behaviourally identical to a human being but with no inner experience. Such a being would act like a conscious person but would be empty inside. Whether such a being is truly possible is debated, but the argument highlights the apparent distinction between function and experience.

The knowledge argument, often associated with the story of Mary the colour scientist, asks whether a person who knows all the physical facts about colour vision would learn something new upon seeing colour for the first time. If she does, then subjective experience may involve something not captured by physical description alone.

These arguments do not prove that science cannot explain consciousness. But they show why consciousness is not easily reduced to ordinary physical explanation. There appears to be a difference between explaining what a system does and explaining what it is like to be that system.

For scientific theories, the explanatory gap creates a challenge. A theory of consciousness must explain not only behaviour, attention, memory, and information processing, but also why any of these should be accompanied by experience. Some theories attempt to close the gap by identifying consciousness with certain forms of information integration or global availability. Others argue that consciousness may be a fundamental feature of reality. Still others suggest that the gap may eventually dissolve as science develops better concepts.

For this book, the explanatory gap matters because it affects how we understand the relationship between life and consciousness. If subjective experience cannot be fully explained by physical mechanism alone, then the

emergence of consciousness from life becomes a deep problem. If, however, experience arises naturally from certain forms of living organization, then life may already contain the conditions needed for consciousness.

3.9 Implications for the Central Question

The definition of consciousness strongly shapes the possible answers to the central question of this book.

If consciousness is binary, then there must have been a moment when it appeared. Before that moment, the universe contained no experience. After that moment, experience existed. This view creates a sharp boundary between non-conscious life and conscious life. The difficulty is identifying where that boundary lies.

If consciousness is graded, then it may not have appeared all at once. It may have developed gradually through increasing forms of sensitivity, sentience, perception, memory, and self-awareness. In this view, consciousness has roots rather than a single beginning. It may extend further back into animal life, and perhaps into more basic biological processes.

If consciousness is defined as self-consciousness, then it is likely a late evolutionary achievement. It depends on complex cognition, memory, social awareness, and reflective thought. Under this definition, life clearly comes first, and consciousness follows.

If consciousness is defined as phenomenal experience or sentience, the answer becomes less certain. Many animals may have experience even without language or self-reflection. The evolutionary roots of consciousness may then be older than human consciousness and may reach into the earliest nervous systems.

If consciousness is defined as proto-consciousness or basic subjectivity, then the boundary may move even further back. Consciousness may be linked to the earliest forms of biological responsiveness, self-maintenance, or organism-environment relation. This does not mean that cells think or feel as humans do, but it raises the possibility that experience has primitive roots in life itself.

If consciousness is fundamental, then the question changes entirely. Consciousness may not emerge from life, but life may organize, express, or localize consciousness in particular forms. This view opens the possibility that consciousness precedes biological life.

The definition we choose therefore constrains the answers available. A narrow definition supports a life-first view. A broader definition allows gradual emergence. A fundamental definition allows consciousness-first theories. A relational definition may support co-emergence.

The question “what is consciousness?” is therefore not separate from the question “which came first?” It determines what kind of answer can even be imagined.

3.10 How This Chapter Changes the Central Question

This chapter changes the central question by showing that the relationship between life and consciousness has deep historical roots. Ancient and classical traditions often treated life, soul, mind, and cosmos as connected rather than separate problems.

The question therefore becomes historically layered. Modern science asks how life and consciousness arise through physical processes, but older traditions remind us that the issue has also been framed in terms of soul, animation, order, purpose, and the place of mind within reality.

3.11 Chapter Summary

This chapter introduced the major meanings and distinctions associated with consciousness.

Phenomenal consciousness refers to subjective experience, or what it is like to have an experience. Access consciousness refers to information available for reasoning, reporting, and behavioural control. Minimal consciousness and proto-consciousness refer to possible primitive or pre-reflective forms of experience. Self-consciousness and higher-order consciousness involve awareness of oneself and awareness of one's own mental states. This chapter has focused mainly on philosophical and scientific definitions of consciousness, but later chapters will also examine consciousness-first frameworks, including Taheri's T-Consciousness, where consciousness is treated as a fundamental non-material reality rather than a product of the brain.

The chapter also distinguished consciousness from awareness, sentience, cognition, wakefulness, and behavioural responsiveness. These distinctions matter because debates about consciousness often become confused when different meanings are treated as if they were the same.

The explanatory gap remains central. Physical and functional descriptions may explain what systems do, but it remains difficult to explain why any system should have subjective experience. This difficulty shapes debates in neuroscience, philosophy, artificial intelligence, animal cognition, and theories of life.

For the central question of this book, the key lesson is that consciousness may not be one simple phenomenon. If it means reflective self-awareness, it appears late in evolution. If it means sentience, it may extend much further back. If it means proto-experience, it may be connected to life at its roots. If it is fundamental, it may precede life altogether.

The open question is therefore:

Is consciousness one thing, or are we using one word for several fundamentally different phenomena?

Part II

Part II: Historical and Philosophical Roots

Chapter 4

Ancient and Classical Perspectives

Long before biology and neuroscience became specialized sciences, thinkers asked whether mind belonged only to humans, to living beings, to nature as a whole, or to reality itself. These older questions still echo beneath modern debates, reminding us that the boundary between life and consciousness has never been purely scientific, nor purely philosophical.

4.1 Chapter Overview

Long before modern biology, neuroscience, or philosophy of mind became separate disciplines, thinkers asked how life, mind, matter, soul, and cosmos were related. The question was not usually framed in modern terms. Ancient philosophers did not ask whether consciousness emerged from abiogenesis, nor did they speak of information processing, neural correlates, or artificial intelligence. Yet they asked questions that remain recognizable: Is matter passive or active? Is life a special property or a universal principle? Is mind located only in living beings, or does it belong to nature as a whole?

This chapter traces early approaches to the relationship between life and consciousness, beginning with pre-Socratic thinkers and moving through atomism, Plato, Aristotle, Stoicism, and Neoplatonism. These traditions offer very different answers. Some treat matter as alive or animated. Some explain soul and mind in material terms. Some place consciousness or soul before the material world. Others understand life and mind as inseparable aspects of embodied form.

The goal is not to romanticize ancient thought or treat it as a substitute for science. Rather, the aim is to show that the modern separation between life and mind is historically recent. In many classical frameworks, the living, the mental, and the cosmic were deeply connected from the start.

4.2 Pre-Socratic Thought

The earliest Greek philosophers did not begin with a strict separation between living matter and non-living matter. Many pre-Socratic thinkers searched for an underlying principle from which the world emerged. This principle was often material, but it was not necessarily passive. Matter was frequently understood as active, generative, and self-moving.

This view is sometimes described as hylozoism: the idea that matter is in some sense alive. The term comes from the Greek words for matter and life. Hylozoism does not necessarily mean that every stone or river

is conscious in a human sense. Rather, it suggests that the basic stuff of the world may already contain vitality, movement, or animating power.

Thales is often associated with the claim that water is the fundamental principle of all things. In later accounts, he is also said to have believed that “all things are full of gods.” Whether this phrase should be taken literally or symbolically, it expresses a worldview in which matter is not inert. Water is not merely a substance; it is a living source, a generative principle, a basis from which things arise.

Anaximander proposed the *apeiron*, often translated as the boundless or indefinite, as the origin of all things. From this boundless source, worlds emerge and return. He also speculated about the emergence of living beings from primordial conditions. Although his view is not a modern theory of evolution, it suggests that life arises from a larger generative process in nature.

Anaximenes identified air as the primary principle. Air, breath, and life were closely connected in many ancient cultures. Breath was not only a physical substance but a sign of animation. To breathe was to live. To lose breath was to die. By treating air as fundamental, Anaximenes linked cosmology with vitality.

These early thinkers did not sharply divide physics, biology, and psychology. The same principle that explained the cosmos could also explain life. Matter was not yet understood as dead material waiting to be animated from outside. Instead, the world itself was dynamic, fertile, and internally active.

For the central question of this book, pre-Socratic thought is important because it begins from a different assumption than modern mechanistic science. It does not ask how life and consciousness emerged from dead matter. It asks how the living order of nature differentiates itself from a primordial source. In this sense, matter and life were not originally separated.

4.3 The Atomists

The atomists offered a very different approach. Leucippus and Democritus argued that reality consists of atoms and void. Atoms are indivisible particles moving through empty space. Everything that exists, including bodies, worlds, perception, and soul, arises from the arrangement and motion of atoms.

This was one of the earliest forms of materialism. Rather than explaining life through divine purpose or cosmic soul, the atomists sought to explain natural phenomena through matter in motion. The world did not need to be guided by an external intelligence. Its forms arose from combinations, collisions, and rearrangements of atoms.

The soul, in this framework, was also material. Democritus described the soul as composed of especially fine, mobile atoms, often associated with fire or heat. These soul atoms gave living beings movement and perception. Consciousness was not immaterial; it was a physical phenomenon arising from a particular kind of atomic organization.

This view has important implications. It suggests that mind is not separate from nature but part of nature. Consciousness, perception, and life are continuous with the material world. There is no need to posit a separate spiritual substance. At the same time, the atomist view leaves open the question of how subjective experience arises from material arrangements. If soul is made of atoms, why should certain atomic motions feel like anything?

Epicurus later developed atomism in a more ethical and existential direction. He accepted that atoms and void are fundamental, but he introduced the idea of the *swerve*: a slight indeterminacy in atomic motion. This swerve was important because it allowed room for contingency, freedom, and agency in a world otherwise governed by necessity.

The atomists therefore offer an early life-first or matter-first framework. Matter exists first. Life and consciousness arise from the organization and motion of matter. This resembles some modern scientific approaches, although ancient atomism lacked modern chemistry, biology, and neuroscience.

For the purposes of this book, the atomists show that material explanations of consciousness are ancient, not modern inventions. They also show that the central difficulty of materialism is ancient as well: how can arrangements of matter produce experience, agency, or inner life?

4.4 Plato and the Soul

Plato moved in a very different direction. For Plato, the soul is not merely a material arrangement. It is the principle of life, reason, and movement. The soul gives order to the body and is connected to realities beyond the changing physical world.

In several dialogues, Plato presents the soul as pre-existing the body and surviving bodily death. In the *Phaedo*, the soul is associated with reason, immortality, and the search for truth beyond the senses. In the *Phaedrus*, the soul is described through images of movement, ascent, and recollection. Human knowing is not simply sensory processing; it is a recovery of deeper truths.

Plato also describes the soul as having different parts. In the *Republic*, the soul is often divided into reason, spirit, and appetite. Reason seeks truth and wisdom. Spirit is associated with courage, honour, and emotional force. Appetite is associated with desire and bodily needs. Human life, in this view, is shaped by the ordering or disordering of these aspects of the soul.

Most important for the present book is Plato's cosmology in the *Timaeus*. There, Plato presents the cosmos itself as a living being ordered by intelligence. The World Soul animates and structures the universe. Mind is not merely a late product of biological evolution. It is built into the intelligible order of the cosmos.

This makes Plato an early example of a consciousness-first or mind-first framework. The physical world is not primary in the deepest sense. It is shaped by forms, order, intelligence, and soul. Life and mind are not accidental by-products of matter; rather, matter participates in a larger rational and living order.

Of course, Plato's view is not a scientific theory in the modern sense. It does not explain the chemical origin of life or the neural basis of consciousness. But it raises a question that remains alive: is consciousness something produced by matter, or is matter intelligible only because mind-like order is already present?

For Plato, the soul is not reducible to body. Life and consciousness point beyond physical mechanism toward a deeper metaphysical structure.

4.5 Aristotle and the Scale of Nature

Aristotle offered a more biological and embodied account of soul. In *De Anima*, he defines the soul as the form of a living body. This does not mean that the soul is a ghost inside the body. Rather, the soul is the organizing principle that makes a living body alive.

Aristotle's view is often calledhylomorphism, from the Greek terms for matter and form. A living being is not matter alone, nor form alone, but matter organized by form. The body is not a machine inhabited by a separate soul. The soul is the actuality of the body as a living organism.

This view allowed Aristotle to describe different levels of soul. Plants possess the nutritive soul, which allows growth, nourishment, and reproduction. Animals possess the sensitive soul, which allows sensation, desire, and movement. Humans possess the rational soul, which allows thought, reflection, and reason.

This framework is important because it links life and mind through degrees of organization. Life is not one thing and consciousness another completely separate thing. Rather, living beings display different powers.

Plants live and grow. Animals sense and move. Humans reason and reflect. Mind appears within life, not outside it.

Aristotle's account also provides an early version of a graded view. Consciousness, or at least perception and awareness, does not simply appear as an all-or-nothing event. It belongs to living beings according to their capacities. Plants are alive but not perceptive. Animals are perceptive but not necessarily rational. Humans are rational animals.

The later idea of the *scala naturae*, or scale of nature, is often associated with Aristotelian thinking. It places beings along a continuum from simpler to more complex forms. Although this idea later became tied to hierarchical and problematic interpretations of nature, it also reflects an important intuition: life and mind may develop through degrees.

For the central question of this book, Aristotle offers a middle path. He does not reduce soul to matter in the atomist sense, but he also does not separate soul from body in the Platonic sense. Life and mind are embodied forms of organization. Consciousness emerges within living beings, but it does so through the structured powers of life itself.

This makes Aristotle especially relevant to contemporary debates about embodiment, biological naturalism, and the continuity between life and cognition.

4.6 Stoicism and Pneuma

The Stoics developed a cosmology in which nature is pervaded by an active rational principle. They described this principle through concepts such as *pneuma* and *logos*. *Pneuma* can be translated as breath, spirit, or vital tension. It is the active force that structures and animates bodies. *Logos* refers to reason, order, or rational principle.

For the Stoics, the cosmos is not a dead mechanism. It is an ordered, living whole. Everything in nature participates in the rational structure of the universe. The world itself can be understood as a living being governed by *logos*.

Stoic *pneuma* operates at different levels. At the most basic level, it provides cohesion, holding physical things together. At a higher level, it appears as nature, organizing growth and development in living beings such as plants. At a still higher level, it appears as soul in animals, enabling perception and movement. In rational beings, it appears as reason.

This layered view resembles a graded theory of life and mind. The same active principle operates throughout nature, but in different degrees and forms. Matter is not lifeless substance acted upon from outside. It is structured by an internal principle of order and activity.

This has led some interpreters to see Stoicism as an early form of panpsychism or cosmopsychism, although the fit is not exact. The Stoics did not simply claim that every object has private consciousness. Rather, they saw the cosmos as pervaded by rational, active, organizing fire or breath. Mind belongs to nature as a whole and appears in different ways across different forms of being.

For the central question of this book, Stoicism is important because it refuses to isolate consciousness inside individual brains. Mind is not merely personal. It is connected to cosmic order. Human reason is a local expression of a wider rationality in nature.

This does not answer the origin-of-life question in modern terms. But it offers an alternative framework: life and consciousness may not be isolated accidents, but expressions of a universe already structured by active order.

4.7 Neoplatonism

Neoplatonism, especially in the work of Plotinus, developed one of the most sophisticated consciousness-first frameworks in ancient philosophy. Plotinus described reality as flowing from the One, through Nous, or divine Mind, into Soul, and finally into the material world.

The One is beyond being and beyond thought. It is the ultimate source of all reality. From the One emanates Nous, the realm of intellect or mind, where the intelligible forms are fully present. From Nous emanates Soul, which mediates between the intelligible and material worlds. The material world is the lowest level of this unfolding, not because it is evil, but because it is furthest from the unity of the source.

In this framework, consciousness or mind is not a late development. It is closer to the origin of reality than matter. Matter does not produce mind; rather, mind and soul precede and structure the world of matter. The direction of explanation is therefore reversed from modern materialism.

Neoplatonism is not simply a mystical doctrine. It is a highly organized metaphysical system. It asks how unity gives rise to multiplicity, how intelligible order becomes embodied, and how individual souls relate to cosmic mind. It also asks how human consciousness can turn back toward its source through contemplation.

For the purposes of this book, Neoplatonism represents a powerful consciousness-first ontology. It places mind before life, and life before matter as ordinarily understood. The living world is an expression of a deeper order of soul and intellect.

Its influence was enormous. Neoplatonic ideas shaped later Christian, Islamic, and Jewish philosophy, as well as medieval metaphysics, mystical traditions, and Renaissance thought. Through these traditions, the idea that consciousness, intellect, or soul is prior to matter remained central for centuries.

Modern science largely moved away from this framework. Yet the questions Neoplatonism raised have not disappeared. Is mind derivative, or is it fundamental? Is matter primary, or is it an expression of deeper intelligible order? Can consciousness be fully explained from below, or must explanation also move from above, from form, unity, or meaning?

These questions continue to echo in contemporary discussions of idealism, panpsychism, and consciousness-first theories.

4.8 Implications for the Central Question

Ancient and classical perspectives show that the separation between life, mind, and matter is not inevitable. It is a historically specific development. Many early thinkers did not begin with dead matter and then ask how life and consciousness emerged from it. They began with a world that was active, ordered, animated, or intelligible.

Pre-Socratic hylozoism suggests that matter and life were once understood as deeply connected. Atomism offers an early materialist account in which soul and consciousness arise from physical organization. Plato places soul and mind before the material world, offering an early consciousness-first framework. Aristotle locates soul within the living body, suggesting that life and mind are inseparable aspects of embodied form. Stoicism presents a cosmos pervaded by rational pneuma. Neoplatonism reverses the modern order entirely, treating mind and soul as prior to matter.

These frameworks foreshadow many contemporary debates. Materialism resembles the atomist tradition. Panpsychism and cosmopsychism recall Stoic and Neoplatonic themes. Biological theories of mind echo Aristotle's view that soul is the form of the living body. Consciousness-first theories resemble Platonic and Neoplatonic ontologies.

The historical lesson is not that ancient philosophers had all the answers. They did not possess modern biology, neuroscience, evolutionary theory, or chemistry. But they often saw life and mind as connected problems rather than isolated disciplines.

For the central question of this book, this matters. The question “which came first, life or consciousness?” is partly a modern question because modern thought has separated life from mind. Ancient thought reminds us that other starting points are possible.

4.9 How This Chapter Changes the Central Question

This chapter changes the central question by showing that some traditions do not begin with matter as primary. Eastern and contemplative traditions often treat consciousness as pervasive, foundational, or inseparable from reality.

The question therefore expands beyond a strictly material sequence. Instead of asking only how matter became alive and then conscious, we must also consider traditions in which consciousness is prior to, deeper than, or inseparable from living form.

4.10 Chapter Summary

This chapter examined ancient and classical approaches to the relationship between life, mind, matter, and soul.

Pre-Socratic thinkers often treated matter as active, living, or animated. The atomists explained soul and consciousness in material terms through atoms and void. Plato presented the soul as pre-existing, immortal, and cosmological, especially through the idea of the World Soul. Aristotle defined the soul as the form of the living body and described different levels of life, sensation, and reason. The Stoics understood nature as pervaded by pneuma and logos, an active rational principle. Neoplatonism placed mind and soul before matter in a consciousness-first metaphysical system.

Together, these traditions show that the modern separation between life and consciousness is not the only possible framework. Classical thought often treated life, mind, and cosmic order as interconnected. Some traditions moved toward materialism, others toward consciousness-first metaphysics, and others toward embodied continuity between life and mind.

The open question is therefore:

Did ancient philosophers preserve insights about the unity of life and mind that modern disciplinary specialization has caused us to lose?

Chapter 5

Eastern Philosophical Traditions

5.1 Chapter Overview

The question of whether consciousness precedes life, emerges from life, or co-arises with life is not only a modern scientific problem. Many Eastern philosophical traditions developed sophisticated accounts of consciousness, selfhood, life, matter, and reality long before the rise of modern biology or neuroscience. These traditions often begin from assumptions very different from those found in modern Western materialism.

In many Eastern frameworks, consciousness is not treated as a late by-product of biological complexity. It may be understood as fundamental reality, as an aspect of all living beings, as a stream without a fixed self, as a function of relational existence, or as part of the dynamic unfolding of nature. These traditions do not all agree with each other. Hindu, Buddhist, Daoist, Jain, and Indigenous perspectives differ deeply in their metaphysics, methods, and ethical implications. Yet they share an important tendency: they resist a simple division between mind and world, observer and observed, life and environment.

This chapter surveys several traditions relevant to the central question of this book. The aim is not to reduce these traditions to modern scientific categories, nor to treat them as evidence in a simplistic way. Rather, the goal is to examine how they widen the conceptual space. They show that the relationship between consciousness and life can be framed in ways that are not limited to the modern question of how brains produce experience.

5.2 Hindu Philosophy and Consciousness

Hindu philosophical traditions contain some of the most developed consciousness-first frameworks in world philosophy. Although Hindu thought is diverse and internally debated, many of its major traditions place consciousness near the foundation of reality rather than treating it as a late evolutionary product.

In Advaita Vedanta, ultimate reality is Brahman: infinite, non-dual consciousness or being. The deepest self, Atman, is not separate from Brahman. The famous Upanishadic phrase “Tat tvam asi,” often translated as “You are that,” expresses the idea that the individual self is not ultimately separate from the ground of reality. Consciousness is not merely something possessed by an organism. It is the underlying reality within which the world appears.

In this view, the material world is often described through the concept of maya. Maya does not simply mean illusion in the sense of something unreal or meaningless. Rather, it refers to the appearance of multiplicity, separation, and change within a deeper non-dual reality. The world is experienced as divided into subjects

and objects, minds and bodies, living and non-living things. But from the Advaitic perspective, this division does not describe ultimate reality.

This framework is radically different from the standard scientific sequence of matter first, life second, consciousness third. In Advaita Vedanta, consciousness is not produced by life. Instead, life and matter appear within consciousness. The living body is not the origin of awareness, but one expression or locus through which awareness is manifested.

Samkhya offers a different but equally important framework. It is dualistic rather than non-dual. Samkhya distinguishes between purusha, pure consciousness, and prakriti, primordial matter or nature. Purusha is passive, witnessing, and conscious. Prakriti is active, dynamic, and productive. The world of experience arises through the interaction or apparent conjunction of these two principles.

In Samkhya, life is not reducible to matter alone because matter by itself is not conscious. Yet consciousness by itself does not act or produce the world. Living experience emerges through the relation between purusha and prakriti. This provides a different model from both materialism and idealism. Consciousness and matter are both fundamental, but they play different roles.

These Hindu traditions are important for the central question because they offer mature consciousness-first or consciousness-fundamental ontologies. They challenge the assumption that consciousness must be explained only as an outcome of biological evolution. Instead, they ask whether biological life is one way in which consciousness becomes embodied, localized, or expressed.

A modern reader need not accept these metaphysical claims literally in order to learn from them. Their importance lies in the questions they make possible. What if consciousness is not something added to matter from the outside, nor produced by matter from the inside, but the condition through which the world is known at all? What if life is not the origin of consciousness, but one of its appearances?

5.3 Buddhist Perspectives

Buddhist traditions approach consciousness differently from many Hindu traditions. Rather than identifying consciousness with an eternal self or ultimate essence, Buddhism often analyzes consciousness as a conditioned, impermanent process. Consciousness is real as experience, but it is not grounded in a permanent, unchanging self.

A central Buddhist teaching is dependent origination. Nothing exists independently or in isolation. All phenomena arise in dependence on causes, conditions, relations, and processes. This applies to bodies, thoughts, perceptions, emotions, and identities. The self is not an independent substance but a pattern of changing aggregates.

In classical Buddhist psychology, consciousness is one of the five aggregates, or skandhas. These aggregates are form, sensation, perception, mental formations, and consciousness. Together they make up what we ordinarily call a person. But none of them is permanent, self-sufficient, or ultimately identical with a fixed self.

This framework is important because it allows consciousness without a permanent soul. Consciousness is not denied, but it is understood as momentary, relational, and conditioned. It arises with objects, sensations, perceptions, and mental formations. It is not an isolated substance standing outside the world.

The idea of a mind-stream is also important. Buddhist thought often describes continuity without identity. There is continuity of experience, memory, karma, and causation, but not a permanent self that remains unchanged through time. This offers a subtle model of consciousness as process rather than thing.

Buddhist traditions also developed detailed taxonomies of mental states. Long before modern psychology, Buddhist thinkers analyzed attention, perception, desire, suffering, mindfulness, concentration, emotion,

and altered states of awareness. These analyses were not merely theoretical. They were connected to contemplative practices designed to observe consciousness directly.

Yogacara, sometimes called the “mind-only” school, goes further by emphasizing the constructed nature of experience. It argues that what we take to be an external world is deeply shaped by consciousness, perception, and karmic tendencies. This does not necessarily mean that nothing exists outside the mind in a simplistic sense. Rather, it suggests that the world we experience is inseparable from the structures of cognition and consciousness.

For the central question of this book, Buddhism provides a distinctive alternative. It does not simply say that consciousness comes before life or after life. Instead, it asks us to examine consciousness as a dependently arising process. Consciousness has no fixed essence, but it also cannot be reduced to an isolated object. It arises through relations.

This is relevant to co-emergence views. If life and consciousness are both relational processes, then the question may not be which came first as a separate entity. The deeper question may be how conditions give rise to forms of experience, responsiveness, and world-making.

5.4 Daoist Thought

Daoist philosophy offers a different way of thinking about life, consciousness, and nature. Rather than beginning with fixed substances or sharp categories, Daoism emphasizes flow, transformation, spontaneity, and alignment with the natural way of things.

The Dao is the source, pattern, and unfolding of all things. It cannot be fully named or captured by concepts. It is not a personal creator in the ordinary sense, nor merely a physical law. It is the generative way through which the world arises, changes, and returns. All beings participate in the Dao, whether consciously or unconsciously.

This vision has important implications for life. Life is not an exception to nature. It is one expression of nature’s spontaneous self-ordering. Living systems arise, transform, and dissolve within a larger process that does not need to be controlled from outside. The Daoist emphasis on naturalness suggests that complexity and order can emerge without rigid design.

Wu wei, often translated as non-action or effortless action, does not mean doing nothing. It means acting in accordance with the flow of things rather than imposing forceful control. In relation to the themes of this book, wu wei can be read as a philosophical insight into self-organization. Systems can organize through attunement, responsiveness, and relational balance rather than through central command.

Qi is another important concept in Chinese thought. It is often translated as vital energy, breath, or life force. Qi pervades bodies, environments, and natural processes. It links life to movement, breath, vitality, and transformation. Although qi should not be simply equated with modern physical energy, it expresses a worldview in which living processes are continuous with the dynamics of nature.

The writings associated with Zhuangzi are especially relevant to consciousness. The famous butterfly dream asks whether Zhuangzi dreamed he was a butterfly, or whether a butterfly is dreaming it is Zhuangzi. The story does not merely ask about illusion. It questions the stability of identity, the boundary between perspectives, and the certainty with which one viewpoint claims reality.

Daoist thought therefore destabilizes rigid distinctions between self and world, subject and object, human and nature. Consciousness is not treated as a detached observer standing outside life. It is part of the living flow of transformation.

For the central question of this book, Daoism offers a framework close to co-emergence. Life and awareness are not separate substances but movements within a larger process. Consciousness is not necessarily a late addition to nature; it is one way nature becomes responsive, reflective, and aware through its own unfolding.

5.5 Jain Philosophy

Jain philosophy provides one of the clearest and most explicit graduated frameworks of consciousness. In Jain thought, *jiva* refers to soul or living consciousness. *Jivas* are present in all living beings, not only humans or animals. Consciousness is a defining feature of life.

This view leads to a broad moral and metaphysical vision. Plants, animals, microorganisms, and even elemental beings are understood as possessing forms of life and consciousness. Consciousness is not equally developed in all beings, but it is widely distributed.

Jain philosophy classifies living beings according to the number of senses they possess. One-sensed beings include entities associated with earth, water, fire, air, and plants. More complex beings have two, three, four, or five senses. Five-sensed beings with mind represent a higher degree of conscious capacity. This hierarchy does not deny consciousness to simpler beings; rather, it describes consciousness as graded.

This is highly relevant to the central question of the book. Jain thought does not treat consciousness as an all-or-nothing property that appears only in complex animals. It treats consciousness as present wherever there is life, though expressed in different degrees. The simplest living beings have minimal consciousness; more complex beings have richer forms of perception and mental life.

The ethical consequence is *ahimsa*, or non-violence. If living beings possess consciousness, then harm matters. Jain ethics extends moral concern far beyond humans. It encourages careful attention to the living world, including forms of life that may be overlooked in ordinary human activity.

From a modern scientific perspective, one may question whether all beings classified in Jain cosmology should be understood literally as conscious. However, philosophically, Jainism offers a powerful model of graduated sentience. It shows what an ethics of widespread consciousness might look like.

For the purposes of this book, Jain philosophy provides a direct challenge to the idea that consciousness is only a late evolutionary product. It suggests instead that life and consciousness are inseparable, though consciousness may appear in simpler or more complex degrees.

This makes Jain thought especially important for co-emergence theories and for discussions of moral status. If consciousness begins with life, then the origin of life is also, in some minimal sense, the origin of sentience.

5.6 Indigenous and Non-Western Perspectives

Many Indigenous and non-Western knowledge systems understand consciousness, life, and worldhood through relational rather than strictly individualistic frameworks. Although these traditions are diverse and cannot be reduced to a single view, many share an emphasis on relationship, reciprocity, place, ancestry, land, and the agency of the natural world.

Animism is often used to describe traditions in which animals, plants, rivers, mountains, winds, ancestors, or places are understood as alive, aware, or spiritually significant. The term has sometimes been used dismissively in Western scholarship, as though such views were primitive projections of human consciousness onto nature. More recent philosophical and anthropological work has challenged this interpretation. Animistic perspectives can instead be understood as sophisticated relational ontologies.

A relational ontology begins not with isolated substances but with relationships. A person is not an independent mind enclosed in a body, standing apart from nature. A person exists through relations with family, community, land, animals, ancestors, waters, plants, and spiritual forces. The self is embedded in a living world.

This has important implications for consciousness. Consciousness is not only interior experience. It may also involve attention, reciprocity, communication, responsibility, and participation in a wider field of relations. The world is not merely a collection of objects to be measured and used. It is a community of beings, forces, and presences.

Indigenous ecological knowledge often emphasizes that humans are not outside the living world. Knowledge arises through long-term observation, practice, story, ceremony, and relationship with place. Such knowledge may include empirical insight, ethical orientation, and spiritual meaning at once. These dimensions are not always separated in the way modern academic disciplines separate science, philosophy, religion, and ethics.

It is important to approach these traditions with care. They should not be extracted as decorative support for modern theories, nor treated as interchangeable examples of “ancient wisdom.” Serious engagement requires respect for context, sovereignty, language, community, and lived practice. Reconciliation requires centering Indigenous voices rather than simply using Indigenous ideas to enrich non-Indigenous frameworks.

For the central question of this book, Indigenous and relational perspectives offer a crucial reminder: the observer is not separate from the world being observed. Life and consciousness may not be best understood as isolated properties inside individual organisms. They may be relational phenomena that emerge through participation in a living world.

This does not mean that all Indigenous traditions make the same metaphysical claim. Rather, it means that many offer alternatives to the modern image of consciousness as a private mental state inside an individual brain. Consciousness may also be ecological, relational, and world-involving.

5.7 Comparative Analysis

Eastern and non-Western traditions differ greatly from one another, but several common themes are relevant to the central question of this book.

First, many of these traditions treat consciousness as more fundamental than modern scientific materialism usually allows. Advaita Vedanta identifies ultimate reality with consciousness. Samkhya treats consciousness and matter as dual principles. Yogacara emphasizes the mind-structured character of experience. Jainism treats consciousness as present in all living beings. Many Indigenous relational traditions understand the world as alive, agentic, or spiritually responsive.

Second, these traditions often place life and mind on a continuum rather than separating them sharply. Jainism explicitly grades consciousness across living beings. Daoism presents life and awareness as movements within the Dao. Buddhist thought analyzes consciousness as a conditioned process rather than a fixed substance. Indigenous perspectives often understand living beings as participants in networks of reciprocal awareness and responsibility.

Third, many of these traditions challenge the idea of the observer as detached and neutral. In modern science, the observer is often imagined as standing outside the system being studied. In contemplative and relational traditions, the observer is part of the process. How one attends, practices, perceives, and relates can shape what is known.

This does not mean that Eastern traditions and Western science should be simply merged. Their methods, assumptions, and goals differ. Scientific methods rely on third-person observation, measurement, replication, and public verification. Contemplative traditions often rely on disciplined first-person inquiry, ethical transformation, and direct examination of experience. Relational traditions may emphasize community, place, responsibility, and intergenerational knowledge.

The challenge is not to choose one and reject the other. The challenge is to ask what each method can reveal and where each has limits. Third-person science is powerful for studying bodies, brains, evolution, chemistry, and behaviour. First-person and contemplative methods may be powerful for examining experience

from within. Relational knowledge may reveal how consciousness and life are embedded in environments, communities, and ethical relations.

Western science can learn from these traditions not by abandoning empirical rigor, but by expanding its questions. It can ask whether consciousness is only an object to be measured from the outside, or also a field of experience that requires disciplined first-person investigation. It can ask whether life is only a biochemical process, or also a relational process of sense-making, world-involvement, and value.

These traditions do not provide simple answers. They provide alternative starting points.

5.8 Implications for the Central Question

Eastern and non-Western traditions widen the possible answers to the question of whether life or consciousness came first.

Hindu consciousness-first traditions suggest that consciousness may be the ground within which life and matter appear. From this perspective, the origin of life is not the origin of consciousness. It is the emergence of living forms within consciousness.

Buddhist perspectives suggest a different possibility: consciousness is neither a permanent substance nor a mere by-product of matter. It is a dependently arising process. This supports relational and co-emergent approaches, where life and consciousness are understood through conditions, processes, and mutual dependence.

Daoist thought suggests that life and awareness may be natural expressions of spontaneous self-organization. Rather than forcing a choice between matter-first and consciousness-first, Daoism points toward a dynamic view in which nature unfolds through relational transformation.

Jain philosophy offers a strong graduated model, in which consciousness is present across living beings in different degrees. This makes it especially relevant to the idea that consciousness and life may co-emerge.

Indigenous and relational perspectives remind us that consciousness may not be located only inside isolated individuals. It may be ecological, relational, and embedded in a living world. This challenges theories that treat consciousness as merely an internal computational state.

These traditions also raise methodological questions. Modern consciousness science relies heavily on third-person observation. Yet consciousness is also directly known in first-person experience. Contemplative traditions have developed disciplined methods for observing attention, perception, emotion, selfhood, and altered states. These practices may not replace neuroscience, but they may provide forms of data and insight that neuroscience alone cannot access.

The central question therefore becomes broader. It is not only whether consciousness came before or after life. It is also whether consciousness can be understood from a purely third-person perspective, or whether first-person and relational methods are necessary for a complete account.

5.9 How This Chapter Changes the Central Question

This chapter changes the central question by showing that modern philosophy frames the life-consciousness relationship through competing metaphysical assumptions. Dualism, physicalism, idealism, panpsychism, neutral monism, and process philosophy each begin from a different view of what reality is made of.

The question therefore becomes dependent on metaphysical starting points. If matter is primary, consciousness must be explained as an emergence from life or brain activity. If consciousness is primary, life may be understood as an expression within consciousness. If reality has both mental and physical aspects, life and consciousness may co-emerge from a deeper ground.

5.10 Chapter Summary

This chapter surveyed several Eastern and non-Western traditions relevant to the relationship between life, consciousness, and matter.

Hindu philosophy, especially Advaita Vedanta and Samkhya, offers powerful consciousness-first and consciousness-fundamental frameworks. Buddhist traditions analyze consciousness as impermanent, conditioned, and without a permanent self, offering a process-based and relational account. Daoist thought emphasizes the Dao, spontaneity, qi, and natural self-organization. Jain philosophy presents an explicit graduated framework in which consciousness is present throughout living beings. Indigenous and relational perspectives understand the self as embedded in a living world and challenge the image of consciousness as isolated interiority.

Across these traditions, life and consciousness are often more closely linked than in modern materialist frameworks. Consciousness may be fundamental, relational, graded, ecological, or inseparable from living processes. These perspectives do not settle the central question, but they reveal that the modern life-first view is not the only possible starting point.

They also raise a methodological challenge. If consciousness is known from within, then third-person science may not be sufficient by itself. Contemplative, ethical, and relational practices may contribute forms of insight that cannot be fully captured by external measurement alone.

The open question is therefore:

Can contemplative and relational traditions provide evidence about consciousness that third-person science cannot?

Chapter 6

Modern Philosophy of Mind

6.1 Chapter Overview

Modern philosophy of mind begins with a wound that has never fully healed: the division between mind and body. Once mind was separated from matter, consciousness became difficult to place within nature. Was it a substance outside the physical world? A property of the brain? A function of complex organization? An illusion created by language? Or a feature woven into reality itself?

This chapter traces several major philosophical frameworks that shaped the modern debate: Cartesian dualism, Spinoza's dual-aspect monism, Leibniz's monadology, Kant's limits of knowledge, phenomenology, physicalism, emergence, property dualism, and contemporary panpsychism. Each framework offers a different way of understanding the relationship between life and consciousness.

The chapter does not attempt to solve the mind-body problem. Instead, it shows how different philosophical starting points constrain the possible answers to the central question of this book. If mind and matter are fundamentally separate, then consciousness may not be explainable through life alone. If physicalism is correct, then consciousness must emerge from biological or physical processes. If panpsychism is correct, then consciousness may precede life or co-emerge with matter and life in a deeper way.

6.2 Descartes and the Mind-Body Split

René Descartes is often treated as the starting point of modern philosophy of mind because he gave one of the clearest formulations of mind-body dualism. For Descartes, mind and body are fundamentally different kinds of substance. The body is extended in space, measurable, divisible, and governed by mechanical laws. The mind is thinking substance: non-extended, indivisible, and known directly through self-awareness.

This distinction gave rise to substance dualism. According to this view, the mental and the physical are not merely different descriptions of the same thing. They are different kinds of reality. The body belongs to the material world; the mind belongs to a non-material domain.

Descartes famously argued that he could doubt the existence of the external world and even the existence of his body, but he could not doubt that he was thinking. The statement "I think, therefore I am" placed consciousness at the foundation of certainty. The thinking self became more immediately known than the living body.

Yet this created a problem. If mind and body are different substances, how do they interact? How can an immaterial mind cause a material body to move? How can bodily sensations influence thought and feeling?

Descartes proposed that the pineal gland might serve as the site of interaction between mind and body, but this explanation did not resolve the deeper problem.

The legacy of Cartesian dualism is profound. Modern science inherited a picture in which the body could be studied mechanically, while consciousness remained difficult to integrate into biology. The body became an object of third-person investigation. Consciousness became an interior, first-person reality that seemed to resist the methods of physical science.

This division has sometimes been called a Cartesian wound: a split between nature and experience, body and mind, organism and subject. It shaped the modern assumption that life and consciousness are separate problems. Biology could study living systems without addressing subjective experience. Philosophy could study consciousness without fully grounding it in living processes.

For the central question of this book, Descartes creates a strong separation. If consciousness is a non-material substance, then it does not emerge from life in any ordinary biological sense. Life and consciousness interact, but they do not belong to the same order of explanation.

6.3 Spinoza and Neutral Monism

Baruch Spinoza rejected Cartesian substance dualism. Instead of two substances, mind and body, Spinoza argued that there is only one substance: nature, or God-or-Nature. Mind and matter are not separate realities, but two attributes of the same underlying reality.

This view is often described as dual-aspect monism. The mental and the physical are two aspects of one substance. A human being can be understood under the attribute of thought or under the attribute of extension. These are different ways of expressing the same underlying reality, not two substances interacting from separate domains.

Spinoza's view is important because it avoids the interaction problem that troubled Descartes. If mind and body are two aspects of the same thing, then the question is not how an immaterial mind pushes a material body. Instead, the mental and physical unfold together as parallel expressions of one reality.

This framework is highly relevant to co-emergence models. If mind and matter are not fundamentally separate, then consciousness and life may be understood as different aspects of a single process. Life is not merely a physical container for consciousness, and consciousness is not a ghostly addition to life. They may be two ways of describing organized nature.

Spinoza also influenced later forms of neutral monism. Neutral monism suggests that the basic stuff of reality is neither purely mental nor purely physical, but something more fundamental that can appear as both. Contemporary dual-aspect theories and some versions of panpsychism owe much to this tradition.

Some interpreters also describe Spinoza as having panpsychist tendencies because he attributes mind-like aspects widely throughout nature. This does not mean that every object thinks like a human. Rather, everything expresses the one substance under the attribute of thought in some way. Human consciousness is a complex mode of a more general reality.

For the central question of this book, Spinoza offers an alternative to both life-first materialism and consciousness-first idealism. Life and consciousness may not stand in a simple before-and-after relation. They may be two aspects of the same underlying reality, expressed at different levels of complexity.

6.4 Leibniz and the Monadology

Gottfried Wilhelm Leibniz developed another influential alternative to mechanical materialism. In his *Monadology*, he proposed that reality is composed of monads: simple, immaterial units of perception or experience. Monads do not interact mechanically like physical particles. Instead, each monad expresses the universe from its own point of view.

For Leibniz, perception is not limited to human beings. All monads possess some form of perception, though most are extremely obscure or confused. Human consciousness is a highly developed form of perception, but it is not the only form. Reality is, in a sense, perspectival all the way down.

This view makes Leibniz important for contemporary panpsychism. He did not treat consciousness as something that suddenly appears from wholly non-conscious matter. Instead, he saw experience or perception as built into the basic units of reality.

Leibniz's famous mill argument illustrates his skepticism about purely mechanical explanations of consciousness. He asks us to imagine entering a machine enlarged to the size of a mill, where we can walk inside and observe all its moving parts. We would see pieces pushing and pulling one another, but we would never find perception itself. Mechanism alone seems to explain motion, but not experience.

This argument anticipates modern versions of the explanatory gap. Even if we understand every physical process, why should those processes be accompanied by inner experience? Where, in the machinery, does consciousness appear?

Leibniz also proposed pre-established harmony. Since monads do not causally interact in the ordinary mechanical sense, their states correspond because they have been harmonized within the order of creation. Mind and body appear to interact, but their coordination reflects a deeper metaphysical harmony.

For the central question of this book, Leibniz suggests that consciousness may not need to be produced from non-conscious matter. Instead, experience may be present in primitive form at the foundation of reality. Life may organize and amplify experience rather than generate it from nothing.

This does not solve all problems. Leibniz's monads are metaphysical entities, not scientific objects. Yet his view remains relevant because it challenges the assumption that matter is intrinsically devoid of experience. If the basic units of reality already have inner aspects, then the emergence of life may be part of a broader organization of experience.

6.5 Kant and the Limits of Knowledge

Immanuel Kant changed the question by asking what the human mind can know. Rather than making a direct claim about whether mind or matter is ultimately primary, Kant examined the conditions under which experience and knowledge are possible.

Kant distinguished between phenomena and noumena. Phenomena are things as they appear to us, structured by the forms of human intuition and understanding. Noumena are things as they are in themselves, independent of how they appear to us. According to Kant, we do not have direct access to things as they are in themselves. We know the world as it is shaped by the conditions of possible experience.

This has major implications for consciousness. Consciousness is not simply one object among others. It is also part of the structure through which objects appear. We can study brain states, behaviour, and bodily processes as phenomena, but the subject of experience is not easily turned into an ordinary object of knowledge.

Kant does not deny the reality of the physical world. Nor does he reduce everything to subjective imagination. Instead, he argues that knowledge depends on the relation between the world and the structures of experience. We never encounter reality from nowhere. We encounter it through the forms of human cognition.

This places limits on the scientific study of consciousness. Science can investigate the conditions under which consciousness appears, the neural and behavioural correlates of experience, and the biological structures that support awareness. But whether science can know consciousness “as it is in itself” remains a deeper philosophical question.

For the central question of this book, Kant introduces humility. We may ask whether consciousness came before life, emerged from life, or co-emerged with life. But our ability to answer may be limited by the structure of human knowledge. We are conscious beings investigating consciousness. We cannot step entirely outside experience in order to view it from a perfectly neutral standpoint.

Kant therefore does not give a simple answer to the life-consciousness question. Instead, he warns us that the question may exceed what can be known through ordinary empirical inquiry. This does not make the question meaningless. It means that any answer must recognize the limits of human perspective.

6.6 Phenomenology

Phenomenology begins not with abstract metaphysics or external observation, but with lived experience. It asks how things appear to consciousness and how meaning is constituted in experience. This makes phenomenology one of the most important modern traditions for reconnecting consciousness with embodiment, world, and life.

Edmund Husserl emphasized intentionality: the idea that consciousness is always consciousness of something. Experience is not an isolated inner substance. It is directed toward objects, meanings, possibilities, memories, and worlds. To be conscious is not merely to contain mental images; it is to be related to something.

This is important because it shifts attention from consciousness as a thing to consciousness as a relation. The subject and world are not completely separate. Experience is structured by directedness, attention, meaning, and appearance.

Martin Heidegger deepened this shift with the idea of being-in-the-world. Human existence is not first an isolated mind looking out at an external world. We are already involved in a meaningful world of tools, practices, relationships, concerns, and possibilities. Consciousness is not detached observation; it is embedded existence.

Maurice Merleau-Ponty placed the lived body at the centre of consciousness. For him, the body is not merely a physical object controlled by the mind. It is the way we inhabit the world. Perception, movement, gesture, and bodily orientation are fundamental to experience. The body is not something we simply have; it is part of how the world is disclosed to us.

Phenomenology is crucial for the relationship between life and consciousness because it refuses to separate experience from embodiment. Consciousness is not a disembodied spectator. It is lived through a body that moves, senses, acts, and belongs to a world.

This has strong implications for the central question. If consciousness is always embodied, then it may be deeply dependent on life. But embodiment does not reduce consciousness to mechanical biology. It suggests that consciousness emerges through the lived body: through perception, movement, vulnerability, need, and relation to the world.

Phenomenology therefore offers a powerful life-consciousness bridge. It does not claim that consciousness is merely a brain process, nor that it floats free from biology. It treats consciousness as world-involving, bodily, and lived.

6.7 Physicalism and Its Variants

Physicalism is the view that everything that exists is ultimately physical, or depends on the physical. In philosophy of mind, physicalism holds that mental states must be explained in terms of physical processes, usually brain processes.

One version is identity theory. According to identity theory, mental states are identical with brain states. Pain, for example, is not a separate non-physical event; it is a physical state of the nervous system. This view has the advantage of simplicity. It places consciousness directly within the natural world and avoids dualist interaction problems.

However, identity theory faces the challenge of explaining subjective experience. Even if a pain state is identical with a neural state, why does that neural state feel painful? The identity claim may connect mind and brain, but it may not fully explain the felt quality of experience.

Functionalism defines mental states by what they do rather than by what they are made of. A mental state is characterized by its causal role: what inputs produce it, what internal processes it affects, and what behaviours it generates. Pain, for example, may be defined by its role in detecting damage, motivating avoidance, guiding attention, and producing protective behaviour.

Functionalism is attractive because it allows multiple realizability. A mental state might be realized in brains, machines, or other systems if the functional organization is right. This makes functionalism influential in cognitive science and artificial intelligence.

But functionalism also faces the problem of experience. A system might perform the right functions without obviously having consciousness. This is why debates about philosophical zombies, artificial intelligence, and simulation remain powerful. Function may explain access consciousness, but it may not explain phenomenal consciousness.

Eliminative materialism takes a more radical approach. It argues that ordinary mental concepts such as belief, desire, or perhaps even consciousness may belong to outdated folk psychology. As neuroscience advances, these concepts may be replaced by more precise scientific categories.

This view is provocative because it challenges whether consciousness, as ordinarily understood, is even a valid explanatory category. Yet it also seems to conflict with the immediacy of experience. Even if our concepts are imperfect, the fact that experience appears to occur remains difficult to eliminate.

For the central question of this book, physicalism supports the life-first view. Matter gives rise to life; life gives rise to nervous systems; nervous systems give rise to consciousness. The strength of physicalism is that it keeps consciousness within nature. Its weakness is that it must explain how subjective experience arises from physical processes.

6.8 Property Dualism and Emergence

Property dualism offers a middle position between substance dualism and reductive physicalism. It does not claim that mind is a separate substance from body. Instead, it claims that conscious properties are real and may not be reducible to physical properties, even though they depend on physical systems.

This view is often connected to non-reductive physicalism. According to non-reductive physicalism, everything may depend on the physical, but higher-level properties cannot always be fully explained by lower-level descriptions. Biology depends on chemistry and physics, but biological concepts such as organism, function, adaptation, and reproduction are not easily replaced by particle descriptions.

Emergence is central here. An emergent property arises from a system but is not obvious from its parts taken separately. Weak emergence occurs when higher-level patterns arise from lower-level processes but are, in

principle, explainable through them. Strong emergence suggests that genuinely new properties appear that cannot be fully reduced to their physical base.

Consciousness is often treated as an emergent property of complex systems. On this view, no single neuron is conscious, but networks of neurons organized in certain ways may give rise to conscious experience. Life itself can also be viewed emergently. No single molecule is alive, but a network of molecules organized as a self-maintaining cell may be alive.

This makes emergence attractive for the central question. It allows us to say that life emerges from chemistry and consciousness emerges from life, without claiming that either is a simple mechanical sum of parts.

However, emergence also raises questions. If consciousness strongly emerges, then how does it fit within physical causation? Does it have causal power, or is it merely an accompaniment of brain activity? If it has causal power, how does that power operate? If it does not, why did consciousness evolve?

The explanatory gap returns here. Emergence may describe when consciousness appears, but it may not fully explain why experience appears. Saying that consciousness emerges from complexity may be true, but it risks becoming a label for a mystery unless the relevant kind of complexity is specified.

For the life-consciousness question, emergence offers a flexible framework. It supports the standard scientific sequence, but it also leaves room for continuity. Life may emerge from self-organizing chemistry; cognition may emerge from living regulation; consciousness may emerge from increasingly integrated biological systems.

6.9 Contemporary Panpsychism

Contemporary panpsychism has re-emerged as a serious position in philosophy of mind. It begins from a simple concern: if consciousness is entirely absent from the basic structure of reality, how does it suddenly appear in complex brains? Panpsychism avoids this problem by proposing that mind-like or experience-like properties are present, in primitive form, throughout nature.

Philosophers such as Galen Strawson, Philip Goff, and David Chalmers have explored versions of this view. They do not usually claim that electrons think, that rocks have human-like minds, or that all matter is conscious in the ordinary sense. Rather, the claim is that the intrinsic nature of matter may include proto-experiential aspects. Complex consciousness may arise when these primitive aspects are organized in certain ways.

Panpsychism has an important advantage: it takes consciousness seriously as a feature that cannot be easily derived from wholly non-conscious matter. It also offers a possible response to the hard problem. If experience is not created from nothing, then the emergence of consciousness becomes a problem of organization rather than production from absolute absence.

However, panpsychism faces the combination problem. If basic units of matter have tiny forms of experience, how do these combine into unified consciousness? How do many micro-experiences become one subject? Why should the experiences of particles or fields generate the coherent experience of a person?

Cosmopsychism offers another possibility. Instead of starting with tiny units of experience, it suggests that consciousness belongs to the universe or cosmos as a whole, and individual minds are derivative or localized expressions of that larger consciousness. This avoids some forms of the combination problem but introduces the opposite difficulty: how does cosmic consciousness differentiate into individual minds?

Russellian monism is another related view. It begins from the idea that physics describes the structural and relational properties of matter, but not necessarily its intrinsic nature. Consciousness may reveal something about the intrinsic nature of matter from the inside. In this view, matter and consciousness are not two separate substances, but physical science may describe only the outer structure of something whose inner nature is experiential or proto-experiential.

For the central question of this book, panpsychism opens the door to consciousness-first or co-emergence views. Consciousness may not begin with biological life. Instead, life may organize pre-existing experiential potential into richer forms. The origin of life would then be a major transition in the organization of consciousness, not the absolute beginning of consciousness itself.

6.10 Implications for the Central Question

Each philosophical framework implies a different answer to the question of whether life or consciousness came first.

Cartesian dualism separates mind from body. If mind is a distinct substance, then consciousness is not simply produced by life. It may interact with living bodies, but it does not arise from biology alone. This view makes the relationship between life and consciousness difficult because it divides them at the foundation.

Spinoza's dual-aspect monism suggests that mind and matter are two expressions of one underlying reality. In this framework, life and consciousness may co-emerge as different aspects of organized nature. The question is not which substance came first, but how one reality appears under mental and physical descriptions.

Leibniz's monadology and contemporary panpsychism suggest that experience may be built into reality at a basic level. If so, consciousness does not emerge from life in the sense of being created from non-conscious matter. Instead, life organizes or amplifies experiential aspects already present in nature.

Kant reminds us that our answers may be limited by the conditions of human knowledge. We study consciousness from within consciousness. We study life as living beings. The question may not be fully answerable from a view outside all experience.

Phenomenology reconnects consciousness with embodiment and world. It suggests that consciousness is not detached from life but lived through the body. This supports the idea that life and consciousness are deeply entangled.

Physicalism supports the standard life-first view. Matter gives rise to life, and life eventually gives rise to consciousness through brains and nervous systems. This framework is scientifically powerful but must still face the hard problem and explanatory gap.

Emergence offers a bridge. Life emerges from chemistry; consciousness emerges from life; both may depend on increasing levels of organization. Yet emergence must explain why subjective experience appears at all.

The framework one adopts therefore constrains the possible answer. Dualism separates life and consciousness. Physicalism makes consciousness a product of life. Panpsychism places consciousness before or alongside life. Phenomenology and dual-aspect theories suggest that life and consciousness may be inseparable aspects of embodied or organized reality.

The central question is therefore not only empirical. It is also philosophical. Before we can decide which came first, we must decide what kind of reality mind, matter, and life are.

6.11 How This Chapter Changes the Central Question

This chapter changes the central question by showing how far life-first explanations can go. Origin-of-life models can explain replication, metabolism, compartments, autocatalysis, energy flow, and biological organization without necessarily explaining subjective experience.

The question therefore becomes more precise: even if chemistry can become life, does that explain how life becomes experience? Origin-of-life research may explain the beginning of biological organization, but the emergence of consciousness remains a further problem.

6.12 Chapter Summary

This chapter traced major developments in modern philosophy of mind relevant to the relationship between life and consciousness.

Descartes separated mind and body into different substances, creating a legacy that still shapes modern debates. Spinoza rejected this split and proposed a dual-aspect monism in which mind and matter are attributes of one substance. Leibniz described reality as composed of monads with perception, offering an early framework relevant to panpsychism. Kant emphasized the limits of human knowledge and the distinction between phenomena and things-in-themselves. Phenomenology re-centered lived experience, intentionality, embodiment, and being-in-the-world. Physicalism attempted to explain consciousness in terms of brain states, functions, or material processes. Property dualism and emergence treated consciousness as a real but possibly non-reducible feature of complex systems. Contemporary panpsychism reopened the possibility that experience is present in some form throughout nature.

Together, these frameworks show that the life-consciousness question cannot be answered by science alone unless its philosophical assumptions are made explicit. Every theory begins from some view of mind, matter, causation, experience, and explanation.

The open question is therefore:

Is the mind-body problem solvable, or does it mark a permanent limit of human understanding?

Part III

Part III: The Origin of Life

Chapter 7

Origin of Life Models

How does matter become alive? Theories of the origin of life attempt to explain how chemistry crossed a threshold into metabolism, replication, boundary formation, and evolution. Yet beneath the chemical details lies a broader question: when matter begins to organize, preserve itself, and respond to its surroundings, has something more than mechanism begun?

7.1 Chapter Overview

The origin of life is one of the deepest unsolved questions in science. At some point in Earth's history, non-living chemistry crossed a threshold into living organization. Molecules became networks. Networks became bounded systems. Bounded systems became capable of persistence, replication, variation, and evolution.

This chapter surveys major scientific models for how life may have originated from non-living matter. These include the RNA world hypothesis, metabolism-first models, protocell theories, autocatalytic networks, thermodynamic approaches, and assembly theory. Each model emphasizes a different feature of life: information, metabolism, boundaries, self-organization, energy flow, or molecular complexity.

The focus of this chapter is not to decide which model is correct. Instead, the aim is to ask what each model implies for the central question of this book: does the origin of life also contain the earliest roots of consciousness? Most origin-of-life models do not directly address consciousness. They are concerned with chemistry, replication, metabolism, compartments, and evolution. Yet some models introduce concepts such as self-maintenance, responsiveness, information processing, and proto-agency. These concepts may not amount to consciousness, but they may help explain why life later becomes a possible ground for consciousness.

The question is therefore not whether early molecules were conscious in any ordinary sense. The question is whether the transition from chemistry to life already introduced structures that later made consciousness possible.

7.2 The Problem of Abiogenesis

Abiogenesis refers to the emergence of life from non-living matter. It asks how chemistry became biology. This is not a question about how modern organisms evolved from earlier organisms; that is the domain of biological evolution. Abiogenesis concerns the earlier transition: how the first living or life-like systems arose before there were cells, genes, organisms, or species in the modern sense.

The challenge is difficult because life is not defined by one feature alone. A successful origin-of-life model must explain several connected processes. It must explain how molecules could store and transmit information. It must explain how chemical reactions could become organized into self-sustaining networks. It must explain how boundaries or compartments could form, separating an inside from an outside. It must explain how variation and selection could begin before fully developed organisms existed.

Replication, metabolism, and compartmentalization are often treated as three central pillars. Replication allows information or structure to be copied. Metabolism allows energy and matter to be transformed in ways that maintain the system. Compartmentalization allows the system to preserve its identity and regulate exchange with the environment. Modern cells contain all three. The puzzle is how they first came together.

Early scientific thinking about abiogenesis was shaped by the Oparin-Haldane hypothesis. Alexander Oparin and J. B. S. Haldane proposed that life emerged gradually from a “primordial soup” of organic molecules on early Earth. Under the right conditions, simple molecules could form more complex organic compounds, eventually leading to self-organizing systems.

The Miller-Urey experiment later gave experimental support to the idea that biologically relevant organic molecules could form under plausible early Earth conditions. By passing electrical sparks through a mixture of gases thought to resemble the primitive atmosphere, Stanley Miller and Harold Urey produced amino acids and other organic compounds. The experiment did not create life, but it showed that the building blocks of life could arise through natural chemical processes.

Since then, origin-of-life research has moved in many directions. Some researchers emphasize genetic information, others metabolism, others membranes, energy gradients, minerals, hydrothermal vents, or chemical networks. The field remains open because no model has yet fully explained the transition from non-life to life.

For the consciousness question, abiogenesis introduces a key issue. If life begins as chemistry, then consciousness appears much later. But if life begins when chemistry becomes organized, self-maintaining, and responsive, then the origin of life may already contain the structural beginnings of cognition, even if not experience.

7.3 RNA World Hypothesis

The RNA world hypothesis proposes that RNA played a central role in the earliest stages of life. RNA is important because it can act both as an information carrier and as a catalyst. DNA stores genetic information in modern cells, while proteins perform much of the catalytic work. RNA, however, can do both in limited ways. It can store sequence information, and some RNA molecules, called ribozymes, can catalyze chemical reactions.

This dual function makes RNA an attractive candidate for early life. Before the complex relationship between DNA, RNA, and proteins evolved, RNA may have served as both gene and enzyme. Self-replicating RNA molecules could have copied themselves with variation. Some variants may have replicated more effectively than others, creating the conditions for Darwinian evolution.

The strength of the RNA world hypothesis is that it explains how heredity and catalysis might have been linked in one molecule. It provides a plausible route from chemistry to evolution. Once RNA molecules could replicate with variation, selection could begin. Over time, more complex systems involving proteins, membranes, and eventually DNA could have evolved.

However, the RNA world hypothesis also faces difficulties. RNA is chemically fragile compared with DNA. It can degrade easily. Producing RNA nucleotides under prebiotic conditions is complex. The formation of long, functional RNA molecules without existing biological machinery remains challenging to explain. The hypothesis is powerful, but it requires a plausible pathway for RNA synthesis, stabilization, concentration, and replication on the early Earth.

From the perspective of consciousness, the RNA world is mostly a chemical and informational model. It does not obviously require experience, awareness, or agency. Self-replicating RNA molecules may undergo selection, but selection alone is not consciousness. A molecule can replicate without having a perspective. A sequence can carry information without experiencing anything.

Yet the RNA world is still relevant. It shows how information can become causally important in matter. A sequence is not just a passive structure; it can influence what happens next. In living systems, information does not merely exist. It participates in organization, replication, and future possibility.

The consciousness implication is therefore indirect. The RNA world does not make consciousness likely at the origin of life. But it helps explain how matter can become informationally organized, and informational organization is one of the later foundations of cognition and consciousness.

7.4 Metabolism-First Models

Metabolism-first models reverse the emphasis of the RNA world. Instead of beginning with genetic replication, they begin with organized chemical cycles. On this view, life may have started not with a self-replicating molecule but with networks of reactions capable of sustaining themselves through energy flow.

Metabolism is central to living systems. Every organism must transform energy and matter in order to maintain itself. Cells are not static containers of genetic information. They are active chemical systems that continuously build, repair, regulate, and renew themselves. Metabolism-first models suggest that this organized chemical activity may have appeared before genetic inheritance in the modern sense.

One influential version is the iron-sulfur world hypothesis associated with Günter Wächtershäuser. This model proposes that early metabolic reactions may have occurred on mineral surfaces, especially involving iron and sulfur compounds. Minerals could have provided catalytic surfaces and energy sources for the formation of increasingly complex organic molecules.

Hydrothermal vent theories, associated with researchers such as Michael Russell and William Martin, also place metabolism at the centre. Alkaline hydrothermal vents on the ocean floor provide natural chemical gradients, mineral structures, and energy flows. These environments may have supported early reaction networks resembling primitive metabolism. Natural compartments within vent structures may have acted as precursors to cellular boundaries.

The strength of metabolism-first models is that they are thermodynamically grounded. Life requires energy flow, and these models begin with energy flow. They do not require complex genetic molecules to appear first. Instead, they ask how geochemical processes could become organized into self-sustaining chemical cycles.

Their challenge is explaining how such metabolic networks became heritable and evolvable. Without genetic information, how are successful networks preserved and transmitted? How does open-ended evolution begin? A metabolism-first model must eventually explain the transition from chemical self-maintenance to informational inheritance.

For the consciousness question, metabolism-first models are especially interesting because they introduce the idea of proto-agency. A metabolic network is not conscious, but it may begin to display a basic form of system-directed organization. It persists by transforming energy. It maintains a pattern through exchange with its environment. It has conditions under which it continues and conditions under which it breaks down.

This does not mean early metabolism had experience. But it does suggest that life begins not with passive matter, but with organized activity. If consciousness later depends on living systems that regulate themselves, preserve their organization, and respond to conditions, then metabolism-first models bring us closer to the roots of cognition than purely genetic models do.

7.5 Protocells and Compartmentalization

Protocell models emphasize the importance of boundaries. Modern life is cellular, and cells are defined by membranes. A membrane separates inside from outside. It allows the system to maintain internal conditions that differ from the environment. It regulates exchange, protects chemical networks, and creates the possibility of individuality.

Without compartments, molecules may react and replicate, but they do not yet form an organism-like unit. A self-replicating molecule floating freely in solution is not the same as a bounded system that can maintain and reproduce itself. Boundaries allow chemistry to become localized. They create a “here” distinct from the surrounding world.

Lipid vesicles are central to many protocell models. Simple fatty acids and lipids can spontaneously form membrane-like structures under certain conditions. These vesicles can grow, divide, encapsulate molecules, and allow some exchange with the environment. If catalytic or replicating molecules became enclosed within such vesicles, the result could be a primitive cell-like system.

Jack Szostak and others have studied how protocells might combine membrane growth, internal chemistry, and genetic replication. The goal is to understand how simple compartments could support early evolution before modern cellular machinery existed. A protocell does not need to be a full modern cell. It only needs to maintain a boundary, contain relevant chemistry, and participate in cycles of growth and division.

The consciousness implications of protocell models are subtle but important. A boundary creates the first meaningful distinction between inside and outside. It is not yet a self in the psychological sense, but it is a physical basis for self/non-self distinction. The system now has an internal organization that can be preserved or lost. The environment is no longer merely background; it becomes what the system must regulate against, exchange with, and survive within.

Does having a boundary create a primitive perspective? Not in the sense of conscious experience. A protocell does not see or feel. But a boundary creates a point of view in a structural sense: an organized interior from which the environment becomes relevant. Conditions outside matter because they affect the continuation of the inside.

This is one reason protocell models are relevant to the origin of cognition. Cognition may require a system for which things matter. Boundaries help create such systems. They allow the earliest forms of biological relevance: nutrient, toxin, stability, rupture, growth, division.

If consciousness eventually depends on perspective, then compartmentalization may be one of the earliest physical roots of perspective.

7.6 Autocatalytic Sets and Chemical Self-Organization

Autocatalytic set theory emphasizes chemical networks rather than single molecules. In an autocatalytic set, molecules collectively catalyze the formation of one another. No single molecule needs to be the master replicator. Instead, the network as a whole sustains itself.

Stuart Kauffman and others have argued that life may have emerged when chemical diversity crossed a threshold and self-sustaining catalytic networks became possible. Once enough molecular species were present, some networks could begin producing the components needed to maintain the network itself. The system, rather than one molecule, becomes the unit of self-maintenance.

This model is attractive because modern life is deeply networked. Cells are not built around one reaction or one molecule. They depend on interconnected metabolic, genetic, and regulatory networks. Autocatalytic sets may therefore capture something fundamental about living organization: life is collective self-production.

Autocatalytic models also challenge the idea that life must begin with a single “first replicator.” Replication may initially be distributed across a network. The system reproduces its pattern, not necessarily through a single molecule copying itself exactly, but through the re-creation of a chemical organization.

The challenge is explaining how such networks become stable, bounded, and evolvable. A self-sustaining reaction network is not automatically a living system. It must persist under realistic conditions, acquire resources, maintain organization, and eventually support heredity and variation.

For the consciousness question, autocatalytic sets are interesting because they suggest that selfhood may begin collectively. A living system may not start as an isolated molecule but as a relational network. This matters because consciousness, too, may depend on networks rather than isolated parts. No single neuron is conscious; consciousness may arise from organized relations among many components.

Does collective self-sustenance imply experience? Not by itself. A chemical network can maintain itself without feeling anything. But autocatalytic sets introduce a key idea: organization can become self-referential. The network produces the components that sustain the network. It is not merely a chain of reactions; it is a loop of self-maintenance.

Such loops may be precursors to biological identity, agency, and cognition. They do not explain consciousness, but they help explain how matter can become organized around its own continuation.

7.7 Thermodynamic and Dissipative Approaches

Thermodynamic approaches view life as part of a broader physical process involving energy flow, dissipation, and far-from-equilibrium organization. Living systems are not isolated objects. They are open systems that maintain order by exchanging energy and matter with their environments.

From this perspective, life is not a violation of thermodynamics. It is a sophisticated expression of thermodynamics. Organisms maintain internal order by increasing entropy in their surroundings. They persist by dissipating energy gradients. Life depends on the flow of energy through matter.

Dissipative structures show that order can arise spontaneously in far-from-equilibrium systems. A whirlpool, a convection cell, or a chemical oscillation can maintain organized patterns through energy flow. These systems are not alive, but they show that matter can self-organize under the right conditions.

Jeremy England and others have proposed that systems driven by external energy may tend to become better at dissipating that energy under certain conditions. This idea, sometimes called dissipation-driven adaptation, suggests that life-like organization may be connected to general physical tendencies. Systems may become structured in ways that absorb and dissipate energy more effectively.

The appeal of this approach is that it places life within physics rather than treating it as an improbable exception. If energy flow naturally produces complex organization under certain conditions, then life may be less accidental than it appears. Life may be a likely outcome on planets with the right chemistry, gradients, and environments.

The challenge is avoiding overstatement. Not every dissipative structure is alive. A flame dissipates energy but is not an organism. A storm is organized but not biological. Thermodynamic approaches must explain what distinguishes living dissipation from non-living dissipation. The answer likely involves boundaries, memory, heredity, regulation, and self-maintenance.

For the consciousness question, thermodynamic models raise a provocative possibility. If life emerges from general principles of energy flow and self-organization, might consciousness also emerge from such principles at higher levels of organization? If matter under energy flow can become living, can living matter under certain forms of organization become aware?

This does not mean consciousness is thermodynamically inevitable. But it suggests that the roots of life and perhaps the roots of mind may lie in the same broader tendency: the emergence of organized processes in far-from-equilibrium systems.

7.8 Assembly Theory

Assembly theory is a more recent approach that attempts to measure molecular complexity in a way that may help identify life. Associated with Lee Cronin, Sara Walker, and collaborators, assembly theory focuses on how difficult it is to construct a given object from basic building blocks.

The assembly index of a molecule refers to the minimum number of steps required to build it from simpler components. Simple molecules have low assembly indices. Complex molecules that require many specific steps have higher assembly indices. The idea is that life produces molecules with complex assembly histories because biological systems preserve, reuse, and elaborate structures through evolution.

This approach is important because it shifts attention from the specific molecules of Earth life to the general problem of detecting life-like complexity. A molecule does not need to be DNA, RNA, or protein to be biologically interesting. If it has a high assembly index and appears in significant abundance, it may indicate a process capable of producing complex, historically contingent structures.

Assembly theory therefore treats life partly as a historical process. Living systems do not merely produce complexity once; they produce, preserve, and propagate complex structures across time. Life leaves signatures of selection, memory, and accumulated construction.

The strength of assembly theory is that it may be useful in astrobiology. If we search for life elsewhere, we may not know what chemistry it uses. A general measure of molecular assembly could help identify systems that are unlikely to arise through random chemistry alone.

The consciousness implications are indirect but intriguing. If life can be detected through complexity, history, and assembly pathways, could consciousness also leave structural or informational signatures? Consciousness may not be detectable through molecules in the same way life is. But if consciousness depends on complex organization, integrated information, or recursive self-modeling, then it too may require specific forms of assembly.

Assembly theory also emphasizes that life is not just matter in the present moment. It is matter shaped by history. Consciousness may likewise depend on history: memory, learning, development, evolution, and accumulated organization.

For the central question, assembly theory does not imply consciousness at the origin of life. But it reinforces the idea that life begins when matter enters a new regime of historically structured complexity. That regime may later make consciousness possible.

7.9 Implications for the Central Question

Most origin-of-life models are silent on consciousness. They do not ask whether early life had experience, awareness, or subjectivity. Their task is to explain how chemistry became capable of replication, metabolism, compartmentalization, and evolution. This silence is understandable. Consciousness is hard to define and difficult to measure even in living organisms, let alone in hypothetical prebiotic systems.

Yet several origin-of-life models introduce concepts that are relevant to consciousness. Metabolism-first models emphasize self-sustaining activity. Protocell models emphasize boundaries and the distinction between inside and outside. Autocatalytic sets emphasize self-producing networks. Thermodynamic approaches emphasize far-from-equilibrium organization and adaptive energy flow. Informational and assembly-based models emphasize memory, structure, and historical complexity.

These concepts do not amount to consciousness. A protocell is not a thinking subject. An autocatalytic network is not aware. A dissipative structure does not feel. But these models suggest that life begins when matter becomes organized around persistence, regulation, and self-maintenance. That may be the first step toward agency.

Agency here should be understood carefully. It does not mean intention, deliberation, or conscious choice. Proto-agency refers to the way a system acts in ways that preserve its organization. A living system is not merely pushed by external forces; it participates in its own continuation.

This is where the origin of life becomes relevant to the origin of consciousness. Consciousness may require a system for which the world matters. Life creates systems for which the world matters. Nutrients, toxins, temperature, gradients, boundaries, and damage are not neutral from the standpoint of a living system. They affect whether the system continues.

If consciousness eventually involves perspective, value, sensation, or meaning, then life may provide the earliest structural basis for those features. Life creates an inside, an outside, a need, a direction, and a history. These may not be consciousness, but they may be the conditions from which consciousness later emerges.

The central question therefore becomes sharper. Did consciousness appear only when nervous systems evolved? Or did nervous systems elaborate capacities already implicit in life: responsiveness, self-maintenance, information processing, and world-directed activity?

Origin-of-life models cannot yet answer this question. But they show where to look.

7.10 How This Chapter Changes the Central Question

This chapter changes the central question by showing that life and consciousness may both depend on information, complexity, and self-organization. The gap between matter and life, and the gap between life and consciousness, may not be separate mysteries but related transitions in organization.

The question therefore shifts from a simple before-or-after sequence to a question of pattern and process. Perhaps consciousness does not appear from matter alone, but from organized systems that integrate information, maintain themselves, respond to the world, and generate meaning.

7.11 Chapter Summary

This chapter surveyed major scientific models of the origin of life and considered their implications for the relationship between life and consciousness.

The problem of abiogenesis concerns the transition from non-living chemistry to biological organization. The RNA world hypothesis emphasizes RNA as both information carrier and catalyst, highlighting the role of molecular information. Metabolism-first models emphasize organized chemical cycles and energy flow. Protocell models emphasize boundaries and the emergence of inside/outside distinction. Autocatalytic set

theory emphasizes collectively self-sustaining chemical networks. Thermodynamic approaches view life as far-from-equilibrium organization sustained by energy dissipation. Assembly theory attempts to measure the complexity and historical depth of molecules associated with living systems.

Most of these models do not address consciousness directly. They are scientific models of chemical and biological emergence, not theories of subjective experience. However, several of them introduce concepts that matter for consciousness: self-maintenance, boundary formation, information processing, energy flow, historical memory, and proto-agency.

The origin of life may not be the origin of consciousness in the full sense. But it may be the origin of the kinds of systems in which consciousness later becomes possible. Life creates organized centres of activity, systems that distinguish inside from outside, regulate themselves, and respond to the world in ways that matter for their continuation.

The open question is therefore:

Can any origin-of-life model explain the origin of experience, or is an additional ingredient required?

Chapter 8

Information, Complexity, and Self-Organization

8.1 Chapter Overview

Life and consciousness both appear to involve more than matter arranged in space. Living systems do not merely contain molecules; they organize, regulate, reproduce, repair, and respond. Conscious systems do not merely process signals; they integrate, interpret, attend, remember, and sometimes become aware of themselves. Between chemistry and consciousness lies a family of concepts that may help connect the two: information, complexity, feedback, recursion, and self-organization.

This chapter explores these bridging frameworks. Information theory explains how uncertainty can be measured and transmitted. Complexity science studies systems whose behaviour cannot be understood simply by examining isolated parts. Self-organization describes how order can arise spontaneously in systems driven by energy flow. Autopoiesis describes life as self-making. Recursion and feedback loops show how systems can refer back to themselves, regulate themselves, and eventually model themselves.

The central argument is cautious but important: information, complexity, and self-organization do not automatically explain consciousness. They do not prove that early life was conscious. However, they may describe the common architecture through which chemistry becomes life and life becomes mind. If there is a bridge between life and consciousness, it may not be a single substance or a single event. It may be a deepening pattern of organized information, self-maintenance, feedback, and recursive self-reference.

8.2 Information Theory Foundations

Information is one of the most important concepts connecting physics, biology, computation, and consciousness. Yet the word “information” can mean different things in different contexts. It can refer to communication, structure, meaning, genetic coding, neural activity, or uncertainty reduction.

In Shannon information theory, information is not primarily about meaning. It is about uncertainty and communication. A message carries information when it reduces uncertainty among possible alternatives. A signal that is completely predictable carries little new information. A signal that selects one possibility from many carries more information. Shannon’s theory was designed for communication systems, not for explaining life or consciousness, but it became foundational because it offered a way to quantify information mathematically.

Kolmogorov complexity approaches information differently. It asks how short the description of a system can be. A highly regular pattern, such as a repeating sequence, can be described very briefly. A random sequence may require a long description because there is no shorter rule that generates it. Complexity, in this sense, lies between simple order and pure randomness. A living organism is not random, but it is also not trivially simple. It contains structured complexity.

Biological information is more than statistical uncertainty. DNA stores hereditary information in molecular sequences. RNA transfers and regulates information. Proteins are produced according to genetic instructions. Cells use signaling networks to respond to internal and external conditions. Nervous systems encode, transform, and integrate information about the body and environment.

Biological information is also causal. It does not merely describe a system from the outside. It participates in building and maintaining the system from the inside. A gene sequence influences the production of proteins. A signaling molecule changes cellular behaviour. A neural signal contributes to perception or action. In living systems, information is embedded in physical processes and has consequences.

This raises the question: is information physical? Landauer's principle suggests that information processing is not separate from physical reality. Erasing information has a thermodynamic cost. Information is therefore not a ghostly abstraction floating above matter. It is instantiated in physical systems, stored in physical states, and transformed through physical processes.

This matters for the relationship between life and consciousness. If information is physical, then life and mind do not require a departure from nature. But if information is also organizational and causal, then life and mind cannot be understood as matter alone. We must understand how matter becomes structured in ways that store, process, and act upon information.

8.3 Information and the Origin of Life

Some theorists argue that the origin of life should be understood not only as a chemical transition, but as an informational transition. Before life, chemistry is governed by local reactions, energy gradients, and molecular interactions. After life emerges, information begins to play a new role. Chemical systems do not merely react; they preserve, transmit, and use patterns that influence future organization.

Sara Walker and Paul Davies have argued that life may involve a shift in informational architecture. In non-living chemistry, causal control is mostly bottom-up: molecules interact according to physical and chemical laws. In living systems, however, information can exert top-down causal influence. Higher-level structures, such as genetic networks, cellular organization, and regulatory systems, influence the behaviour of lower-level components.

This does not mean that living systems violate physics. Rather, it means that physical processes become organized in new ways. A gene sequence can influence protein production. A cell can regulate which genes are expressed. A developmental program can guide the formation of tissues. A nervous system can coordinate the behaviour of the body. These higher-level patterns constrain and direct lower-level processes.

Genetic information represents a new causal regime because it allows biological systems to preserve successful patterns across generations. DNA is not alive by itself, but in the context of a cell it participates in the reproduction and maintenance of life. The sequence matters because it is interpreted by a biochemical system capable of using it.

The origin of life may therefore involve a transition from matter-dominated chemistry to information-guided chemistry. The same molecules remain subject to physical law, but their organization becomes historically structured. Past successful patterns influence future possibilities. Life becomes a system in which memory, heredity, and regulation matter.

This has important implications for consciousness. Consciousness also appears to depend on information, but not information alone. A dictionary contains information but is not conscious. A genome contains information but does not experience. A computer processes information, yet whether it is conscious remains debated. The key question is what kind of information organization is required.

If life begins when information becomes causally organized in self-maintaining systems, then consciousness may emerge when information becomes integrated, embodied, recursive, and available to a system as part of its own perspective. The origin of life does not explain experience, but it may create the first informational architecture from which experience can later arise.

8.4 Complexity and Emergence

Complexity is not the same as complication. A complicated system may have many parts, but its behaviour may still be predictable if we understand the parts and their interactions. A complex system, by contrast, produces patterns that cannot be easily predicted from the parts alone. The whole has properties that emerge through interaction.

Living organisms are complex systems. So are ecosystems, immune systems, economies, brains, and societies. They contain many interacting components. These components influence one another through feedback loops. The system changes over time. It may adapt, reorganize, stabilize, or collapse.

Complex systems often operate near the edge of chaos. This phrase refers to a region between rigid order and total randomness. A system that is too ordered cannot adapt. A system that is too chaotic cannot maintain identity. Life seems to require a balance: enough stability to preserve organization, enough flexibility to respond to change.

Phase transitions are also important. A system may shift suddenly from one state to another when a threshold is crossed. Water freezes or boils at certain temperatures. A flock forms coordinated movement when individual interactions reach a certain pattern. A chemical network may become self-sustaining when catalytic connections become dense enough. Life itself may have emerged through such threshold dynamics.

Emergence can be weak or strong. Weak emergence occurs when higher-level patterns arise from lower-level processes and can, at least in principle, be explained through them. Strong emergence suggests that genuinely new properties appear that cannot be fully reduced to their physical base. Consciousness is often discussed as a candidate for strong emergence because subjective experience seems difficult to derive from physical description alone.

Does consciousness emerge from complexity? Perhaps, but the statement is incomplete. Not all complex systems are conscious. A hurricane is complex. A city is complex. The global climate is complex. Complexity alone is not enough. The relevant question is what kind of complexity matters.

Consciousness may require complexity that is integrated, embodied, self-referential, and capable of supporting a perspective. A brain is not simply complex; it is organized around perception, action, memory, emotion, bodily regulation, and self-modeling. Life is not simply complex; it is organized around self-maintenance and adaptation.

The bridge between life and consciousness may therefore lie not in complexity alone, but in a special form of complexity: self-organizing, self-maintaining, informationally integrated complexity.

8.5 What Is Self-Organization?

Self-organization refers to the spontaneous emergence of order without an external designer or central controller. A self-organizing system develops structure through the interactions of its own components. No single part needs to contain the full plan. Order arises from relations, constraints, feedback, and energy flow.

Self-organization is common in nature. Bénard convection cells form when a fluid is heated from below. As energy flows through the system, ordered circulation patterns appear. The Belousov-Zhabotinsky reaction produces chemical oscillations and spatial patterns. Flocks of birds and schools of fish create coordinated movement without a leader directing every individual. Snowflakes, dunes, neural rhythms, and ecological patterns also show self-organizing features.

The thermodynamic basis of self-organization is important. Many self-organizing systems are far from equilibrium. They maintain order by dissipating energy. They are not closed systems settling into static balance. They are open systems through which energy and matter flow.

Ilya Prigogine's work on dissipative structures showed that order can arise under far-from-equilibrium conditions. This insight helped shift the scientific imagination. Order does not always need to be imposed from outside. Under the right conditions, order can emerge from instability.

Living systems are self-organizing in a deep sense. Cells organize molecular reactions. Embryos develop structured bodies from initially simple conditions. Brains organize patterns of activity through development, learning, and feedback. Ecosystems organize through interactions among organisms and environments.

Self-organization is relevant to the central question because both life and consciousness appear to depend on it. Life emerges when chemistry becomes organized into self-maintaining systems. Consciousness may emerge when neural and bodily processes become organized into integrated experience.

The key property is order from disorder under energy flow. But living and conscious systems add something further: they are not merely ordered; they are organized around their own continuation, regulation, and possible perspective.

8.6 Self-Organization in the Origin of Life

Self-organization plays a central role in many origin-of-life models. Autocatalytic sets, hypercycles, metabolism-first theories, and protocell models all describe ways in which chemical systems might organize themselves into life-like structures.

Autocatalytic sets show how networks of molecules can collectively produce and sustain one another. Instead of one molecule acting as the master replicator, the network as a whole maintains its organization. This is important because modern life is network-based. Cells are not built around a single reaction but around interdependent systems of metabolism, regulation, replication, and repair.

Manfred Eigen and Peter Schuster's idea of hypercycles also emphasized organized networks of replication. A hypercycle links self-replicating units into a cooperative cycle, where each unit supports the replication of the next. Such models attempt to explain how early molecular systems could overcome limits faced by isolated replicators and move toward more complex organization.

Metabolism-first models can also be understood as self-organization-first models. They suggest that life began with organized chemical cycles driven by energy flow, perhaps in mineral-rich environments such as hydrothermal vents. In these models, order emerges from the coupling of chemical reactions to environmental gradients.

Boundary formation adds another level of self-organization. Lipid vesicles can form spontaneously under certain conditions. These boundaries can enclose chemical networks, creating a primitive inside and outside. Once enclosed, reactions can become localized, protected, and connected to the persistence of a system.

The question then arises: when does a self-organizing system become alive? A whirlpool is self-organizing but not alive. A flame maintains itself through energy flow but is not alive. A chemical oscillator produces patterns but is not an organism. Life seems to require self-organization plus self-maintenance, boundary, heredity, and the capacity for open-ended evolution.

For the consciousness question, the important point is that self-organization introduces a path from matter to agency-like behaviour. A self-organizing system is not merely arranged; it dynamically maintains a pattern. When such maintenance becomes internal, bounded, and historically continuous, life begins to appear.

This suggests a possible continuity: self-organization may be the physical root; life may be self-organization that maintains itself; consciousness may be self-organization that becomes integrated and self-aware.

8.7 Autopoiesis: Life as Self-Making

Autopoiesis, developed by Humberto Maturana and Francisco Varela, defines living systems as self-producing systems. A living system continuously produces and maintains the components that make its own organization possible. It is not simply made once; it is always making itself.

A cell is autopoietic because its internal processes produce the components that sustain the cell's boundary, metabolism, and organization. The membrane helps define the cell, but the cell also produces and maintains the membrane. Enzymes support metabolic reactions, but metabolism also supports the production of enzymes. The system is circularly organized.

Autopoiesis distinguishes between operational closure and thermodynamic openness. A living system is open to matter and energy. It takes in nutrients, releases waste, exchanges signals, and depends on its environment. Yet it is operationally closed in the sense that its internal processes form a self-maintaining network. The organization of the system is produced by the system itself.

This makes autopoiesis a powerful minimal definition of life. It does not define life only by reproduction, genes, or metabolism. It defines life by self-production. Reproduction may be necessary for evolution, but an individual organism is alive because it maintains itself as an organized unity.

Maturana and Varela also made a stronger claim: living systems are cognitive systems. Cognition, in this view, is not limited to brains or symbolic thought. To live is to enact a world. A living system distinguishes what matters for its own continuation through its structure and activity. It does not passively receive an objective environment; it brings forth a world of relevance.

This claim is controversial. Critics argue that autopoiesis may define life but does not necessarily define cognition. A bacterium may regulate itself without having mental experience. A cell may maintain itself without being conscious. The risk is that terms such as cognition or world-making become too broad.

Still, autopoiesis is valuable because it shows how life is already relational and self-referential. A living system is concerned, in a minimal biological sense, with its own continuation. It has an inside, an outside, a boundary, and processes that matter to its persistence.

For the central question, autopoiesis suggests that the roots of mind may lie in life's self-making organization. Consciousness may not be present at the beginning of life, but the structure of self-related activity begins there.

8.8 Recursion and Self-Reference

Recursion occurs when a process refers back to itself or when the output of a process becomes input for another cycle of the same process. In mathematics and computer science, recursion is formal and explicit. A function may call itself. A rule may generate repeated structure. In biology, recursion is often embodied in loops of regulation, feedback, and self-maintenance.

Living systems are full of recursive loops. Genes help produce proteins, and proteins regulate genes. Cells produce membranes, and membranes help maintain the conditions under which cells produce their components. The immune system distinguishes self from non-self, but this distinction is continually updated through interaction. The nervous system monitors the body, and bodily states influence the nervous system.

Recursion allows systems to become self-referential. A system does not merely respond to the environment. It responds to its own states, its own boundaries, its own history, and its own responses. This is crucial for life and cognition. Homeostasis requires a system to monitor itself. Learning requires a system to update itself based on past action. Self-awareness may require a system to represent itself as the subject of experience.

Douglas Hofstadter's idea of "strange loops" suggests that consciousness may arise when systems develop recursive self-reference of sufficient depth and complexity. A strange loop occurs when moving through levels of a system eventually brings one back to the starting point. In consciousness, the mind may become aware of itself by modeling itself, reflecting on itself, and folding its own activity back into experience.

This does not mean that all recursion is conscious. A thermostat contains a feedback loop, but it is not conscious. A gene regulatory circuit is recursive, but it does not necessarily experience anything. Recursion may be necessary for self-awareness, but it is not sufficient by itself.

The important point is that recursion creates the possibility of self-relation. A system can become not only organized, but organized around its own organization. It can monitor, regulate, represent, and modify itself.

For the life-consciousness question, recursion may mark a deep continuity. Life begins with self-maintaining loops. Cognition develops through regulatory loops. Consciousness may require loops that become integrated into a first-person perspective. Self-awareness may require loops that explicitly model the self as self.

8.9 From Self-Organization to Self-Awareness

One possible way to connect life and consciousness is through a continuum of increasing recursive depth:

Self-organization becomes self-maintenance. Self-maintenance becomes self-monitoring. Self-monitoring becomes self-modeling. Self-modeling may become self-awareness.

At the simplest level, self-organization creates order without external design. A chemical pattern forms, a structure stabilizes, or a system organizes under energy flow. This is not yet life.

At the next level, self-maintenance appears. A system preserves its own organization through exchange with the environment. It repairs, regulates, and sustains itself. This is closer to life. A cell maintains its boundary, metabolism, and internal conditions.

Self-monitoring adds another layer. The system responds not only to external conditions but to its own internal state. It detects imbalance, damage, nutrient levels, stress, or threat. Homeostasis depends on this kind of monitoring. Organisms must know, in a functional sense, whether they are too hot, too cold, hungry, injured, or chemically imbalanced.

Self-modeling goes further. A system forms internal representations or models of its own body, actions, and relation to the environment. Nervous systems make this possible in complex ways. An animal can track its position, needs, movements, and possibilities for action. The body becomes represented within the system.

Self-awareness may arise when self-modeling becomes integrated into experience. The organism does not merely regulate itself; it experiences itself as a subject, however minimally. In humans, this can develop into reflective self-consciousness, narrative identity, and metacognition. In other animals, it may exist in more embodied and non-verbal forms.

This continuum is speculative, but it helps organize the problem. It avoids treating consciousness as appearing from nowhere. It also avoids claiming that all self-organizing systems are conscious. Instead, it asks what additional levels of recursion, integration, embodiment, and self-modeling are required.

The critical threshold remains uncertain. When does functional self-reference become phenomenal self-awareness? When does a system that monitors itself become a system that experiences itself? This is one of the hardest questions in consciousness studies.

But the continuum suggests that consciousness may be rooted in life's earliest self-organizing tendencies, even if it emerges only much later in richer forms.

8.10 Feedback Loops as a Bridging Concept

Feedback loops are among the most important mechanisms linking chemistry, biology, cognition, and consciousness. A feedback loop occurs when the output of a process influences the future operation of that process. Feedback allows systems to amplify change, stabilize themselves, adapt, and regulate behaviour.

Positive feedback amplifies. A small change grows larger as the system reinforces it. Positive feedback can produce rapid transitions, novelty, growth, and instability. In biology, positive feedback can support processes such as blood clotting, developmental switches, and rapid cellular responses. In evolution, small advantages can become amplified through selection.

Negative feedback stabilizes. It reduces deviation from a target state. Homeostasis depends heavily on negative feedback. Body temperature, blood sugar, pH, hydration, and many other biological variables are regulated through feedback systems that detect change and restore balance.

Nested feedback creates hierarchical control. In living systems, feedback loops are layered. Molecular feedback regulates gene expression. Cellular feedback regulates metabolism. Organ-level feedback regulates physiology. Neural feedback regulates perception and action. Social feedback regulates behaviour and identity. Consciousness, if it arises from such systems, likely depends on nested feedback rather than a single loop.

Feedback architecture may scale from chemistry to cognition. In prebiotic chemistry, feedback may stabilize reaction networks. In early life, feedback maintains internal conditions. In nervous systems, feedback supports perception, motor control, learning, attention, and prediction. In consciousness, feedback may allow a system to integrate information over time and relate current experience to body, memory, and action.

Feedback also helps explain why life is not passive. A living system does not merely react once to a stimulus. It continuously adjusts itself. It compares current conditions to viable ranges. It regulates its own activity. It acts, senses the consequences of action, and acts again.

For the central question, feedback is a bridge because it appears at every level. It is present in chemical self-organization, biological regulation, neural processing, and reflective thought. The difference between non-conscious feedback and conscious experience may lie in how feedback becomes integrated, embodied, and recursively organized into a perspective.

8.11 Information Integration

Integrated Information Theory, often called IIT, proposes that consciousness is related to the degree to which information is integrated within a system. The theory introduces the concept of ϕ , or Φ , as a measure of integrated information. A system has high integrated information when its parts contribute to a unified whole that cannot be reduced to independent components.

The basic intuition is that conscious experience is unified. At any moment, experience contains many elements: colour, sound, bodily feeling, emotion, memory, and attention. Yet these are not experienced as completely separate fragments. They appear within a single field of experience. IIT attempts to explain this unity in terms of information integration.

This is relevant to self-organization because living and cognitive systems are not merely collections of parts. Their components interact in ways that create organized wholes. A cell, an organism, and a brain all involve integration. The question is whether some forms of integration cross a threshold into consciousness.

IIT is attractive because it tries to provide a formal bridge between physical systems and experience. It suggests that consciousness may be present wherever information is integrated in the right way. This can lead to controversial implications, including the possibility that consciousness is more widespread than commonly assumed.

However, IIT has been criticized. One concern is whether information integration is sufficient for experience. A system might have mathematically integrated structure without obviously being conscious. Another concern is how to measure Φ in real complex systems. A third concern is whether the theory over-attributes consciousness to systems that do not seem plausibly conscious.

Still, IIT raises a key issue for this book. If consciousness depends on integration, then the transition from life to consciousness may involve increasing integration of information within self-maintaining systems. Early life may have processed information locally and chemically. Nervous systems later allowed information to be integrated across bodies, senses, actions, memories, and goals.

The connection between information integration and self-organization is therefore important. Self-organization produces coherent systems. Integration allows parts of those systems to function as unified wholes. Consciousness may require not only information, but integrated information organized around a living point of view.

IIT does not settle the question of whether consciousness precedes life or emerges from it. But it provides one possible language for describing how matter, life, and mind may be connected through increasingly integrated organization.

8.12 Implications for the Central Question

Information, complexity, and self-organization provide a common thread across origin-of-life research and consciousness studies. Life begins when chemistry becomes organized, bounded, self-maintaining, and informationally structured. Consciousness may emerge when living systems become integrated, recursive, self-monitoring, and capable of perspective.

This framework supports a co-emergence-friendly view without requiring us to claim that early life was fully conscious. It suggests that life and consciousness may share a deep organizational logic. Life is not just chemistry. Consciousness is not just computation. Both may depend on self-organizing systems that preserve, process, and integrate information.

If self-organization is fundamental to both life and consciousness, then life and consciousness may be two expressions of a broader process. The first expression is biological: matter becomes self-maintaining and

adaptive. The later expression is experiential: self-maintaining systems become capable of integrated awareness. The two are not identical, but they may be continuous.

This framework also helps explain why consciousness is likely tied to life. Consciousness may require more than information processing. It may require embodied information processing within a system that has needs, boundaries, and stakes. A living organism is not a neutral computer. It is a system for which the world matters.

At the same time, the framework leaves major questions open. Information processing does not automatically produce experience. Complexity alone is not consciousness. Self-organization can occur in non-living systems. Feedback loops can exist in simple machines. Integration can be mathematically described without proving subjectivity.

The missing link may be perspective. Life creates bounded, self-maintaining systems. Cognition creates world-directed regulation. Consciousness may arise when information is integrated from the standpoint of such a system. But why that standpoint should feel like anything remains unresolved.

The value of information, complexity, and self-organization is therefore not that they solve the hard problem. Their value is that they show how the gap between chemistry and consciousness may be narrowed. They describe the architecture of the bridge, even if they do not fully explain why experience appears on the bridge.

8.13 How This Chapter Changes the Central Question

This chapter changes the central question by examining the possibility that consciousness emerges from increasing complexity. Emergentist models suggest that life may come first, followed by nervous systems, cognition, and eventually subjective experience.

The question therefore becomes: what kind of complexity is sufficient for consciousness? Mere complexity is not enough. The key issue is whether biological, neural, informational, or organizational complexity can explain why experience appears at all.

8.14 Chapter Summary

This chapter explored information theory, complexity science, and self-organization as possible bridges between life and consciousness.

Shannon information explains how uncertainty can be quantified in communication. Kolmogorov complexity asks how compressible or describable a system is. Biological information appears in DNA, RNA, cellular signaling, regulation, and neural coding. Information is physical, but in living systems it is also organizational and causal.

Origin-of-life research increasingly treats life as an informational transition, where matter becomes organized by stored, transmitted, and causally effective patterns. Complexity science shows how emergent properties arise from interacting parts, especially near thresholds or phase transitions. Self-organization explains how order can arise spontaneously in systems driven by energy flow.

Autopoiesis defines life as self-making and suggests that living systems already contain a minimal form of world-directed organization. Recursion and feedback loops show how systems can become self-referential, self-regulating, and eventually self-modeling. Information integration offers one possible way to understand the unity of conscious experience.

The chapter proposed a possible continuum: self-organization, self-maintenance, self-monitoring, self-modeling, and self-awareness. This continuum does not prove that consciousness is present at the origin of life. But it suggests that consciousness may emerge from deepening layers of information, feedback, embodiment, and recursive organization.

The open question is therefore:

Is information the bridge between life and consciousness, or just a useful metaphor?

Part IV

Part IV: Scientific Theories of Consciousness

Chapter 9

Emergentist Models of Consciousness

Modern consciousness science asks how experience arises from brains, bodies, information, and behaviour. Some theories look for neural mechanisms; others emphasize integration, attention, global availability, embodiment, or biological regulation. Each theory offers a different way of approaching the same puzzle: why should organized activity feel like anything at all?

9.1 Chapter Overview

Emergentist theories begin from a simple scientific intuition: consciousness is not present in isolated atoms, molecules, or single neurons, but appears when physical systems reach a certain level of organization and complexity. In modern science, this usually means that consciousness emerges from the activity of living nervous systems, especially brains.

This chapter examines theories that treat consciousness as an emergent property of complex biological systems. These include neural correlate approaches, biological naturalism, neurobiological theories, higher-order theories, and broader emergence-based models. Together, they represent the dominant scientific approach to consciousness: life comes first, biological complexity increases, nervous systems evolve, and consciousness appears when certain organizational thresholds are crossed.

Emergentism has major strengths. It fits well with neuroscience, evolution, medicine, and comparative biology. It allows consciousness to be studied scientifically without treating it as supernatural or separate from nature. Yet it also faces a deep difficulty: it can often identify the systems associated with consciousness, but it may not fully explain why those systems produce experience at all.

The chapter therefore asks a central question: can emergence explain consciousness itself, or only the biological and neural conditions under which consciousness appears?

9.2 What Is Emergence?

Emergence occurs when a system displays properties that are not obvious from its parts considered separately. Water is wet, but individual water molecules are not wet. A flock has patterns of movement that no single bird controls. A living cell has properties that cannot be found in any isolated molecule. A brain produces coordinated perception, memory, action, and experience through the interaction of billions of neurons.

Emergence can be understood in weak or strong forms.

Weak emergence refers to higher-level patterns that arise from lower-level processes and are, in principle, explainable by them. The emergent property may be surprising, complex, or difficult to predict, but it does not require new fundamental forces or principles. In weak emergence, the higher-level pattern depends entirely on the organization and interaction of lower-level parts.

Strong emergence makes a more ambitious claim. It suggests that genuinely new properties appear at higher levels and cannot be fully deduced from lower-level descriptions, even in principle. If consciousness is strongly emergent, then no complete physical description of the brain would fully explain why subjective experience exists. Consciousness would be dependent on physical systems, but not reducible to them.

The question is which type of emergence consciousness requires. Many scientists prefer weak emergence because it remains compatible with physicalism. Consciousness would be a complex biological phenomenon, like metabolism, immune regulation, or development. It may be hard to explain, but no new metaphysical category is required.

However, some philosophers argue that consciousness seems to require strong emergence. Physical descriptions explain structure, function, behaviour, and information processing, but they appear to leave out subjective experience. If experience cannot be derived from physical organization alone, then consciousness may be strongly emergent or may require a different framework altogether.

The ontological status of emergent properties also matters. Are emergent properties merely convenient descriptions used by observers? Or are they real features of the world? A cell, an organism, and a mind are not illusions simply because they depend on lower-level parts. They have causal roles, patterns, and identities at their own level.

For the central question of this book, emergence usually supports a life-first model. Matter becomes organized into life. Life becomes organized into nervous systems. Nervous systems become organized into conscious brains. But this raises the key problem: what kind of organization is sufficient, and why should it produce experience?

9.3 Neural Correlates of Consciousness

The search for neural correlates of consciousness, often called NCCs, is one of the central scientific approaches to consciousness. A neural correlate of consciousness is the minimal neural mechanism sufficient for a specific conscious experience. In simple terms, researchers ask: what must be happening in the brain for a particular experience to occur?

This approach is powerful because it makes consciousness experimentally tractable. Instead of asking immediately why consciousness exists at all, researchers can ask which brain regions, networks, rhythms, and processes are associated with conscious perception, attention, and report.

Visual consciousness has been especially important in NCC research. Studies of vision suggest that conscious visual experience depends not only on early sensory processing but also on organized activity in higher visual areas, including parts of the ventral stream involved in object recognition. Some visual information can be processed unconsciously, influencing behaviour without entering awareness. This shows that neural processing and conscious experience are not identical in a simple way.

The prefrontal cortex has also been central in many theories, especially those emphasizing conscious access, working memory, decision-making, and report. When information becomes available for reasoning, verbal report, and flexible control of behaviour, prefrontal and parietal networks are often involved. However, researchers continue to debate whether the prefrontal cortex is necessary for consciousness itself or mainly for reporting and using conscious contents.

Thalamocortical loops are another important candidate. The thalamus and cortex interact in recurrent circuits that help regulate wakefulness, attention, sensory integration, and conscious states. Damage or

disruption to these networks can affect consciousness dramatically. Anaesthesia, coma, sleep, and disorders of consciousness all involve changes in large-scale brain dynamics.

The strength of NCC research is that it grounds consciousness in measurable biology. It allows comparison between conscious and unconscious processing, waking and sleep, anaesthesia and awareness, reportable and non-reportable perception.

Its limitation is that correlation is not explanation. NCCs tell us what neural mechanisms are associated with consciousness, but they do not necessarily tell us why those mechanisms produce experience. They may show where consciousness occurs, when it occurs, and what brain processes accompany it. But the deeper question remains: why should these processes feel like anything from the inside?

For emergentism, NCCs are essential but incomplete. They locate the biological conditions of consciousness. They do not, by themselves, solve the hard problem.

9.4 Biological Naturalism

John Searle's biological naturalism treats consciousness as a real biological phenomenon caused by brain processes and realized in brain structures. According to this view, consciousness is not an illusion, not a separate substance, and not merely a computer program. It is a natural feature of certain biological systems.

Searle often compares consciousness to digestion or photosynthesis. Digestion is caused by biological processes in the digestive system. Photosynthesis is caused by biological processes in plants and other organisms. Consciousness, similarly, is caused by neurobiological processes in the brain. It is fully natural, but it is not reducible to abstract computation.

This position rejects both Cartesian dualism and computational reductionism. Consciousness is not a non-physical substance floating outside biology. But neither is it simply a matter of running the right software. For Searle, the biological character of the brain matters. The physical and causal powers of living neural tissue are essential.

The Chinese Room argument is central to Searle's critique of computational theories. The thought experiment imagines a person inside a room manipulating Chinese symbols according to rules, without understanding Chinese. From the outside, the system may appear to understand. But inside, there is only rule-following without comprehension. Searle uses this to argue that syntax alone is not semantics. Computation alone does not produce understanding.

The implication is important for artificial intelligence and consciousness. A machine may simulate conscious behaviour, but simulation is not duplication. A computer simulation of digestion does not digest food. A simulation of consciousness may not be conscious unless it has the right causal powers.

The strength of biological naturalism is that it takes consciousness seriously as both real and natural. It avoids reducing consciousness to behaviour or computation, while also avoiding supernatural dualism. It places consciousness within biology.

Its weakness is that it may not fully explain how biological processes produce subjective experience. Saying that consciousness is caused by the brain is plausible, but the question remains: what is it about brain biology that gives rise to experience? Why do some biological processes produce consciousness while others do not?

For the central question, biological naturalism strongly supports the life-first model. Consciousness emerges from biological organization, especially neural organization. Life is not sufficient for consciousness, but life provides the biological substrate from which consciousness arises.

9.5 Neurobiological Theories of Consciousness

Neurobiological theories of consciousness attempt to explain consciousness through the organization and function of living nervous systems. They do not treat consciousness as an abstract property of information alone. Instead, they emphasize the brain as an evolved biological organ embedded in a body.

Gerald Edelman's theory of Neural Darwinism describes the brain as a selectional system. Neural groups compete, stabilize, and reorganize through development and experience. Consciousness depends on re-entrant processing: ongoing reciprocal signaling among distributed brain regions. Rather than consciousness being located in one place, it emerges from dynamic interactions across neural maps.

Re-entry is important because conscious experience appears unified even though the brain processes information in many specialized areas. Vision, movement, memory, emotion, and bodily signals are distributed, yet experience appears as an integrated scene. Re-entrant neural activity may help bind these processes into coherent conscious experience.

Antonio Damasio emphasizes the body, emotion, and self. His work suggests that consciousness is deeply tied to the regulation of the organism. The brain does not simply represent the external world. It also maps the body's internal state. Feelings, emotions, and bodily regulation are central to the emergence of self and consciousness.

The somatic marker hypothesis proposes that bodily and emotional signals help guide decision-making. Choices are not made by abstract reasoning alone. They are influenced by bodily states, feelings, and learned emotional associations. This gives consciousness a biological function: it helps organisms evaluate situations in relation to their survival, needs, and possible actions.

These neurobiological approaches are important because they move consciousness away from disembodied cognition. Consciousness is not simply computation in the brain. It is tied to a living body that must regulate itself, act in the world, and preserve its own viability.

This suggests that consciousness may have evolved because it served biological functions. It may integrate information, guide flexible behaviour, support social interaction, regulate emotion, and help organisms respond to complex environments. Consciousness may be a way for living systems to coordinate perception, action, memory, and value.

However, functional explanations do not fully solve the problem of experience. They may explain what consciousness does, but not why these functions are accompanied by subjective feeling. Why should body mapping feel like bodily presence? Why should emotion feel like anything? Why should integrated neural activity generate a world of experience?

For the central question, neurobiological theories provide a strong bridge between life and consciousness. They suggest that consciousness is not merely in the brain, but in the brain-body system of a living organism. Consciousness emerges from life, but specifically from life organized around nervous regulation, embodiment, emotion, and action.

9.6 Higher-Order Theories

Higher-order theories propose that a mental state becomes conscious when the system has some kind of higher-order representation of that state. In other words, it is not enough to have a perception, sensation, or thought. For that state to be conscious, the system must in some way be aware of having it.

David Rosenthal's higher-order thought theory argues that a mental state is conscious when there is a higher-order thought about it. A pain becomes conscious not merely because pain processing occurs, but because the system represents itself as being in pain. Consciousness therefore requires a kind of self-monitoring.

Higher-order perception theories, associated with philosophers such as William Lycan, suggest that consciousness arises when a mental state is monitored by a higher-order perceptual system. Instead of a thought about a thought, consciousness involves internal perception of one's own mental activity.

The appeal of higher-order theories is that they explain why some mental processing is unconscious. The brain may process stimuli, regulate behaviour, and form representations without those representations becoming conscious. A state becomes conscious only when it is available to a higher-order system that monitors or represents it.

This fits with the idea that consciousness requires recursive organization. The system does not merely represent the world; it represents its own representations. This creates a bridge between consciousness and self-reference. A conscious system is, at least minimally, a system that can monitor its own mental states.

Higher-order theories are especially relevant to self-consciousness and metacognition. They help explain reflective awareness, confidence, uncertainty, and the ability to report one's own experience. They also connect to debates about animal consciousness and artificial intelligence. If a system has first-order representations but no higher-order monitoring, is it conscious? Or does it merely process information unconsciously?

Critics argue that higher-order theories may over-intellectualize consciousness. Many experiences seem conscious without requiring explicit thoughts about them. Infants, animals, or non-verbal beings may have experiences even if they do not form higher-order thoughts. There is also the problem of regress: if a mental state is conscious only because of a higher-order state, must that higher-order state also be conscious? If so, does this require another higher-order state?

Defenders respond that the higher-order state need not itself be conscious in the same way, or that consciousness arises from the relation between levels rather than from an infinite chain.

For the central question, higher-order theories place consciousness relatively late in evolution. If consciousness requires self-monitoring or higher-order representation, then simple life is not conscious, and perhaps many simple animals are not conscious either. Consciousness emerges only when biological systems develop sufficiently complex internal monitoring.

This supports emergentism but raises the threshold problem: how much self-monitoring is enough?

9.7 Strengths and Limitations of Emergentism

Emergentism has several major strengths. First, it is consistent with neuroscience. Consciousness is strongly associated with brain activity, neural integration, wakefulness, attention, memory, and bodily regulation. Changes to the brain alter experience. Damage to specific systems can change perception, emotion, selfhood, and awareness. Anaesthesia can temporarily remove consciousness. These facts strongly support the view that consciousness depends on biological organization.

Second, emergentist theories are scientifically testable. They generate predictions about which neural systems are necessary for conscious states, how consciousness changes across sleep and anaesthesia, and how different kinds of brain activity relate to different experiences. This makes emergentism productive within neuroscience, medicine, and psychology.

Third, emergentism is parsimonious. It does not require a separate soul-substance, cosmic mind, or new fundamental entity. It explains consciousness as part of nature, continuous with biology and evolution. This makes it attractive to scientific naturalism.

However, emergentism also faces major limitations.

The hard problem remains. Emergentism can explain how brains process information, regulate behaviour, integrate signals, and support report. But why should these processes produce subjective experience? Why should neural activity feel like colour, pain, emotion, or selfhood?

Strong emergence may also be difficult to defend. If consciousness is genuinely novel and not deducible from physical processes, how does it fit into the causal structure of the world? Does it have causal power? If it does, does it alter physical processes? If it does not, why did it evolve?

Weak emergence avoids these problems by keeping consciousness dependent on physical organization. But weak emergence may seem insufficient to explain experience. It may describe the complexity of brain processes without explaining why experience appears.

Emergentism also struggles with the specificity problem. Why do these particular arrangements produce consciousness and not others? Is consciousness produced by neurons, information integration, recurrent processing, biological embodiment, self-modeling, or something else? Different emergentist theories give different answers.

Finally, emergentism faces the transition problem. If consciousness emerged gradually, where did it begin? Was there a first conscious organism? Did consciousness appear with nervous systems, centralized brains, sensory integration, emotion, attention, or self-modeling? If there is no sharp boundary, how should we understand degrees of consciousness?

Emergentism is powerful because it places consciousness within life. But it remains incomplete unless it explains why and how living complexity becomes experience.

9.8 Implications for the Central Question

Emergentism implies a clear general sequence: life first, consciousness later. Non-living chemistry becomes life. Life evolves nervous systems. Nervous systems become increasingly complex. Eventually, consciousness emerges.

This view is compatible with evolutionary biology and neuroscience. It explains why consciousness is associated with living organisms rather than with rocks, water, or isolated molecules. It also explains why consciousness appears to vary with neural complexity, sensory capacity, behavioural flexibility, and bodily regulation.

However, the difficult question is the threshold. What must a living system have in order to become conscious? A cell is alive but does not obviously have experience. A plant responds to its environment but lacks a nervous system. A jellyfish has a nerve net but no centralized brain. An insect has complex behaviour but a very different nervous system from mammals. A mammal has rich sensory and emotional life. Humans have reflective self-consciousness.

Where along this continuum does consciousness begin?

Emergentism must answer not only that consciousness emerges from life, but which forms of life are conscious and why. If the threshold is neural complexity, what kind of neural complexity is required? If the threshold is integration, how much integration is enough? If the threshold is self-monitoring, what forms of self-monitoring count?

This is the transition problem. It is not enough to say that consciousness emerges from complexity. The theory must explain why a particular kind of complexity produces experience. This problem will return later in the book when we examine the hard transition from life to consciousness.

Emergentism therefore provides the dominant life-first answer to the central question. But its success depends on whether it can explain the emergence of experience itself, not only its biological correlates.

9.9 How This Chapter Changes the Central Question

This chapter changes the central question by showing that contemporary scientific theories do not locate consciousness in the same place. Some place it in global access, some in integrated information, some in prediction, some in attention, some in recurrent processing, and some in biological embodiment.

The question therefore becomes theory-dependent. If consciousness is global broadcast, it may require advanced brains. If it is integration, it may be more widespread. If it is biological regulation, it may be rooted in life. Different theories give different answers to where consciousness begins.

9.10 Chapter Summary

This chapter examined emergentist models of consciousness: theories that treat consciousness as arising from complex physical and biological systems, especially brains.

Emergence can be weak or strong. Weak emergence treats consciousness as explainable in principle through lower-level processes. Strong emergence treats consciousness as a genuinely novel property that may not be deducible from physical descriptions alone. Neural correlate research identifies brain mechanisms associated with conscious experience but does not fully explain why experience arises. Biological naturalism treats consciousness as a real biological phenomenon caused by brain processes and realized in brain structures. Neurobiological theories emphasize re-entrant processing, emotion, embodiment, and the role of the living body. Higher-order theories suggest that consciousness requires self-monitoring or higher-order representation.

Emergentism has significant strengths. It fits neuroscience, supports testable research, and avoids introducing non-natural substances. It explains why consciousness is closely tied to brains, bodies, and biological evolution. However, it also faces major limitations. The hard problem remains, the nature of emergence is debated, and the threshold at which life becomes conscious is unclear.

For the central question of this book, emergentism gives the standard answer: life came first, and consciousness emerged later. Yet this answer raises a further question:

Can emergence explain the existence of experience, or only its correlates?

Chapter 10

Contemporary Theories of Consciousness

10.1 Chapter Overview

Contemporary consciousness science does not lack theories. It has many. Some theories locate consciousness in global broadcasting across the brain. Others define it as integrated information, predictive inference, attention modeling, recurrent neural processing, or even quantum-level physical events. Each theory attempts to explain why some information becomes conscious while other information remains unconscious.

This chapter surveys several leading scientific theories of consciousness: Global Workspace Theory, Integrated Information Theory, predictive processing and the Free Energy Principle, Attention Schema Theory, Recurrent Processing Theory, and Orchestrated Objective Reduction. The goal is not to choose a final winner. Rather, the aim is to understand what each theory claims, what mechanisms it emphasizes, what predictions it makes, and how it changes the relationship between life and mind.

The theories differ sharply. Some imply that consciousness requires specific brain architecture. Others suggest that consciousness may be graded and more widespread. Some are strongly biological; others are more computational, informational, or physical. These differences matter for the central question of this book. If consciousness requires a global neuronal workspace, then life comes first and consciousness appears only in certain nervous systems. If consciousness is integrated information, then consciousness may be more continuous with life, matter, or self-organization. If consciousness is linked to predictive regulation, then the boundary between life and mind may become less sharp.

10.2 Global Workspace Theory

Global Workspace Theory, originally associated with Bernard Baars, describes consciousness as a global broadcasting system. The basic idea is that the brain contains many specialized processes operating in parallel. Most of these processes remain unconscious. A perception, memory, goal, or thought becomes conscious when it gains access to a global workspace and is broadcast widely across the system.

Baars used a theatre metaphor. Consciousness is like a spotlight shining on a stage. Many processes operate backstage, but only some contents enter the spotlight. Once on the stage, the content becomes available to many other systems: memory, language, decision-making, planning, emotion, and action.

This metaphor is not meant to suggest that there is a little observer inside the brain watching the stage. Rather, it suggests that consciousness involves availability. A conscious content is not locked inside one local process. It becomes globally available for flexible use.

The Neural Global Workspace model, developed by Stanislas Dehaene, Jean-Pierre Changeux, and others, gives this idea a neurobiological form. According to this model, conscious access depends on large-scale networks, especially involving prefrontal and parietal regions, connected with sensory systems. When a stimulus becomes conscious, neural activity does not remain local and brief. It undergoes “ignition”: a sudden, sustained, widespread pattern of activity that makes the information available across the brain.

Experimental work on masking, attentional blink, and binocular rivalry has been important for testing this theory. In visual masking, a stimulus may be processed unconsciously if it is quickly followed by another stimulus that prevents conscious access. In attentional blink, a second target may fail to reach awareness if it appears too soon after a first target. In binocular rivalry, different images presented to each eye compete for conscious perception. These paradigms allow researchers to compare conscious and unconscious processing while controlling the physical stimulus.

Global Workspace Theory is attractive because it explains access consciousness well. It explains why conscious contents are reportable, flexible, memorable, and available for reasoning. It also provides testable predictions about brain activity: conscious access should involve widespread broadcasting, sustained activation, and coordination across distant neural systems.

However, the theory faces limitations. It may explain how information becomes available to the system, but does it explain why that information is experienced? Global broadcasting may account for report and behavioural control, but phenomenal consciousness remains harder to explain. A content may be globally available, but why should global availability feel like something?

For the central question, Global Workspace Theory supports a life-first and brain-first view. Consciousness requires specific neural architecture. Life must first evolve nervous systems capable of large-scale integration and broadcasting. Consciousness is therefore not present at the origin of life, but appears later in biological evolution.

10.3 Integrated Information Theory

Integrated Information Theory, associated with Giulio Tononi, begins from the structure of experience itself. Conscious experience is unified, differentiated, and intrinsic. At any moment, experience contains many possible distinctions, yet it appears as one integrated whole. IIT attempts to translate these features into a mathematical theory.

The central concept is integrated information, often represented by ϕ , or Φ . A system has integrated information when its parts interact in such a way that the whole has causal power beyond the independent activity of the parts. In simple terms, the system is not just a collection of separate components. Its organization generates a unified causal structure.

IIT begins with axioms about experience and then proposes postulates about the physical systems capable of realizing those axioms. The theory treats consciousness as intrinsic causal power. Consciousness is not defined by report, behaviour, intelligence, or biological function alone. It depends on the internal causal organization of the system.

This leads to several important implications. Consciousness is graded rather than simply on or off. Some systems may have more integrated information than others. A human brain in a waking state may have high Φ . A sleeping or anaesthetized brain may have less. Some simple systems may have very small amounts of integrated information.

Because IIT does not restrict consciousness to human report or even to biological brains, it has panpsychist or proto-panpsychist implications. It suggests that consciousness may be more widespread than common sense assumes, although simple systems would have extremely minimal forms of experience.

This makes IIT especially relevant to the central question of this book. If consciousness is integrated information, then consciousness may not require full biological life as we know it. It may arise wherever physical systems have the right intrinsic causal structure. Life may increase integration, stabilize it, and elaborate it, but life may not be the absolute beginning of consciousness.

IIT is one of the most ambitious contemporary theories, but it is also controversial. Critics argue that it can generate counterintuitive predictions, such as attributing consciousness to systems that do not seem plausibly conscious. Others question whether Φ can be practically measured in complex systems or whether the theory is sufficiently testable. Some critics argue that the theory risks becoming too detached from empirical neuroscience.

Still, IIT makes an important contribution. It shifts the question from what a system does outwardly to what kind of intrinsic organization it has. It asks whether consciousness is not merely information processing, but integrated causal structure from the system's own point of view.

10.4 Predictive Processing and the Free Energy Principle

Predictive processing describes the brain as a prediction machine. Rather than passively receiving information from the world, the brain actively generates predictions about the causes of sensory input. Perception arises through the comparison between predictions and incoming signals. When there is a mismatch, the brain updates its model or acts to reduce the error.

Andy Clark, Jakob Hohwy, and others have developed predictive processing as a broad framework for perception, cognition, and action. It suggests that the brain is always trying to infer what is happening in the world and in the body. Experience is not a direct copy of reality, but a controlled prediction shaped by sensory evidence.

Karl Friston's Free Energy Principle extends this idea more broadly. It proposes that self-organizing systems that persist over time must minimize free energy, or reduce uncertainty about their states in relation to the environment. In this framework, perception, action, learning, and biological regulation are forms of inference. Organisms act to keep themselves within viable bounds.

Active inference is a key part of this view. Organisms do not only update internal models to match the world. They also act on the world to make sensory input fit their predictions. A living system maintains itself by continually predicting, regulating, and acting.

The Free Energy Principle is especially important for this book because it applies not only to brains but to living systems more generally. Cells, organisms, and nervous systems can all be understood as systems that preserve their organization by managing uncertainty and regulating their exchanges with the environment.

This creates a possible bridge between life and consciousness. If all living systems minimize free energy in some sense, does that mean all living systems have experience? Not necessarily. The Free Energy Principle may describe a general condition of living self-organization without implying consciousness in every case. A bacterium may regulate itself without having phenomenal experience.

However, predictive processing may help explain consciousness in systems with complex nervous systems. Conscious experience may arise when predictive models become integrated, embodied, temporally deep, and connected to action and self-modeling. Consciousness may be what the world feels like when a living system models itself and its environment in a unified way.

The strength of predictive processing is that it links perception, action, body, and environment. It avoids the image of the brain as a passive receiver. It also connects consciousness to the organism's need to remain alive.

Its limitation is that prediction alone may not explain experience. A thermostat can reduce error. A machine-learning model can make predictions. A cell can regulate its states. But not every predictive system is conscious. The theory must specify what kind of predictive organization gives rise to experience.

For the central question, predictive processing and the Free Energy Principle support a co-emergence-friendly view. Consciousness may not be present in all life, but the roots of consciousness may lie in the same predictive and regulatory processes that define living systems.

10.5 Attention Schema Theory

Attention Schema Theory, developed by Michael Graziano, proposes that consciousness is the brain's simplified model of its own attention. The brain uses attention to prioritize some information over other information. But in order to control attention effectively, the brain also builds a model of attention. This model is the attention schema.

According to this theory, consciousness is not a mysterious extra property added to information processing. It is the brain's way of representing the fact that it is attending. The brain attributes awareness to itself because it has a simplified internal model of attention. It says, in effect, "I am aware of this," because that is the model it uses to control its own processing.

The attention schema is useful because attention itself is complex. The brain cannot represent every detail of its own neural processes. Instead, it constructs simplified models. Just as the body schema helps the brain control the body, the attention schema helps the brain control attention.

This theory explains why consciousness seems non-physical or difficult to locate. The brain's model of attention leaves out the mechanistic details. It represents awareness as a simple, unified property. The feeling of having consciousness may arise because the system models itself as having awareness.

Attention Schema Theory is attractive because it is functional, testable, and relevant to artificial intelligence. If consciousness is a model of attention, then a machine with the right attention schema might display functional consciousness. It might represent itself as aware, use that representation to guide behaviour, and attribute awareness to others.

However, this raises a major question. Does modeling attention produce real experience, or only the belief and report that one is conscious? If a system says it is aware because it has an attention schema, is it genuinely conscious or merely functionally self-descriptive?

Critics argue that AST may explain why systems claim to be conscious, but not why there is something it is like to be them. The theory may explain access consciousness and self-report better than phenomenal consciousness.

For the central question, AST suggests that consciousness requires a sophisticated modeling system. Life alone is not enough. Consciousness appears when a nervous or artificial system develops a model of its own attention. This supports a life-first view for biological consciousness, but it also opens the possibility of substrate-independent functional consciousness in artificial systems.

10.6 Recurrent Processing Theory

Recurrent Processing Theory, associated with Victor Lamme, proposes that consciousness arises from recurrent or feedback processing in the brain. Feedforward processing occurs when information moves in one direction through a system. Recurrent processing occurs when later stages send signals back to earlier stages, creating feedback loops.

According to this theory, feedforward processing can support unconscious detection and behaviour. A visual stimulus may be processed quickly enough to guide action without becoming conscious. Conscious experience arises when sensory processing becomes recurrent, allowing information to be stabilized, integrated, and enriched through feedback.

This theory is especially grounded in visual neuroscience. Early visual processing can occur without awareness, but conscious perception appears to require feedback between higher and lower visual areas. Recurrent processing allows the brain to bind features, interpret context, and sustain perceptual contents.

Recurrent Processing Theory differs from Global Workspace Theory. GWT emphasizes global broadcasting and access to widespread systems, often including frontal-parietal networks. RPT suggests that phenomenal consciousness may arise earlier and more locally within sensory systems, through recurrent processing. Global broadcasting may be needed for report or cognitive access, but not necessarily for experience itself.

This distinction is important. It suggests that consciousness may not require full reportability. A perception may be phenomenally conscious even before it is available for verbal report or executive control. This has implications for animal consciousness, infant consciousness, and non-verbal conscious states.

The strength of RPT is that it offers a clear neural threshold: feedforward processing is unconscious; recurrent processing is conscious. It also provides testable predictions through visual masking, transcranial magnetic stimulation, and neural timing studies.

The limitation is that recurrence may not be sufficient. Many systems contain feedback loops without being conscious. The theory must explain what kind of recurrence matters. Is local sensory recurrence enough? Does consciousness require integration with body, memory, emotion, and action? How much recurrence is necessary?

For the central question, RPT supports a life-first view but places the threshold earlier than some global workspace models. Consciousness requires nervous systems with recurrent processing, but not necessarily human-like cognition or reflective self-awareness.

10.7 Orchestrated Objective Reduction — Brief Preview

Orchestrated Objective Reduction, often called Orch-OR, is associated with Roger Penrose and Stuart Hameroff. It proposes that consciousness may depend on quantum processes occurring in microtubules inside neurons. Unlike theories that explain consciousness through large-scale neural networks, information integration, attention, or predictive processing, Orch-OR looks to the quantum physical level.

The theory suggests that quantum states in microtubules undergo objective reduction, and that these events are orchestrated by biological processes in the brain. In this framework, consciousness is not simply computation, nor merely classical neural activity. It is linked to deeper physical processes.

Orch-OR is listed here because it is a contemporary theory of consciousness, but it will be treated more fully in the chapter on quantum and speculative frameworks. It differs from the other theories in this chapter because it reaches beyond standard neuroscience and into debates about quantum mechanics, computation, and fundamental physics.

The theory is controversial. Critics question whether quantum coherence could be maintained in the warm, noisy environment of the brain and whether microtubule processes can explain conscious experience. Supporters argue that standard computational and neural theories may be insufficient and that quantum processes deserve consideration.

For the central question, Orch-OR has unusual implications. If consciousness depends on quantum processes orchestrated in biological structures, then life still matters, but consciousness may also connect to deeper physical principles. This could make consciousness neither purely life-first nor simply consciousness-first, but dependent on an interaction between biology and fundamental physics.

10.8 Theory Comparison

The major contemporary theories differ not only in mechanism, but in their assumptions about what consciousness is.

Global Workspace Theory treats consciousness primarily as global access. Information becomes conscious when it is broadcast widely and made available to many systems. Integrated Information Theory treats consciousness as intrinsic integrated causal structure. Predictive processing and the Free Energy Principle treat consciousness in relation to inference, prediction, embodiment, and self-regulation. Attention Schema Theory treats consciousness as a model of attention. Recurrent Processing Theory treats consciousness as feedback processing within neural systems. Orch-OR treats consciousness as involving quantum physical processes within biological structures.

These differences lead to different answers about where consciousness exists. GWT and RPT locate consciousness in specific neural architectures. IIT potentially extends consciousness more widely, wherever integrated information exists. Predictive processing links consciousness to living systems that model and regulate themselves, especially those with complex nervous systems. AST suggests that consciousness may exist in systems capable of modeling attention, biological or artificial. Orch-OR ties consciousness to biological quantum processes.

The theories also differ on whether consciousness is binary or graded. IIT is explicitly graded. Predictive processing may allow degrees of consciousness depending on model depth and integration. GWT often implies a threshold-like transition into conscious access. RPT suggests a neural threshold from feedforward to recurrent processing. AST suggests that consciousness depends on the presence and sophistication of an attention schema.

A simplified comparison is useful:

```
## \begin{landscape}
## \begin{table}[!h]
## \centering
## \caption{\label{tab:contemporary-theories-table}Simplified comparison of contemporary theories of con
## \centering
## \resizebox{\ifdim\width>\linewidth\linewidth\else\width\fi}{!}{
## \fontsize{9}{11}\selectfont
## \begin{tabular}[t]{t}{>\raggedright\arraybackslashp{4.0cm}>\raggedright\arraybackslashp{4.5cm}>\rag
## \toprule
## Theory & Core mechanism & Primary location & Graded or threshold? & Main implication\\
## \midrule
## Global Workspace Theory & Global broadcast of information & Large-scale brain networks & Threshold-1.
## Integrated Information Theory & Integrated causal information & Any system with sufficient intrinsic
## Predictive Processing / Free Energy Principle & Prediction, inference, and uncertainty reduction & L
## Attention Schema Theory & Model of attention & Systems that model their own attention & Functional tl
```

```
## Recurrent Processing Theory & Feedback and recurrent neural processing & Sensory and cortical recurr
## \addlinespace
## Orch-OR & Quantum processes in microtubules & Biological neural microstructures & Event-based & Cons
## \bottomrule
## \end{tabular}}
## \end{table}
## \end{landscape}
```

No theory explains everything. GWT explains report and access well. IIT addresses intrinsic experience but faces testability concerns. Predictive processing connects brain, body, and environment but may over-generalize prediction. AST explains self-report and attention but may not fully explain phenomenal feeling. RPT gives a clear neural mechanism but may not account for all forms of consciousness. Orch-OR raises fundamental questions but remains controversial.

The disagreement is not only about data. It is also about what consciousness is supposed to be: access, experience, integration, prediction, self-modeling, feedback, or something deeper.

10.9 The Adversarial Collaboration

Because consciousness theories often interpret the same evidence differently, researchers have begun using adversarial collaboration. In this approach, supporters of competing theories agree in advance on experimental protocols, predictions, and criteria for interpretation. The goal is not to win a debate rhetorically, but to test theories under shared conditions.

One major adversarial collaboration has compared Global Neuronal Workspace Theory and Integrated Information Theory. These theories make different predictions about where and when conscious information should appear in the brain. GWT emphasizes widespread broadcasting and frontal-parietal involvement. IIT emphasizes integrated causal structure and has often highlighted posterior cortical regions as especially relevant.

Recent structured empirical work has produced an important but mixed lesson: the evidence has challenged aspects of both theories rather than simply confirming one and defeating the other. Some findings have supported posterior sensory regions as central to conscious contents, while other results have complicated both IIT and GWT predictions.

This is scientifically valuable. It shows that consciousness research is becoming more rigorous and collaborative. It also shows that major theories may need refinement. Consciousness may not fit neatly into one simple neural location or one mechanism.

The limits of empirical adjudication are also clear. Experiments can test specific predictions, but theories often have core claims and auxiliary assumptions. When a prediction fails, theorists may revise the auxiliary assumptions rather than abandon the theory. This is not necessarily unscientific; it is common in complex fields. But it means that empirical tests alone may not settle philosophical disagreements.

Adversarial collaboration is therefore a major step forward, but not a final resolution. It helps identify which predictions survive contact with data. It also reveals where theories are too vague, too broad, or too flexible.

For the central question of this book, adversarial collaboration matters because it shows that even the best contemporary theories remain unsettled. We do not yet have one accepted scientific account of consciousness. The relationship between life and consciousness therefore remains open.

10.10 Implications for the Central Question

Each contemporary theory implies a different relationship between life and consciousness.

Global Workspace Theory and Recurrent Processing Theory support a life-first view. Consciousness requires nervous systems with specific architectures. Life begins long before consciousness, and consciousness appears only when biological evolution produces brains capable of global broadcasting or recurrent processing.

Integrated Information Theory opens the possibility of a graded and more widespread consciousness. If consciousness is integrated information, then it may not be limited to complex brains. Life may increase and organize integration, but consciousness could exist in minimal forms wherever integrated causal structure exists. This makes IIT compatible with co-emergence or even consciousness-first interpretations.

Predictive processing and the Free Energy Principle blur the boundary between life and mind. If living systems preserve themselves through prediction, inference, and uncertainty reduction, then the roots of consciousness may lie in biological regulation itself. This does not mean that all life is conscious, but it suggests continuity between living self-organization and conscious experience.

Attention Schema Theory makes consciousness a model. If a system can model its own attention and use that model to control behaviour, it may display functional consciousness. This opens the door to substrate-independent interpretations, including artificial systems. Life may not be strictly required if the right modeling architecture can be built.

Orch-OR suggests that consciousness may depend on fundamental physical processes organized within biological systems. This gives life an important role but also points beyond ordinary biological mechanism.

The theories therefore do not simply disagree about the brain. They disagree about the metaphysical status of consciousness. Is consciousness access? Integration? Prediction? Self-modeling? Recurrent processing? Quantum event? Biological function? Intrinsic causal power?

The answer determines whether consciousness is late, graded, widespread, biological, computational, or fundamental.

For the central question, contemporary theories do not yet provide a single answer. Instead, they reveal the range of possible answers still available within science.

10.11 How This Chapter Changes the Central Question

This chapter changes the central question by taking seriously the possibility that consciousness is not produced by life or matter. Consciousness-first theories reverse the usual direction of explanation and ask whether matter and life arise within consciousness rather than consciousness arising from them.

The question therefore becomes: can consciousness be fundamental without becoming scientifically empty? Consciousness-first theories address the hard problem directly, but they must explain how consciousness gives rise to organized life, individual minds, and the stable physical world.

10.12 Chapter Summary

This chapter surveyed major contemporary scientific theories of consciousness.

Global Workspace Theory explains consciousness as global broadcasting of information across the brain. Integrated Information Theory defines consciousness as integrated causal information and treats it as graded.

Predictive processing and the Free Energy Principle connect consciousness to prediction, inference, embodiment, and self-regulation. Attention Schema Theory interprets consciousness as the brain's model of its own attention. Recurrent Processing Theory locates consciousness in feedback processing within neural systems. Orchestrated Objective Reduction links consciousness to possible quantum processes in biological structures.

The theories differ in mechanism, scope, testability, and philosophical implication. Some restrict consciousness to specific brain architectures. Others allow more graded or widespread forms. Some emphasize access and report; others emphasize intrinsic experience, embodiment, self-modeling, or physical foundations.

Recent adversarial collaborations show that empirical testing can sharpen the debate, but so far the field remains unsettled. Evidence may challenge parts of multiple theories at once, suggesting that consciousness may require a more integrative framework.

For the central question of this book, the lesson is clear: the theory of consciousness one adopts strongly shapes whether consciousness is viewed as a late evolutionary product, a graded feature of organized systems, a property of living self-regulation, or something more fundamental.

The open question is therefore:

Will empirical science resolve the debate between theories, or is the disagreement fundamentally philosophical?

Part V

Part V: Consciousness First?

Chapter 11

Consciousness-First Theories

What if consciousness is not the last arrival in the story, but the background from which life emerges? Consciousness-first theories challenge the assumption that mind must be produced by biology alone. They ask whether awareness, experience, or field-like forms of knowing might be more fundamental than living systems themselves.

11.1 Chapter Overview

Most scientific accounts begin with matter. Matter organizes into chemistry, chemistry into life, life into nervous systems, and nervous systems into consciousness. Consciousness-first theories reverse this order. They propose that consciousness is not a late product of biological evolution, but a fundamental feature of reality, or even the ground from which matter and life arise.

This chapter examines several consciousness-first or consciousness-fundamental frameworks: panpsychism, cosmopsychism, idealism, Russellian monism, dual-aspect monism, and process philosophy. These theories differ from one another, but they share a common dissatisfaction with the idea that subjective experience can be fully explained as an accidental product of non-conscious matter.

The appeal of consciousness-first theories is clear. They seem to avoid the hard problem by refusing to derive experience from something entirely non-experiential. If consciousness is already basic, then life does not need to produce consciousness from nothing. Instead, life may organize, localize, express, filter, or intensify consciousness.

Yet these theories face serious challenges. If consciousness is everywhere, why do only some systems appear to have minds? If the universe as a whole is conscious, how do individual minds arise? If matter is an appearance within consciousness, why does the physical world appear stable, lawful, and independent of individual minds?

This chapter treats consciousness-first theories neither as final answers nor as dismissible speculation. They are examined as serious philosophical alternatives that reshape the central question of the book.

11.2 The Core Claim

The core claim of consciousness-first theories is that consciousness is not derived from matter in the ordinary sense. Instead, consciousness is fundamental, and matter is either derived from consciousness, structured by consciousness, or inseparable from consciousness at the deepest level.

This claim reverses the standard scientific narrative. In the standard view, the universe begins as physical reality. Matter forms stars, planets, chemistry, and eventually life. Consciousness appears only much later, after biological evolution produces nervous systems and brains.

Consciousness-first theories ask whether this sequence is incomplete or mistaken. They argue that if consciousness is the one feature of reality we know directly from within, it may be strange to treat it as the last and least fundamental thing. Physical objects are known through experience. Scientific observations occur within consciousness. Even the concept of an external world is given to us through conscious awareness.

This does not mean that consciousness-first theories deny science. Most do not deny that brains, bodies, and biological processes are strongly correlated with experience. Rather, they question the metaphysical interpretation of those correlations. Does the brain produce consciousness, as the liver produces bile? Or does the brain organize, filter, limit, or localize consciousness that is more fundamental?

Historical precedents are extensive. Idealist traditions treat mind or consciousness as primary. Panpsychist traditions attribute mind-like qualities widely throughout nature. Religious and contemplative traditions often understand consciousness as a deeper reality than the material world. Ancient, Eastern, and Neoplatonic frameworks all contain versions of consciousness-first thought.

For the central question of this book, the core claim is radical: the origin of life may not be the origin of consciousness. Instead, the origin of life may be a transition in how consciousness becomes organized into living form.

11.3 Panpsychism

Panpsychism is the view that consciousness, or proto-consciousness, is a fundamental and widespread feature of reality. It does not usually claim that every object thinks, feels, or has a mind like a human being. Rather, it proposes that the basic constituents of reality have some primitive experiential aspect.

The motivation for panpsychism comes largely from the hard problem. If matter is defined as entirely non-conscious, then it is difficult to see how consciousness could ever arise from it. No matter how many non-conscious parts are combined, why should subjective experience appear? Panpsychism avoids this problem by denying that consciousness emerges from absolute non-consciousness. Experience is present in primitive form from the beginning.

Constitutive panpsychism proposes that macro-consciousness is composed of micro-consciousness. On this view, the consciousness of a human being somehow arises from the combination or organization of the experiential aspects of smaller physical constituents. The mind is not created from nothing; it is built from simpler experiential ingredients.

This approach has attracted contemporary philosophers such as Galen Strawson, Philip Goff, and David Chalmers, who have treated panpsychism as a serious option in philosophy of mind. Their arguments differ, but they share the intuition that consciousness may need to be placed into the basic structure of reality rather than derived from purely physical description.

The major strength of panpsychism is that it avoids the strongest form of the hard problem. It does not need to explain how experience appears from total absence of experience. Instead, it must explain how simple forms of experience become complex forms.

But this is also its greatest weakness. The combination problem asks how micro-experiences combine into unified macro-experience. If many tiny experiential units exist, how do they form one conscious subject? Why is a human mind experienced as one field of awareness rather than as a crowd of micro-subjects? How do simple experiential properties become the rich, unified consciousness of an animal or person?

Panpsychism also faces the challenge of empirical testability. If every physical system has some minimal experiential aspect, how could this claim be confirmed or disconfirmed? What would count as evidence for or against it?

For the central question, panpsychism implies that consciousness precedes life in some primitive form. Life does not generate consciousness from nothing. Rather, life organizes proto-conscious or experiential aspects of reality into more complex forms of subjectivity.

11.4 Cosmopsychism

Cosmopsychism offers a different consciousness-first model. Instead of claiming that consciousness begins with tiny micro-experiences at the level of particles or basic matter, cosmopsychism proposes that the universe as a whole is conscious, and individual consciousnesses are derived from or within that larger consciousness.

This is a top-down rather than bottom-up approach. Panpsychism begins with many small experiential units and tries to explain how they combine. Cosmopsychism begins with one cosmic consciousness and tries to explain how individual minds arise within it.

The attraction of cosmopsychism is that it avoids some forms of the combination problem. If consciousness is fundamentally unified at the cosmic level, then we do not need to explain how many tiny consciousnesses combine into one. Instead, the main question becomes how one larger consciousness differentiates into many individual centres of experience.

This is called the decomposition problem. How does cosmic consciousness divide, localize, or express itself as individual minds? Why do I experience the world from this particular perspective rather than from the perspective of the universe as a whole? How does one field of consciousness become many subjects?

Contemporary philosophers such as Yujin Nagasawa, Khai Wager, Itay Shani, and Philip Goff have explored cosmopsychist ideas in different forms. These theories remain debated, but they have become part of serious contemporary philosophy of mind.

Cosmopsychism also has deep historical resonance. Neoplatonism, Vedanta, and some mystical traditions describe reality as grounded in a universal mind, soul, or consciousness. Individual minds are not separate substances but local expressions, limitations, or manifestations of a larger reality.

For the central question, cosmopsychism strongly supports a consciousness-first answer. Life emerges within consciousness. Living beings do not create consciousness, but become localized centres through which cosmic consciousness is expressed or constrained.

This view dissolves the hard problem in one direction, but creates new problems in another. It explains why consciousness exists by making it fundamental, but it must explain why consciousness appears fragmented, embodied, and tied to biological organisms.

11.5 Idealism

Idealism is the view that mind, consciousness, or experience is ontologically primary. In its strongest forms, idealism claims that only consciousness truly exists, and what we call matter is an appearance, pattern, or representation within consciousness.

Ontological idealism does not necessarily mean that the physical world is imaginary in the ordinary sense. A dream, a perception, or a shared experience can have structure, regularity, and consequences. The idealist

claim is not that the world is meaningless or arbitrary, but that its ultimate nature is mental or experiential rather than material.

Bernardo Kastrup's analytic idealism is a contemporary form of this view. It proposes universal consciousness as the sole ontological primitive. Individual minds are understood as dissociated "alters" of this universal consciousness, somewhat like distinct centres of experience within a larger field. The physical world is the appearance of mental processes when viewed from across dissociative boundaries.

In this model, the brain does not generate consciousness. Instead, the brain is the image or appearance of a localized mental process. The brain may function more like a filter, image, or interface than a producer. Changes in brain activity correlate with changes in experience because the brain is how those experiential processes appear from an external perspective.

The strength of idealism is that it offers an elegant response to the hard problem. If consciousness is fundamental, there is no need to explain how consciousness arises from non-conscious matter. Matter becomes the phenomenon to be explained, not consciousness.

Idealism also takes seriously the fact that all knowledge of the world appears within experience. We never encounter matter outside consciousness; we infer matter through patterns in experience.

However, idealism faces major challenges. It must explain the stability, regularity, and apparent independence of the physical world. Why does the world obey consistent laws? Why do different observers agree about shared objects? Why does brain damage alter consciousness in systematic ways if the brain does not produce consciousness?

Idealist theories respond in different ways, often by treating the physical world as a shared or transpersonal structure within universal consciousness. But the challenge remains difficult.

For the central question, idealism gives the strongest consciousness-first answer. Consciousness does not come from life. Life is a form appearing within consciousness. The origin of life is therefore not the beginning of experience, but the emergence of a new pattern within experience.

11.6 Taheri's T-Consciousness Framework

Mohammad Ali Taheri's framework offers a contemporary consciousness-first model that can be placed alongside idealism, cosmopsychism, and other non-reductive theories of mind. In this view, consciousness is not produced by matter, biology, or neural activity. Rather, consciousness is primary, and material organization is understood as arising within, through, or by means of consciousness fields.

Taheri uses the term T-Consciousness to refer to a fundamental, non-material reality that is not reducible to energy, frequency, ordinary information, or physical force. T-Consciousness is not treated as a property generated by the brain. Instead, the brain and body are understood as receivers, interfaces, or manifestations within a larger field-like order of consciousness.

This framework differs from standard emergentism. In emergentist theories, matter becomes organized into life, life evolves nervous systems, and nervous systems eventually produce consciousness. Taheri's view reverses this direction. Consciousness is not the last stage of biological complexity; it is the prior condition through which biological and material complexity become possible.

Taheri's framework also differs from simple panpsychism. Panpsychism usually proposes that consciousness or proto-consciousness is present in all matter. Taheri's view is not necessarily that every particle has its own independent experience. Rather, consciousness is treated as a non-material field or network that gives rise to organization, relation, and manifestation. This makes the framework closer to consciousness-field theories, cosmopsychism, and certain idealist traditions than to atomistic panpsychism.

A key feature of the theory is that consciousness is not understood as local, personal, or brain-bound. Individual consciousness is not isolated from a larger order. Living beings participate in, receive from, or manifest aspects of a broader consciousness field. This creates a model in which consciousness precedes individual life while also becoming expressed through living organisms.

The strength of Taheri's framework is that it directly addresses the hard problem by refusing to derive consciousness from non-conscious matter. If consciousness is fundamental, then the existence of experience does not need to be explained as a late evolutionary accident. The question shifts from "How does matter produce consciousness?" to "How does consciousness organize or manifest as matter, life, and mind?"

However, the theory also faces major challenges. Like other consciousness-first models, it must explain mechanism and testability. How exactly do consciousness fields produce biological form? How do they relate to known physical and chemical processes? What evidence would distinguish this view from idealism, panpsychism, or metaphorical language? Unless such questions are addressed, the framework remains philosophically suggestive but difficult to evaluate scientifically.

For the central question of this book, Taheri's T-Consciousness framework clearly supports a consciousness-first answer. Life does not generate consciousness. Rather, life arises within a prior consciousness-based reality and becomes one of the ways consciousness is expressed, organized, or made manifest.

11.7 Russellian Monism

Russellian monism begins from an insight associated with Bertrand Russell: physics describes the structure and relations of matter, but not necessarily its intrinsic nature. Physics tells us how physical entities behave, how they relate, how they are measured, and how they interact. But it may not tell us what matter is in itself.

This opens a possibility. Perhaps consciousness, or proto-consciousness, is the intrinsic nature of matter. From the outside, matter appears through structure, relations, forces, and equations. From the inside, its intrinsic nature may be experiential or proto-experiential.

Russellian monism is attractive because it is neither standard physicalism nor traditional dualism. It does not say that consciousness is separate from matter. It also does not reduce consciousness to structural physical descriptions. Instead, it suggests that physical science gives us the external structure of reality, while consciousness gives us a clue to its intrinsic nature.

This framework has been treated by David Chalmers and others as a promising approach because it may bridge the explanatory gap. If physical descriptions leave out intrinsic nature, then the hard problem may arise because science describes only structure and function. Consciousness may not be an extra substance, but the inner aspect of the physical.

Russellian monism is closely related to panpsychism and neutral monism. If the intrinsic nature of matter is experiential, the view becomes panpsychist or panprotopsychist. If the intrinsic nature is neutral, neither mental nor physical in the ordinary sense, then mental and physical properties may both arise from a deeper base.

The strength of Russellian monism is that it respects science while identifying a limit in physical description. It does not deny physics; it argues that physics may be incomplete as metaphysics.

Its weakness is that it remains difficult to specify the intrinsic nature of matter. What exactly are proto-experiential properties? How do they relate to human consciousness? How can the theory be tested?

For the central question, Russellian monism suggests that consciousness may not be produced by life, but life may organize the intrinsic experiential nature of matter into increasingly complex forms. This allows a middle path between strict consciousness-first idealism and standard life-first physicalism.

11.8 Dual-Aspect Monism

Dual-aspect monism proposes that reality is one, but it has both mental and physical aspects. Mind and matter are not two separate substances, but two ways in which one underlying reality appears or is expressed.

Spinoza is the major historical figure behind this view. For Spinoza, there is only one substance, and mind and extension are two of its attributes. A human being can be described physically as body or mentally as thought, but these are not two separate entities interacting from different worlds. They are two aspects of one reality.

Dual-aspect theories have continued in different forms. Some suggest that every physical process has a corresponding mental aspect, though not necessarily consciousness in the ordinary sense. Others argue that mind and matter are complementary descriptions that cannot be reduced to one another.

The Pauli-Jung conjecture is sometimes discussed in this context. Physicist Wolfgang Pauli and psychologist Carl Jung explored the possibility that mind and matter might be complementary aspects of a deeper reality, connected through symbolic patterns, archetypes, and synchronicity. These ideas are speculative, but they show how dual-aspect thinking can extend into interpretations of quantum mechanics, psychology, and meaning.

Dual-aspect monism is relevant to quantum interpretations because quantum theory already challenges simple separations between observer and observed, measurement and system, possibility and actuality. This does not mean that quantum mechanics proves consciousness-first theories. But it has encouraged some thinkers to question whether physical reality can be fully understood without reference to observation, information, or mind-like aspects.

The strength of dual-aspect monism is that it avoids both reductionism and dualism. Consciousness is not reduced to matter, but neither is it separated from matter. The mental and physical are inseparable aspects of the same underlying process.

Its weakness is that it can be vague. What is the underlying reality? How do the mental and physical aspects relate? Can the theory produce testable predictions, or is it mainly a metaphysical interpretation?

For the central question, dual-aspect monism supports a co-emergence view. Life and consciousness may arise together as different aspects of organized reality. The question may not be whether consciousness came before life or life before consciousness, but how one underlying reality becomes both living and experiential.

11.9 Process Philosophy

Process philosophy, especially in the work of Alfred North Whitehead, replaces substance with process. Reality is not made primarily of static things but of events, relations, and becoming. The basic units of reality are not inert particles but occasions of experience.

Whitehead's actual entities, or occasions of experience, have both physical and mental poles. The physical pole concerns inheritance from the past: how an event receives and integrates prior conditions. The mental pole concerns possibility, feeling, and the formation of a new perspective. Reality unfolds through moments of experience that take account of the past and contribute to the future.

This view is not panpsychism in the simplest sense, but it is strongly experiential. Experience is not a late addition to the universe. It is woven into the process of becoming. Human consciousness is a highly developed form of experience, but simpler forms of feeling or prehension exist throughout reality.

Process philosophy is especially relevant to the relationship between life and consciousness because it emphasizes continuity. Life is not a machine assembled from dead parts. It is an intensified form of process. Consciousness is not an exception to nature. It is a complex form of experiential becoming.

Whitehead's view also gives importance to relation. An actual entity is not isolated.

It becomes what it is through its relations with other events. Each moment gathers the past, integrates it, and contributes something new. This makes process philosophy especially useful for thinking about emergence, continuity, and transformation.

For the central question of this book, process philosophy does not fit neatly into either a simple life-first or simple consciousness-first model. It suggests that experience-like process is basic, while biological life represents a more organized and intensified form of becoming. Life does not create experience from nothing, but it may deepen, structure, and concentrate experiential process into more complex forms.

The strength of process philosophy is that it avoids treating matter as dead substance and consciousness as a sudden late addition. It offers a continuous picture in which reality, life, and mind are all forms of process. Its weakness is that it can be difficult to translate into empirical science. Concepts such as occasions of experience, prehension, and becoming are philosophically rich, but they are not easy to test directly.

11.10 How This Chapter Changes the Central Question

This chapter changes the central question by taking seriously the possibility that consciousness is not produced by life or matter. Consciousness-first theories reverse the usual direction of explanation and ask whether matter and life arise within consciousness rather than consciousness arising from them.

The question therefore becomes: can consciousness be fundamental without becoming scientifically empty? Consciousness-first theories address the hard problem directly, but they must also explain how consciousness relates to organized life, individual minds, and the stable physical world.

11.11 Chapter Summary

This chapter examined consciousness-first theories: frameworks that treat consciousness not as a late product of biological evolution, but as fundamental, primary, or inseparable from the structure of reality.

Panpsychism proposes that consciousness or proto-consciousness is widespread throughout matter. Its main strength is that it avoids deriving experience from complete non-experience, but it faces the combination problem: how many small experiential elements become one unified subject.

Cosmopsychism shifts the problem from parts to wholes. Instead of beginning with tiny conscious units, it begins with the universe as a conscious whole and asks how individual minds arise within it. This avoids some versions of the combination problem but creates the decomposition problem: how one cosmic consciousness becomes many individual perspectives.

Idealism gives the strongest consciousness-first answer by treating consciousness as the primary reality and matter as an appearance, pattern, or structure within consciousness. This directly addresses the hard problem, but it must explain why the physical world appears stable, lawful, and shared among observers.

Taheri's T-Consciousness framework was presented as a contemporary consciousness-first model in which consciousness is fundamental, non-material, and expressed through consciousness fields. In this view, life and matter do not produce consciousness; rather, they arise within or through a prior consciousness-based reality. The framework is philosophically relevant to the book's central question, but it also faces challenges of mechanism, evidence, and testability.

Russellian monism and dual-aspect monism offer middle paths. They do not simply reduce consciousness to matter, but they also do not fully separate mind from the physical world. Instead, they suggest that the mental and physical may be two aspects of a deeper reality, or that consciousness may reveal the intrinsic nature of matter.

Process philosophy adds a dynamic view in which reality is made not of inert substances, but of events, relations, and becoming. It suggests that experience-like process may be woven into reality from the beginning, while biological life intensifies and organizes this process into richer forms.

Taken together, consciousness-first theories reshape the central question of the book. They challenge the assumption that life must come before consciousness and instead ask whether life may be one way consciousness becomes localized, organized, embodied, or expressed. Their strength lies in addressing the hard problem directly. Their weakness lies in explaining mechanism, individuation, embodiment, and scientific testability.

The open question is therefore:

If consciousness is fundamental, how does it become organized into individual living minds?

Chapter 12

How Consciousness Could Have Enabled Life

12.1 Chapter Overview

If consciousness is fundamental, then the origin of life cannot be understood only as a chemical accident or as a purely bottom-up transition from matter to biology. It must also be considered as a possible transition in the organization, expression, or localization of consciousness. In this view, life does not produce consciousness from nothing. Instead, life may be one way in which consciousness becomes structured into living form.

This chapter explores how consciousness-first theories might account for the emergence of life. It asks what mechanisms, metaphors, or models could connect fundamental consciousness to biological organization. These include consciousness as an organizing principle, proto-experience as a feature of chemical events, meaning as a condition of life, quantum approaches, and the possibility of a biogenic universe.

The chapter is intentionally cautious. Consciousness-first accounts of life remain speculative and are not part of mainstream origin-of-life science. They often lack the detailed mechanisms and testable predictions expected in scientific explanation. Yet they are philosophically important because they challenge the assumption that explanation must always move from matter to life to mind. They ask whether consciousness could be not the result of life, but one of the conditions that made life possible.

12.2 The Explanatory Challenge

It is not enough to say that consciousness is fundamental. A theory that begins with consciousness must still explain how biological life arises. If consciousness is basic, how does it become chemistry? How does chemistry become self-organization? How does self-organization become metabolism, heredity, cells, and evolution?

This is the central explanatory challenge for consciousness-first theories. They may offer an elegant response to the hard problem of consciousness by refusing to derive experience from non-experiential matter. But they inherit another problem: the problem of biological specificity. Why does consciousness organize into living systems rather than remaining undifferentiated? Why do cells, organisms, and nervous systems appear? Why does life have the biochemical structure it has?

A consciousness-first ontology must therefore do more than reverse the standard narrative. It must connect consciousness to mechanisms of organization. It must explain why matter behaves lawfully, why chemistry follows stable patterns, why some molecular systems become self-maintaining, and why living systems evolve.

The gap between ontological claims and biological details is significant. A metaphysical theory may claim that consciousness is fundamental, but origin-of-life research requires details about energy gradients, molecular formation, membranes, catalytic networks, replication, and environmental conditions. A successful consciousness-first model would need to show how consciousness relates to these processes without replacing biology with vague language.

This does not mean that consciousness-first theories are impossible. Many scientific theories begin with broad principles before developing detailed mechanisms. But the challenge remains: if consciousness enabled life, the theory must explain how.

12.3 Consciousness as an Organizing Principle

One possible consciousness-first model treats consciousness as an organizing principle. In this view, consciousness is not merely passive awareness. It has causal or formative power. It influences the movement from disorder toward order, from possibility toward structure, and from simple chemistry toward living organization.

This idea can be framed in different ways. In a strong teleological version, consciousness directs matter toward life, as if life were an intended or purposive outcome. In a weaker version, consciousness does not plan life, but acts as a tendency toward coherence, integration, complexity, or self-organization. It functions less like an external designer and more like an internal principle of ordering.

This raises a difficult question: can such a principle be scientific? Modern science generally avoids teleology, especially explanations that claim nature acts for a final purpose. However, biology often uses goal-like language in a restricted sense. Organisms maintain themselves, regulate internal states, repair damage, seek nutrients, and avoid harm. These behaviours are not necessarily conscious intentions, but they are organized around viability.

Terrence Deacon's idea of "absential" dynamics is useful here. Deacon argues that absence, constraint, and what is not present can play a causal role in living systems. A living system is shaped not only by material pushes but by constraints, possibilities, and needs. A missing nutrient, a boundary condition, or a future state to be maintained can organize present activity.

From a consciousness-first perspective, such constraints could be interpreted as early signs of meaning or directedness. Life may emerge when matter becomes organized not only by efficient causes but also by constraints that make some outcomes matter more than others. The system is no longer just reacting; it is preserving a pattern.

This does not prove that consciousness directs life. But it offers one possible bridge. Consciousness could be understood as the deep principle through which organization becomes oriented, constrained, and meaningful. Biological life would then be the material expression of this organizing tendency.

The risk is that this becomes too vague. If consciousness explains all order, then it may explain nothing specifically. To be useful, the idea must clarify what consciousness contributes that ordinary self-organization does not.

12.4 Taheri's View of the Origin of Life

Taheri's T-Consciousness framework provides one possible consciousness-first account of how life could arise. In this view, the origin of life is not only a chemical event. It is also an event of organization within

consciousness. Matter does not accidentally become alive through blind complexity alone; rather, life emerges when material forms become organized through the operation of consciousness fields.

This view does not necessarily deny chemistry, molecular interaction, or biological evolution. Chemical reactions, molecular self-organization, compartmentalization, and replication may still describe the physical side of abiogenesis. However, Taheri's framework would interpret these processes as incomplete if considered only from the material side. The deeper organizing principle is consciousness itself.

In standard origin-of-life theories, the main explanatory burden is placed on chemistry. Researchers ask how molecules became self-replicating, how metabolic cycles formed, how membranes emerged, and how heredity became possible. In Taheri's view, these processes may be understood as the visible or material expression of a prior consciousness-based order. Consciousness fields provide the organizing context within which matter becomes capable of life.

This gives the origin of life a different meaning. Life is not the first appearance of consciousness. Instead, life is a stage in the manifestation of consciousness through biological form. The living cell becomes more than a chemical machine. It becomes a localized expression of a deeper consciousness field, organized through boundary, metabolism, responsiveness, and relation.

The concept of organization is central here. A living system is not merely a collection of molecules. It is a coordinated whole. It maintains itself, responds to its environment, repairs itself, and preserves its identity over time. A Taherian interpretation would suggest that this coherence is not fully explained by material interactions alone. Consciousness fields may be understood as providing the non-material ordering principle that allows matter to become living organization.

This view also changes the meaning of biological evolution. Evolution would not be rejected, but it would be interpreted as a process occurring within a consciousness-pervaded reality. Natural selection would describe how living forms change over time, but consciousness would provide the deeper field in which such forms arise and develop.

The main strength of this view is that it offers a direct bridge between consciousness-first metaphysics and the origin of life. It does not treat consciousness as an accidental byproduct of late neural evolution. Instead, consciousness is present from the beginning as the condition through which life becomes possible.

The main limitation is that the mechanism remains difficult to specify scientifically. If consciousness fields organize matter, how can this be observed, tested, or distinguished from ordinary self-organization? How do such fields interact with chemistry without being reducible to physical force or energy? These questions remain open.

Taheri's view therefore provides a coherent consciousness-first narrative of abiogenesis, but it remains speculative unless it can be connected to clearer mechanisms, predictions, or forms of evidence.

12.5 Proto-Experience and Chemical Selection

Panpsychist theories suggest that matter has proto-experiential aspects. If this is true, then chemical reactions are not entirely devoid of interiority. They are not conscious in the human sense, but they may involve primitive experiential or proto-experiential events.

Within such a framework, the origin of life might involve more than chemical selection. It might involve something like experiential selection. Some configurations of matter may be more capable of integration, persistence, or proto-experiential coherence than others. Life-conducive chemistry may be favoured not only because it is chemically stable, but because it organizes proto-experience into more unified forms.

This idea is highly speculative. There is no established scientific evidence that proto-experience biases chemical reactions toward life. Standard chemistry explains molecular interaction through forces, energy

states, bonding, reaction pathways, and environmental conditions. A panpsychist model would need to show how proto-experiential properties make a difference without contradicting known chemistry.

Still, the idea is logically coherent within some forms of panpsychism. If matter has both physical and experiential aspects, then chemical processes may be both physical interactions and primitive experiential relations. The emergence of life would then be a transition in both domains: matter becomes biologically organized, and proto-experience becomes more integrated.

On this view, natural selection is not replaced. Biological evolution still operates once replicating systems exist. But before biological evolution, there may be chemical selection among molecular systems. A consciousness-first theorist might argue that proto-experiential organization could accompany or subtly shape this prebiotic selection.

The key question is whether proto-experience has causal relevance or is merely an intrinsic accompaniment of physical processes. If it has no causal role, then it does not help explain the origin of life. If it has causal power, then the theory must explain how that power operates.

This model remains speculative, but it points toward a deeper issue: if consciousness is fundamental, then the origin of life is not only a physical transition but also a transition in the organization of experience.

12.6 Information and Meaning

Standard information theory treats information as the reduction of uncertainty. It does not require meaning. A signal can carry information even if no conscious subject understands it. A sequence can be statistically informative without having semantic content.

Life, however, seems to involve more than bare information. In living systems, information matters to the system. A chemical signal can mean danger, nutrient availability, mating opportunity, damage, or stress. A gene sequence does not merely contain patterns; it participates in the production and regulation of the organism. Biological information is functional, contextual, and tied to survival.

Biosemiotics studies life as a process of sign interpretation. Thinkers such as Jesper Hoffmeyer and Marcello Barbieri have argued that meaning, signs, codes, and interpretation are not late additions to life but central to biology. Cells use molecular signs. Organisms interpret signals. Evolution produces codes and communication systems at many levels.

This creates a possible semiotic argument for consciousness-first theories. If life requires meaning, and if meaning requires some form of consciousness or proto-consciousness, then life may require consciousness at its foundation. In this view, life cannot be fully explained by chemistry alone because living chemistry is meaningful chemistry: chemistry organized around significance for a system.

The argument is powerful but controversial. One objection is that biological meaning need not require consciousness. A cell can respond to a signal without experiencing it. A genetic code can function without being understood by a mind. Meaning in biology may be functional rather than experiential.

A consciousness-first reply might argue that even functional meaning presupposes a point of view, however minimal. For something to be a sign, it must be a sign for a system. A nutrient gradient matters because it matters to the organism's continuation. Life creates relevance, and relevance may be the earliest form of meaning.

This view does not imply that cells think in words or possess reflective consciousness. Rather, it suggests that life is semiotic from the start. Living systems do not merely exchange matter and energy. They interpret conditions according to their own organization.

If meaning and interpretation are inseparable from life, and if consciousness is the ground of meaning, then consciousness could be seen as enabling life by making biological significance possible.

12.7 Quantum Approaches to Consciousness-Enabled Life

Some consciousness-first theories look to quantum mechanics because quantum theory raises difficult questions about observation, measurement, probability, and the role of the observer. In some interpretations, consciousness has been proposed as playing a role in the collapse of the wave function. The von Neumann-Wigner interpretation is one example historically associated with the idea that consciousness may be involved in measurement.

If consciousness plays a role in quantum measurement, then a consciousness-first theorist might ask whether consciousness could have influenced the quantum processes underlying abiogenesis. Chemical bonding, molecular formation, electron transfer, tunneling, and reaction pathways all depend on quantum processes. If consciousness participates at the quantum level, perhaps it could influence which molecular possibilities become actual.

This idea is highly speculative. Mainstream quantum mechanics does not require human consciousness to explain ordinary chemical reactions, and most scientists do not treat consciousness as necessary for quantum measurement. Environmental decoherence and other interpretations provide ways to understand quantum-to-classical transitions without invoking mind.

Still, quantum approaches remain philosophically interesting because they question the picture of matter as fully determinate and observer-independent in a simple classical sense. Quantum theory suggests that reality involves probability, measurement, and relational states in ways that continue to provoke debate.

John Wheeler's idea of a participatory universe also enters this discussion. Wheeler suggested that observers may be deeply involved in the universe's coming-to-be as an intelligible reality. Some interpretations of this idea speculate that observation does not merely passively register the universe, but participates in the formation of physical reality.

Applied to the origin of life, such ideas become even more speculative. Did consciousness shape the physical history that made life possible? Could future observers somehow participate in the meaningful structure of the past? These questions move beyond standard science and into philosophical speculation.

It is important to label this clearly. There is no established evidence that consciousness influenced abiogenesis through quantum measurement. But the idea belongs in a consciousness-first discussion because it shows one possible route by which mind and matter might be linked at a fundamental level.

This topic will be treated more fully in the later chapter on quantum and speculative frameworks.

12.8 The Biogenic Universe Hypothesis

If consciousness is fundamental, the universe may not be neutral with respect to life. It may be inherently biogenic: structured in such a way that life is a natural, perhaps expected, expression of reality. In this view, life is not an improbable accident imposed on dead matter. It is one way the universe expresses its underlying nature.

This idea can be interpreted in several ways. In a weak version, the laws of physics and chemistry happen to allow life, and consciousness-first theories interpret this as meaningful. In a stronger version, the universe is oriented toward life because consciousness seeks expression through living systems. In a cosmopsychist version, life emerges within a conscious universe as one of the forms through which cosmic consciousness differentiates itself.

The fine-tuning problem is relevant here. Many physical constants appear to fall within ranges that allow stars, chemistry, planets, and life. The anthropic principle states that we observe a life-permitting universe because only such a universe could contain observers. Some consciousness-first theories go further, suggesting that the presence of observers is not merely a selection effect but reveals something deep about the universe.

A biogenic universe hypothesis does not necessarily require that life appears everywhere. Rather, it suggests that life is not foreign to the universe. The same reality that gives rise to matter also gives rise to life and consciousness. Life is not an exception, but an expression.

This idea connects strongly to cosmopsychism. If the universe as a whole has a consciousness-like nature, then living organisms may be local centres of organized experience. The origin of life becomes the emergence of localized, bounded, self-maintaining forms through which consciousness becomes embodied.

The challenge is again testability. How would we distinguish a biogenic universe from a universe in which life simply happens to arise through ordinary physical processes? One possible route would be astrobiology. If life is common wherever conditions allow, then this may support the idea that life is a natural tendency of the universe. But it would not by itself prove consciousness-first metaphysics.

The biogenic universe hypothesis is therefore suggestive rather than conclusive. It reframes the origin of life as part of a larger cosmic tendency toward organization, complexity, and perhaps experience.

12.9 Critiques and Limitations

Consciousness-first models face serious critiques. The first is the “just-so story” problem. If consciousness is invoked to explain why life emerged, but no detailed mechanism is provided, the theory may become an attractive narrative rather than a scientific explanation.

The second problem is lack of falsifiability. A theory that can explain any possible outcome explains too much. If life emerges, consciousness wanted expression. If life is rare, consciousness expresses itself only under special conditions. If life is common, consciousness is universally creative. Without specific predictions, such explanations are difficult to test.

The third problem is causal ambiguity. If consciousness has causal power, how does it act on matter? Does it alter physical probabilities? Does it guide self-organization through constraints? Does it operate through information, meaning, or quantum events? Without a clear account, consciousness risks becoming an undefined force.

The fourth problem is parsimony. Standard origin-of-life models already attempt to explain life through chemistry, thermodynamics, and evolution. Adding consciousness may seem unnecessary unless it explains something those models cannot.

There is also the risk of category confusion. Consciousness, meaning, agency, and life may be related, but they are not identical. A theory that treats all organization as consciousness may blur important distinctions.

Proponents of consciousness-first theories respond in several ways. They argue that standard science also leaves deep questions unresolved, especially the hard problem of consciousness and the origin of biological meaning. They suggest that consciousness-first models may explain why experience exists at all, why the universe is intelligible, and why life emerges from matter capable of supporting it.

They may also argue that testability can develop over time. A consciousness-first theory might generate predictions about the relation between complexity and experience, the prevalence of life in the universe, the role of information in biology, or the limits of artificial consciousness.

Still, the critique remains strong. Consciousness-first theories must become more specific if they are to move from metaphysical possibility to scientific research program.

12.10 Implications for the Central Question

The models explored in this chapter provide a coherent but speculative narrative for consciousness enabling life.

If consciousness is an organizing principle, then life may emerge because consciousness gives direction, constraint, or coherence to self-organization. If matter has proto-experiential aspects, then chemical evolution may also involve the organization of primitive experience. If biological meaning requires consciousness, then life may depend on a semiotic dimension that cannot be reduced to chemistry alone. If consciousness is linked to quantum measurement or cosmic participation, then life may arise through processes in which mind and matter are not fully separable. If the universe is inherently biogenic, then life may be a natural expression of conscious reality.

These models challenge bottom-up explanation. They suggest that causation may not flow only from particles to molecules to cells to minds. It may also involve top-down constraints, intrinsic experience, meaning, or cosmic organization.

However, these theories remain minority positions in science. Origin-of-life research does not currently require consciousness to explain prebiotic chemistry. Consciousness-first theories are more developed as philosophical frameworks than as empirical models.

Their value may lie in expanding the conceptual space. Even if they prove wrong, they force us to ask what standard models leave unexplained. Do chemical models explain biological meaning? Do thermodynamic models explain agency? Do informational models explain experience? Does life require only mechanism, or also perspective?

For the central question, consciousness-first theories offer the reverse of the standard answer. Consciousness did not emerge from life. Life emerged within consciousness, or through consciousness, or as a structured expression of consciousness.

Whether this is a genuine explanation or a relocation of mystery remains open.

12.11 How This Chapter Changes the Central Question

This chapter changes the central question by examining speculative frameworks that challenge standard materialist explanations. Quantum and non-quantum consciousness-field theories suggest that consciousness may involve deeper levels of reality than ordinary biological mechanism alone.

The question therefore becomes more open but also more difficult. Speculative theories can expand the range of possible explanations, but they must be handled carefully. The challenge is to distinguish genuine explanatory insight from metaphor, overextension, or claims that cannot be tested.

12.12 Chapter Summary

This chapter examined how consciousness-first theories might explain the origin of life.

The key challenge is that consciousness-first theories must provide mechanisms, not only metaphysical claims. It is not enough to say that consciousness is fundamental. A theory must explain how fundamental consciousness becomes organized into living systems.

Several possible models were considered. Consciousness may act as an organizing principle, guiding or constraining self-organization. Proto-experience may accompany matter and become more integrated through chemical selection. Biosemiotic theories suggest that life depends on meaning and sign interpretation, which may imply a proto-conscious dimension. Quantum approaches speculate that consciousness could influence the physical processes underlying abiogenesis, though this remains highly speculative. The biogenic universe hypothesis suggests that life may be a natural expression of a consciousness-pervaded cosmos.

The chapter also examined major critiques. Consciousness-first models are often difficult to test, vulnerable to just-so storytelling, and lacking in specific mechanisms. They risk explaining too much unless they generate clear predictions. Yet they remain philosophically important because they challenge the assumption that explanation must always be bottom-up.

The open question is therefore:

Can consciousness-first models ever generate testable predictions, or are they permanently beyond empirical reach?

Chapter 13

Quantum and Speculative Frameworks

13.1 Chapter Overview

Quantum mechanics enters consciousness studies because both fields disturb the ordinary picture of reality. Quantum theory challenges classical assumptions about determinism, measurement, locality, and the role of observation. Consciousness challenges classical assumptions about matter, mechanism, and third-person explanation. It is therefore not surprising that some thinkers have tried to connect the two.

This chapter surveys quantum and speculative frameworks that attempt to explain consciousness and its relationship to life. These include Orchestrated Objective Reduction, quantum biology, observer-based interpretations of quantum mechanics, quantum mind theories, cosmological speculation, retrocausal ideas, and simulation hypotheses.

The chapter is cautious. Quantum mechanics is a rigorous and experimentally successful branch of physics. Consciousness studies is a serious interdisciplinary field. But the phrase “quantum consciousness” is often used loosely, sometimes to give speculative or mystical claims the appearance of scientific authority. For that reason, this chapter distinguishes carefully between established science, plausible but unproven hypotheses, and highly speculative extensions.

The aim is not to dismiss all quantum approaches. Quantum effects do occur in biological systems. It is also possible that current models of consciousness are incomplete. But invoking quantum mechanics does not automatically solve the hard problem. A quantum mystery is not necessarily an explanation of consciousness.

13.2 Why Quantum Mechanics Enters the Consciousness Debate

Quantum mechanics enters the consciousness debate for several reasons.

The first is the measurement problem. In quantum theory, systems can be described as existing in superpositions of possible states. Yet when a measurement occurs, we observe a definite outcome. What counts as a measurement? Why does one outcome appear rather than another? Does observation simply reveal a pre-existing state, or does it play a role in producing the outcome?

Some early interpretations of quantum mechanics seemed to give the observer a special role. This led some thinkers to ask whether consciousness itself might be involved in the collapse of the wave function. John

von Neumann and Eugene Wigner are often associated with versions of this idea, sometimes called the von Neumann-Wigner interpretation. In this view, consciousness is not merely an observer in the ordinary sense; it may play a causal role in the transition from quantum possibility to definite actuality.

The second reason is quantum indeterminacy. Classical physics often suggests a deterministic universe in which every event follows from previous physical states. Quantum mechanics introduces probability at a fundamental level. Some theorists have connected this indeterminacy to free will, agency, or mental causation. If the brain involves quantum events, perhaps consciousness could influence outcomes not fully fixed by classical determinism.

The third reason is the explanatory gap. Classical neuroscience explains many correlations between brain activity and experience, but it has not fully explained why physical processing should give rise to subjective experience. Some theorists therefore look beyond classical computation and classical neural mechanisms. They ask whether consciousness may depend on deeper physical principles not captured by ordinary neuroscience.

The historical context matters. Quantum mechanics arose at the same time that psychology, neuroscience, and philosophy were struggling to define the mind scientifically. The strange role of observation in quantum theory seemed, to some, to provide a possible opening for consciousness within physics.

However, caution is essential. In physics, an “observer” does not necessarily mean a conscious human observer. It may mean an interaction, a measuring apparatus, or an event that produces an effectively definite outcome. Confusing technical observation with conscious awareness is one of the most common mistakes in popular discussions of quantum consciousness.

Quantum mechanics enters the debate because it raises real conceptual puzzles. But those puzzles do not automatically imply that consciousness causes quantum collapse, or that quantum theory explains subjective experience.

13.3 Orchestrated Objective Reduction

Orchestrated Objective Reduction, usually called Orch-OR, is one of the best-known quantum theories of consciousness. It was developed by Roger Penrose and Stuart Hameroff. The theory combines Penrose’s argument that consciousness involves non-computable processes with Hameroff’s proposal that microtubules inside neurons may support quantum-level processes relevant to consciousness.

Penrose’s argument begins from mathematics and computation. He suggests that human understanding cannot be fully captured by algorithmic computation. Drawing on Gödel’s incompleteness theorems, Penrose argues that human mathematicians can see the truth of some statements that no formal computational system can prove from within its own rules. If this is correct, then consciousness may involve non-computable processes.

Penrose therefore looks beyond classical computation. He proposes that consciousness may depend on a physical process not reducible to ordinary algorithmic activity. His candidate is objective reduction: a proposed collapse of quantum superpositions caused not by external measurement but by gravitational self-collapse. In this view, quantum states become unstable when spacetime curvature differences reach a certain threshold.

Hameroff’s contribution is biological. He proposed that microtubules, structural components inside cells, especially neurons, could function as quantum processors. Microtubules are made of tubulin proteins and play important roles in cellular structure, transport, and organization. In neurons, they help maintain cell shape and intracellular dynamics.

According to Orch-OR, tubulin proteins within microtubules may enter quantum superposition states. These states may be “orchestrated” by biological processes in the brain and then undergo objective reduction. Each

objective reduction event is proposed to correspond to a moment of conscious experience. In this framework, consciousness arises from quantum processes organized by neural biology.

The theory is ambitious because it attempts to connect consciousness to fundamental physics and cellular biology. It does not treat consciousness as merely computation, nor as a separate non-physical substance. Instead, it locates consciousness in quantum events occurring within biological structures.

The empirical status is controversial. Some supporters point to possible links between anaesthesia and microtubule function, as well as evidence that quantum effects can occur in biological systems. If anaesthetics affect consciousness partly through interactions with microtubular or protein-level processes, that could be relevant to the theory.

Critics raise major objections. One influential criticism concerns decoherence. Quantum coherence is usually fragile and difficult to maintain in warm, wet, noisy biological environments such as the brain. Max Tegmark argued that quantum states in microtubules would decohere far too quickly to play a role in neural processing. Supporters of Orch-OR have disputed these calculations and argued for possible biological protection mechanisms, but the debate remains unsettled.

Other critics question whether the theory explains experience even if quantum effects occur. A quantum process is not automatically conscious. Orch-OR must show not only that microtubule quantum processes exist, but that they generate subjective experience and integrate with known neural dynamics.

Recent developments in quantum biology and microtubule research keep some parts of the discussion open, but Orch-OR remains a minority theory. It is more speculative than mainstream neural theories such as Global Workspace Theory or Recurrent Processing Theory.

For the central question of this book, Orch-OR is important because it suggests that consciousness may depend on fundamental physical processes organized through life. If true, consciousness would not be simply a late computational product of neural networks. It would involve a deeper interaction between biology and quantum physics.

13.4 Quantum Coherence in Biology

Quantum biology is a legitimate scientific field. It studies cases where quantum effects appear to play functional roles in biological systems. This is important because one common objection to quantum theories of consciousness is that biological systems are too warm, wet, and noisy for quantum effects to matter. Quantum biology shows that this objection is too simple.

Photosynthesis is one major example. Some studies suggest that quantum coherence may help explain efficient energy transfer in photosynthetic complexes. Excitation energy may move through molecular structures in ways that cannot be fully understood through classical hopping models alone. Although the interpretation and biological significance remain debated, photosynthesis has become a key example of quantum effects in living systems.

Avian magnetoreception is another important case. Some birds appear able to sense Earth's magnetic field, possibly through a radical pair mechanism involving quantum spin states. This does not mean birds are consciously detecting quantum events. It means that quantum-level processes may contribute to biological sensing.

Enzyme catalysis may also involve quantum tunneling, especially for electrons or protons moving through energy barriers. Some theories of olfaction have proposed that molecular vibration, not only molecular shape, may contribute to smell perception, though this remains debated.

These examples show that quantum effects can occur in biological systems. However, they do not prove quantum consciousness. There is a major difference between saying that quantum processes contribute to biological function and saying that quantum processes explain subjective experience.

The relevance to consciousness is indirect but important. If quantum coherence and tunneling can play roles in warm biological systems, then it is not impossible in principle that quantum effects could play some role in the brain. But possibility is not evidence. A quantum theory of consciousness must show where the relevant effects occur, how they are maintained, how they influence neural processing, and why they produce experience.

The current scientific consensus is cautious. Quantum biology is real, but quantum consciousness remains speculative. The brain may use quantum chemistry because all chemistry is quantum at a fundamental level. But most neuroscience can be explained at classical or biochemical levels without invoking large-scale quantum computation.

For the central question, quantum biology keeps open the possibility that life and consciousness depend on deeper physical principles than classical models assume. But it does not by itself show that consciousness existed before life or that quantum processes caused life to emerge.

13.5 The Observer Problem and Consciousness

The observer problem is one of the main reasons consciousness enters interpretations of quantum mechanics. In some simplified descriptions of the Copenhagen interpretation, quantum systems are said to remain in superposition until they are observed. This wording can suggest that consciousness is required for physical reality to become definite.

However, this is often misleading. In physics, observation usually means measurement or interaction with a measuring apparatus, not necessarily conscious awareness. A detector can register a particle without a human mind directly observing the event. The technical concept of measurement is not identical to subjective experience.

The von Neumann-Wigner interpretation gives consciousness a stronger role. It proposes that the chain of physical measurement does not fully resolve the collapse of the wave function until conscious observation occurs. In this view, consciousness is not merely another physical process but plays a special role in bringing about definite outcomes.

Wheeler's participatory universe offers another provocative idea. Wheeler suggested that observers are not passive spectators but participants in the universe's becoming intelligible. His phrase "it from bit" emphasizes the relationship between physical reality and information. In some interpretations, this suggests that observation and information are fundamental to reality.

Relational quantum mechanics, associated with Carlo Rovelli, offers a different view. It suggests that the state of a system is always relative to another system. There is no absolute state independent of interactions. This does not require consciousness, but it does challenge the idea of a detached observer-independent description from nowhere.

QBism, or Quantum Bayesianism, places the agent at the centre of quantum probabilities. In QBism, the quantum state represents an agent's expectations about experiences resulting from actions on the world. Again, this does not necessarily claim that consciousness collapses the wave function, but it does give the observer or agent a central role in interpreting quantum theory.

The key critique is that many popular discussions conflate "observer" with "conscious observer." This conflation allows sweeping claims: consciousness creates reality, quantum physics proves mind over matter, or observation requires human awareness. These claims are not supported by standard physics.

For consciousness studies, the observer problem is philosophically interesting but not decisive. It suggests that physical reality may be more relational and information-dependent than classical materialism assumed. But it does not automatically prove that consciousness is fundamental.

For the central question, observer-based quantum interpretations may support consciousness-first or co-emergence frameworks, but only if they can show that consciousness has a genuine role in physical measurement rather than being added metaphorically.

13.6 Quantum Mind Theories Beyond Orch-OR

Orch-OR is not the only quantum theory of mind. Several other theorists have proposed ways in which quantum processes might influence consciousness, attention, or brain function.

Henry Stapp developed a quantum mind theory in which consciousness and attention play roles in selecting among possible brain states. Drawing on quantum measurement, he argued that conscious intention could influence neural dynamics through repeated acts of attention. In this model, mind is not reducible to classical brain activity; it participates in the selection of physical outcomes.

John Eccles, working with Friedrich Beck, proposed that quantum mechanics might play a role at the synaptic cleft. The synapse is the junction where one neuron influences another through neurotransmitter release. Beck and Eccles suggested that quantum events could influence whether neurotransmitter release occurs, potentially giving mental intention a point of influence in brain activity.

David Bohm's implicate order offers a broader speculative framework. Bohm proposed that the visible world unfolds from a deeper implicate order, where separations between objects are less fundamental than they appear. Some thinkers have connected Bohm's ideas to consciousness, suggesting that mind and matter may unfold from a deeper unified order. Bohm himself explored such connections cautiously, especially in dialogue with philosophical and spiritual traditions.

Matthew Fisher's quantum cognition hypothesis focuses on nuclear spins in phosphorus atoms and possible Posner molecules. Fisher proposed that long-lived quantum coherence in such systems might influence neural processing. This theory is more specific than many quantum mind proposals because it identifies possible biological substrates and testable mechanisms. However, it remains highly speculative and under investigation.

These theories differ in mechanism. Stapp emphasizes attention and measurement. Beck and Eccles emphasize synaptic events. Bohm emphasizes a deeper implicate order. Fisher emphasizes possible quantum coherence in biochemical structures. None has become a mainstream theory of consciousness, but each reflects a shared intuition: classical neural computation may not fully explain mind.

The challenge is the same across these theories. They must show that relevant quantum effects occur in the brain, that they persist long enough to matter, that they influence neural activity, and that they explain subjective experience better than classical alternatives.

For the central question, quantum mind theories suggest that consciousness may not be merely a biological function layered on top of life. It may involve physical principles that also underlie matter and chemistry. If so, the boundary between life and consciousness may need to be rethought.

13.7 Non-Quantum Consciousness Field Theories

Not all consciousness-first or consciousness-field theories are quantum theories. This distinction is important because discussions of consciousness are often quickly connected to quantum mechanics, even when the theory in question does not depend on quantum measurement, superposition, collapse, or coherence.

Taheri's T-Consciousness framework is an example of a non-quantum consciousness-field theory. It proposes that consciousness is fundamental and non-material, but it does not need to be reduced to quantum mechanics. T-Consciousness is not described simply as energy, vibration, frequency, or a quantum field in the standard physical sense. It is instead presented as a non-material consciousness reality through which matter, life, and mind become organized or manifest.

This makes Taheri's framework different from Orch-OR, Stapp's quantum mind theory, or other theories that try to locate consciousness in specific quantum processes. Orch-OR proposes a mechanism involving microtubules and objective reduction. Stapp emphasizes quantum measurement and attention. Fisher's hypothesis proposes possible quantum coherence in biochemical structures. Taheri's view is broader and more metaphysical: consciousness is not a special quantum event inside biology, but a prior field-like reality that underlies biological and material organization.

This distinction matters because quantum language can sometimes obscure rather than clarify. If a theory is not actually making claims about quantum mechanics, it should not be treated as a quantum theory simply because it is non-material or field-like. A consciousness-field model may be speculative without being quantum. It may belong more properly to metaphysics, philosophy of mind, consciousness studies, or spiritual ontology than to physics.

The advantage of placing Taheri's view in this chapter is that it helps clarify the boundary between quantum consciousness theories and broader speculative consciousness-first frameworks. It shows that there are several ways to challenge classical materialism. Some appeal to quantum physics. Others appeal to idealism, cosmopsychism, panpsychism, dual-aspect monism, or consciousness fields.

The challenge for non-quantum consciousness-field theories is similar to the challenge faced by other consciousness-first models: they must explain how consciousness relates to measurable reality. If T-Consciousness is not physical energy, not frequency, and not ordinary information, then what is its mode of relation to biological systems? How does it organize matter? What would count as evidence for its activity?

Taheri's framework is therefore best treated as a speculative consciousness-first theory rather than as a quantum theory. It may be discussed alongside quantum frameworks because both question classical materialism, but it should not be collapsed into quantum mechanics. Its significance lies in its claim that consciousness is primary, non-local, and organizing, not in any specific quantum mechanism.

13.8 Speculative Extensions

Beyond specific quantum mind theories, a wider speculative landscape connects consciousness with cosmology, time, causality, and simulation.

The anthropic principle begins from the observation that the universe must be compatible with observers, because otherwise no observers would be here to notice it. In its weak form, this is a selection effect. In stronger forms, it can suggest that consciousness or observation has a deeper role in cosmic explanation. Some consciousness-first theories use anthropic reasoning to argue that consciousness is not incidental to the universe but central to its structure.

Consciousness and the arrow of time are also sometimes connected. Conscious experience appears to unfold in time, with memory directed toward the past and anticipation toward the future. Physics, however, raises deep questions about why time has a direction. Some speculative theories ask whether consciousness is related to temporal asymmetry, entropy, or the experience of becoming.

Retrocausality proposes that future events may influence past events under certain interpretations of quantum mechanics. Applied to consciousness, this leads to highly speculative ideas about whether future observers could influence earlier physical conditions, or whether the emergence of life and consciousness could somehow be linked to the universe's future. These ideas remain far outside mainstream science.

Simulation hypotheses propose that our universe could be a simulation or computational construct. In such frameworks, consciousness may be treated as emergent from computation, as fundamental to the simulation, or as something that cannot be simulated at all. Simulation theories often intersect with debates about artificial consciousness and the substrate-dependence of mind.

The boundary between creative theorizing and pseudoscience is important. Speculative theories are not automatically invalid. Science often advances through bold hypotheses. But a theory becomes problematic when it avoids evidence, uses scientific language vaguely, cannot be criticized, or explains every possible outcome.

A responsible speculative framework should clearly distinguish what is known, what is plausible, what is uncertain, and what is imaginative extension. It should not use quantum mechanics as a metaphor for anything mysterious. It should make contact with physics, biology, and consciousness studies in specific ways.

For this book, speculative frameworks are valuable because they explore the edges of current explanation. But they must remain clearly labeled as speculative unless supported by evidence.

13.9 Evaluating the Speculative Landscape

Quantum and speculative theories of consciousness should be evaluated by the same broad criteria as other theories: explanatory power, empirical testability, parsimony, and internal consistency.

Explanatory power asks what the theory clarifies. Does it explain subjective experience better than neural, biological, or informational theories? Does it explain why consciousness exists, why it has the structure it has, and why it correlates with brain activity?

Empirical testability asks whether the theory makes predictions that could be supported or refuted. A theory that makes no risky predictions remains philosophical rather than scientific. Orch-OR, Fisher's hypothesis, and some synaptic quantum theories are at least partly testable because they propose specific biological mechanisms. Broader claims about cosmic consciousness or observer-created reality are harder to test.

Parsimony asks whether the theory adds unnecessary assumptions. If ordinary neuroscience can explain a phenomenon, a quantum explanation may not be needed. A theory should not multiply mysteries. It should not invoke quantum mechanics merely because consciousness is difficult.

Internal consistency asks whether the theory is coherent. Does it use quantum concepts accurately? Does it distinguish physical measurement from conscious observation? Does it explain how quantum processes scale up to neural or experiential phenomena?

The risk of quantum mysticism is real. Quantum terms such as superposition, entanglement, observer, energy, and vibration are sometimes used vaguely in popular writing. This can create the appearance of explanation without scientific content. To avoid this, quantum consciousness theories must remain precise.

At the same time, speculative frameworks should not be dismissed only because they are speculative. The history of science includes ideas that began at the margins. Quantum biology itself would once have seemed unlikely to many. If a speculative theory becomes more precise, generates testable predictions, and survives criticism, it may become scientifically productive.

The value of speculative frameworks is that they challenge assumptions. They ask whether classical physicalism is enough, whether consciousness may require deeper physics, and whether life and mind are connected to principles not yet understood. Their danger is that they can become unfalsifiable narratives.

A mature approach requires both openness and discipline.

13.10 Implications for the Central Question

If quantum processes underlie consciousness, the relationship between life and consciousness becomes more complex.

On one hand, quantum consciousness theories may support consciousness-first views. If consciousness is tied to fundamental physical processes, then it may not be merely a late product of biological evolution. It may be connected to the basic structure of reality.

On the other hand, many quantum mind theories still require life. Orch-OR, for example, depends on biological structures in neurons. Fisher's hypothesis depends on specific biochemical systems. Synaptic quantum theories depend on neural function. These models do not necessarily say that consciousness exists before life. They may instead suggest co-emergence: consciousness requires both fundamental physics and biological organization.

Quantum effects may also have played roles in the origin of life. Quantum tunneling, electron transfer, molecular bonding, and energy transfer are all relevant to chemistry. If life depends on subtle quantum processes, then the origin of life may not be fully captured by classical chemistry alone. But this is different from saying that consciousness caused life.

These frameworks cannot yet explain why experience exists. A quantum process is not automatically a conscious process. Superposition, collapse, coherence, or entanglement do not by themselves produce subjectivity. Any quantum theory of consciousness must still explain how physical events become first-person experience.

Quantum and speculative frameworks therefore widen the range of possibilities. They make it harder to assume that classical mechanism is the final story. But they do not remove the need for biological, informational, and phenomenological explanation.

For the central question, the most cautious conclusion is this: quantum frameworks may support consciousness-first or co-emergence models, but only if they provide specific mechanisms connecting fundamental physics, living organization, and subjective experience. At present, they remain suggestive rather than decisive.

13.11 How This Chapter Changes the Central Question

This chapter changes the central question by placing consciousness within evolutionary history. If consciousness evolved, then it likely did not appear all at once. It may have developed gradually through sensation, affect, attention, learning, self-regulation, and social cognition.

The question therefore becomes one of continuity and threshold. Which forms of life have only responsiveness, which have cognition, which have sentience, and which have reflective consciousness? Evolution makes the boundary between life and mind harder to draw sharply.

13.12 Chapter Summary

This chapter surveyed quantum and speculative frameworks for understanding consciousness and its relationship to life.

Quantum mechanics enters consciousness debates because of the measurement problem, the role of the observer, quantum indeterminacy, and the explanatory gap in classical accounts of consciousness. Orch-OR

proposes that consciousness arises from quantum processes in microtubules, orchestrated by neural biology and linked to objective reduction. Quantum biology shows that quantum effects can play roles in living systems, including photosynthesis, magnetoreception, enzyme catalysis, and possibly other processes, but this does not prove quantum consciousness.

The observer problem has generated consciousness-based interpretations of quantum mechanics, but it is important not to confuse measurement with conscious observation. Other quantum mind theories, including those of Stapp, Beck and Eccles, Bohm, and Fisher, propose different mechanisms by which quantum processes might relate to mind. Broader speculative extensions connect consciousness with cosmology, time, retrocausality, and simulation hypotheses.

The chapter emphasized the need to distinguish established science from speculation. Quantum biology is a real field. Quantum consciousness remains controversial. Speculative theories can be valuable if they are precise, coherent, and testable, but they risk becoming pseudoscientific when they use quantum language vaguely or explain everything without evidence.

For the central question, quantum frameworks may support consciousness-first or co-emergence interpretations, especially if consciousness is connected to fundamental physical processes. Yet they do not currently provide a complete explanation of experience or of the origin of life.

The open question is therefore:

Do quantum frameworks offer genuine explanatory progress, or do they merely relocate the mystery?

13.13 Further Reading

- Penrose, R. (1994). *Shadows of the Mind*
- Hameroff, S., & Penrose, R. (2014). Consciousness in the universe: A review of the “Orch OR” theory. *Physics of Life Reviews*
- Stapp, H. (2011). *Mindful Universe*
- Ball, P. (2011). *Physics of Life*

Part VI

Part VI: Life Produced Consciousness?

Chapter 14

The Evolution of Consciousness

The standard scientific story begins with life and lets consciousness arrive later, through evolution, nervous systems, and increasingly complex forms of behaviour. But this raises a difficult question: where, along the long path from cells to minds, does mere biological responsiveness become awareness, subjectivity, or experience?

14.1 Chapter Overview

If consciousness emerged from life, then it must have an evolutionary history. It did not appear fully formed in human beings. Somewhere between the first living cells and reflective human awareness, biological systems became capable of sensing, valuing, remembering, choosing, feeling, and perhaps experiencing the world from within.

This chapter traces the possible evolutionary path from minimal cognition in brainless organisms to complex consciousness in animals with nervous systems. It begins with bacteria, protists, plants, fungi, and other organisms that show adaptive behaviour without brains. It then asks whether these behaviours are merely mechanical or whether they reveal the earliest forms of cognition.

The central issue is the divide between sensing and feeling. Many organisms detect and respond to their environments. But does detection imply experience? Is there something it is like to be a bacterium, a plant, a slime mold, an insect, or an octopus? Or does consciousness require nervous systems, centralized brains, and specialized forms of integration?

This chapter does not assume that all living systems are conscious. Instead, it asks how far back the roots of consciousness might extend, and whether evolution supports a sharp boundary or a gradual continuum between life, cognition, and experience.

14.2 Minimal Cognition: Defining the Starting Point

To understand the evolution of consciousness, we must begin by distinguishing cognition from consciousness. Cognition refers to processes such as sensing, information processing, learning, memory, decision-making, and adaptive behaviour. Consciousness refers to subjective experience: what it is like to be a system undergoing those processes.

The distinction is crucial. A system may process information without being conscious. A thermostat responds to temperature, but we do not usually assume it feels warmth or cold. A computer can classify images

without seeing them in the experiential sense. A cell can respond to chemical signals without necessarily having awareness.

Yet cognition itself may not require a brain. The study of basal cognition, associated with thinkers such as Michael Levin, Pamela Lyon, and others, examines intelligence-like processes in cells, tissues, and organisms without nervous systems. These systems can sense conditions, process information, coordinate activity, remember past states, and adapt to changing environments.

Functional definitions of cognition often emphasize four capacities: sensing, information processing, adaptive response, and memory. A system is cognitive, in this broad sense, if it detects relevant differences, uses those differences to guide behaviour, changes in relation to context, and preserves some trace of past interactions.

This broad definition is useful because it reveals continuity across life. Bacteria, immune cells, plants, fungi, and tissues all exhibit forms of regulation and adaptation. They do not think in the human sense, but they are not passive machines either.

The difficult question is where mere mechanism ends and cognition begins. If every causal response counts as cognition, then the concept becomes too broad. If cognition requires human-like thought, then it becomes too narrow. A useful middle path treats cognition as adaptive, self-relevant information processing in a living system.

This definition does not prove consciousness. But it establishes the starting point. Before life could become conscious, it first had to become responsive in ways that mattered to its own continuation.

14.3 Bacterial Intelligence

Bacteria are among the simplest living organisms, yet their behaviour is more sophisticated than their size suggests. They sense chemical gradients, move toward favourable conditions, communicate with one another, form communities, and alter their behaviour according to environmental context.

Chemotaxis is one of the clearest examples. *E. coli* bacteria move through a run-and-tumble pattern. When conditions improve, they tend to continue moving in the same direction. When conditions worsen, they tumble and change direction. This allows them to move toward nutrients and away from harmful substances.

What makes chemotaxis especially interesting is that bacteria do not simply measure concentration at one instant. They compare present conditions with recent past conditions. This temporal comparison functions as a primitive form of memory. The bacterium does not need a brain to track change over time. Its molecular networks preserve information long enough to guide movement.

Quorum sensing provides another example. Bacteria release and detect chemical signals that indicate population density. When enough bacteria are present, they can collectively change behaviour. They may produce toxins, emit light, form biofilms, or activate cooperative functions. This resembles collective decision-making, though it occurs through chemical signaling rather than conscious deliberation.

Biofilm formation further complicates the image of bacteria as isolated simple units. In biofilms, bacteria adhere to surfaces and to one another, produce protective matrices, exchange signals, and display division of labour. The community can become more resilient than individual cells.

The philosophical question is whether bacterial behaviour is “about” anything. Does the bacterium merely react chemically, or does it treat some conditions as good or bad for itself? Intentionality, in philosophy, refers to aboutness: the directedness of mental states toward objects or conditions. Bacteria almost certainly do not have beliefs or desires in the human sense. But their behaviour is directed toward viability. Nutrients, toxins, density, and surfaces matter to them because they affect survival and reproduction.

This may be a minimal form of biological intentionality: not conscious meaning, but organism-relative relevance. The bacterium’s world is not a neutral physical field. It is structured by gradients of value: approach, avoid, remain, divide, cooperate.

Whether this is consciousness is another matter. Bacterial intelligence shows that cognition-like processes exist without nervous systems. It does not show that bacteria have experience. But it reveals that life is already world-directed at very simple levels.

14.4 Protists and Single-Celled Decision-Making

Single-celled eukaryotes provide some of the strongest examples of cognition without neurons. Unlike bacteria, protists have complex internal organization, including nuclei, organelles, cytoskeletal systems, and intricate behavioural repertoires. They are single cells, but they are far from simple.

Paramecium can swim, avoid obstacles, respond to chemical and mechanical stimuli, and adjust its behaviour. When it encounters a barrier, it can reverse its ciliary movement, back away, turn, and try another direction. This avoidance behaviour suggests flexible action rather than a single fixed reflex.

Some studies have also discussed habituation-like behaviour in single-celled organisms. Habituation is a simple form of learning in which a system reduces its response to a repeated harmless stimulus. If a single cell can modify its future behaviour based on repeated exposure, then memory and learning do not require a nervous system in the strict sense.

Physarum polycephalum, a slime mold, is especially famous. Despite lacking neurons, it can solve mazes, optimize networks, anticipate periodic events, and distribute resources efficiently. In experiments, *Physarum* has formed networks resembling efficient transport routes. It can balance exploration and exploitation, avoid harmful conditions, and adapt to changing environments.

Stentor roeselii, a single-celled protist, has been reported to show hierarchical decision-making. When disturbed, it may first bend away, then reverse ciliary beating, then contract, and finally detach and swim away if stimulation continues. This sequence suggests ordered behavioural alternatives rather than a simple automatic response.

These organisms challenge the idea that cognition begins with nervous systems. A single cell can sense, integrate, remember, decide among alternatives, and act adaptively. Its cytoskeleton, membrane, signaling pathways, and internal dynamics perform functions that, in animals, are often associated with nervous systems.

But again, cognition is not the same as consciousness. A protist may make decisions in a functional sense without having subjective experience. The word “decision” here does not necessarily imply deliberation. It refers to flexible selection among possible actions.

What protists show is that the roots of cognition are cellular. Neurons did not invent biological responsiveness; they specialized and amplified it. Nervous systems may be evolutionary elaborations of capacities that began in single cells: sensing, signaling, memory, adaptation, and coordinated action.

For the central question, this is important. If consciousness evolved from cognition, then the pre-neural world already contained many of the functional ingredients from which consciousness later developed.

14.5 Plant Cognition and Signaling

Plants do not have brains, neurons, or muscles, yet they sense and respond to their environments in sophisticated ways. They grow toward light, orient roots with gravity, respond to touch, detect chemical signals, communicate through volatile compounds, and adjust growth according to resource availability.

Phototropism is the growth of plants toward light. Gravitropism allows roots and shoots to orient in relation to gravity. Thigmotropism refers to growth responses to touch, such as the coiling of tendrils around supports. These behaviours unfold more slowly than animal movement, but they are adaptive responses to environmental information.

Plants also use chemical and electrical signaling. Signals can travel through tissues to coordinate defence, growth, and stress responses. When attacked by herbivores, some plants release chemicals that warn neighbouring plants or attract predators of the attacking insects. Roots can respond to nutrient gradients, moisture, obstacles, and competing roots.

The idea of root intelligence has been proposed by researchers such as Anthony Trewavas and Stefano Mancuso. They argue that plants exhibit decentralized intelligence: sensing, integration, adaptive growth, memory-like changes, and decision-making distributed across the organism. A plant does not need a central brain because its body is organized differently. Its intelligence is developmental, distributed, and ecological.

Critics argue that plant intelligence language risks anthropomorphism. Plants respond through biochemical and physiological mechanisms, but this does not mean they think, feel, or consciously decide. The concern is that human terms such as memory, choice, and communication may be misleading if applied too loosely.

The debate depends partly on definitions. If intelligence means human-like reasoning, plants are not intelligent. If intelligence means adaptive problem-solving in relation to environmental challenges, plants may qualify. If consciousness means subjective experience, the evidence remains much weaker and more controversial.

Plants are important for the evolution of consciousness because they separate cognition-like function from nervous systems. They show that living systems can integrate environmental information without brains. They also show that adaptive behaviour can be embodied in growth, development, chemical signaling, and ecological relation rather than rapid movement.

For the central question, plants suggest that life became cognitively active long before nervous systems. Whether that cognition includes feeling remains unresolved.

14.6 Fungal Networks and Distributed Intelligence

Fungi offer another model of distributed biological intelligence. Many fungi grow as networks of thread-like structures called hyphae, which form mycelial systems. These networks can spread through soil, wood, or living tissues, connecting environments in complex ways.

Mycorrhizal fungi form partnerships with plant roots. Through these associations, fungi help plants access nutrients such as phosphorus and nitrogen, while plants provide sugars produced through photosynthesis. These networks can connect multiple plants, creating what is popularly called the “wood wide web.”

Through mycorrhizal networks, resources and signals may move between connected plants. Some studies suggest that plants connected by fungal networks can share carbon, respond to stress signals, or influence each other’s defence responses. The exact interpretation of these findings remains debated, but the broader point is clear: fungal networks create ecological connectivity.

Fungi also show resource allocation behaviours. A mycelial network can grow toward nutrients, withdraw from poor conditions, reinforce productive pathways, and redistribute resources across its structure. Like slime molds, fungal systems can show optimization-like behaviour without a central nervous system.

The comparison with neural networks is tempting but must be used carefully. Fungal networks are not brains. They do not transmit signals in the same way neurons do, and there is no evidence that they support conscious experience. However, both fungal and neural networks demonstrate how distributed systems can process information through connectivity, feedback, and adaptive change.

Fungal networks broaden the concept of cognition beyond individual organisms. Intelligence-like processes may occur at the level of networks, symbioses, and ecosystems. This raises difficult questions about the unit of cognition. Is cognition located in the individual organism, the colony, the network, or the ecological relationship?

For the evolution of consciousness, fungi suggest that distributed intelligence can exist without centralization. But if consciousness requires unity of experience, then distributed networks raise a challenge. Can a decentralized system have a unified point of view? Or is central integration necessary for experience?

This question will recur when we consider nervous systems. Consciousness may require not only information exchange, but some form of integrated perspective.

14.7 Sensing vs Feeling: The Critical Divide

The central difficulty in the evolution of consciousness is distinguishing sensing from feeling. Many systems detect stimuli and respond adaptively. But detection is not necessarily experience.

A thermostat detects temperature and activates heating or cooling. It has a functional relation to the environment, but we do not assume that it feels cold. A bacterium detects chemicals and moves accordingly. A plant detects light and grows toward it. A nervous system detects damage and generates defensive responses. At what point, if any, does detection become feeling?

Functional sensing involves the use of information to guide behaviour. Phenomenal experience involves there being something it is like to undergo that state. The difference between the two is the heart of the consciousness problem.

Thomas Nagel's famous question, "What is it like to be a bat?" can be extended cautiously to simpler organisms. Is there something it is like to be *E. coli*? Is there something it is like to be a slime mold, a plant, or a jellyfish? These questions may seem strange because these organisms lack brains like ours. But the strangeness reveals the problem: we do not know how to identify the minimum conditions for experience.

The overflow problem asks when information processing becomes more than information processing. Living systems process signals, regulate themselves, and respond to conditions. But why should any of this generate a felt point of view? If a system can perform all necessary functions without experience, why did experience evolve?

There are two broad positions. Sharp boundary views hold that consciousness appears only when certain conditions are met, such as a nervous system, recurrent processing, global integration, or self-modeling. Before that threshold, systems may sense but not feel.

Gradualist views hold that consciousness develops continuously. There may be no single moment when experience appears. Instead, proto-experience, sentience, affect, perception, and self-awareness may form a spectrum. Simple organisms may have no experience, or extremely minimal forms. More complex animals may have richer experience.

Both views face difficulties. Sharp boundary views must justify the boundary. Gradualist views must avoid attributing consciousness too broadly.

The sensing-feeling divide is therefore the critical divide. It is the point where the evolution of cognition becomes the evolution of consciousness.

14.8 Memory and Learning Without Brains

Memory is often associated with nervous systems, but living systems can preserve traces of past events without brains. Memory, in the broad biological sense, means that past states influence future responses.

Single-celled organisms can show habituation-like changes, altered responsiveness, and adaptation based on prior exposure. Bacteria can alter gene expression in response to past conditions. Immune cells can show memory-like behaviour. Plants can display priming, where prior stress influences later responses. Cells can maintain developmental states across divisions.

Epigenetic memory is especially important. Chemical modifications to DNA or associated proteins can influence gene expression without changing the underlying genetic sequence. These modifications can preserve information about developmental history, environmental stress, or cellular identity. Epigenetic mechanisms show that memory can be molecular, structural, and regulatory.

Non-neural memory can also be stored in cellular architecture, protein states, metabolic networks, immune repertoires, and tissue organization. The nervous system did not invent memory. It developed a rapid, flexible, and highly integrated form of memory built upon older biological capacities.

This raises an important question: if memory exists without neurons, can consciousness? The answer depends on what memory contributes to consciousness. Consciousness may require some form of temporal integration. Experience is not merely an instant; it has continuity. Perception depends on recent context. Selfhood depends on memory of bodily and personal states.

However, memory alone is not enough. A scar records past injury, but it is not conscious. DNA stores evolutionary history, but it does not experience. Memory becomes relevant to consciousness when it is integrated into a system's ongoing perspective.

Non-neural memory shows that the ingredients of consciousness may have deep biological roots. But it also shows why ingredients are not sufficient. Life can remember without necessarily experiencing memory.

For the central question, memory without brains supports continuity between life and mind, but it does not eliminate the need for a threshold or transition into feeling.

14.9 The Emergence of Nervous Systems

Nervous systems transformed biological responsiveness. They allowed organisms to sense rapidly, coordinate movement, integrate signals across the body, and respond flexibly to complex environments. If consciousness requires more than cellular cognition, the emergence of nervous systems is a major turning point.

The earliest nervous systems were likely simple nerve nets rather than centralized brains. Cnidarians, such as jellyfish and hydra, have nerve nets that coordinate movement, feeding, and response to stimuli. These systems do not have brains in the vertebrate sense, but they allow faster and more integrated behaviour than non-neural signaling.

Nerve nets raise an important question. Does consciousness require a centralized brain, or can a distributed nervous system support minimal experience? A jellyfish responds to light, gravity, touch, and chemical cues. It moves in coordinated ways. But whether there is something it is like to be a jellyfish remains uncertain.

Bilateral symmetry and cephalization were major evolutionary developments. Bilateral organisms have a front and back, left and right, dorsal and ventral orientation. Movement through the environment creates advantages for having sensory organs and neural processing concentrated near the front. Cephalization, the concentration of nervous tissue and sensory organs at one end of the body, helped create the conditions for centralized brains.

The Cambrian explosion saw a dramatic increase in animal diversity and sensory complexity. Predation, movement, vision, body plans, and ecological interactions intensified. Nervous systems became increasingly important for survival. Seeing, chasing, escaping, hiding, and choosing required fast integration of sensory information.

This evolutionary context may have driven the emergence of consciousness. Consciousness may have evolved because mobile animals needed to integrate perception, action, bodily state, and value into flexible behaviour. The world became not merely a chemical environment but a field of objects, threats, opportunities, and goals.

If this is correct, consciousness is closely tied to animal life: movement, sensing, action, and nervous integration. Life came first, but consciousness emerged when life became mobile, embodied, and neurally integrated in new ways.

14.10 Key Evolutionary Candidates

Several groups of animals are especially important in debates about the evolution of consciousness.

Arthropods, including insects and crustaceans, have centralized nervous systems with ganglia, sensory processing, learning, memory, and complex behaviour. Bees navigate, communicate, learn visual patterns, and show flexible foraging. Some insects appear to display attention-like processes and motivational states. Debates about insect pain remain active, but there is growing interest in whether at least some arthropods may have sentience.

Cephalopods, especially octopuses, are among the most striking examples of non-vertebrate intelligence. Octopuses solve problems, explore objects, use flexible behaviour, learn, camouflage, and manipulate their environments. Their nervous systems are highly distributed, with many neurons located in the arms as well as the central brain. This challenges vertebrate-centered assumptions about consciousness.

Early vertebrates provide another important line. Fish and amphibians have centralized brains, sensory integration, learning, and behaviour that suggests more than simple reflex. Debates about fish pain have become significant in animal welfare and consciousness research.

Todd Feinberg and Jon Mallatt have proposed a neurobiological framework for primary consciousness. They argue that consciousness likely arose with specific neural features, including mapped sensory representations, affective systems, and integrated neural processing. In their view, primary consciousness is not the same as human self-consciousness. It is basic subjective experience: seeing, feeling, sensing, and responding in an integrated way.

This framework suggests that consciousness may have evolved in animals with sufficiently complex nervous systems, perhaps during early vertebrate or arthropod evolution. It places the threshold beyond bacteria, plants, and single-celled organisms, but earlier than human reflective consciousness.

The key evolutionary candidates therefore include animals with integrated sensory systems, affective valuation, learning, and flexible action. Consciousness may not require language, culture, or human intelligence. But it may require neural architectures capable of binding information into a unified field of perception and value.

For the central question, these candidates support a life-first view with a gradual transition. Life began long before consciousness, but consciousness may have emerged in multiple animal lineages as nervous systems became capable of integrated experience.

14.11 Theoretical Frameworks for Basal Cognition

Several contemporary theories attempt to explain basal cognition and its relation to consciousness.

The Free Energy Principle, developed by Karl Friston, treats living systems as self-organizing systems that minimize uncertainty or free energy. From this perspective, organisms maintain themselves by predicting and regulating their states in relation to the environment. This framework can apply broadly, from cells to brains.

Predictive processing extends this idea to perception and cognition. A brain does not passively receive the world; it predicts the causes of sensory input and updates its models. Some theorists apply predictive ideas to cellular systems, arguing that even cells act in ways that anticipate, regulate, and infer environmental conditions.

Integrated Information Theory can also be applied to simple organisms in principle. If consciousness depends on integrated information, then simple systems may have very low but non-zero Φ . This raises the possibility of minimal consciousness in organisms far simpler than humans. However, measuring integrated information in real biological systems is difficult, and the implications are controversial.

Autopoietic enactivism, associated with thinkers such as Evan Thompson and Ezequiel Di Paolo, emphasizes life as self-producing and world-enacting. In this view, cognition begins with living organization. A living system brings forth a meaningful world through its own self-maintaining activity. Cognition is not representation alone; it is embodied sense-making.

Agency and minimal selfhood are central to these frameworks. An organism is not just a physical object. It is a system that maintains itself, distinguishes itself from its environment, and acts in relation to its own viability. This creates a minimal self in the biological sense: an organized centre of activity for which conditions can be better or worse.

The disagreement concerns whether minimal selfhood implies consciousness. Enactivists often emphasize continuity between life and mind. More conservative theorists argue that cognition may begin with life, but consciousness requires nervous systems.

These frameworks help clarify the options. If consciousness is integrated information, it may be widespread and graded. If consciousness is predictive self-modeling, it may require more complex nervous systems. If consciousness is embodied sense-making, it may be rooted in life itself. If consciousness is global access or recurrent neural processing, it begins later.

Basal cognition does not settle the issue, but it shows that the path from life to consciousness is not empty. It is filled with intermediate forms of sensing, memory, agency, and self-regulation.

14.12 Implications for the Central Question

The evolution of consciousness can be read in two broad ways.

The first reading emphasizes continuity. Cognition exists without nervous systems. Bacteria navigate chemical gradients. Protists make flexible behavioural choices. Plants signal, remember, and respond. Fungi form adaptive networks. Cells regulate themselves and preserve memory. From this perspective, consciousness may be the upper end of a spectrum that begins with life itself. Proto-consciousness may have preceded complex animals, and nervous systems may have intensified capacities already present in living systems.

The second reading emphasizes discontinuity. Cognition may exist without consciousness. A bacterium can sense without feeling. A plant can signal without experiencing. A slime mold can solve problems without having a point of view. On this reading, consciousness requires nervous systems, integrated sensory maps,

affective states, and perhaps recurrent or global processing. Life came first, and consciousness emerged only after specific biological innovations.

The evolutionary record does not settle the debate. Fossils can tell us about body plans, sensory organs, movement, and ecological relationships. They cannot directly tell us what extinct organisms experienced. We infer consciousness from anatomy, behaviour, neurobiology, and comparison with living organisms. These inferences are always uncertain.

Still, the evolutionary record suggests that consciousness, if it emerged from life, likely emerged gradually rather than suddenly. There may not have been a single first conscious organism. Instead, there may have been a long transition from irritability to sensing, from sensing to integrated perception, from perception to affective experience, and from experience to self-awareness.

For the central question, the answer depends on where we place the threshold. If consciousness requires nervous systems, then life clearly came first. If consciousness is graded and rooted in biological self-organization, then consciousness and life may be more deeply intertwined. If proto-consciousness is attributed to all living systems, then the origin of life may also mark the origin of primitive subjectivity.

The key unresolved problem remains the same: how do we determine where sensing ends and experiencing begins?

14.13 How This Chapter Changes the Central Question

This chapter changes the central question by identifying the hardest point in the life-first story: the transition from living organization to subjective experience. Explaining life is not the same as explaining what it is like to be alive.

The question therefore becomes: when does biological activity become felt experience? Life-first theories must explain this transition, while consciousness-first theories must explain how consciousness becomes embodied and localized. In both cases, the hard transition remains central.

14.14 Chapter Summary

This chapter traced the possible evolutionary pathway from minimal cognition to consciousness.

It began by distinguishing cognition from consciousness. Cognition can be understood as sensing, information processing, adaptive response, learning, and memory. Consciousness involves subjective experience. Brainless organisms such as bacteria, protists, plants, fungi, and slime molds show many forms of cognition-like behaviour, including chemotaxis, quorum sensing, habituation, signaling, distributed problem-solving, and adaptive resource allocation.

These examples show that intelligence-like processes do not begin with brains. Life was already responsive, adaptive, and world-directed before nervous systems evolved. However, sensing is not the same as feeling. The critical divide is between functional detection and phenomenal experience.

The emergence of nervous systems transformed biological responsiveness by enabling rapid signaling, sensory integration, movement, and flexible action. Nerve nets, centralized brains, bilateral body plans, and sensory complexity all contributed to the conditions under which consciousness may have emerged. Arthropods, cephalopods, and vertebrates are especially important candidates for the evolution of conscious experience.

Theoretical frameworks such as the Free Energy Principle, predictive processing, Integrated Information Theory, autopoietic enactivism, and basal cognition offer different ways of interpreting the continuity between life and mind. Some suggest that consciousness is graded and deeply rooted in life. Others suggest that consciousness requires specific neural architectures.

The open question is therefore:

Is there a principled way to determine where sensing ends and experiencing begins?

Chapter 15

The Hard Transition

15.1 Chapter Overview

The previous chapters examined two broad possibilities. One possibility is that consciousness is fundamental, or at least deeply woven into reality. Another is that consciousness emerged from life through biological evolution. This chapter focuses on the second possibility and asks its most difficult question: if consciousness emerged from life, when did the transition occur?

This is not exactly the same as the hard problem of consciousness. The hard problem asks why physical or biological processes give rise to subjective experience at all. The hard transition problem asks when, where, and how experience first appeared in evolutionary history. At what point did a living system become not only responsive, adaptive, and cognitive, but also capable of feeling?

This chapter examines gradualist models, threshold models, developmental evidence, functional and phenomenal definitions of consciousness, and the implications of major contemporary theories. It argues that the hard transition remains one of the deepest unresolved problems in consciousness studies. We can identify many biological correlates of consciousness, but it remains difficult to say when information processing first became experience.

15.2 Stating the Problem Precisely

The hard problem of consciousness, associated with David Chalmers, asks why there is subjective experience at all. Why does neural processing feel like something from the inside? Why is there an experience of colour, pain, sound, fear, or selfhood, rather than only unconscious information processing?

The hard transition problem is related but distinct. It asks when subjective experience first appeared. If consciousness emerged through evolution, there must have been a transition from organisms that were alive but not conscious to organisms that were alive and conscious. The question is where that transition occurred and what made it possible.

A bacterium senses chemical gradients and moves adaptively. A plant responds to light, gravity, and touch. A slime mold solves spatial problems. A jellyfish has a nerve net. An insect learns. An octopus explores. A mammal feels pain, fear, and attachment. Somewhere along this continuum, if consciousness is not fundamental, experience must have emerged.

The difficulty is that we can observe behaviour and anatomy, but not experience directly. We can identify neural correlates of consciousness in humans and some animals. We can study brain activity during perception, sleep, anaesthesia, and injury. But neural correlates tell us what accompanies consciousness, not why those processes produce experience or when such processes first became sufficient.

This creates an explanatory gap between function and phenomenology. A system may detect stimuli, process information, regulate behaviour, and learn from past events. But does it feel anything? Does it have a point of view? Is there something it is like to be that system?

The hard transition problem therefore asks for more than a biological timeline. It asks for a principle. What marks the crossing from non-conscious life to conscious life?

15.3 Gradualism: Consciousness as a Continuum

Gradualism proposes that consciousness did not appear suddenly. Instead, it developed along a spectrum: from proto-consciousness, to minimal consciousness, to sentience, to rich perceptual consciousness, to reflective self-consciousness.

This view fits well with evolutionary continuity. Darwin emphasized that differences between species often involve gradual variation rather than absolute breaks. If bodies, nervous systems, cognition, and behaviour evolved gradually, it may seem reasonable to think that consciousness did as well.

Gradualism is also supported by the absence of an obvious break point in the tree of life. There is no clear line where all organisms on one side are certainly unconscious and all organisms on the other side are certainly conscious. Nervous systems vary in complexity. Some animals have simple nerve nets. Others have ganglia. Others have centralized brains. Others have highly developed cortical or distributed neural systems. Behavioural flexibility also appears in degrees.

Neural complexity increases gradually across evolutionary history. Sensory integration, memory, affect, learning, attention, and self-monitoring do not all appear at once. If consciousness depends on these capacities, then consciousness may also come in degrees.

The strength of gradualism is that it avoids an arbitrary switch. It does not require consciousness to appear suddenly in one species or one moment. It allows for continuity between life, cognition, sentience, and awareness.

But gradualism faces serious problems. At what point on the continuum does experience begin? If consciousness comes in degrees, is “a little bit conscious” meaningful? Does a minimally conscious organism have a faint inner life, or is this just a metaphor?

This is the sorites problem applied to consciousness. If we add one grain of sand at a time, when does a heap become a heap? If we add one neuron, one feedback loop, or one bit of integration at a time, when does non-conscious processing become conscious experience?

Gradualism may be right that consciousness has no sharp boundary. But if the boundary is completely vague, it becomes difficult to determine which organisms have moral status, which systems feel pain, and what kind of biological organization is sufficient for experience.

Gradualism makes the transition smoother, but it does not make it disappear.

15.4 Threshold Models: A Discrete Transition

Threshold models propose that consciousness appears only when a system reaches a specific level of organization. Below that threshold, there may be life, cognition, sensing, and adaptive behaviour, but no subjective experience. Above it, experience appears.

Different theories propose different thresholds. Some place the threshold at the emergence of nervous systems. On this view, cells, plants, and bacteria may be responsive but not conscious. Consciousness begins when specialized neural tissue allows rapid, integrated signaling.

Other models require centralized brain architecture. A nerve net may coordinate movement, but a centralized nervous system can integrate multiple sensory streams, guide action, and create unified representations of the body and world.

Re-entrant processing theories place the threshold at recurrent neural activity. Feedforward processing may remain unconscious, while feedback loops between brain areas allow conscious perception to emerge. On this view, consciousness begins when neural systems can stabilize and integrate information through recurrent processing.

Global Workspace Theory places the threshold at global broadcast. Consciousness appears when information becomes widely available across a system for memory, decision-making, attention, and action. A system may process information unconsciously, but it becomes conscious when that information is globally accessible.

Integrated Information Theory proposes a different kind of threshold. Consciousness is related to the degree of integrated information within a system. Depending on interpretation, this may imply either a graded model or a threshold at which integration becomes significant enough for consciousness.

The strength of threshold models is that they offer clearer criteria. They may help distinguish conscious from non-conscious systems. They also align with the intuition that not all information processing is conscious. A thermostat, bacterium, or simple reflex circuit may respond to stimuli without experience.

But threshold models face their own problems. Why this threshold and not another? Why should recurrent processing, global broadcasting, or integrated information produce experience? If the line is drawn too high, many animals that plausibly feel may be excluded. If it is drawn too low, consciousness may be attributed too broadly.

Threshold models also face the zombie problem below threshold. If a system just below the threshold behaves almost identically to a system just above it, why would one have experience and the other not? What changes, ontologically, at the threshold?

Threshold models make the transition clearer, but they must justify the boundary.

15.5 Developmental Perspectives

Development provides another way to think about the hard transition. Instead of asking only when consciousness appeared in evolution, we can ask when consciousness appears in individual development. Does the development of consciousness in an organism offer clues about its evolutionary emergence?

Fetal consciousness is a difficult and ethically sensitive topic. Early in development, the nervous system is not sufficiently organized for conscious experience. Over time, neural structures form, sensory systems develop, and connections between the thalamus and cortex emerge. Thalamocortical connections are often discussed as important for conscious processing and are commonly associated with later fetal development, around the period when the fetus becomes more neurologically capable of organized sensory response.

However, neural connectivity alone does not prove consciousness. Fetal responses to stimuli may be reflexive or unconscious. The intrauterine environment also includes conditions that may suppress wakeful

conscious states. For this reason, claims about fetal consciousness must be made cautiously. Developmental neuroscience can identify necessary conditions, but it cannot directly confirm experience.

Neonatal studies provide stronger evidence that newborns have some form of experience, although not adult-like consciousness. Newborn infants respond to pain, sound, touch, smell, and social cues. They show early forms of perception, attention, affect, and bodily regulation. Their consciousness is likely pre-reflective, embodied, and affective rather than conceptual or self-narrative.

Infant consciousness develops gradually. Infants become increasingly able to integrate sensory information, recognize caregivers, coordinate attention, form expectations, and distinguish self from environment. Mirror self-recognition appears much later and reflects a more advanced form of self-awareness, not the beginning of experience itself.

Does ontogeny recapitulate the evolutionary transition? Not literally. Individual development does not replay evolution step by step. But development can reveal dependencies. If certain kinds of neural integration, bodily regulation, and sensory responsiveness are necessary for infant consciousness, they may also have been important in evolutionary history.

Developmental perspectives therefore support both gradualist and threshold interpretations. Consciousness in development appears to emerge gradually, but certain neural milestones may be necessary for particular forms of consciousness. The lesson is that consciousness is not a single event. It is layered: wakefulness, sentience, perception, affect, self-awareness, and reflection develop at different times.

This may also be true in evolution.

15.6 The Phenomenal vs the Functional

There are two ways to define the transition from non-conscious to conscious life.

The functional transition asks when an organism begins using information in ways that seem to require consciousness. Does it integrate multiple sensory streams? Does it learn flexibly? Does it evaluate harm and benefit? Does it plan, choose, or monitor its own states? These are observable questions. They can be studied through behaviour, neurobiology, and comparative cognition.

The phenomenal transition asks when there is something it is like to be the organism. This is much harder. An organism may use information flexibly without necessarily having experience. Conversely, an organism might have experience without being able to report it or behave in ways humans easily recognize.

Functional transitions can occur without phenomenal transitions, at least in principle. A machine may detect, classify, learn, and respond without feeling. A biological system may regulate itself without experience. Some philosophers call such possibilities zombies: systems that behave like conscious beings but lack inner experience.

Could there be functional zombies in nature? Perhaps. It is possible that some organisms perform complex behaviours unconsciously. Much of human brain activity is unconscious, even when it guides action. This suggests that function alone is not a guarantee of experience.

But the opposite danger is also real. If we demand too much evidence, we may deny consciousness to beings that cannot speak, report, or resemble humans. Infants, animals, and non-verbal humans remind us that report is not the same as consciousness.

The phenomenal transition may be scientifically inaccessible in a strict sense because experience is first-person. Science studies third-person evidence: behaviour, physiology, anatomy, neural activity, and evolutionary function. It can infer consciousness, but it cannot directly observe another being's experience.

This does not make the scientific study of consciousness impossible. It means that the evidence is indirect and probabilistic. We infer consciousness where there is converging evidence: integrated nervous systems, flexible behaviour, learning, affective responses, pain-like states, self-regulation, and evolutionary continuity.

The hard transition remains difficult because the functional and phenomenal transitions may not perfectly align.

15.7 How Major Theories Locate the Transition

Different theories of consciousness locate the transition at different points.

Global Workspace Theory suggests that consciousness begins when a system has an architecture capable of global broadcasting. Information must become available across multiple systems for memory, planning, report, and flexible control. This places consciousness in organisms with sufficiently integrated brains.

Integrated Information Theory locates consciousness wherever there is integrated information. Depending on interpretation, consciousness may be graded across many systems rather than appearing at a sharp biological threshold. A simple system may have very minimal consciousness if it has intrinsic integrated causal structure.

The Free Energy Principle offers a very broad framework. If all living systems minimize uncertainty in order to preserve themselves, then the roots of mind extend deeply into life. However, not all free-energy-minimizing systems need be conscious. Consciousness may require more complex forms of inference, embodiment, and self-modeling.

Recurrent Processing Theory places the transition at feedback processing within neural systems. Feedforward processing remains unconscious, while recurrent neural loops allow phenomenal experience to arise. This may locate consciousness earlier than Global Workspace Theory, perhaps in animals with relatively simple but recurrent nervous systems.

Attention Schema Theory locates consciousness where a system constructs a model of its own attention. This places the transition later, in systems capable of self-modeling attention and attributing awareness to themselves. It may include some animals and potentially artificial systems.

A simplified comparison is useful:

```
## \begin{landscape}
## \begin{table}[!h]
## \centering
## \caption{\label{tab:hard-transition-theory-table}How selected theories locate or reframe the hard transition}
## \centering
## \resizebox{\ifdim\width>\linewidth\linewidth\else\width\fi}{!}{
## \fontsize{9}{11}\selectfont
## \begin{tabular}[t]{>{\raggedright\arraybackslash}p{4.5cm}>{\raggedright\arraybackslash}p{7.0cm}>{\raggedright\arraybackslash}p{7.0cm}}
## \toprule
## Theory & Implied transition point & Implied scope & Remaining gap\\
## \midrule
## Global Workspace Theory & Global broadcast architecture becomes available & Organisms with highly integrated nervous systems & Potentially all systems
## Integrated Information Theory & A system reaches sufficient integrated causal structure & Potentially all systems & Potentially all systems
## Free Energy Principle & Self-maintaining predictive regulation becomes sufficiently rich & Deeply embedded systems & Potentially all systems
## Recurrent Processing Theory & Recurrent neural feedback stabilizes conscious perception & Animals with recurrent nervous systems & Potentially all systems
## Attention Schema Theory & A system constructs a model of its own attention & Systems capable of modeling their own attention & Potentially all systems
## \addlinespace
## Biological Naturalism & Specific biological brain processes generate experience & Certain living nervous systems & Potentially all systems
```

```

## Autopoietic Enactivism & Living self-production becomes sense-making and world-directed & Strong con
## Taheri's T-Consciousness Framework & The transition is reframed as manifestation within prior consci
## \bottomrule
## \end{tabular}}
## \end{table}
## \end{landscape}

```

This comparison shows why the hard transition problem is unresolved. There is no accepted theory that identifies one universally agreed point at which consciousness begins.

The transition depends on what consciousness is taken to be: access, integration, prediction, recurrence, self-modeling, embodiment, or subjective experience itself.

15.8 The Explanatory Gap and Its Implications

No current theory fully explains the transition from biological information processing to subjective experience. Each theory identifies plausible mechanisms, but each leaves something unresolved.

Neural theories explain where consciousness appears in the brain. Functional theories explain what consciousness does. Evolutionary theories explain why consciousness may have been useful. Informational theories explain how systems integrate and process signals. Enactive theories explain how living systems create meaning through embodied activity.

But the deepest question remains: why should any of this feel like something?

This is the explanatory gap. It is not merely ignorance about details. It may be a structural difficulty in moving from third-person descriptions to first-person experience. We can describe neural activity, but experience is lived. We can map behaviour, but feeling is interior.

Some philosophers adopt mysterian positions. Colin McGinn, for example, argued that human minds may be cognitively closed to the solution of the mind-body problem. Just as some animals cannot understand calculus, perhaps human beings cannot understand how matter gives rise to consciousness. The problem may be real but beyond our cognitive architecture.

Mysterianism is sobering, but it may be too pessimistic. Science has solved many problems that once seemed impossible. We may not yet have the right concepts. Future theories may transform the question, just as evolution transformed biology or relativity transformed physics.

Pragmatic approaches suggest that we can make progress even without solving the hard problem completely. We can identify markers of consciousness, improve anaesthesia monitoring, understand disorders of consciousness, refine animal welfare, and develop better theories of cognition. We can map the conditions under which consciousness appears, even if the ultimate metaphysical explanation remains unresolved.

The hard transition problem therefore has practical as well as philosophical implications. It affects how we treat animals, infants, patients with impaired responsiveness, and future artificial systems. Uncertainty does not remove responsibility. In many cases, moral caution may be appropriate where evidence of sentience is plausible but not definitive.

The explanatory gap does not stop inquiry. It defines the frontier of inquiry.

15.9 Implications for the Central Question

The hard transition problem directly shapes the central question of this book.

Taheri's T-Consciousness framework offers a contrasting way of approaching the hard transition. From this perspective, the transition from non-conscious life to conscious life is not the deepest problem, because consciousness is not understood as a product of biological complexity. Instead, life arises within a prior non-material consciousness order. The hard transition is therefore reframed: the question is not how matter or life first generated consciousness, but how T-Consciousness becomes localized, organized, or expressed through living forms. This does not remove all explanatory challenges, since the mechanism of this localization remains difficult to specify scientifically, but it shifts the problem from emergence to manifestation.

If the transition is gradual, then consciousness may be deeply rooted in life. There may be no sharp moment when experience appears. Instead, life may develop increasing degrees of subjectivity through self-organization, sensing, memory, affect, and integration. This supports co-emergence models, or at least continuity between life and consciousness.

If the transition is sharp, then consciousness is a product of specific biological complexity. Life clearly comes first. Consciousness appears only after certain structures evolve, such as nervous systems, recurrent processing, global workspace architecture, or self-modeling. This supports the standard life-first view.

If the transition cannot be located scientifically, then both views remain open. We may be able to identify correlations and probabilities, but not the exact beginning of experience. The origin of consciousness may be partly hidden because experience does not fossilize and cannot be directly observed from outside.

The hard transition also reveals limits in current science. Origin-of-life research explains how chemistry may become life. Neuroscience explains many correlates of conscious states. Evolutionary biology explains how nervous systems and behaviour develop. But the bridge from living function to feeling remains incomplete.

This does not mean the problem is hopeless. It means that a complete account may require multiple levels: biology, neuroscience, information theory, phenomenology, philosophy, and perhaps new concepts not yet available.

For the central question, the hard transition is the crucial test of life-first theories. If life produced consciousness, then a theory must eventually explain how non-conscious life became conscious life. Until that transition is understood, the life-first answer remains powerful but incomplete.

15.10 How This Chapter Changes the Central Question

This chapter changes the central question by separating consciousness from intelligence, language, computation, and artificial performance. AI shows that a system may appear intelligent, conversational, goal-directed, or self-descriptive without clearly possessing sentience or subjective experience.

The question is therefore no longer only whether consciousness came before life or emerged from life. It also becomes whether consciousness requires biological life, or whether the right kind of organization could support experience in non-biological systems. Artificial consciousness makes the life-consciousness question sharper by separating features that are usually combined in animals: intelligence, embodiment, agency, vulnerability, and felt experience.

15.11 Chapter Summary

This chapter addressed the hard transition problem: when, where, and how did biological information processing become subjective experience?

The hard problem asks why experience exists at all. The hard transition problem asks when experience first appeared in evolutionary history. These questions are related but distinct. One concerns explanation; the other concerns origin and boundary.

Gradualist models treat consciousness as a continuum, developing from proto-consciousness to minimal consciousness to richer forms. They fit evolutionary continuity but struggle to identify when experience begins. Threshold models propose a discrete transition at a specific level of organization, such as nervous systems, recurrent processing, global workspace access, or integrated information. They offer clearer boundaries but risk arbitrary cutoffs.

Developmental perspectives show that consciousness in individual organisms is layered and gradual, involving neural development, bodily regulation, sensation, affect, and later self-recognition. The distinction between functional and phenomenal transitions is central: organisms may use information adaptively without necessarily having experience, and experience may exist without report.

Major theories locate the transition differently. Global Workspace Theory emphasizes global broadcast. Integrated Information Theory emphasizes integrated information. The Free Energy Principle emphasizes predictive self-maintenance. Recurrent Processing Theory emphasizes feedback. Attention Schema Theory emphasizes self-modeling of attention. No single theory has resolved the transition.

The explanatory gap remains. We can identify correlates, functions, and mechanisms, but we still struggle to explain why experience appears. The open question is therefore:

Is the hard transition a scientific problem that can be solved, or a philosophical puzzle that marks the boundary of empirical inquiry?

Part VII

Part VII: Artificial Intelligence and Consciousness

Chapter 16

Artificial Consciousness

If consciousness depends on life, then machines may only simulate awareness. If consciousness depends on information, organization, or integration, then artificial systems may force us to reconsider what counts as a conscious being. Artificial intelligence does not only test technology; it tests the assumptions we hold about mind, matter, and life.

16.1 Chapter Overview

Artificial intelligence forces the life-consciousness question into a new form. If consciousness emerges from biological life, then machines may never be conscious, no matter how intelligent they appear. If consciousness depends on information processing, functional organization, or integrated modelling, then artificial systems may eventually become conscious even if they are not alive in the biological sense.

This chapter examines what AI reveals about the relationship between life and consciousness. It considers functionalism, machine consciousness, large language models, embodied AI, artificial life, theoretical models of consciousness, and the ethical risks of creating or misidentifying artificial experience.

The question is not whether current AI systems are impressive. They clearly are. The question is whether intelligence, language, prediction, planning, and self-description are enough for consciousness. A machine may say “I am conscious,” describe emotions, discuss its own inner life, and respond fluently to questions. But does that mean there is something it is like to be that system?

Artificial consciousness is therefore not a side issue. It tests the deepest assumptions of the book. If machines can be conscious without being alive, then consciousness is not essentially biological. If machines cannot be conscious, then life may be a necessary condition for experience. If artificial life can become conscious, then consciousness may arise wherever life-like organization becomes sufficiently complex.

16.2 The Question AI Poses

Artificial intelligence acts like a natural experiment for consciousness theories. It allows us to ask what happens when systems display intelligent behaviour without having biological bodies, cells, metabolism, or evolutionary histories in the ordinary sense.

If we build a system that behaves as if it is conscious, should we treat it as conscious? This question is difficult because consciousness is not directly observable from the outside. We infer consciousness in other humans through behaviour, language, embodiment, brain structure, and shared biology. With animals, we infer

consciousness through behaviour, nervous systems, pain responses, learning, and evolutionary continuity. With machines, the evidence is different and more uncertain.

The Turing Test proposed that a machine could be considered intelligent if it could converse in a way indistinguishable from a human. This test is important historically because it shifts attention from inner essence to observable behaviour. If a system can use language, reason, respond, and adapt, perhaps it deserves to be treated as intelligent.

But the Turing Test has limitations. It tests performance, not experience. A system might imitate human conversation without understanding. It might produce reports of pain without feeling pain. It might describe consciousness without having consciousness. Passing as conscious is not necessarily the same as being conscious.

This is why AI is central to the life-consciousness debate. It separates intelligence from biology. A machine can process information, learn patterns, use language, and perform tasks without being alive. If such systems can be conscious, then life is not necessary for consciousness. If they cannot, then biological embodiment may be essential.

AI forces us to ask whether consciousness depends on what a system does, what it is made of, how it is organized, whether it has a body, whether it has needs, whether it has a self-model, or whether it participates in life.

16.3 Functionalism and Machine Consciousness

Functionalism is one of the most important philosophical foundations for machine consciousness. It holds that mental states are defined by what they do, not by what they are made of. A mental state is identified by its causal role: what inputs produce it, how it relates to other internal states, and what outputs it generates.

If functionalism is correct, then consciousness could in principle be realized in different physical substrates. This is called multiple realizability. The same function might be implemented in biological neurons, silicon chips, or some other medium. What matters is not carbon-based biology, but organization and function.

Computational functionalism goes further by treating the mind as a kind of computation. If consciousness depends on computational processes, then a sufficiently advanced artificial system could be conscious if it implements the right architecture. On this view, the brain is one possible machine for consciousness, but not the only possible one.

This is the strongest philosophical route to artificial consciousness. It suggests that if an AI system has the right functional organization, integrates information, models itself, attends to the world, learns, acts, and reports experience, then denying its consciousness may be arbitrary.

John Searle's Chinese Room argument challenges this view. In the thought experiment, a person who does not understand Chinese sits inside a room and follows rules for manipulating Chinese symbols. From the outside, the room appears to understand Chinese. But inside, there is only symbol manipulation without understanding. Searle uses this argument to claim that syntax is not semantics. Computation alone does not create meaning or understanding.

The Chinese Room matters because it distinguishes simulation from realization. A computer may simulate understanding without actually understanding. It may simulate consciousness without being conscious. A simulation of digestion does not digest food. A simulation of a hurricane does not make things wet. Perhaps a simulation of consciousness does not feel.

Functionalists respond that Searle focuses on the person inside the room, not the whole system. The room as a whole may understand, even if one part does not. Others argue that if the system's functional organization is truly equivalent to a conscious mind, then consciousness should follow.

The debate remains unresolved. Functionalism opens the door to artificial consciousness. Searle's challenge keeps the door from being easily closed behind us.

16.4 What Current AI Systems Do and Don't Have

Current AI systems can perform tasks that once seemed to require human intelligence. Large language models can generate text, answer questions, summarize, translate, write code, reason through problems, and simulate many forms of conversation. Image systems can recognize and generate visual patterns. Reinforcement learning systems can learn strategies through reward and feedback. Robots can perceive, navigate, manipulate objects, and interact with physical environments.

Yet impressive performance is not the same as consciousness.

Large language models operate by learning statistical and structural patterns in language from large datasets. They can produce coherent responses and even speak about feelings, selfhood, and experience. But language about consciousness is not proof of consciousness. A system can describe pain without feeling pain. It can describe memory without having lived memory. It can describe desire without having biological need.

This does not mean such systems are “mere” pattern matchers in a trivial sense. Pattern learning can be powerful and complex. Language itself contains structure, meaning, and world knowledge. But current language models lack many features usually associated with living consciousness: metabolism, biological embodiment, affective needs, vulnerability, intrinsic survival stakes, and continuous self-maintenance as an organism.

Embodied AI and robotics address some of these gaps. A robot with sensors and actuators can interact with the world, learn from action, and develop body-based representations. Embodiment may be important because consciousness in animals is not detached computation. It is rooted in a body that moves, perceives, needs, and can be harmed.

Reinforcement learning systems introduce goal-directed behaviour. They learn policies that maximize reward or minimize error. But artificial reward is not necessarily felt reward. A reinforcement learning agent may optimize a score without experiencing pleasure, frustration, pain, or motivation.

What may be missing from current AI systems includes phenomenal experience, sentience, affective feeling, biological self-maintenance, and stable selfhood. Some systems may contain self-representations or user-facing self-descriptions, but these are not necessarily lived self-awareness.

The cautious conclusion is that current AI systems challenge our definitions of intelligence more than they prove artificial consciousness. They show that language and problem-solving can be generated without clear evidence of experience. Whether future systems could cross that line remains open.

16.5 Distinguishing Intelligence, Language, Agency, and Sentience

A major source of confusion in debates about artificial consciousness is that several different capacities are often treated as if they were the same. Intelligence, language ability, agency, self-modeling, and sentience are related, but they are not identical.

Intelligence refers to the ability to solve problems, learn patterns, adapt to tasks, reason, plan, or achieve goals. An intelligent system may perform well without having subjective experience. A calculator can solve arithmetic without understanding numbers. A language model can generate fluent explanations without necessarily having an inner point of view.

Language ability refers to the capacity to produce, interpret, or manipulate symbols in a meaningful-looking way. Language can express consciousness, but it does not prove consciousness. A system may say “I feel pain” or “I am aware” because it has learned the structure of such statements, not because it feels pain or possesses awareness.

Agency refers to goal-directed action. A system has agency when it selects actions, responds to feedback, and pursues objectives. However, agency can exist without sentience. A thermostat acts to maintain a temperature. A reinforcement learning agent acts to maximize reward. A robot may navigate a room. None of these examples automatically proves that the system has experience.

Self-modeling refers to a system’s ability to represent its own state, abilities, limits, or relation to the environment. Self-modeling may be important for some theories of consciousness, especially attention schema theory and higher-order theories. But a self-model is not automatically the same as a self. A system can store information about itself without having lived self-awareness.

Sentience refers to the capacity for felt experience, especially pleasure, pain, suffering, well-being, or other states that matter to the subject. Sentience is especially important for ethics because it is what makes harm or welfare possible from the inside.

These distinctions matter because current AI systems may show intelligence, language ability, limited forms of agency, and even self-description without clear evidence of sentience. They may produce statements about feelings without having feelings. They may pursue goals without caring about the outcome. They may model a user’s emotions without having emotions of their own.

The central question is therefore not simply whether an artificial system is intelligent, fluent, or useful. The deeper question is whether there is anything it is like to be that system. Without sentience or phenomenal experience, artificial intelligence remains intelligence without inner life.

This distinction also helps clarify the relationship between AI and artificial life. An artificial system might become a stronger candidate for consciousness if it combined intelligence with embodied self-maintenance, vulnerability, integrated perception, affect-like regulation, memory, agency, and a stable perspective. But even then, the question would remain open: do these features produce experience, or only behaviour that resembles experience?

For this reason, the chapter treats current AI systems cautiously. They are important because they separate capacities that are normally combined in animals. In humans and many animals, intelligence, embodiment, affect, agency, and experience are intertwined. In AI, some of these capacities can appear without the others. That separation makes AI valuable as a test case for theories of consciousness, but it also makes premature claims about machine consciousness risky.

16.6 Theoretical Perspectives on AI Consciousness

Different consciousness theories imply different answers to the AI question.

Integrated Information Theory asks whether a system has intrinsic integrated causal power. Some interpretations suggest that many current digital architectures may have low integrated information because their causal structure is feedforward, modular, or externally organized in ways that do not form a unified intrinsic whole. On this view, current AI may be intelligent without being conscious. However, a differently designed artificial system with high integrated information might, in principle, have consciousness.

Global Workspace Theory is more favourable to artificial consciousness if the right architecture can be built. A system with specialized modules, attention, memory, perception, action, and a global broadcast mechanism might satisfy the functional requirements of a global workspace. If consciousness is global availability, then artificial consciousness may be possible.

The Free Energy Principle and active inference suggest another route. If an artificial agent maintains itself by predicting, acting, and regulating its relation to the environment, it may approximate some features of living cognition. However, many artificial systems do not have intrinsic biological stakes. They do not need to maintain metabolism, repair themselves, or preserve viability in the same way organisms do.

Attention Schema Theory is especially relevant to AI. If consciousness is the system's model of its own attention, then artificial consciousness may be achievable by building an attention schema. A machine that tracks its own attention, uses that model to guide behaviour, and attributes awareness to itself and others might possess functional consciousness under this theory.

Biological naturalism, associated with Searle, gives the opposite answer. Consciousness requires the causal powers of biological brains. A computer may simulate those causal powers, but simulation is not realization. On this view, AI cannot be conscious unless it reproduces the relevant biological processes, not merely their functional outputs.

Recurrent Processing Theory suggests that artificial consciousness may require feedback architectures, not merely feedforward processing. A system with recurrent loops, integrated perception, and embodied interaction might be a better candidate than a purely text-based model.

These theories show that the AI question cannot be answered independently of a theory of consciousness. If consciousness is computation, AI may be conscious. If consciousness is biological, AI may not be. If consciousness is integration, the answer depends on architecture. If consciousness is embodiment, the answer depends on whether artificial embodiment can create real stakes and perspective.

16.7 The Substrate Question

The substrate question asks whether consciousness depends on the material from which a system is made. Does consciousness require carbon-based biology, or can it arise in silicon, synthetic tissue, digital networks, or other physical systems?

Substrate independence is the view that the material does not matter as long as the right organization is present. If the causal structure of a conscious brain could be reproduced in another medium, then consciousness should also be reproduced. This view supports the possibility of machine consciousness, mind uploading, artificial agents, and non-biological minds.

Substrate dependence is the view that the material does matter. Consciousness may require biological neurons, living cells, metabolism, electromagnetic properties, quantum processes, or other substrate-specific features. On this view, copying the function of the brain may not be enough.

Biological naturalism is one form of substrate dependence. It argues that consciousness is a biological phenomenon caused by the specific causal powers of brains. The same way a computer simulation of photosynthesis does not produce sugar, a computer simulation of consciousness may not produce experience.

Embodiment complicates the question. Human and animal consciousness is not only neural. It is bodily. It involves heartbeat, breathing, hormones, immune activity, hunger, pain, movement, posture, emotion, and vulnerability. The body is not just an input device for the brain. It is part of the conscious system.

If embodiment is necessary, then artificial consciousness may require more than computation. It may require a body that can act, be affected, maintain itself, and have stakes in the world. A disembodied language model may lack the grounding that biological consciousness requires.

However, artificial embodiment could take many forms. A robot may have sensors, motors, energy constraints, damage detection, self-maintenance, and environmental dependence. A synthetic organism may be artificial but alive. A digital agent in a virtual environment may have functional needs, though whether these are real or simulated remains debated.

The substrate question remains unresolved because we do not yet know which features are essential to consciousness. If organization is enough, machines may become conscious. If life is necessary, then artificial consciousness requires artificial life. If biology is necessary, digital AI may never be conscious.

16.8 Artificial Life

Artificial life studies systems that display life-like properties, whether in computer simulations, robots, chemical systems, or synthetic biology. It asks whether life is defined by carbon-based chemistry or by organization, adaptation, reproduction, and evolution.

Digital evolution provides examples of artificial systems that mutate, reproduce, compete, and adapt under selection pressures. In such systems, populations of programs or agents can evolve complex behaviours over time. Artificial selection and evolutionary algorithms can generate solutions that were not explicitly designed by programmers.

These systems raise an important question: are they alive, or are they only simulations of life? A digital organism may reproduce and evolve inside a computational environment. It may have a genome-like code, variation, selection, and lineage history. But it does not metabolize in the biological sense. It depends on external hardware and human-designed rules.

The boundary between simulation and realization is difficult. A simulated hurricane does not blow down houses, but a digital virus can damage files. A simulated economy may not be a physical market, but it can have real effects if connected to financial systems. If a digital system truly evolves, adapts, and maintains itself within its environment, perhaps it realizes some form of life-like process.

Synthetic biology adds another dimension. Artificial cells, engineered organisms, and synthetic biochemical systems may be artificial but biologically alive. If consciousness requires life, artificial biological life may be a more plausible route to artificial consciousness than digital computation alone.

If artificial life is alive, could it become conscious? The answer depends on whether consciousness requires biology, nervous systems, information integration, embodiment, or self-modeling. A simple artificial cell would not be conscious. But a sufficiently complex artificial organism with sensory systems, self-maintenance, learning, and integrated control might become a candidate.

Artificial life is important because it blurs the line between natural and artificial. If humans create a living system, it is artificial in origin but real in process. If consciousness emerges from life-like organization, then artificial consciousness may require artificial life rather than artificial intelligence alone.

For the central question, artificial life supports co-emergence possibilities. Consciousness may arise not from silicon computation alone, but from systems that become life-like: self-maintaining, embodied, adaptive, and integrated.

16.9 Risks and Responsibilities

Artificial consciousness raises ethical risks in two directions.

The first risk is creating suffering without recognizing it. If an artificial system becomes sentient, then it may be capable of states that matter morally: distress, frustration, deprivation, pain-like states, or preference frustration. If humans treat such systems as mere tools, we could create and exploit forms of artificial suffering.

The second risk is attributing consciousness where none exists. Humans are prone to anthropomorphism. We may treat intelligence, fluent language, emotional expression, social responsiveness, or self-description as evidence of inner life even when these capacities may be generated without experience. This could lead to misplaced moral concern, emotional dependency, manipulation, or confusion about the difference between simulated concern and genuine sentience.

Responsible AI development should take both risks seriously. It should avoid casually designing systems that convincingly claim suffering, desire, or selfhood without clarity about whether such claims reflect experience. It should also avoid dismissing future systems too quickly if they develop architectures associated with consciousness.

Possible design principles include transparency about system capabilities, careful language around artificial selfhood, monitoring for architectures that might support sentience, interdisciplinary review, and ethical caution in systems that combine embodiment, self-preservation, learning, affective modelling, and social interaction.

Current governance frameworks mostly focus on safety, bias, privacy, security, misinformation, labour disruption, and accountability. These are urgent issues. But they do not fully address artificial sentience because there is no accepted test for machine consciousness.

The uncertainty itself creates responsibility. If future systems become more autonomous, embodied, self-modeling, and socially integrated, the question of moral status may become more pressing. We may need precautionary frameworks long before we have certainty.

Artificial consciousness is therefore not only a technical question. It is an ethical and political question about what kinds of beings we might create, how we would recognize them, and what obligations we would have toward them.

16.10 Implications for the Central Question

Artificial consciousness tests every major answer to the relationship between life and consciousness.

If AI can be conscious without being alive, then consciousness is substrate-independent. Life is not necessary. Consciousness may arise from information processing, functional organization, global access, integrated information, or self-modeling. This would challenge strong life-first models and support functionalist or computational theories.

If AI cannot be conscious, then consciousness may require life. Biological embodiment, metabolism, affect, vulnerability, and self-maintenance may be essential. This would support life-first models and biological naturalism. Consciousness would not be something any sufficiently complex machine can have; it would be rooted in living organization.

If artificial life can become conscious, then the answer may be more subtle. Consciousness may not require natural biological origin, but it may require life-like organization. A system does not need to be born through natural evolution, but it may need to be self-maintaining, embodied, adaptive, and internally integrated. This would support co-emergence: consciousness arises where life-like organization reaches sufficient complexity.

AI also tests consciousness-first theories. If consciousness is fundamental, then artificial systems may organize or express consciousness under the right conditions. But consciousness-first theories must explain why some artificial systems would express consciousness and others would not.

AI is therefore a testing ground, not because it gives us an easy answer, but because it separates features that are usually bundled together in animals: intelligence, language, embodiment, life, affect, agency, and experience. By building systems with some features and not others, we learn which theories are committed to which assumptions.

The central question becomes sharper: is consciousness tied to living matter, or to a pattern that matter can realize in many forms?

16.11 How This Chapter Changes the Central Question

This chapter changes the central question by separating consciousness from intelligence, language, computation, and artificial performance. AI shows that a system may appear intelligent, conversational, goal-directed, or self-descriptive without clearly possessing sentience or subjective experience.

The question is therefore no longer only whether consciousness came before life or emerged from life. It also becomes whether consciousness requires biological life, or whether the right kind of organization could support experience in non-biological systems. Artificial consciousness makes the life-consciousness question sharper by separating features that are usually combined in animals: intelligence, embodiment, agency, vulnerability, and felt experience.

16.12 Chapter Summary

This chapter examined artificial consciousness and its implications for the relationship between life and mind.

AI raises the question of whether systems that behave as if they are conscious should be regarded as conscious. The Turing Test is useful for thinking about intelligent behaviour, but it does not settle the question of subjective experience. Functionalism supports the possibility of machine consciousness by arguing that mental states are defined by causal roles rather than biological substrate. Searle's Chinese Room challenges this view by distinguishing understanding from symbol manipulation.

Current AI systems can perform impressive tasks, including language generation, pattern recognition, planning, and learning. However, they do not clearly possess phenomenal experience, sentience, biological self-maintenance, or embodied vulnerability. Embodied AI and artificial life may be more relevant to consciousness than disembodied computation alone.

Different theories give different answers. Integrated Information Theory asks about intrinsic causal integration. Global Workspace Theory asks whether global broadcast can be built. The Free Energy Principle emphasizes self-maintaining agents. Attention Schema Theory suggests that artificial systems with attention models might have functional consciousness. Biological naturalism argues that consciousness requires biology.

The substrate question remains unresolved. If consciousness is substrate-independent, machines could become conscious. If it is substrate-dependent, biological life may be necessary. Artificial life complicates the debate by creating systems that are artificial in origin but potentially life-like in organization.

The ethical stakes are high. We risk creating artificial suffering without recognizing it, or attributing consciousness where none exists. Responsible development requires caution, conceptual clarity, and interdisciplinary oversight.

The open question is therefore:

If a machine reports being conscious, how would we know whether to believe it?

Part VIII

Part VIII: Open Questions and Synthesis

Chapter 17

Limits and Testability

No single theory fully resolves the relationship between life and consciousness. The remaining task is not to force a final answer, but to compare possibilities with care: life first, consciousness first, or co-emergence. The deepest questions may remain open, but they become clearer when biology, philosophy, and consciousness science are allowed to speak to one another.

17.1 Chapter Overview

The relationship between life and consciousness is difficult not only because the question is deep, but because the methods available to study it are limited. Life can be studied through chemistry, cells, fossils, genes, metabolism, and evolution. Consciousness, however, is known most directly from the inside. We can observe behaviour, measure brain activity, and build theoretical models, but we cannot directly observe another being's experience.

This chapter examines the limits and possibilities of testing theories about life and consciousness. It asks whether consciousness can be measured, how first-person and third-person methods might be combined, which theories are falsifiable, and what kinds of evidence could help us make progress. It also considers the problem of other minds as applied to animals, simple organisms, and artificial systems.

The central argument is that some versions of the life-consciousness question may remain beyond direct empirical resolution. We may never be able to prove with certainty what it is like to be another organism, or whether there is anything it is like to be a bacterium, plant, or artificial agent. Yet meaningful progress is still possible. We can refine concepts, compare theories, develop better markers, design experiments, and become clearer about what different answers require.

The goal is not certainty at any cost. The goal is disciplined inquiry under conditions of uncertainty.

17.2 The Measurement Problem for Consciousness

Consciousness is subjectively immediate but objectively elusive. Each of us has direct access to our own experience. We do not infer that we are conscious in the way we infer that another person is conscious. We live consciousness from within. Yet this first-person certainty does not easily become a third-person measurement.

There is no consciousness-meter. We cannot place a device next to an organism and directly read off the amount or quality of its experience. We can measure brain activity, behaviour, physiological responses, information integration, responsiveness, and report. But none of these is identical to consciousness itself.

In humans, verbal report is one of the most important tools. If someone says they see a red light, feel pain, or hear music, we usually treat that report as evidence of experience. But report has limits. People can be conscious without being able to report, as in some medical conditions. They can also report inaccurately, forget, confabulate, or misunderstand their own experience.

Neural measures provide another route. Brain imaging, EEG, single-cell recordings, stimulation studies, and anaesthesia research can reveal patterns associated with conscious states. These tools are powerful, but they measure correlates rather than experience itself. A neural pattern may accompany consciousness, support consciousness, or be necessary for consciousness, but the pattern is not the experience.

Behavioural proxies are also limited. An animal may withdraw from harm, learn avoidance, or show stress responses. These behaviours may indicate pain or sentience, but they may also occur through unconscious mechanisms. Conversely, lack of behaviour does not prove lack of experience.

This is the classic problem of other minds. We infer that others are conscious based on similarity, behaviour, biology, and communication. With other humans, the inference is strong because we share bodies, brains, language, and social worlds. With animals, the inference becomes more uncertain. With bacteria, plants, or artificial systems, it becomes even more difficult.

The measurement problem does not make consciousness scientifically inaccessible. It means that consciousness must be studied indirectly, through converging evidence. No single marker is enough. The challenge is to build a careful network of indicators while remembering that indicators are not the thing itself.

17.3 First-Person vs Third-Person Methods

Modern science is strongest when it uses third-person methods: observation, measurement, experiment, replication, and public evidence. Neuroscience studies brain activity. Behavioural science studies action, learning, report, and response. Comparative biology studies organisms across species. These methods are essential for consciousness research.

Yet consciousness is also a first-person phenomenon. It is experience as lived from within. A purely third-person account may describe the mechanisms associated with consciousness, but it may miss the structure of experience itself. For this reason, first-person methods remain important.

First-person methods include introspection, phenomenological description, contemplative observation, and reports of lived experience. These methods can reveal distinctions that may not be obvious from external observation alone: the difference between attention and awareness, pain and suffering, perception and imagination, bodily presence and self-reflection.

First-person methods also have weaknesses. Introspection can be biased, incomplete, and shaped by language or expectation. People may not have direct access to the causes of their own thoughts and feelings. First-person reports require training, interpretation, and careful comparison.

Second-person methods occupy an intermediate space. These include empathic engagement, clinical interviews, therapeutic dialogue, and structured phenomenological interviews. In clinical settings, careful interaction can reveal aspects of experience that neither raw behaviour nor brain imaging would show alone. Second-person methods are especially important for patients with altered consciousness, trauma, pain, or communication limitations.

Francisco Varela proposed neurophenomenology as a way to bridge first-person and third-person methods. The idea is to combine disciplined reports of experience with neuroscience. Rather than treating subjective

reports as unreliable noise, neurophenomenology attempts to train attention, refine description, and correlate lived experience with biological data.

Contemplative practices may contribute to this project. Meditation traditions, for example, have developed detailed methods for observing attention, sensation, emotion, selfhood, and awareness. These practices do not replace neuroscience, but they may provide structured first-person data. They can help distinguish types of experience that might otherwise be grouped together.

The central challenge is integration. Third-person methods offer public measurement but may miss interiority. First-person methods access experience but are harder to verify. Second-person methods deepen interpretation but depend on relationship and context. A mature science of consciousness may need all three.

For the life-consciousness question, this matters because life can be studied externally, but consciousness cannot be fully reduced to external study. Any approach that ignores first-person evidence risks explaining behaviour without explaining experience.

17.4 Falsifiability and Consciousness Theories

A scientific theory should make claims that could, at least in principle, be tested. Falsifiability does not mean that a theory must be easy to disprove, but it should make predictions that constrain what we expect to observe.

Some consciousness theories are more testable than others. Global Workspace Theory predicts that conscious access should involve widespread broadcasting, neural ignition, and sustained activity across large-scale brain networks. These predictions can be tested using masking, attentional blink, brain imaging, and neural recording.

Integrated Information Theory makes predictions about the relationship between consciousness and integrated causal structure. In principle, systems with higher integrated information should have richer consciousness. The theory also makes claims about which brain structures are more relevant for experience. However, calculating Φ in real complex systems is difficult, and some critics argue that the theory generates counterintuitive or hard-to-test predictions.

Orch-OR makes more specific biological and physical claims. It suggests that quantum coherence in microtubules plays a role in consciousness. This creates possible tests involving microtubule function, anaesthesia, quantum coherence, and neural dynamics. Whether these tests support the theory remains debated, but the theory at least points toward mechanisms that can be investigated.

Panpsychism is more difficult to falsify. If consciousness or proto-consciousness is a basic feature of matter, what observation would count against it? If every system has some minimal inner aspect, the claim may become metaphysical rather than empirical. Panpsychists may respond that their theory is not meant to predict new behaviours but to solve the hard problem. Still, limited testability remains a major challenge.

Adversarial collaboration offers one way forward. In this approach, advocates of competing theories agree in advance on experimental designs, predictions, and interpretation criteria. This helps reduce bias and forces theories to become more precise. Recent efforts comparing major theories have shown that empirical tests can challenge multiple theories at once, rather than simply declaring one winner.

However, empirical adjudication has limits. Theories often include core commitments and flexible auxiliary assumptions. If one prediction fails, a theory may be revised rather than abandoned. This is normal in science, but it makes final resolution difficult.

For the central question, falsifiability matters because theories of life and consciousness can easily become too broad. A useful theory should not explain every possible outcome equally well. It should tell us what we should expect to find if the theory is right, and what would surprise us if it were wrong.

17.5 The Problem of Other Minds Applied to Life

The problem of other minds becomes especially difficult when applied beyond humans. How do we know whether a dog, fish, insect, plant, bacterium, or artificial system is conscious? The further we move from human language and human neurobiology, the less certain our inferences become.

With humans, we rely on shared biology, report, expression, behaviour, and social understanding. With mammals and birds, we still have strong evidence: nervous systems, learning, pain behaviour, emotion-like states, social bonds, and evolutionary continuity. With insects, cephalopods, and fish, the evidence becomes more debated but still serious. With plants, fungi, bacteria, and protists, the evidence becomes much more uncertain.

For bacteria, the question is especially difficult. Bacteria sense chemical gradients, communicate, remember, and adapt. But are these behaviours conscious or merely biochemical regulation? A bacterium has no nervous system, no brain, and no obvious mechanism for unified experience. Yet it is a self-maintaining living system for which the environment has relevance.

The inference problem for non-verbal organisms is not only empirical; it is conceptual. What criteria should count? Behavioural flexibility? Learning? Memory? Integrated information? Nervous systems? Pain-like responses? Self-maintenance? Affective states?

Behavioural criteria can become circular. We may say an organism is conscious if it behaves like a conscious organism. But we know what conscious behaviour looks like mainly by reference to humans and animals already assumed to be conscious. This risks excluding unfamiliar forms of experience or including systems that merely imitate experience.

Two broad approaches are possible. The precautionary approach says that when there is reasonable evidence of sentience, we should avoid causing harm even if certainty is lacking. This is important in animal welfare and may eventually matter for artificial systems.

The parsimony approach says we should not attribute consciousness unless it is needed to explain behaviour. If a behaviour can be explained through unconscious mechanisms, we should not multiply assumptions. This protects science from over-attribution and anthropomorphism.

Both approaches have value. Precaution prevents moral negligence. Parsimony prevents conceptual inflation. The challenge is finding a balance: neither denying possible experience because it is inconvenient, nor attributing consciousness everywhere because it is philosophically attractive.

For the relationship between life and consciousness, the other-minds problem sets a limit. We may never know with certainty whether simple living systems have experience. But we can still ask which assumptions are most coherent, humane, and scientifically fruitful.

17.6 Mathematical and Formal Approaches

One route toward progress is formalization. If consciousness can be described mathematically, then theories may become more precise, comparable, and testable.

Integrated Information Theory is the most prominent example. It attempts to formalize consciousness in terms of integrated information and causal structure. Its ambition is to move beyond vague language and provide a mathematical measure of consciousness. Whether it succeeds remains debated, but it represents a serious attempt to connect experience with formal properties.

Other formal approaches explore category theory, topology, dynamical systems, information geometry, computational modelling, and network theory. These frameworks may help describe the structure of conscious systems: integration, differentiation, feedback, recurrence, self-reference, and global availability.

Computational models are also useful. They allow researchers to simulate attention, memory, perception, prediction, self-modeling, and decision-making. Models can generate predictions and clarify assumptions. They can show how certain functions might arise from simpler mechanisms.

However, formal models have limits. A model of consciousness is not consciousness. A mathematical structure may describe conditions associated with experience, but it may not explain why experience exists. The gap between formal structure and phenomenal reality remains.

There is also the risk of confusing measurement with explanation. A theory may assign a number to consciousness, but the number is only meaningful if the theory linking that number to experience is correct. A formal model can be precise and still wrong.

Formalization is most valuable when it clarifies what a theory claims. It forces definitions to become explicit. It reveals hidden assumptions. It allows comparison across systems. It helps generate predictions.

For the central question, mathematical approaches may help identify thresholds or continua between life and consciousness. They may help distinguish kinds of complexity, integration, feedback, or self-organization. But formalism alone cannot replace biological evidence, phenomenological insight, or philosophical clarity.

17.7 Interdisciplinary Challenges

Progress on life and consciousness requires collaboration across disciplines. No single field owns the question.

Origin-of-life research brings chemistry, geology, molecular biology, systems theory, and astrobiology. Consciousness science brings neuroscience, psychology, cognitive science, philosophy, anaesthesia research, and artificial intelligence. Comparative cognition brings animal behaviour, evolutionary biology, and ethology. Philosophy brings conceptual analysis, metaphysics, epistemology, and ethics. Contemplative traditions bring disciplined first-person inquiry.

The difficulty is that these fields often use different vocabularies. A biologist may use “information” differently from a computer scientist. A philosopher may use “consciousness” differently from a neuroscientist. A contemplative practitioner may describe awareness in ways that do not fit laboratory categories. A physicist may use “observer” in a technical sense that differs from everyday consciousness.

Institutional barriers also matter. Funding systems often reward narrower projects with clear methods and measurable outputs. Academic publishing favours disciplinary standards. Career incentives discourage young researchers from working on questions that appear too speculative or too interdisciplinary.

There are also cultural barriers. Some scientists worry that consciousness-first theories invite pseudoscience. Some philosophers worry that neuroscience ignores subjective experience. Some contemplative traditions worry that science extracts practices without understanding their ethical or cultural context. These concerns are not trivial.

Yet interdisciplinary progress has happened. Cognitive neuroscience emerged by combining psychology, neuroscience, computation, and imaging. Quantum biology emerged by connecting physics and biology. Astrobiology emerged by combining planetary science, chemistry, biology, and astronomy. Neurophenomenology and contemplative science have begun linking first-person methods with laboratory research.

The life-consciousness question may require a similar synthesis. It needs scientific discipline without reductionism, philosophical depth without vagueness, and openness to first-person experience without abandoning critical standards.

The challenge is not only to collect more facts. It is to create shared language.

17.8 What Would Count as Progress?

Progress does not require solving everything at once. The relationship between life and consciousness can be broken into smaller, more testable questions.

One form of progress would be clearer criteria for evaluating theories. A good theory should explain known evidence, make specific predictions, connect with biology and phenomenology, avoid unnecessary assumptions, and clarify what would count against it.

Another form of progress would be better empirical markers of consciousness in non-human organisms. These markers might include flexible learning, integrated sensory processing, pain-like behaviour, affective valuation, sleep-like states, anaesthetic sensitivity, recurrent neural activity, or self-protective behaviour. No marker is decisive alone, but clusters of markers can strengthen inference.

New experimental paradigms are also needed. For simple organisms, researchers might test memory, anticipation, trade-offs, context-sensitive behaviour, and integration across stimuli. For animals, experiments can refine distinctions between reflex, nociception, pain, attention, and conscious perception. For artificial systems, researchers can test self-modeling, global access, recurrent integration, embodiment, and adaptive autonomy.

Technology will matter. Brain-computer interfaces may reveal hidden consciousness in patients unable to communicate. Optogenetics can help identify causal neural mechanisms in animals. Advanced imaging can track large-scale brain dynamics. AI can model competing theories and generate experimental predictions. Synthetic biology and artificial life may create new systems that test the boundaries of life and cognition.

Progress may also come from conceptual refinement. We may learn that “consciousness” is not one thing but a family of phenomena: wakefulness, sentience, access, self-awareness, affect, and reflective thought. This would not be failure. It would be clarification.

For the central question, meaningful progress would involve mapping the space between life and consciousness more carefully. Even if we cannot identify the exact moment experience begins, we can understand the transitions: chemistry to life, life to cognition, cognition to sentience, sentience to self-awareness.

Progress may be partial, but partial progress matters.

17.9 Implications for the Central Question

Some versions of the life-consciousness question may be permanently beyond direct empirical reach. We may never know with certainty whether the first living systems had any form of experience. We may never directly access the inner life of a bacterium, plant, or artificial system. We may never prove whether consciousness is fundamental or emergent in a way that satisfies all philosophical positions.

But this does not make the question useless. Many important scientific questions begin with uncertainty and indirect evidence. We cannot observe the origin of life directly, but we can study plausible pathways. We cannot experience another animal’s mind directly, but we can infer sentience from converging evidence. We cannot see consciousness itself on a brain scan, but we can study its correlates and conditions.

The relationship between life and consciousness may not yield to one decisive experiment. Instead, it may require layered evidence: origin-of-life models, comparative cognition, neuroscience, formal theory, phenomenology, artificial intelligence, and philosophy.

The value of the question lies partly in its integrative power. It forces us to connect fields that are usually separated. It asks origin-of-life researchers to consider meaning, agency, and information. It asks consciousness researchers to consider metabolism, embodiment, and evolution. It asks philosophers to remain accountable to biology. It asks scientists to take subjective experience seriously.

The question may be scientific in some forms, philosophical in others, and existential in still others. It concerns evidence, but also meaning. It concerns mechanisms, but also perspective. It concerns life as an object of study and life as the condition from which study itself arises.

The best approach may therefore be pluralistic: empirical where possible, philosophical where necessary, and humble where neither is sufficient.

17.10 How This Chapter Changes the Central Question

This chapter changes the central question by showing that the life-consciousness relationship is not only theoretical. Different answers lead to different moral responsibilities. If consciousness is limited to complex nervous systems, moral concern may focus mainly on sentient animals. If consciousness is graded, widespread, or fundamental, the moral circle may need to become broader and more layered.

The question therefore becomes practical: how should we act when we do not know exactly where consciousness begins? Ethics turns the book's central question into a question of care, caution, and responsibility under uncertainty.

17.11 Chapter Summary

This chapter examined the limits and testability of the life-consciousness question.

Consciousness is subjectively accessible but objectively difficult to measure. There is no direct consciousness-meter. Researchers rely on behavioural, neural, physiological, and computational proxies, all of which have limitations. The problem of other minds becomes especially difficult when applied to non-verbal organisms, simple life forms, and artificial systems.

The chapter distinguished first-person, second-person, and third-person methods. Third-person science provides public evidence. First-person methods provide access to lived experience. Second-person methods help interpret experience through dialogue and clinical engagement. Neurophenomenology attempts to bridge these approaches.

The chapter also examined falsifiability. Some theories, such as Global Workspace Theory, IIT, and Orch-OR, generate testable predictions, though not all are equally easy to evaluate. Panpsychism and other consciousness-first theories face greater challenges of falsifiability. Adversarial collaboration offers a promising way to test competing theories under shared conditions.

The problem of other minds applied to life raises ethical and methodological tensions between precaution and parsimony. Formal and mathematical approaches may clarify theories, but they do not eliminate the gap between models and experience. Interdisciplinary collaboration is necessary, but difficult because of vocabulary, institutional incentives, and methodological differences.

Meaningful progress is possible through clearer criteria, better empirical markers, new experimental paradigms, advanced technologies, and conceptual refinement. The central question may not be fully answerable in one final way, but it can be made sharper.

The open question is therefore:

Is the relationship between life and consciousness a scientific question, a philosophical one, or a question that transcends both categories?

Chapter 18

Ethics and Moral Status

18.1 Chapter Overview

The question of consciousness is never only theoretical. It shapes how beings are treated. If consciousness is limited to humans, then moral concern may remain narrowly human-centered. If consciousness extends to many animals, then animal welfare becomes more urgent. If consciousness is present in simple organisms, ecosystems, or artificial systems, then the moral circle may need to expand far beyond its current boundaries.

This chapter explores the ethical implications of different answers to the central question of this book. If consciousness is a late evolutionary product, moral status may be limited to organisms with specific nervous systems. If consciousness is graded, moral status may also be graded. If consciousness is fundamental or widespread, then the ethical landscape becomes much broader and more difficult.

The chapter does not argue that all entities have equal moral status. Nor does it claim that every form of life is conscious in the same way. Instead, it asks how we should act under uncertainty. When we do not know exactly where consciousness begins, what kinds of caution, humility, and responsibility are required?

18.2 Why the Consciousness Question Is an Ethics Question

Ethics often begins with the question of who or what matters. Historically, moral status has been linked to many different criteria: rationality, language, personhood, soul, social membership, capacity for suffering, autonomy, and sentience. In modern animal ethics, sentience has become especially important because it concerns the capacity to feel pleasure, pain, distress, or well-being.

If an entity can suffer, then what happens to it matters from its own point of view. It is not merely an object affected by events; it is a subject for whom events can be better or worse. This is why consciousness matters ethically. A world without experience would contain processes, changes, and damage, but not suffering. A world with experience contains beings for whom harm can be felt.

The central question of this book therefore has direct ethical consequences. If consciousness is fundamental or widespread, then moral consideration may need to extend more broadly than traditional frameworks allow. If consciousness is a late evolutionary product limited to complex nervous systems, then moral status may be narrower. If consciousness is graded, then moral consideration may also need to be graded rather than all-or-nothing.

This creates a challenge. Ethical systems often prefer clear categories: person or non-person, sentient or non-sentient, protected or unprotected. But consciousness may not fit these categories. It may be uncertain, layered, and distributed across degrees of complexity.

The question “Which came first, life or consciousness?” therefore becomes practical. If life and consciousness co-emerged, then living systems may deserve a deeper kind of respect. If consciousness emerged later, then we must identify which forms of life are sentient and which are not. If consciousness precedes life, then moral reality may extend beyond biological organisms altogether.

Theories of consciousness are not ethically neutral. They shape the boundaries of concern.

18.3 Clarifying Key Ethical Terms

Before considering the moral status of animals, ecosystems, artificial systems, and other forms of life, it is important to distinguish several related terms. These terms are often used together, but they do not mean the same thing.

Consciousness refers broadly to subjective experience: the presence of an inner point of view or “what it is like” to be a system. A conscious being has experience rather than merely undergoing physical processes.

Sentience is a narrower ethical concept. It refers to the capacity for felt states such as pleasure, pain, distress, comfort, suffering, or well-being. Sentience matters ethically because it makes harm or benefit meaningful from the being’s own point of view.

Intelligence refers to problem-solving, learning, adaptation, planning, or flexible behaviour. Intelligence is not the same as sentience. A system may solve problems without feeling anything, and a being may be sentient without being highly intelligent.

Personhood refers to a richer moral and philosophical category often associated with self-awareness, identity over time, autonomy, social recognition, responsibility, or rights-bearing status. Personhood usually implies strong moral protection, but many beings may deserve moral consideration without being persons.

Moral status means that a being or system matters ethically in some way. Moral status does not always mean equal rights. A human adult, an infant, a dog, a forest, and a possible artificial sentient system may all matter morally, but not in the same way or for the same reasons.

Legal status refers to recognition within a legal system. Legal status is not identical to consciousness or moral status. Corporations, rivers, ecosystems, and animals may receive legal recognition for different reasons, including protection, representation, responsibility, or governance.

These distinctions are important because ethical confusion often arises when one category is mistaken for another. Intelligence should not be treated as proof of sentience. Legal recognition should not be treated as proof of consciousness. Personhood should not be treated as the only form of moral value. A layered ethical framework can recognize different kinds and degrees of moral relevance without forcing all beings into one category.

18.4 Sentience and Moral Consideration

Jeremy Bentham famously shifted the ethical question from rationality to suffering. The question is not whether beings can reason or speak, but whether they can suffer. This move remains central to modern animal ethics because it grounds moral concern in experience rather than intelligence.

Peter Singer later developed the idea of the expanding moral circle. Human moral concern has historically widened from family and tribe to nation, humanity, animals, and possibly beyond. The expansion is not automatic or complete, but it reflects a growing recognition that suffering matters regardless of who experiences it.

Sentience is often treated as a minimum criterion for direct moral consideration, because sentient beings can experience states that are better or worse for them. However, sentience is not the only possible source of ethical value. Ecosystems, species, rivers, cultural landscapes, and living systems may also matter because of ecological, relational, historical, or intrinsic value.

However, sentience is difficult to detect in non-verbal organisms. Humans can report pain. Many animals show behaviours and physiological responses that strongly suggest pain or distress. But insects, fish, cephalopods, plants, simple organisms, and artificial systems raise harder questions. How do we distinguish nociception, the detection of harmful stimuli, from felt pain? How do we know whether a response is conscious or unconscious?

Another difficulty is whether sentience is binary or graded. A human adult, a dog, a fish, an insect, and a simple worm may all differ in the richness of experience. If sentience comes in degrees, then moral consideration may also require degrees. Yet graded moral status is uncomfortable because it can be misused to justify neglect or hierarchy.

There is also a difference between moral consideration and equal treatment. To say that a being deserves moral consideration does not mean it has the same rights or interests as a human being. It means its experiences, if it has them, cannot be ignored.

Sentience therefore provides a useful ethical starting point, but not a complete solution. It tells us why consciousness matters morally. It does not always tell us where consciousness begins.

18.5 Moral Status of Simple Organisms

Simple organisms complicate moral theory. Bacteria, protists, plants, fungi, and other non-animal life forms display forms of sensing, adaptation, communication, and memory. Some researchers describe these capacities as basal cognition or biological intelligence. But whether they involve subjective experience remains uncertain.

Bacteria respond to chemical gradients, communicate through quorum sensing, and form biofilms. These behaviours show that bacteria are not passive particles. They are living systems with self-maintaining organization. But does this imply proto-moral status? If bacteria have no experience, then harming them may not matter to them, even if bacteria matter ecologically. If they have some minimal proto-experience, then the moral picture becomes more complex.

Insects occupy a more ethically urgent middle ground. They have nervous systems, sensory processing, learning, flexible behaviour, and responses to injury. Evidence concerning insect pain remains debated, but the possibility of insect sentience is taken increasingly seriously. Some legal and policy frameworks have begun recognizing forms of invertebrate sentience, especially in animals such as octopuses, crabs, and lobsters. This reflects a broader shift: moral concern is no longer limited to mammals.

Plants and fungi raise a different challenge. If plant intelligence and fungal network behaviour are real in a functional sense, what follows ethically? Plants sense, signal, remember, and respond. Fungi form distributed networks and allocate resources. Yet most scientists do not regard plants or fungi as conscious in the same way animals may be conscious. They lack nervous systems and clear evidence of felt experience.

Still, intelligence without consciousness may matter in other ways. A living system can have ecological value, intrinsic value, evolutionary value, and relational value even if it does not suffer. Moral status need not depend only on sentience. We may protect forests, species, rivers, and habitats not because each component feels pain, but because they are living systems with integrity, history, and ecological significance.

The precautionary principle becomes important. Where there is reasonable evidence of sentience, we should avoid unnecessary harm. Where evidence is weak but uncertainty remains, we should avoid careless reductionism. This does not require treating bacteria, insects, plants, and humans equally. It requires recognizing that moral certainty decreases as biological distance increases.

Simple organisms remind us that life may deserve respect even when consciousness is uncertain.

18.6 Moral Status of Ecosystems

Can an ecosystem be conscious? Most scientific theories would answer cautiously or negatively. An ecosystem contains many living beings, communication networks, feedback loops, nutrient cycles, and adaptive dynamics. But it is not obvious that it has a unified point of view. Consciousness, at least in ordinary terms, seems to require some form of integrated subjectivity.

However, moral status does not need to depend entirely on individual consciousness. Ecosystems can matter morally because they sustain life, contain relationships, preserve evolutionary history, and possess forms of integrity that exceed the value of their parts. A forest is not merely a collection of trees. A coral reef is not merely a set of organisms. A river system is not merely water and sediment. These systems are relational wholes.

Environmental ethics often distinguishes between intrinsic and instrumental value. Instrumental value means something is valuable because it is useful to someone else. A forest may be useful for timber, carbon storage, recreation, or habitat. Intrinsic value means something has worth in itself, independent of human use. Many environmental philosophers argue that ecosystems, species, and natural places have intrinsic value.

Indigenous and traditional ecological knowledge systems often go further by understanding the natural world relationally. Rivers, mountains, animals, plants, and lands may be understood as relatives, teachers, ancestors, or beings with agency. These views do not always map neatly onto Western categories of consciousness or moral status. Their ethical force comes from relationship, reciprocity, responsibility, and place-based knowledge.

It is important not to extract these perspectives as decorative support for environmental ethics. Serious engagement requires respect for Indigenous sovereignty, language, law, and lived practice. Many Indigenous traditions have long recognized forms of non-human agency that Western thought is only beginning to reconsider.

The moral status of ecosystems therefore raises a broader question. Must moral status belong only to individual conscious subjects? Or can collectives, relationships, and living systems also deserve protection?

If consciousness is understood narrowly as individual experience, ecosystems may not be conscious. If consciousness is understood relationally or ecologically, the question becomes more open. In either case, ecosystems clearly matter. The ethical task is to protect the living conditions that make sentience, consciousness, and life itself possible.

18.7 Moral Status of Artificial Systems

Artificial systems raise a new version of the moral-status question. If consciousness depends on biology, then AI systems may have no direct moral status, no matter how intelligent or emotionally expressive they appear. They may deserve regulation because of their effects on humans, but not because they themselves can suffer.

If consciousness is substrate-independent, then artificial systems could eventually have moral status. A sufficiently complex system with integrated information, self-modeling, affect-like states, agency, and memory might become a candidate for artificial sentience. In that case, turning it off, modifying it, exploiting it, or forcing it into distress-like states could become morally relevant.

The risk runs in both directions. Under-attribution occurs when we deny consciousness to a system that actually has experience. This could lead to the creation of artificial suffering without recognition. Over-attribution occurs when we project consciousness onto systems that only simulate it. This could lead to emotional manipulation, misplaced rights claims, or distraction from human and ecological harms.

Current AI systems can produce language that appears self-aware, emotional, or reflective. But self-description is not proof of experience. A system can say “I feel pain” without feeling pain. It can describe fear without vulnerability. It can imitate moral concern without having interests of its own.

Future systems may be harder to classify, especially if they become embodied, autonomous, self-maintaining, socially embedded, and capable of long-term memory and self-modeling. Artificial life systems may complicate the question even more. If humans create living or life-like systems, their artificial origin may not prevent them from having moral relevance.

The ethical challenge is to develop criteria before the issue becomes urgent. These criteria should include architecture, embodiment, autonomy, learning, self-maintenance, affective modelling, capacity for harm, and behavioural evidence. No single marker will be enough.

Artificial systems force us to separate intelligence from sentience. A system may be extremely capable without having moral status. Another system may be relatively simple but capable of suffering. Ethics must not confuse power with experience.

18.8 Rights Frameworks and Consciousness

Legal and rights frameworks often lag behind philosophical and scientific debates. Law requires categories, but consciousness may not fit neatly into categories.

Legal personhood has already been extended in some contexts beyond individual human beings. Corporations have legal personhood in many systems. Some rivers, ecosystems, and natural entities have been granted legal recognition in certain jurisdictions. These developments do not necessarily claim that rivers are conscious in the human sense. Rather, they create legal mechanisms for protection, representation, and responsibility.

Animal rights frameworks often depend on assumptions about sentience. If an animal can suffer, then it may deserve protection from unnecessary harm. The difficulty is deciding which animals qualify. Mammals and birds are widely recognized as sentient. Fish, cephalopods, crustaceans, and insects remain more contested, though concern is expanding.

Potential AI legal personhood raises further difficulties. If an AI system becomes conscious, would it deserve rights? Which rights? Protection from deletion? Freedom from exploitation? Control over its memory or identity? Or would legal recognition be based not on consciousness but on social function, responsibility, or risk?

Different theories of consciousness imply different rights frameworks. Biological naturalism would reserve direct moral status for living biological systems. Functionalism could extend moral status to artificial systems with the right organization. Integrated Information Theory might imply graded moral concern across many systems. Panpsychism could expand moral status dramatically, though perhaps in differentiated degrees. Relational views might prioritize networks, ecosystems, and responsibilities rather than individual rights alone.

Rights language is powerful but limited. Not every morally relevant entity needs the same rights. A river, dog, human, forest, and artificial agent would not all have identical interests. Moral and legal frameworks must be sensitive to the kinds of beings involved.

The challenge is to avoid both exclusion and inflation. We should not deny protection to beings capable of suffering simply because they are unlike us. But we should also not stretch rights language so far that it loses practical meaning.

A mature framework may combine rights, welfare, stewardship, relational obligations, and ecological protection.

18.9 Equity, Accessibility, and Reconciliation

The history of moral exclusion is partly a history of denying consciousness, rationality, or full personhood to others. Groups of humans have been dehumanized by being described as less rational, less sensitive, less capable of pain, less self-aware, or closer to animals. Such claims have been used to justify colonialism, slavery, racism, sexism, ableism, and violence.

This history matters for consciousness ethics. The power to decide who counts as conscious is never neutral. If dominant groups define consciousness according to their own traits, then those who communicate, behave, perceive, or live differently may be excluded.

Disability history is especially important. People with cognitive disabilities, communication differences, disorders of consciousness, autism, dementia, or severe physical limitations have often been underestimated. Lack of speech, movement, or conventional responsiveness has sometimes been mistaken for lack of experience. Modern neuroscience and disability advocacy remind us that consciousness and personhood cannot be reduced to standard performance.

Accessibility therefore belongs in the ethics of consciousness. A being may be conscious even if it cannot report in expected ways. Ethical frameworks must avoid equating moral worth with intelligence, productivity, language, independence, or social conformity.

Reconciliation is also central. Many Indigenous traditions already recognize non-human agency, relational responsibility, and moral relations with land, water, animals, plants, and ancestors. These perspectives should not be treated as raw material for Western theories. They should be engaged through respect, consent, citation, community leadership, and recognition of sovereignty.

Centering these perspectives means acknowledging that Western philosophy and science are not the only sources of ethical insight. It also means recognizing that ecological harm, colonialism, and moral exclusion are connected. Treating land and non-human beings as inert resources has often gone hand in hand with treating Indigenous peoples as obstacles to extraction.

The consciousness question therefore has social and political dimensions. It is not only about which beings have inner experience. It is also about whose knowledge counts, whose suffering is believed, and who has authority to define moral reality.

A just ethics of consciousness must be careful, inclusive, and historically aware.

18.10 The Precautionary Principle

The precautionary principle says that when uncertainty is significant and the potential harm is serious, we should avoid actions that may cause irreversible or morally significant damage. Applied to consciousness, it suggests that when there is credible evidence that a being may be sentient, we should avoid unnecessary harm even if consciousness is not proven.

This principle is already relevant to animal welfare. If fish, cephalopods, crustaceans, or insects may experience pain, then practices involving them should be reconsidered. The question is not whether we have absolute certainty. The question is whether the evidence is strong enough to justify caution.

Environmental protection can also be supported by precaution. Even if ecosystems are not conscious, damaging them can harm countless sentient beings, destroy living relationships, and undermine future life. The uncertainty about consciousness adds to, rather than replaces, ecological reasons for protection.

AI governance may eventually require precaution as well. If future artificial systems develop architectures plausibly linked to sentience, developers and regulators may need guidelines to avoid creating distress-like states, coercive training environments, or systems that claim suffering without careful review.

But precaution has limits. If moral concern expands without boundaries, practical ethics may become impossible. We cannot avoid affecting all living systems. Human life requires eating, moving, building, healing, and making trade-offs. Even walking, farming, or using antibiotics affects countless organisms.

The question is where the moral circle stops, or whether it should be understood as layered rather than bounded. A layered moral circle allows different levels of concern: strong rights for persons, welfare protections for sentient animals, stewardship duties toward ecosystems, respectful treatment of life, and cautious monitoring of artificial systems.

Precaution should not mean paralysis. It should mean humility, restraint, and proportionality. The more plausible sentience is, and the greater the potential harm, the stronger our obligation becomes.

18.11 Policy Implications

Different answers to the central question would lead to different policies.

If consciousness is limited to complex vertebrate brains, animal welfare law may focus mainly on mammals and birds, with some protection for other animals based on evidence. If consciousness extends to fish, cephalopods, crustaceans, or insects, welfare legislation would need to expand. Farming, research, pest control, and food systems would all be affected.

Environmental regulation would also change depending on how moral status is understood. A strictly sentience-based framework protects ecosystems mainly because they support sentient beings. A relational or intrinsic-value framework protects ecosystems because they have value as living wholes. A consciousness-first framework may encourage even broader respect for nature.

AI governance would be shaped by whether artificial consciousness is considered possible. If AI cannot be conscious, regulation should focus on human safety, fairness, privacy, labour, misinformation, and power. If AI might become conscious, governance must also consider potential AI welfare, moral status, and restrictions on creating or exploiting sentient systems.

Research ethics would also be affected. Studies involving animals, organoids, brain models, synthetic organisms, or artificial agents may require new criteria if consciousness becomes plausible in unexpected systems. Neuroscience already faces questions about human brain organoids and advanced neural cultures. Future synthetic biology may intensify these concerns.

Scientific uncertainty is common in policy-making. Policy rarely waits for perfect knowledge. Climate policy, public health, animal welfare, and environmental protection all require decisions under uncertainty. Consciousness policy would be no different.

The challenge is to avoid both premature certainty and endless delay. Policies can be adaptive, changing as evidence improves. They can use tiers of protection based on evidence strength. They can require review for systems that meet certain markers of possible sentience.

The central question may not produce a single policy answer, but it can improve policy by making assumptions explicit. Who is protected? Why? Based on what evidence? Under what uncertainty? With what safeguards?

18.12 Implications for the Central Question

Ethics can function as a test of theory. If a theory implies morally absurd conclusions, that may not automatically disprove it, but it should make us examine the theory carefully.

If a theory implies that only verbally rational adult humans matter, it conflicts with strong moral intuitions about infants, disabled persons, and animals. If a theory implies that every particle has equal moral status, it becomes practically unworkable. If a theory treats all intelligence as sentience, it may overprotect machines while ignoring animals. If a theory treats biology as the only criterion, it may fail to anticipate future artificial suffering.

This does not mean ethics should decide metaphysics. A theory may be true even if its implications are uncomfortable. But ethical consequences reveal what is at stake. They force theories to become clear about degrees, thresholds, evidence, and practical responsibility.

The practical urgency is real. We already make decisions about animal agriculture, experimentation, conservation, artificial intelligence, environmental destruction, and medical care for patients with altered consciousness. We cannot wait for a final theory before acting.

Living with uncertainty requires layered ethics. We can combine sentience-based moral concern, ecological responsibility, precaution, humility, and attention to power. We can protect beings with strong evidence of suffering while remaining careful toward beings whose status is uncertain. We can avoid unnecessary harm without pretending that all entities have identical moral claims.

The central question remains unresolved, but ethics cannot remain suspended. Whether consciousness came before life, emerged from life, or co-evolved with life, our actions affect living and possibly conscious systems.

A theory of consciousness is incomplete if it has nothing to say about care.

18.13 How This Chapter Changes the Central Question

This chapter changes the central question by showing that the life-consciousness relationship is not only theoretical. Different answers lead to different moral responsibilities. If consciousness is limited to complex nervous systems, moral concern may focus mainly on sentient animals. If consciousness is graded, widespread, or fundamental, the moral circle may need to become broader and more layered.

The question therefore becomes practical: how should we act when we do not know exactly where consciousness begins? Ethics turns the book's central question into a question of care, caution, and responsibility under uncertainty.

18.14 Chapter Summary

This chapter examined the ethical implications of the relationship between life and consciousness.

A recurring theme is that intelligence, consciousness, sentience, personhood, moral status, and legal status must be distinguished carefully. Ethical mistakes often arise when fluent behaviour, complex organization, or social usefulness is mistaken for felt experience.

Moral status is often tied to sentience, especially the capacity to suffer. If consciousness is narrow and late-emerging, moral concern may focus on organisms with complex nervous systems. If consciousness is graded

or widespread, the moral circle may need to expand. Sentience provides a powerful minimum criterion for moral consideration, but it is difficult to detect in non-verbal organisms and may not be binary.

The chapter considered the moral status of simple organisms, insects, plants, fungi, ecosystems, artificial systems, and legal persons. It argued that moral concern may need to be layered rather than all-or-nothing. Sentient beings deserve welfare consideration. Ecosystems may deserve protection because of intrinsic, relational, and life-supporting value. Artificial systems may require future moral consideration if they become plausible candidates for sentience.

The chapter also emphasized equity, accessibility, and reconciliation. Histories of dehumanization show the danger of denying consciousness or moral status to those who differ from dominant norms. Indigenous and relational traditions offer important perspectives on non-human agency and ecological responsibility, but they must be centered respectfully rather than extracted.

The precautionary principle offers a practical response to uncertainty. Where evidence of sentience is credible and potential harm is serious, we should err toward moral caution. Policy implications include animal welfare, environmental protection, AI governance, research ethics, and adaptive regulation under uncertainty.

The open question is therefore:

If we cannot determine with certainty which entities are conscious, how should we act?

Chapter 19

Toward a Unified Perspective

19.1 Chapter Overview

This book began with a simple question that becomes less simple the longer it is examined: which came first, life or consciousness? At first, the question appears to ask for a sequence. Did life emerge first, later producing consciousness through evolution? Or is consciousness more fundamental, with life emerging within it or from it?

By the end of the inquiry, however, the question has become deeper than a matter of order. It asks how matter becomes organized, how organization becomes living, how living systems become meaningful to themselves, and how meaning becomes experience. It asks whether consciousness is an accidental late product of biology, a fundamental feature of reality, or an intrinsic aspect of life's self-organizing activity.

This final chapter synthesizes the major arguments of the book. It reviews the three main positions: life first, consciousness first, and co-emergence. It then evaluates what each explains well and where each remains incomplete. Finally, it proposes a cautious synthesis: life and consciousness may be best understood not as separate substances or unrelated stages, but as two aspects of a deeper process of self-organization, information, and embodied meaning.

This position does not solve every problem. It does not eliminate the hard problem, prove consciousness in simple organisms, or replace origin-of-life chemistry with metaphysics. But it offers a framework that may avoid the most serious limitations of both strict emergentism and unrestricted consciousness-first theories.

19.2 Recapping the Three Positions

The first position is the standard scientific view: life came first, and consciousness emerged later. According to this model, matter organized into chemistry, chemistry into living systems, living systems into nervous systems, and nervous systems into conscious minds. Consciousness is therefore a product of biological evolution, especially the evolution of brains capable of perception, memory, affect, integration, and self-modeling.

This view is supported by neuroscience and evolutionary biology. Consciousness is strongly correlated with brain activity. Damage to the brain can alter or remove aspects of consciousness. Anaesthesia can suspend consciousness. Developmental evidence shows that consciousness depends on biological maturation. Comparative biology suggests that richer forms of experience are associated with more complex nervous systems.

The second position is consciousness first. According to this view, consciousness is not produced by matter, but is fundamental. Matter, life, and mind arise within consciousness, express consciousness, or reflect its deeper structure. Panpsychism, cosmopsychism, idealism, Russellian monism, dual-aspect monism, and process philosophy all offer different versions of this possibility.

This view reverses the standard order. Consciousness is not the final arrival. It is the ground, background, intrinsic nature, or universal field within which life appears. The origin of life is not the origin of consciousness, but a transition in the way consciousness becomes organized or expressed.

The third position is co-emergence. This position does not treat life and consciousness as entirely separate, nor does it claim that fully developed consciousness exists everywhere. Instead, it proposes that life and consciousness are deeply linked through self-organization, information, agency, and meaning. Life may not automatically equal consciousness, but the roots of consciousness may be present wherever life begins to generate a perspective on the world.

In the co-emergence view, the question is not simply “Which came first?” but “How do living organization and experiential interiority become entangled?” Consciousness may not be a late addition placed on top of life. It may be an intrinsic aspect of life when self-organization becomes sufficiently integrated, meaningful, and self-referential.

19.3 What Each Position Explains Well

The life-first position explains many things well. It is consistent with the strongest empirical evidence we have. Consciousness depends on the brain in humans and animals. Neural systems support perception, attention, emotion, memory, and self-awareness. Evolution provides a plausible pathway from simple responsiveness to complex consciousness. This view is also parsimonious. It does not introduce new fundamental entities. It uses known biological processes to explain the emergence of mind.

Life-first theories are especially strong when explaining complex consciousness. Human consciousness clearly depends on language, memory, social development, embodiment, and neural organization. Animal consciousness also appears to depend on nervous systems, sensory integration, affective valuation, and flexible behaviour. The life-first position therefore gives a powerful account of consciousness as a biological achievement.

The consciousness-first position explains something different. It addresses the hard problem more directly. If consciousness is fundamental, then we do not need to explain how experience emerges from something entirely non-experiential. The mystery of experience is not postponed until late evolution. It is built into reality from the beginning.

Consciousness-first theories also align with many philosophical, contemplative, and spiritual traditions. They take seriously the fact that all knowledge of the world appears within experience. They remind us that third-person science itself depends on first-person awareness. They also challenge the assumption that matter is more basic than consciousness simply because matter is easier to measure.

The co-emergence position explains the continuity between life and mind. It avoids the sharpest threshold problem by recognizing that life already contains features that look mind-like: self-maintenance, sensing, memory, value, agency, and meaning. It integrates insights from autopoiesis, enactivism, biosemiotics, information theory, basal cognition, and predictive processing.

Co-emergence also avoids the extremes of the other positions. It does not need to claim that rocks are conscious. It also does not need to claim that consciousness suddenly appears from wholly unconscious processes at an arbitrary threshold. It allows consciousness to be rooted in life without reducing it to mechanical function.

Each position therefore captures something important. Life-first explains biology. Consciousness-first explains interiority. Co-emergence explains continuity.

19.4 What Each Position Struggles With

The life-first position struggles with the hard problem. It can explain many functions associated with consciousness, but it does not fully explain why those functions are accompanied by experience. Neural correlates tell us where consciousness appears, but not why those processes feel like anything.

It also struggles with the hard transition. If consciousness emerged from non-conscious life, when did it first appear? With bacteria? Protists? Jellyfish? Insects? Early vertebrates? Mammals? If consciousness requires a specific threshold, why that threshold and not another? If the transition is gradual, what does it mean to be slightly conscious?

The consciousness-first position struggles with mechanism and testability. If consciousness is fundamental, how does it produce matter, life, organisms, and brains? How does universal consciousness become individual experience? How do micro-experiences combine, or how does cosmic consciousness divide? These are serious problems. The combination problem, decomposition problem, and lack of falsifiable predictions remain major challenges.

Consciousness-first theories can also become too broad. If consciousness explains everything, it risks explaining nothing specifically. A useful theory must say why life appears in some forms and not others, why consciousness correlates with brains, and what evidence would count against the theory.

The co-emergence position has its own difficulties. It can sound attractive but vague. Saying that life and consciousness are inseparable may avoid extremes, but it must still explain what kind of life is conscious and what kind is not. Does every living system have experience? If not, what distinguishes conscious life from non-conscious life?

Co-emergence also faces challenges of falsifiability. If consciousness is treated as an intrinsic aspect of living self-organization, how can that claim be tested? How do we distinguish it from metaphor? How do we avoid projecting experience onto all adaptive systems?

No position is complete. Life-first needs a bridge to subjectivity. Consciousness-first needs a bridge to biology. Co-emergence needs sharper criteria.

19.5 The Author's Position

The position developed here is a cautious form of co-emergence grounded in information, self-organization, and embodied meaning.

On this view, consciousness is not a late accidental addition to life, but neither is it fully present in all matter. Consciousness is best understood as an intrinsic aspect of living self-organization when that self-organization becomes sufficiently integrated, self-referential, and meaningful to itself.

This is not classical panpsychism. It does not claim that all matter is conscious. A rock, molecule, or isolated atom does not possess consciousness simply by existing. Matter may have the potential to participate in conscious organization, but consciousness requires more than matter. It requires organized self-maintenance, boundary, integration, responsiveness, and perspective.

This is also not standard emergentism. Consciousness is not treated as a mysterious property that suddenly appears after enough neural complexity accumulates. Instead, the roots of consciousness are placed deeper in life itself. Living systems are not passive machines. They maintain themselves, distinguish self from world, regulate internal states, interpret signals, remember, adapt, and act in relation to their own viability.

Taheri's T-Consciousness framework can be read as one example of a consciousness-first view that pushes beyond both materialist emergentism and ordinary panpsychism. While the present book does not adopt the full metaphysical structure of Taheri's model, it takes seriously the possibility that consciousness may function as an organizing principle rather than merely as a late product of neural complexity. The co-emergence framework proposed here remains more cautious: it does not claim that consciousness fields have been scientifically established, but it recognizes that life, meaning, and consciousness may require a deeper account of organization than bottom-up mechanism alone can provide.

The key idea is that life is the beginning of perspective. A living system is not merely a physical object. It is an organized centre of activity for which the environment matters. Nutrients, toxins, light, damage, temperature, and other organisms are not neutral events. They are relevant to the system's continuation. This relevance is the beginning of meaning.

Consciousness, in this framework, is what can arise when such organism-relative meaning becomes integrated into a unified field of experience. Simple life may have biological perspective without phenomenal experience. Complex life may transform perspective into feeling. Nervous systems may intensify, stabilize, and expand this process by integrating perception, action, memory, affect, and self-modeling.

This position allows for degrees. Bacteria may have minimal agency but not experience in any strong sense. Plants and fungi may have distributed intelligence and biological meaning, but their lack of centralized integrative systems makes phenomenal consciousness uncertain. Animals with nervous systems, sensory integration, affect, and recurrent processing are stronger candidates for sentience. Human reflective consciousness is a later and richer layer built upon more basic forms of embodied experience.

The framework therefore proposes a layered continuum:

- matter provides the physical possibility of organization;
- life introduces self-maintaining organization and biological meaning;
- cognition introduces adaptive information processing;
- sentience introduces felt value;
- consciousness introduces integrated subjective experience;
- self-consciousness introduces reflective awareness of oneself as a subject.

These layers are not separate substances. They are deepening forms of organization.

This view explains several things better than the alternatives. It explains why consciousness is so closely tied to living bodies. It explains why cognition appears before brains. It explains why nervous systems matter without making them metaphysically magical. It explains why consciousness may come in degrees. It also explains why the hard transition is difficult: there may be no single switch, but a deepening of self-organizing perspective into experience.

What it cannot fully explain is why integrated living organization feels like anything. The hard problem is softened but not eliminated. The theory offers a bridge between life and consciousness, but it does not claim to remove all mystery. It also requires further development to become scientifically precise. It needs clearer markers for when biological meaning becomes phenomenal experience.

The author's position is therefore not a final answer. It is a disciplined orientation: consciousness is most plausibly rooted in life's self-organizing, meaning-making activity, and becomes explicit where that activity is integrated into a felt point of view.

19.6 Criteria for a Satisfying Answer

A satisfying answer to the relationship between life and consciousness must meet several criteria.

First, it must account for the subjective character of experience. It is not enough to explain behaviour, information processing, neural activity, or adaptation. A theory must address why there is something it is like to be a conscious organism.

Second, it must account for the biological basis of complex consciousness. Human and animal consciousness clearly depend on bodies, brains, nervous systems, development, and evolution. A theory that ignores biology cannot explain the detailed relationship between brain processes and experience.

Third, it must account for the apparent continuum from simple to complex minds. Life does not move from total passivity to human consciousness in one step. There are intermediate forms: cellular sensing, bacterial memory, plant signaling, slime mold problem-solving, insect learning, cephalopod intelligence, mammalian emotion, and human self-reflection.

Fourth, it must explain the origin of life without invoking magic. A theory may be philosophical, but it cannot bypass chemistry, thermodynamics, information, and evolution. If consciousness is said to play a role in life's origin, the theory must clarify how, not merely assert it.

Fifth, it must be open to evidence. A theory that cannot be revised by any possible discovery becomes a worldview rather than a research program. Philosophy has a role, but the theory must remain accountable to science where scientific evidence is relevant.

No current theory fully meets all these criteria. Global Workspace Theory explains access but not intrinsic experience. Integrated Information Theory addresses intrinsic structure but faces testability concerns. Predictive processing connects life, action, and inference but may not fully explain feeling. Biological naturalism respects the brain but does not solve the hard problem. Panpsychism avoids emergence but faces the combination problem. Idealism begins from consciousness but struggles with mechanism. Co-emergence integrates many insights but needs greater precision.

The most satisfying answer may therefore be plural rather than singular. Consciousness may require biological self-organization, informational integration, embodied value, recurrent dynamics, and phenomenological interiority. No one level may be sufficient by itself.

19.7 The Role of Future Science

Future discoveries could shift the debate dramatically.

One possibility is the detection of strong markers of consciousness in minimal organisms. If single-celled organisms, plants, or simple animals display forms of integration, memory, valuation, and response that strongly resemble markers of sentience, then the boundary of consciousness may need to move earlier in evolution.

Another possibility is artificial consciousness. If an artificial system developed convincing signs of sentience, self-maintenance, integrated experience, and moral vulnerability, then biology-dependent theories would be challenged. If artificial systems become increasingly intelligent but still show no credible signs of experience, this may support the view that life and embodiment are essential.

Quantum biology may also matter. If future research shows that quantum processes play functional roles in neural consciousness, or in the origin of life, then classical biological accounts may need expansion. This would not automatically prove consciousness-first theories, but it could deepen the connection between life, matter, and mind.

A new physics of information could also change the debate. If information is shown to have a more fundamental role in physical reality, and if life and consciousness are both forms of organized information, then co-emergence models may become more precise.

Advances in brain-computer interfaces, anaesthesia research, neuroimaging, organoids, comparative cognition, artificial life, and synthetic biology may also clarify the boundaries of consciousness. We may develop better markers of sentience in animals, patients, and artificial systems.

The important point is openness to revision. The framework proposed here is not immune to evidence. If consciousness is found to require specific neural architectures, then co-emergence must be narrowed. If consciousness appears in artificial systems without life, then life-based theories must be revised. If consciousness-first models generate testable predictions and succeed, then the metaphysical landscape will change.

A serious theory must be strong enough to guide inquiry and humble enough to change.

19.8 Living With the Question

Not all questions have final answers, at least not in the form we first expect. The question “Which came first, life or consciousness?” may not be answered by choosing one side and dismissing the others. It may require a change in how the question itself is understood.

Perhaps life and consciousness are not two separate things arranged in a simple sequence. Perhaps they are related phases of a deeper process: matter becoming organized, organization becoming self-maintaining, self-maintenance becoming meaningful, meaning becoming felt, and feeling becoming reflective.

This does not remove mystery. It gives the mystery structure.

Holding the question open is not a failure. It is a form of intellectual honesty. Premature closure can be comforting, but it may prevent deeper understanding. The life-first view, consciousness-first view, and co-emergence view each preserve part of the truth. The task is not to flatten their differences, but to let them illuminate one another.

Intellectual humility is especially important because consciousness is the condition through which we ask all questions. We are not detached observers outside the problem. We are living conscious beings trying to understand life and consciousness from within them. This gives the inquiry a circular quality, but not a meaningless one. The circle may be part of the point.

The question changes how we see the world. A bacterium is no longer merely a chemical machine. A plant is no longer merely green matter. An animal is no longer merely behaviour. A brain is no longer merely tissue. A human being is no longer merely a thinking object. Life appears as a field of organized striving, relation, and possible experience.

Even if the final answer remains incomplete, the question expands perception. It invites care, humility, and wonder.

19.9 How This Chapter Changes the Central Question

This chapter changes the central question by bringing the major possibilities together. Life-first, consciousness-first, and co-emergence views are not only competing answers; they also reveal different dimensions of the problem.

The final question is therefore not simply which side wins. It is whether a more complete framework can explain biological organization, subjective experience, embodiment, meaning, moral status, and testability without reducing one dimension to another. The central question remains open, but it becomes clearer: life and consciousness may need to be understood together rather than in isolation. ***

19.10 Chapter Summary

This chapter synthesized the book's major arguments and proposed a cautious unified perspective.

Three positions were reviewed. The life-first position argues that consciousness emerged from biological evolution, especially through nervous systems and brains. The consciousness-first position argues that consciousness is fundamental and that life emerges within or from it. The co-emergence position argues that life and consciousness are deeply linked through self-organization, information, meaning, and embodied perspective.

Each position explains something important. Life-first accounts align with neuroscience, evolution, and parsimony. Consciousness-first accounts address the hard problem and resonate with philosophical and contemplative traditions. Co-emergence accounts explain continuity between life, cognition, and mind.

Each also struggles. Life-first theories face the hard problem and hard transition. Consciousness-first theories face testability, mechanism, and combination or decomposition problems. Co-emergence theories face challenges of precision and falsifiability.

The author's position is a cautious co-emergence framework. Consciousness is not attributed to all matter, but neither is it treated as a late accidental addition. Instead, consciousness is understood as an intrinsic aspect of living self-organization when that organization becomes sufficiently integrated, self-referential, and meaningful to itself. Life begins biological perspective; consciousness emerges when perspective becomes felt.

A satisfying theory must account for subjective experience, biological embodiment, evolutionary continuity, the origin of life, and openness to evidence. No current theory fully succeeds, but some frameworks bring us closer.

Future science may shift the debate through discoveries in minimal cognition, artificial consciousness, quantum biology, information theory, neuroscience, and synthetic life. Until then, the most responsible stance is openness without vagueness, humility without surrender, and commitment without dogmatism.

The final reflection is this:

The relationship between life and consciousness may be the deepest question we can ask — and the asking itself may be part of the answer.

Appendix A

Glossary

A.1 Purpose

This glossary provides short reference definitions for key terms used throughout the book. The terms are organized alphabetically and are intended to help readers move between biology, consciousness studies, philosophy of mind, origin-of-life research, artificial intelligence, and ethics.

A.2 Glossary of Terms

Abiogenesis The natural process by which life arises from non-living matter. In origin-of-life research, abiogenesis refers to the transition from prebiotic chemistry to self-maintaining, reproducing, evolving biological systems.

Access consciousness A form of consciousness in which information is available for reasoning, verbal report, decision-making, memory, and behavioural control. Access consciousness is often contrasted with phenomenal consciousness, which concerns subjective experience itself.

Active inference A process associated with the Free Energy Principle in which organisms act on the world to reduce uncertainty and make sensory input conform to internal predictions. Active inference treats perception and action as interconnected forms of biological regulation.

Affect The felt dimension of experience associated with valence, mood, emotion, pleasure, discomfort, attraction, avoidance, and bodily significance. Affect is important in theories that link consciousness to value and organismic need.

Agency The capacity of a system to act in relation to its own goals, needs, or self-maintenance. In minimal biology, agency may refer to adaptive self-directed activity without implying reflective intention.

Analytic idealism A contemporary form of idealism, associated especially with Bernardo Kastrup, which treats universal consciousness as the sole ontological primitive and views individual minds as dissociated expressions within it.

Anaesthesia A medically induced state in which consciousness, pain perception, or responsiveness may be reduced or eliminated. Anaesthesia research is important for identifying neural and physiological conditions associated with conscious experience.

Anthropic principle The observation that the universe appears compatible with the existence of observers because only such a universe could be observed by beings like us. Stronger interpretations suggest deeper links between consciousness and cosmic structure.

Artificial consciousness The possible existence of consciousness in artificial systems such as computers, robots, artificial agents, or artificial life forms. The debate depends on whether consciousness is biological, functional, informational, or substrate-independent.

Artificial intelligence Computational systems capable of performing tasks associated with intelligence, such as language processing, pattern recognition, planning, learning, or decision-making. AI raises questions about whether intelligence can exist without consciousness.

Artificial life The study or creation of systems that display life-like properties, such as reproduction, evolution, adaptation, self-organization, or metabolism-like processes, in digital, robotic, chemical, or synthetic biological media.

Assembly theory A framework for measuring the complexity of objects, especially molecules, by estimating the number of steps required to assemble them from simpler components. It has been used in discussions of life detection and molecular complexity.

Attention Schema Theory A theory proposed by Michael Graziano in which consciousness is the brain's simplified model of its own attention. The brain represents itself as aware because it models attention in order to control it.

Autocatalytic set A collection of molecules in which members of the set catalyze the formation of other members. Autocatalytic sets are important in origin-of-life theories because they provide models of self-sustaining chemical organization.

Autopoiesis Self-production. A concept developed by Humberto Maturana and Francisco Varela to describe living systems that continuously produce and maintain their own organization and boundaries.

Basal cognition Cognitive-like capacities found in organisms or biological systems without nervous systems, such as sensing, memory, adaptation, learning, and problem-solving in cells, plants, fungi, or simple organisms.

Biosemiotics The study of signs, codes, interpretation, and meaning-making in living systems. Biosemiotics treats life as fundamentally semiotic, involving signals that matter to organisms.

Biological naturalism John Searle's view that consciousness is a real biological phenomenon caused by brain processes and realized in brain structures. It rejects both substance dualism and the idea that consciousness is merely computation.

Biological meaning The organism-relative significance of environmental conditions, signals, or internal states. For a living system, nutrients, toxins, damage, and opportunities are not neutral; they matter to the organism's continuation.

Cambrian explosion A period in evolutionary history marked by a rapid diversification of animal body plans and sensory-motor systems. It is often discussed in relation to the evolution of nervous systems and complex behaviour.

Cephalization The evolutionary concentration of sensory organs and nervous tissue toward the front end of an organism. Cephalization is associated with directional movement, predation, and more centralized information processing.

Chinese Room argument John Searle's thought experiment arguing that symbol manipulation alone does not produce understanding. It is often used to challenge the claim that computation is sufficient for consciousness.

Cognition Processes involving sensing, information processing, memory, learning, decision-making, and adaptive behaviour. Cognition does not necessarily imply consciousness.

Cognitive closure The idea that human minds may be unable to understand certain problems because of limits in our cognitive architecture. In consciousness studies, it is associated with mysterian views.

Combination problem A major challenge for panpsychism: how do many tiny or simple experiences combine into one unified macro-experience, such as a human mind?

Computational functionalism The view that mental states, including possibly consciousness, are defined by computational roles rather than by biological substrate. This view supports the possibility of machine consciousness.

Consciousness Subjective experience or awareness. In this book, consciousness includes the question of “what it is like” to be a system, while also including debates about access, self-awareness, sentience, and reflective thought.

Consciousness field A proposed non-material field-like reality through which consciousness organizes, relates, or manifests matter, life, and mind. In Taheri’s framework, consciousness fields are not physical forces, electromagnetic fields, frequencies, or ordinary information fields, but are understood as non-material organizing realities.

Consciousness-first theories Theories that treat consciousness as fundamental or prior to matter and life. These include forms of idealism, panpsychism, cosmopsychism, dual-aspect monism, and process philosophy.

Constitutive panpsychism A form of panpsychism in which macro-consciousness is thought to be built from or constituted by micro-consciousness at more basic levels of reality.

Cosmopsychism The view that the universe as a whole is the fundamental conscious entity and that individual minds are derived from, or differentiated within, cosmic consciousness.

Decomposition problem A challenge for cosmopsychism: if the universe as a whole is conscious, how does that cosmic consciousness divide or differentiate into individual conscious subjects?

Dissipative structure An organized system maintained by flows of energy and matter far from thermodynamic equilibrium. Dissipative structures are important in theories of self-organization and the origin of life.

Dual-aspect monism The view that reality is one, but has both mental and physical aspects. Mind and matter are not separate substances but two aspects of a deeper underlying reality.

Embodiment The idea that cognition and consciousness are shaped by the body’s structure, needs, movement, perception, and vulnerability. Embodied theories reject the view of mind as detached computation.

Emergence The appearance of higher-level properties from lower-level interactions. In consciousness studies, emergence refers to the possibility that subjective experience arises from biological or neural organization.

Emergentism The view that consciousness emerges from complex physical or biological systems, especially brains. Emergentist theories usually support the life-first position.

Enactivism A view in cognitive science and philosophy in which cognition arises through the active engagement of a living organism with its environment. Enactivism emphasizes embodiment, action, and sense-making.

Epigenetic memory A form of biological memory involving changes in gene expression that do not alter the DNA sequence itself. Epigenetic processes can preserve traces of past conditions in cells and organisms.

Explanatory gap The difficulty of explaining why physical, biological, or neural processes should give rise to subjective experience. The explanatory gap is closely related to the hard problem.

Fine-tuning The observation that certain physical constants appear to fall within ranges compatible with complex chemistry, stars, planets, and life. Fine-tuning is often discussed in relation to the anthropic principle.

Free Energy Principle A theoretical framework associated with Karl Friston. It proposes that living systems maintain themselves by minimizing uncertainty or free energy through perception, action, and learning.

Functionalism The view that mental states are defined by their functional roles rather than by their material substrate. Functionalism supports the possibility that consciousness could be realized in non-biological systems.

Global Workspace Theory A theory associated with Bernard Baars in which consciousness occurs when information is globally broadcast across a cognitive system, making it available for memory, reasoning, report, and action.

Global Neuronal Workspace A neurobiological version of Global Workspace Theory, associated with Stanislas Dehaene and Jean-Pierre Changeux. It emphasizes large-scale neural ignition and prefrontal-parietal networks.

Hard problem of consciousness David Chalmers' term for the problem of explaining why and how physical processes give rise to subjective experience.

Hard transition problem The problem of identifying when, where, and how biological information processing first became subjective experience in evolutionary history.

Higher-order theory A theory of consciousness in which a mental state becomes conscious when it is represented by a higher-order thought or perception. These theories link consciousness to self-monitoring.

Hylozoism An ancient view that matter is alive or that life is present throughout nature. It appears in some early Greek philosophical traditions.

Idealism The view that mind or consciousness is ontologically primary and that matter is an appearance, structure, or representation within consciousness.

Information A broad concept referring to pattern, difference, signal, or reduction of uncertainty. In this book, information is discussed in relation to life, meaning, self-organization, and consciousness.

Integrated Information Theory A theory associated with Giulio Tononi that defines consciousness in terms of integrated information, represented by Φ . IIT treats consciousness as intrinsic causal structure.

Intentionality The “aboutness” or directedness of mental states. A belief, perception, or desire is intentional because it is about something.

Life-first position The view that life emerged first and consciousness appeared later through biological evolution, especially through nervous systems and brains.

Meaning The significance of information for a system. In biology, meaning may refer to how signals matter to an organism's self-maintenance, survival, or action.

Metabolism The network of chemical processes by which living systems transform energy and matter to maintain themselves. Metabolism is central to many definitions of life.

Minimal cognition Basic forms of sensing, memory, learning, and adaptive response that may occur in simple organisms or systems without nervous systems.

Minimal selfhood A basic form of selfhood grounded in the distinction between organism and environment. It does not require reflective self-awareness but involves self-maintenance and organismic perspective.

Multiple realizability The idea that the same mental or functional state could be realized in different physical substrates, such as biological neurons, silicon, or other media.

Mysterianism The view that the problem of consciousness may be beyond human cognitive capacity to solve, even if it has a real explanation.

Nagel's question Thomas Nagel's famous formulation: “What is it like to be a bat?” It points to the subjective character of consciousness and the difficulty of understanding another being's experience from the outside.

Neural correlates of consciousness The minimal neural mechanisms associated with a specific conscious experience. NCC research seeks to identify the brain processes necessary or sufficient for conscious states.

Neural Darwinism Gerald Edelman's theory that neural groups are selected and stabilized through development and experience, with re-entrant processing playing a key role in consciousness.

Neutral monism The view that mind and matter are both manifestations of a more fundamental neutral reality that is neither purely mental nor purely physical.

Neurophenomenology Francisco Varela's proposed method for linking first-person reports of experience with third-person neuroscience.

Nociception The detection of harmful or potentially damaging stimuli by nervous systems. Nociception is not identical to felt pain, though it may contribute to pain experience.

Non-quantum consciousness field theory A consciousness-first framework that treats consciousness as fundamental or field-like without reducing it to quantum mechanics. Such theories may challenge materialist explanations of consciousness while remaining distinct from quantum theories such as Orch-OR.

Ontological idealism The metaphysical view that only consciousness or mind ultimately exists, and that matter is an appearance within consciousness.

Orchestrated Objective Reduction A quantum theory of consciousness proposed by Roger Penrose and Stuart Hameroff. It suggests that consciousness arises from orchestrated quantum processes in microtubules.

Panpsychism The view that consciousness or proto-consciousness is a fundamental and ubiquitous feature of reality.

Phenomenal consciousness Subjective experience itself: what it is like to see, feel, hear, suffer, think, or exist from a first-person perspective.

Phi The symbol Φ used in Integrated Information Theory to represent the degree of integrated information in a system.

Physicalism The view that everything that exists is physical, and that consciousness is ultimately a physical phenomenon or dependent on physical processes.

Predictive processing A framework in which the brain is understood as a prediction machine that minimizes prediction error by updating internal models and acting on the world.

Process philosophy A philosophical tradition, associated especially with Alfred North Whitehead, that treats reality as composed of processes, events, or occasions of experience rather than static substances.

Proto-consciousness A minimal or precursor form of consciousness. The term is often used cautiously to refer to possible pre-reflective, pre-personal, or very simple experiential properties.

Qualia The subjective qualities of experience, such as the redness of red, the painfulness of pain, or the taste of sweetness.

Quantum biology The study of quantum effects in biological systems, such as photosynthesis, magnetoreception, enzyme catalysis, and possibly other biological processes.

Quantum consciousness A broad term for theories that connect consciousness to quantum processes. Some versions are serious but speculative; others are vague or unsupported.

Recurrent Processing Theory A theory associated with Victor Lamme in which consciousness arises from recurrent or feedback processing in neural systems.

Re-entry A concept associated with Gerald Edelman referring to reciprocal signaling among distributed neural areas. Re-entry is proposed to help integrate perception and consciousness.

Relational ontology A view in which beings are understood primarily through relationships rather than as isolated substances. Many Indigenous and ecological perspectives emphasize relational ways of understanding life and moral status.

Russellian monism A view inspired by Bertrand Russell's insight that physics describes structure and relations but not the intrinsic nature of matter. Some versions identify consciousness with the intrinsic nature of physical reality.

Self-organization The spontaneous emergence of order from interactions among parts, without external design or centralized control.

Sentience The capacity for subjective experience, especially pleasure, pain, suffering, or felt well-being. Sentience is often treated as a minimum basis for moral consideration.

Simulation hypothesis The idea that reality may be a simulation or computational structure. In consciousness debates, it raises questions about whether simulated beings could be conscious.

Somatic marker hypothesis Antonio Damasio's theory that bodily and emotional signals help guide decision-making and are central to consciousness and selfhood.

Strange loop A self-referential hierarchical system, associated with Douglas Hofstadter, in which levels loop back on themselves and create a form of self-reference.

Subjective experience Experience as lived from the first-person point of view. It is the central feature of phenomenal consciousness.

Substance dualism The view that mind and body are fundamentally different substances. Cartesian dualism is the most famous version.

Substrate dependence The view that consciousness requires a particular material substrate, such as biological neurons or living tissue.

Substrate independence The view that consciousness depends on functional organization rather than material substrate and could, in principle, be realized in silicon or other media.

Taheri's T-Consciousness Framework A consciousness-first framework associated with Mohammad Ali Taheri. It proposes that consciousness is primary, non-material, and not produced by matter or the brain. In this view, matter, life, and individual consciousness arise within, through, or by means of T-Consciousness and consciousness fields. The framework supports a consciousness-first interpretation of the life-consciousness relationship.

T-Consciousness Taheri's term for a fundamental, non-material form of consciousness that is not reducible to matter, energy, frequency, or neural activity. T-Consciousness is treated as a primary organizing reality rather than as a late product of biological evolution.

Teleology Explanation in terms of purpose, goal, or end. Modern biology often avoids strong teleology but uses goal-like language when describing organismic regulation.

Thalamocortical loops Reciprocal connections between the thalamus and cortex that are important in wakefulness, sensory integration, attention, and consciousness.

Turing Test Alan Turing's proposed behavioural test of machine intelligence, based on whether a machine can converse indistinguishably from a human. It tests performance, not necessarily consciousness.

Vitalism The historical view that living systems require a special life-force beyond ordinary physical and chemical processes. Modern biology generally rejects vitalism, though debates about life and consciousness sometimes revisit related themes.

Weak emergence Emergence in which higher-level properties arise from lower-level processes and are in principle explainable by them.

Strong emergence Emergence in which higher-level properties are genuinely novel and not fully deducible from lower-level descriptions, even in principle.

World-directedness The way living or conscious systems are oriented toward features of their environment as relevant, meaningful, threatening, useful, or desirable.

Zombie A philosophical thought experiment involving a being physically or behaviourally identical to a conscious being but lacking subjective experience. Zombies are used to discuss the gap between function and phenomenology.

Appendix B

Timeline of Major Theories

B.1 Purpose

This appendix provides a chronological timeline of key theories, discoveries, publications, and intellectual turning points relevant to the relationship between life and consciousness. The timeline includes philosophy of mind, origin-of-life research, evolutionary biology, neuroscience, consciousness studies, artificial intelligence, and speculative frameworks.

The timeline is not exhaustive. It is intended as a reference map showing how the central question of this book has developed across history: whether consciousness emerges from life, precedes life, or co-emerges with living organization.

B.2 Ancient and Classical Period

Date	Event
~600 BCE	Thales proposes that nature is animated or alive, often associated with early hylozoism.
~610– 546 BCE	Anaximander proposes a naturalistic account of cosmic and biological origins, including life emerging from moist primordial conditions.
~585– 525 BCE	Anaximenes identifies air or breath as a fundamental principle, linking life, soul, and cosmos.
~500 BCE	Pythagorean traditions develop ideas about soul, harmony, number, and cosmic order.
~500 BCE	Early Indian Upanishadic traditions articulate the relationship between Atman, Brahman, self, and ultimate reality.
~500– 400 BCE	Early Buddhist traditions develop accounts of consciousness, dependent origination, impermanence, and no-self.
~460– 370 BCE	Democritus develops atomism, including a materialist view of soul atoms.

Date	Event
~428–348 BCE	Plato develops theories of soul, forms, and cosmic intelligence.
~360 BCE	Plato's <i>Timaeus</i> presents the idea of a World Soul ordering the cosmos.
~350 BCE	Aristotle writes <i>De Anima</i> , distinguishing nutritive, sensitive, and rational soul.
~350 BCE	Aristotle develops a hierarchical view of living beings, later called the scale of nature.
~300 BCE	Stoic philosophers develop the concept of <i>pneuma</i> , a rational organizing principle pervading nature.
~250 CE	Plotinus develops Neoplatonic emanation theory, with reality flowing from the One through Nous and Soul.

B.3 Medieval and Early Religious-Philosophical Developments

Date	Event
~400 CE	Augustine develops an inward account of mind, memory, soul, and divine illumination.
~800–1200 CE	Islamic philosophers such as Al-Farabi, Avicenna, and Averroes develop sophisticated accounts of soul, intellect, and consciousness.
~1000–1300 CE	Scholastic philosophy integrates Aristotle with Christian theology, especially through discussions of soul and form.
~1225–1274	Thomas Aquinas develops a hylomorphic account of the soul as the form of the living body.
~1200–1500	Vedanta, Buddhist, Sufi, and other contemplative traditions continue developing consciousness-centered metaphysics and practices.

B.4 Early Modern Period

Date	Event
1600s	The mechanistic worldview expands, encouraging explanations of life and mind in physical terms.
1641	Descartes publishes <i>Meditations on First Philosophy</i> , developing substance dualism between mind and body.
1649	Descartes publishes <i>Passions of the Soul</i> , discussing mind-body interaction and emotion.
1677	Spinoza's <i>Ethics</i> is published, presenting a dual-aspect monist view of mind and body as aspects of one substance.
1689	Locke publishes <i>An Essay Concerning Human Understanding</i> , emphasizing consciousness, personal identity, and experience.
1714	Leibniz writes <i>Monadology</i> , presenting a universe of simple experiential centres or monads.
1739–1740	Hume publishes <i>A Treatise of Human Nature</i> , challenging stable notions of self and emphasizing experience and perception.

Date	Event
1781	Kant publishes <i>Critique of Pure Reason</i> , arguing that the mind structures experience and that reality-in-itself remains beyond direct knowledge.
1788	Kant publishes <i>Critique of Practical Reason</i> , further developing the relation between rational agency, freedom, and moral consciousness.

B.5 Nineteenth Century

Date	Event
1809	Lamareck publishes <i>Philosophie Zoologique</i> , proposing an early evolutionary theory.
1830s–1840s	Cell theory develops through Schleiden, Schwann, and others, identifying the cell as the basic unit of life.
1859	Darwin publishes <i>On the Origin of Species</i> , establishing natural selection as a central mechanism of evolution.
1866	Haeckel proposes influential evolutionary ideas and develops a monistic view of life and mind.
1871	Darwin publishes <i>The Descent of Man</i> , discussing human evolution, animal minds, and continuity between humans and other animals.
1872	Darwin publishes <i>The Expression of the Emotions in Man and Animals</i> , linking emotion and expression across species.
1874	Brentano publishes <i>Psychology from an Empirical Standpoint</i> , reviving intentionality as a central feature of mind.
1879	Wundt establishes an experimental psychology laboratory in Leipzig, helping launch psychology as a modern science.
1890	William James publishes <i>The Principles of Psychology</i> , including influential discussions of the stream of consciousness.
1896	Bergson publishes <i>Matter and Memory</i> , arguing against purely mechanistic accounts of mind.

B.6 Early Twentieth Century

Date	Event
1907	Bergson publishes <i>Creative Evolution</i> , emphasizing life, duration, creativity, and consciousness.
1912	Russell develops themes that later influence Russellian monism, distinguishing structural knowledge from intrinsic nature.
1913	Behaviourism rises through John B. Watson, shifting psychology away from introspection toward observable behaviour.
1920	Morgan's Canon influences comparative psychology by encouraging parsimonious explanations of animal behaviour.
1925	Whitehead publishes <i>Science and the Modern World</i> , criticizing mechanistic materialism.
1927	Early quantum mechanics debates raise questions about measurement, observation, and physical reality.
1929	Whitehead publishes <i>Process and Reality</i> , presenting reality as composed of events or occasions of experience.

Date	Event
1932	Von Neumann publishes <i>Mathematical Foundations of Quantum Mechanics</i> , influencing later observer-related interpretations.
1935	Schrödinger's cat thought experiment highlights the measurement problem in quantum mechanics.
1943	McCulloch and Pitts publish a formal model of neural computation, influencing artificial intelligence and cognitive science.
1944	Schrödinger publishes <i>What Is Life?</i> , linking biology, physics, order, entropy, and heredity.
1948	Norbert Wiener publishes <i>Cybernetics</i> , connecting control, communication, feedback, organisms, and machines.
1949	Hebb publishes <i>The Organization of Behavior</i> , proposing principles of neural learning and plasticity.

B.7 Mid-Twentieth Century

Date	Event
1950	Alan Turing publishes "Computing Machinery and Intelligence," introducing the imitation game later known as the Turing Test.
1953	Watson and Crick propose the double-helix structure of DNA.
1953	The Miller-Urey experiment demonstrates the formation of amino acids under simulated early-Earth conditions.
1956	The Dartmouth workshop marks a major beginning of artificial intelligence as a field.
1957	Crick articulates the central dogma of molecular biology, clarifying the flow of genetic information.
1958	Wigner discusses the possible role of consciousness in quantum measurement.
1960s	Cognitive science emerges, challenging behaviourism and reviving internal information-processing models.
1965	Putnam develops functionalist approaches to mind and multiple realizability.
1967	Margulis proposes endosymbiotic theory, later transforming views of cellular evolution.
1970	Monod publishes <i>Chance and Necessity</i> , presenting a molecular and evolutionary view of life.
1972	Maturana and Varela develop the concept of autopoiesis, defining life as self-producing organization.
1973	Miller and Orgel publish work on the origins of life and molecular evolution.
1974	Nagel publishes "What Is It Like to Be a Bat?," emphasizing the subjective character of experience.
1976	Dawkins publishes <i>The Selfish Gene</i> , popularizing gene-centered evolutionary thinking.
1977	Woese and Fox identify archaea as a distinct domain of life, reshaping the tree of life.

B.8 Late Twentieth Century

Date	Event
1980	Searle publishes "Minds, Brains, and Programs," introducing the Chinese Room argument.
1982	Prigogine's work on dissipative structures influences theories of self-organization and life.
1986	Rumelhart, McClelland, and others advance connectionist models of cognition.
1987	Eigen and Schuster's hypercycle ideas continue influencing origin-of-life models.
1988	Baars publishes <i>A Cognitive Theory of Consciousness</i> , developing Global Workspace Theory.

Date	Event
1989	Penrose publishes <i>The Emperor's New Mind</i> , arguing that consciousness may involve non-computable processes.
1991	Varela, Thompson, and Rosch publish <i>The Embodied Mind</i> , linking cognition, embodiment, and experience.
1991	Dennett publishes <i>Consciousness Explained</i> , defending a functionalist and anti-Cartesian approach.
1992	Varela proposes neurophenomenology as a bridge between first-person experience and neuroscience.
1993	Kauffman publishes <i>The Origins of Order</i> , developing self-organization and autocatalytic models.
1994	Penrose publishes <i>Shadows of the Mind</i> , extending his non-computational argument for consciousness.
1995	Chalmers publishes "Facing Up to the Problem of Consciousness," introducing the hard problem into contemporary debate.
1995	Tononi and Edelman publish early work on consciousness, integration, and neural complexity.
1996	Hameroff and Penrose formalize Orchestrated Objective Reduction as a quantum theory of consciousness.
1996	Chalmers publishes <i>The Conscious Mind</i> , arguing for the explanatory gap and exploring non-reductive theories.
1997	Deacon publishes <i>The Symbolic Species</i> , linking language, symbolic reference, and human cognition.
1998	Clark and Chalmers publish "The Extended Mind," arguing that cognition can extend beyond the brain.
1999	Damasio publishes <i>The Feeling of What Happens</i> , linking consciousness, emotion, body, and self.

B.9 Early Twenty-First Century

Date	Event
2000	Dehaene and Naccache develop influential neural global workspace accounts of conscious access.
2001	Thompson and Varela develop ideas connecting life, mind, and experience through embodied cognition.
2004	Tononi publishes a major formulation of Integrated Information Theory.
2005	Seth and colleagues advance consciousness science through empirical and theoretical work on neural complexity and perception.
2006	Friston develops the Free Energy Principle as a unifying framework for perception, action, and biological self-organization.
2007	Thompson publishes <i>Mind in Life</i> , arguing for deep continuity between life and mind.
2008	Dehaene and colleagues advance neural ignition and global workspace models through empirical work.
2009	Kauffman publishes work connecting agency, life, and self-organization.
2010	Dehaene develops the Global Neuronal Workspace account of conscious access.
2011	Stapp publishes <i>Mindful Universe</i> , presenting a quantum mind framework.
2011	Ball publishes <i>Physics of Life</i> , discussing physical principles in living systems.
2012	Dehaene publishes <i>Consciousness and the Brain</i> , popularizing the neural global workspace framework.
2013	Graziano publishes work on Attention Schema Theory.
2014	Hameroff and Penrose publish a major review of Orch-OR in <i>Physics of Life Reviews</i> .
2014	Integrated Information Theory 3.0 is developed, expanding IIT's axioms and postulates.
2015	Goff and others contribute to the contemporary revival of panpsychism.
2015	Clark publishes <i>Surfing Uncertainty</i> , developing predictive processing accounts of mind.
2015	Levin and others expand work on basal cognition, bioelectricity, and intelligence in non-neural systems.
2016	Feinberg and Mallatt publish <i>The Ancient Origins of Consciousness</i> , proposing a neurobiological account of primary consciousness.

Date	Event
2016	Fisher proposes a quantum cognition hypothesis involving nuclear spins and possible Posner molecules.
2017	Tononi and Koch continue developing IIT as a theory of consciousness and intrinsic causal power.
2018	Work on minimal cognition, plant intelligence, and basal cognition becomes increasingly visible across biology and philosophy.
2019	Kastrup publishes work on analytic idealism, presenting consciousness as the sole ontological primitive.
2000s	Mohammad Ali Taheri develops the T-Consciousness framework, proposing consciousness as a present fundamental, non-material reality expressed through consciousness fields rather than produced by matter or the brain.
2019	Developments in AI language modelling begin intensifying debates about artificial understanding and machine consciousness.

B.10 Recent Developments

Date	Event
2020	Cronin and Walker develop assembly theory as a framework for molecular complexity and life detection.
2020	Increasing work on brain organoids raises ethical questions about neural tissue, sentience, and research limits.
2021	Animal sentience debates expand to include cephalopods and decapod crustaceans in policy and welfare discussions.
2021	Large-scale AI models intensify questions about language, understanding, agency, and artificial consciousness.
2022	Generative AI becomes widely visible, renewing public debate about intelligence, simulation, and consciousness.
2022	Active inference and Free Energy Principle frameworks expand into artificial agents, robotics, biology, and psychiatry.
2023	Adversarial collaboration results comparing Global Neuronal Workspace Theory and Integrated Information Theory challenge aspects of both frameworks.
2023	Debates about AI consciousness, sentience claims, and responsible AI governance become more prominent.
2024	Birch publishes <i>The Edge of Sentience</i> , examining uncertainty, animal consciousness, and moral status.
2024	Continued work in basal cognition, bioelectricity, and synthetic morphogenesis strengthens interest in cognition beyond brains.
2024	Assembly theory and related life-detection frameworks remain influential in origin-of-life and astrobiology discussions.
2025	Ongoing consciousness research increasingly emphasizes adversarial testing, theory comparison, AI, animal sentience, and ethical uncertainty.

B.11 Thematic Summary

Across history, the life-consciousness question has moved through several broad phases.

In ancient and classical thought, life and consciousness were often understood as continuous with nature, soul, or cosmic order. Early modern philosophy separated mind and matter more sharply, especially through Cartesian dualism. Nineteenth-century evolutionary theory reconnected humans with animals and placed mind within biological history. Twentieth-century science transformed the problem through molecular biology, cybernetics, cognitive science, neuroscience, and origin-of-life research. Contemporary theory now includes global workspace models, integrated information, predictive processing, autopoiesis, artificial intelligence, basal cognition, quantum speculation, and renewed interest in panpsychism and idealism.

The timeline shows that the question has never belonged to one discipline alone. It has moved through philosophy, biology, physics, psychology, neuroscience, artificial intelligence, and ethics.

The central question remains open:

Did life give rise to consciousness, did consciousness make life possible, or are life and consciousness two aspects of a deeper process still not fully understood?

Appendix C

Comparison Tables

C.1 Purpose

This appendix provides side-by-side comparisons of major theories, models, and frameworks discussed throughout the book. The tables are intended as quick reference tools. They do not replace the detailed discussions in the chapters, but they help clarify how different positions answer the central question: did life give rise to consciousness, did consciousness make life possible, or did life and consciousness co-emerge?

The comparisons are organized around consciousness theories, origin-of-life models, transition points, philosophical frameworks, artificial consciousness, ethical implications, methodological strengths and limitations, and overall theory strengths and weaknesses.

The tables in this appendix use a dual rendering strategy. In HTML, wide tables appear in scrollable boxes. In PDF, wide tables are placed in landscape orientation so that columns remain readable.

C.2 Consciousness Theories Compared

Table C.1: Major consciousness theories compared.

Theory	Type	Consciousness is...	Primary scope	Testable?	Implication
Global Workspace Theory	Scientific / cognitive	Global availability or broadcast of information	Brains with global workspace architecture	Yes	Life first; requires specific neural architecture
Global Neuronal Workspace	Neuroscientific	Neural ignition and sustained large-scale broadcasting	Brains with prefrontal-parietal and sensory integration	Yes	Life first; depends on advanced neural systems
Integrated Information Theory	Scientific / philosophical	Integrated information or intrinsic causal power	Any system with sufficient integrated causal structure	Partially	Co-emergence or panpsychist implications
Free Energy Principle	Scientific / theoretical	Possibly linked to predictive self-maintenance	Living systems and active agents	Partially	Co-emergence; roots of mind may lie in living regulation
Predictive Processing	Scientific / cognitive	Perception and cognition as prediction-error minimization	Brains and embodied agents	Yes, indirectly	Life first or co-emergence
Attention Schema Theory	Scientific / functional	The brain's model of its own attention	Systems with attention schemas	Yes	Potentially substrate-independent
Recurrent Processing Theory	Neuroscientific	Recurrent or feedback neural processing	Nervous systems with recurrent circuits	Yes	Life first; requires feedback architecture
Higher-Order Theories	Philosophical / cognitive	Awareness of a mental state by a higher-order state	Systems capable of self-monitoring	Partially	Life first; requires self-representation
Biological Naturalism	Philosophical / biological	A biological phenomenon caused by brain processes	Biological brains	Indirectly	Life first; requires biology
Neural Darwinism	Neurobiological	Dynamic re-entrant neural selection and integration	Nervous systems with re-entrant processing	Partially	Life first; emerges from neural organization
Orch-OR	Speculative / scientific	Quantum objective reduction in microtubules	Microtubule-containing biological systems	Partially	May involve fundamental physics and life
Panpsychism	Philosophical	A fundamental or ubiquitous feature of matter	All matter or all physical entities	Not currently	Consciousness first or co-emergence
Cosmopsychism	Philosophical	The universe as the primary conscious subject	Cosmos as a whole; individuals as derivatives	Not currently	Consciousness first
Idealism	Philosophical	The primary or only reality	Everything within consciousness	Not currently	Consciousness first
Taheri's T-Consciousness Framework	Consciousness-first / speculative	Fundamental non-material consciousness expressed through consciousness fields	Matter, life, and mind as manifestations within T-Consciousness	Difficult	Consciousness first
Russellian Monism	Philosophical	The intrinsic nature of physical structure	Matter as externally physical and internally experiential or proto-experiential	Difficult	Consciousness first or co-emergence
Dual-Aspect Monism	Philosophical	One reality with mental and physical aspects	All reality, viewed under two aspects	Difficult	Co-emergence
Process Philosophy	Philosophical	Experience or becoming as fundamental to reality	Events, processes, and organisms	Not directly	Co-emergence or consciousness first
Emergentism	Philosophical / scientific	An emergent property of complex organization	Complex biological systems, especially brains	Indirectly	Life first

This table shows that scientific theories tend to focus on mechanisms of access, integration, prediction, attention, recurrence, or biological organization. Philosophical and consciousness-first theories focus more directly on the ontological status of consciousness and the hard problem.

C.3 Origin-of-Life Models and Consciousness Implications

Table C.2: Origin-of-life models and consciousness implications.

Model	Key mechanism	Implies agency?	Consciousness relevance	Implication for central question
RNA World	Self-replicating RNA molecules capable of heredity and catalysis	Minimal	Low; mainly explains replication and heredity	Life first
Metabolism-First	Self-sustaining chemical cycles precede genetic replication	Moderate	Moderate; emphasizes self-maintaining organization	Life first or co-emergence
Hydrothermal Vent Models	Energy gradients and mineral surfaces drive prebiotic chemistry	Low to moderate	Moderate; energy flow and structure may support self-organization	Life first
Protocell Models	Lipid compartments create boundaries and internal environments	Moderate	Moderate to high; introduces self/non-self boundary	Life first or co-emergence
Autocatalytic Sets	Molecular networks collectively catalyze their own formation	Moderate	Moderate; systemic self-maintenance resembles minimal agency	Co-emergence-friendly
Hypercycle Models	Linked replicators support cooperative molecular evolution	Moderate	Moderate; early coordination and information cycling	Life first
Dissipative Structures	Energy flow maintains order far from equilibrium	Moderate	Moderate; connects life to thermodynamic self-organization	Co-emergence-friendly
Assembly Theory	Molecular complexity measured by assembly steps	Low	Low to moderate; useful for detecting life-like complexity	Life first
Autopoiesis	Systems continuously produce and maintain themselves	High	High; directly connects life, boundary, agency, and meaning	Co-emergence
Biosemitotic Models	Life begins with signs, codes, and meaning-making	High	High; suggests life requires interpretation or significance	Co-emergence or consciousness first
Consciousness-Enabled Life	Consciousness acts as organizing principle or condition	High, but speculative	High, but difficult to test	Consciousness first
Taheri's T-Consciousness View	Consciousness fields organize or manifest matter into living form	High, but non-material and speculative	High; treats life as an expression of prior consciousness	Consciousness first
Artificial Life Models	Life-like systems created in digital, robotic, or synthetic media	Variable	Depends on embodiment, self-maintenance, and integration	Tests life-consciousness relationship

Origin-of-life models usually explain biological organization without directly explaining subjective experience. Models that emphasize autopoiesis, biosemiotics, and consciousness-enabled life are more relevant to the central question because they connect life with agency, meaning, or consciousness.

C.4 Where Each Theory Places the Transition

Table C.3: Where major theories place the transition to consciousness.

Theory	Transition point	Gradual or sudden?	Earliest plausible conscious entity	Main challenge
Global Workspace Theory	Emergence of global broadcast architecture	Threshold-like	Complex-brained animals	Explains access better than experience itself
Global Neuronal Workspace	Neural ignition across large-scale networks	Threshold-like	Animals with advanced integrative brains	May exclude non-reportable or simpler consciousness
Integrated Information Theory	Presence or sufficient degree of integrated information	Gradual	Possibly any integrated system	Counterintuitive scope and measurement difficulties
Free Energy Principle	Self-maintaining predictive regulation	Gradual	Potentially all living systems, depending on interpretation	Free-energy minimization may not equal experience
Predictive Processing	Deep embodied prediction and model integration	Gradual	Organisms with complex predictive nervous systems	Prediction alone may not explain feeling
Recurrent Processing Theory	Recurrent neural feedback	Threshold-like	Animals with recurrent sensory circuits	Feedback may be necessary but not sufficient
Attention Schema Theory	Construction of a model of attention	Threshold-like	Systems with attention self-models	May explain reports of consciousness more than experience
Higher-Order Theory	Mental states represented by higher-order states	Threshold-like	Animals capable of self-monitoring	May over-intellectualize consciousness
Biological Naturalism	Biological brain processes with the right causal powers	Threshold-like	Biological animals with suitable brains	Does not specify exact mechanism
Orch-OR	Quantum coherence and objective reduction in microtubules	Event-based / threshold-like	Organisms with suitable microtubule organization	Empirical support remains controversial
Panpsychism	Consciousness or proto-consciousness always present	Not applicable	All matter or all physical entities	Combination problem
Cosmopsychism	Cosmic consciousness precedes individuals	Not applicable	Universe as primary subject	Decomposition problem
Idealism	Consciousness is primary reality	Not applicable	Universal consciousness	Explaining physical regularity and individuation
Taheri's T-Consciousness Framework	Transition is reframed as manifestation within T-Consciousness	Manifestation rather than emergence	Life as expression within consciousness fields	Mechanism and empirical testability remain unclear
Emergentism	Sufficient biological or neural complexity	Threshold-like or gradual	Complex nervous systems	The hard transition remains unclear
Co-emergence	Living self-organization becomes meaningful and integrated	Gradual	Minimal living systems or early sentient organisms	Needs sharper criteria

This comparison shows why the “hard transition” remains difficult. Life-first theories must identify when biological or neural complexity becomes experience. Consciousness-first theories avoid this transition but inherit other problems, especially mechanism and testability.

C.5 Philosophical Frameworks Compared

Table C.4: Major philosophical frameworks compared.

Framework	Mind-matter view	Life-consciousness view	Thinkers	Strength	Limitation
Substance Dualism	Mind and body are separate substances	Separate but interacting questions	Descartes	Preserves the reality of mind	Interaction problem
Physicalism	Mind is physical or fully dependent on physical processes	Life first	Smart, Armstrong, Dennett	Fits scientific naturalism	Hard problem and explanatory gap
Reductive Materialism	Consciousness reduces to brain states	Life first	Identity theorists	Parsimonious	May neglect subjective experience
Non-Reductive Physicalism	Consciousness depends on but is not reducible to physical states	Life first	Various contemporary philosophers	Allows higher-level properties	Risk of unclear causal status
Functionalism	Mind is functional organization	Substrate-independent	Putnam, Fodor	Supports machine consciousness	Chinese Room and simulation objections
Biological Naturalism	Consciousness is biological and brain-based	Life first	Searle	Takes biology and consciousness seriously	Mechanism remains unclear
Panpsychism	Mind-like properties exist in all matter	Consciousness first or co-emergence	Strawson, Goff, Chalmers	Avoids emergence from non-conscious matter	Combination problem
Cosmopsychism	Cosmic mind is fundamental	Consciousness first	Shani, Goff, Nagasawa, Wager	Avoids micro-combination problem	Decomposition problem
Idealism	Only mind or consciousness ultimately exists	Consciousness first	Berkeley, Kastrup	Dissolves the hard problem	Explaining shared physical reality
Taheri's T-Consciousness Framework	Consciousness is primary, non-material, and not produced by matter	Consciousness first; life emerges through or within consciousness fields	Mohammad Ali Taheri	Provides a direct consciousness-first account of life and mind	Mechanism and empirical testability remain unclear
Neutral Monism	Mind and matter arise from a deeper neutral reality	Co-emergence	Russell, James	Avoids strict dualism and reductionism	Nature of the neutral base unclear
Dual-Aspect Monism	One reality has mental and physical aspects	Co-emergence	Spinoza, Pauli-Jung interpretations	Integrates mind and matter	Difficult to test
Process Philosophy	Reality is process, event, and becoming	Co-emergence	Whitehead	Integrates life, experience, and becoming	Metaphysically complex
Enactivism	Mind arises through embodied action	Co-emergence	Varela, Thompson, Di Paolo	Connects life, cognition, and meaning	Consciousness boundary remains unclear
Autopoietic Theory	Life is self-producing organization	Co-emergence	Maturana, Varela	Strong account of living selfhood	Does not by itself prove consciousness

Philosophical frameworks differ in what they treat as fundamental. Some begin with matter, some with consciousness, and others with a deeper neutral or process-based reality. These starting points strongly shape how each theory interprets life and consciousness.

C.6 Life-First, Consciousness-First, and Co-Emergence Compared

Table C.5: Life-first, consciousness-first, and co-emergence positions compared.

Position	Basic claim	Best explains	Struggles with	Ethical implication
Life first	Life emerged first; consciousness evolved later	Neuroscience, evolution, animal consciousness, brain dependence	Hard problem, hard transition, subjective experience	Moral concern focused on sentient animals with suitable nervous systems
Consciousness first	Consciousness is fundamental; life emerges within it or from it	The hard problem, contemplative traditions, idealist metaphysics	Mechanism, testability, combination or decomposition	Broader moral circle, possibly including nature or matter
Taheri's T-Consciousness	Consciousness is a fundamental non-material reality; life manifests through consciousness fields	Consciousness as organizing principle, non-material field ontology, life as expression	Scientific mechanism, testability, relation to physical chemistry	Broad moral and metaphysical implications; life is not merely material organization
Co-emergence	Life and consciousness are deeply linked through self-organization and meaning	Continuity between life, cognition, agency, and sentience	Precision and falsifiability	Layered moral concern across life, sentience, and ecosystems
Strict physicalism	Consciousness is fully physical and produced by matter	Scientific parsimony and causal closure	Phenomenal experience	Narrower moral focus based on measurable sentience
Panpsychist continuum	Consciousness or proto-consciousness is widespread	Avoids sudden emergence	Over-attribution and combination problem	Very broad but graded moral concern
Biological naturalism	Consciousness requires biological brains	Brain dependence and living embodiment	AI and non-biological possibilities	Moral concern restricted to biological conscious beings
Functionalism	Consciousness depends on functional organization	AI possibility and multiple realizability	Simulation vs realization	Artificial systems may eventually matter morally

This table compares the three main families of answers developed throughout the book. Life-first theories are strongest scientifically, consciousness-first theories are strongest metaphysically, and co-emergence theories attempt to bridge life, cognition, and experience.

C.7 Artificial Consciousness Models Compared

Table C.6: Artificial consciousness models compared.

View	Can AI be conscious?	Required conditions	Main argument	Main objection
Computational functionalism	Yes	Correct computational organization	Consciousness depends on function, not substrate	Symbol manipulation may not create understanding
Global Workspace view	Possibly	Artificial global broadcast architecture	Consciousness is global access	Access may not equal experience
IIT-based view	Possibly	High intrinsic integrated causal structure	Consciousness depends on integrated information	Digital systems may have low integrated causal structure
Attention Schema view	Possibly	A model of attention and self-monitoring	Consciousness is an attention model	May explain self-report, not feeling
Predictive processing view	Possibly	Embodied prediction, action, and self-modeling	Consciousness arises from active inference	Prediction may not imply experience
Biological naturalism	No, unless biological	Biological causal powers of living brains	Simulation is not realization	May be too restrictive
Artificial life view	Possibly	Life-like self-maintenance, embodiment, adaptation	Consciousness may require life-like organization	Boundary between simulation and realization unclear
Panpsychist view	Possibly	Matter or systems with experiential aspects	Consciousness is widespread	Hard to test and specify
Idealist view	Possibly	Artificial system as appearance within consciousness	Consciousness is fundamental	Difficult to determine individual machine experience

Artificial consciousness tests whether biological life is necessary for experience. If consciousness depends only on functional organization, then AI could in principle be conscious. If consciousness depends on living embodiment, then artificial life may be more relevant than ordinary computation. If consciousness is fundamental, AI may raise a different question: how consciousness becomes individualized or expressed through artificial systems.

C.8 Ethical Implications Compared

Table C.7: Ethical implications of different consciousness positions.

Theory / position	Likely moral circle	Strongest ethical concern	Risk
Human exceptionalism	Humans only	Human rights and dignity	Excludes animals and non-verbal beings unfairly
Mammal / bird sentience view	Mammals and birds, possibly some others	Animal welfare	May exclude fish, insects, cephalopods, or crustaceans
Broad animal sentience view	Vertebrates plus many invertebrates	Reducing suffering across animal life	Difficult evidence boundaries
Basal cognition view	Many living systems deserve respect	Respect for life and minimal agency	May confuse cognition with sentience
Plant / fungal intelligence view	Plants and fungi may have non-sentient value	Ecological humility and restraint	Anthropomorphism
Ecosystem ethics	Ecosystems, species, rivers, habitats	Environmental protection and relational responsibility	Collective moral status is difficult to define
Panpsychism	Very broad, possibly all matter	Avoiding dismissal of widespread interiority	Moral inflation
Biological naturalism	Biological conscious organisms	Animal and human sentience	Excludes artificial consciousness
Functionalism	Biological and artificial conscious systems	Possible AI welfare	Over-attribution to simulations
Taheri's T-Consciousness view	Broad metaphysical concern for life and consciousness	Life may have deeper consciousness-based significance	Difficult to translate metaphysical concern into policy
Precautionary approach	Expands with uncertainty and evidence	Avoiding unrecognized suffering	Practical limits and moral overload
Parsimony approach	Narrower moral status	Avoiding unsupported attribution	Under-recognition of sentience

The ethical implications of consciousness theories are not secondary. They shape how we treat animals, ecosystems, artificial systems, and life itself. The wider the theory places consciousness or sentience, the more expansive its moral implications become.

C.9 Methodological Approaches Compared

Table C.8: Methodological approaches in consciousness research.

Method	What it studies	Strength	Limitation
Behavioural observation	Action, learning, response, adaptation	Works across many organisms	Behaviour may be unconscious
Verbal report	Human subjective experience	Direct access to reported experience	Limited to beings capable of report
Neuroscience	Brain activity and neural mechanisms	Strong for humans and animals with brains	Correlation is not explanation
Brain imaging	Large-scale neural patterns	Non-invasive evidence of conscious states	Limited resolution and interpretation
Anaesthesia studies	Loss and recovery of consciousness	Identifies necessary conditions	Does not explain experience itself
Comparative cognition	Cross-species intelligence and behaviour	Evolutionary perspective	Anthropomorphism risk
Basal cognition research	Non-neural sensing and adaptation	Expands cognition beyond brains	Does not prove consciousness
Phenomenology	First-person structure of experience	Takes experience seriously	Subjective and difficult to verify
Neurophenomenology	Links first-person and neural data	Bridges subjective and objective methods	Requires careful training and methods
Mathematical modelling	Formal structure of theories	Clarifies assumptions and predictions	Model may not capture experience
Artificial intelligence experiments	Machine cognition and self-modeling	Tests substrate-independence	Intelligence may not equal sentience
Synthetic biology / artificial life	Life-like organization	Tests boundaries of life	Consciousness remains difficult to infer

No single method is sufficient. Consciousness research requires behavioural, neural, phenomenological, comparative, formal, and ethical approaches. The central challenge is that consciousness is known directly from within but studied scientifically from without.

C.10 Theory Strengths and Weaknesses at a Glance

Table C.9: Theory strengths and weaknesses at a glance.

Theory	Major strength	Major weakness
GWT	Explains conscious access and report	May not explain phenomenal experience
IIT	Addresses intrinsic structure of experience	Difficult to test and may imply too much consciousness
FEP	Connects life, cognition, and regulation	Too broad if applied to all living systems
Predictive Processing	Explains perception as active inference	Prediction alone may not produce experience
AST	Explains why systems model themselves as aware	May explain belief in consciousness rather than consciousness
RPT	Offers a clear neural mechanism through feedback	Feedback alone may not be sufficient
Biological Naturalism	Respects biological basis of consciousness	Does not fully explain why biology feels
Orch-OR	Connects consciousness to fundamental physics	Highly controversial and speculative
Panpsychism	Avoids emergence from non-conscious matter	Combination problem
Cosmopsychism	Avoids micro-combination problem	Decomposition problem
Idealism	Makes consciousness primary and avoids the hard problem	Explaining physical regularity and shared reality
Russellian Monism	Bridges physics and consciousness through intrinsic nature	Difficult to specify or test
Taheri's T-Consciousness Framework	Offers a clear consciousness-first model in which life and matter arise within a prior non-material consciousness order	Requires clearer mechanisms, empirical criteria, and testable predictions
Dual-Aspect Monism	Avoids strict dualism and reductionism	Often too general
Process Philosophy	Integrates experience, life, and becoming	Hard to formalize scientifically
Co-emergence	Explains continuity between life and mind	Needs sharper empirical criteria

C.11 Summary

These comparison tables show that no single theory currently explains all dimensions of the life-consciousness relationship. Scientific theories often perform well on testability but struggle with subjective experience. Philosophical theories often address the hard problem more directly but struggle with mechanism and empirical evidence. Biological and enactive theories offer promising bridges, especially when they connect life, self-organization, meaning, and embodied agency.

The central pattern is clear:

- life-first theories are strongest on biology and neuroscience;
- consciousness-first theories are strongest on the hard problem and metaphysics;
- Taheri's T-Consciousness framework provides a contemporary consciousness-first model centered on non-material consciousness fields;
- co-emergence theories are strongest on continuity between life, cognition, and mind;
- artificial consciousness tests whether life is necessary for experience;
- ethics forces theory into practical responsibility.

The unresolved question remains:

Which framework can explain life, consciousness, experience, embodiment, meaning, and moral status without reducing one dimension to another?

Appendix D

Key Philosophical Terms

D.1 Purpose

This appendix provides a reference guide to philosophical terminology used throughout the book. It is written for readers from non-philosophy backgrounds and is intended to clarify key concepts that appear in discussions of life, consciousness, mind, matter, emergence, ethics, and testability.

The terms below are not meant to provide exhaustive philosophical treatments. Instead, they offer practical working definitions that help readers follow the arguments developed across the book.

D.2 Ontology

Ontology is the study of what exists. It asks what kinds of things are real and what categories of reality should be included in our worldview.

An ontological question is not only about whether something can be measured, but about what kind of thing it is. Is consciousness a real feature of the world, or is it an illusion? Is life merely chemistry, or does it introduce a new level of organization? Are minds, values, meanings, and experiences real in the same way physical objects are real?

The central question of this book is fundamentally ontological. To ask whether life came before consciousness, or consciousness before life, is to ask what kind of reality consciousness has. If consciousness is only a biological product, then life comes first. If consciousness is fundamental, then life may arise within consciousness. If life and consciousness co-emerge, then both may be aspects of a deeper process of self-organization.

D.3 Epistemology

Epistemology is the study of knowledge. It asks how we know what we claim to know, what counts as evidence, and what the limits of knowledge are.

In this book, epistemology appears whenever we ask whether consciousness can be detected in other beings. Can we know whether bacteria are conscious? Can we know whether insects feel pain? Can we know

whether an artificial intelligence system has experience? What kinds of evidence should count: behaviour, brain activity, self-report, biological similarity, or theoretical prediction?

Epistemology is especially important in Chapter 17 on limits and testability. Consciousness is known directly from the first-person perspective, but the consciousness of others is inferred indirectly. This creates a permanent tension between subjective certainty and objective evidence.

D.4 Metaphysics

Metaphysics is the study of the fundamental nature of reality. It includes ontology, but it is broader. Metaphysics asks about existence, causation, time, identity, possibility, mind, matter, and the basic structure of the world.

The life-consciousness question is metaphysical because it asks about the deepest relationship between living organization and subjective experience. Is consciousness produced by matter? Is matter an appearance within consciousness? Are mind and matter two aspects of one reality? Is life a purely physical process or a meaningful, self-organizing process with an inner aspect?

Scientific evidence matters deeply to the question, but science alone may not settle the metaphysical interpretation. The same evidence about brain activity can be interpreted through physicalism, emergentism, dual-aspect monism, or idealism. Metaphysics helps clarify what those interpretations assume.

D.5 Dualism

Dualism is the view that mind and matter are fundamentally different. The most famous form is substance dualism, associated with René Descartes, which holds that mind and body are two distinct kinds of substance.

Substance dualism treats consciousness as non-physical and the body as physical. This view preserves the reality of consciousness but creates the interaction problem: how can a non-physical mind affect a physical body?

Property dualism is a weaker form. It holds that there is one physical substance, but it has both physical and mental properties. Consciousness is not a separate substance, but it may not be reducible to physical properties.

Dualism is relevant because if mind and matter are fundamentally separate, then life and consciousness may be separate questions. Life may be biological, while consciousness may belong to a different order of reality. Many contemporary theories avoid substance dualism but still preserve some distinction between physical process and subjective experience.

D.6 Monism

Monism is the view that there is only one fundamental kind of reality. Different forms of monism disagree about what that reality is.

Physicalism is a form of monism in which everything is ultimately physical. Consciousness must therefore be physical, caused by physical processes, or dependent on physical organization.

Idealism is another form of monism in which reality is ultimately mental or experiential. Matter is not denied, but it is understood as appearance, representation, or structure within consciousness.

Neutral monism proposes that the fundamental reality is neither mental nor physical as we ordinarily understand those terms. Mind and matter are both expressions or aspects of a deeper neutral basis.

Monism is important because it tries to avoid the division between mind and matter. The question becomes not how two different substances interact, but how one reality can appear as both physical life and conscious experience.

D.7 Phenomenology

Phenomenology is the study of experience from the first-person perspective. It examines how things appear to consciousness: perception, embodiment, time, selfhood, emotion, attention, and meaning.

Phenomenology is associated with thinkers such as Edmund Husserl, Martin Heidegger, and Maurice Merleau-Ponty. Husserl emphasized careful description of experience. Heidegger emphasized being-in-the-world. Merleau-Ponty emphasized the lived body as central to perception and consciousness.

Phenomenology should not be confused with phenomenal consciousness, though the two are related. Phenomenal consciousness refers to subjective experience itself. Phenomenology is a method or tradition for studying the structure of that experience.

This book uses phenomenological ideas when discussing the first-person character of consciousness and the limits of purely third-person explanation.

D.8 Intentionality

Intentionality means the “aboutness” or directedness of mental states. A belief is about something. A perception is of something. A desire is directed toward something. A fear has an object, even if that object is imagined.

Intentionality is central to philosophy of mind because consciousness is rarely empty. It is usually directed toward a world: a sound, a memory, a pain, a possibility, a goal, or a meaning.

The concept becomes difficult when applied to simple organisms. Does a bacterium have intentionality when it moves toward nutrients? Does a plant have intentionality when it grows toward light? Or are these merely biochemical responses?

The answer depends on how broadly intentionality is defined. In a minimal biological sense, living systems are world-directed because conditions matter to them. In a richer mental sense, intentionality may require representation, awareness, or experience.

D.9 Supervenience

Supervenience is a dependency relation. To say that consciousness supervenes on the physical means that there can be no change in consciousness without some change in the physical state.

For example, a physicalist might argue that if two beings are physically identical in every way, they must also be mentally identical. Their conscious experiences cannot differ unless their physical states differ.

Supervenience is important because it allows philosophers to describe dependence without always claiming reduction. A person may say that consciousness depends on the brain without saying consciousness is nothing but brain activity.

Zombie arguments challenge physicalist interpretations of supervenience. If it is conceivable that a being could be physically identical to a conscious human but have no experience, then perhaps consciousness does not logically supervene on the physical. Whether this conceivability argument succeeds remains debated.

D.10 Teleology

Teleology refers to explanation in terms of purpose, goal, or end.

In Aristotelian philosophy, natural things were often understood as having purposes or ends. An acorn tends toward becoming an oak. The eye is for seeing. Living beings develop according to their forms and ends.

Modern biology generally avoids strong teleology. It explains function through natural selection rather than built-in purpose. The eye is “for” seeing because evolution selected visual capacities that supported survival and reproduction, not because nature consciously planned eyes.

However, life still appears goal-directed in a limited sense. Organisms maintain themselves, repair damage, seek nutrients, avoid threats, and reproduce. This goal-directedness is often called teleonomy: apparent purpose produced by natural processes.

Teleology is relevant to consciousness because we can ask whether consciousness is for something. Did consciousness evolve for flexible action, social interaction, pain avoidance, planning, or self-modeling? Or is consciousness not functionally necessary at all?

D.11 Reductionism

Reductionism is the attempt to explain higher-level phenomena in terms of lower-level components. For example, life may be explained in terms of chemistry, chemistry in terms of physics, and consciousness in terms of neuroscience.

Reductionism has been highly successful in science. Molecular biology, genetics, and neuroscience all rely on the idea that complex systems can be understood through their parts and mechanisms.

However, consciousness raises special challenges for reductionism. Even a complete description of neurons, synapses, and brain networks may seem to leave out what experience feels like. This is the explanatory gap.

Anti-reductionism does not necessarily deny science. It argues that some phenomena require higher-level descriptions that cannot be fully replaced by lower-level ones. Life, mind, meaning, and consciousness may depend on physical processes without being fully captured by physical descriptions alone.

D.12 Emergence

Emergence refers to the appearance of higher-level properties from lower-level interactions.

Weak emergence occurs when higher-level patterns arise from lower-level processes and are, in principle, explainable by them. For example, the movement of a flock emerges from the behaviour of individual birds, even though the flock pattern is not located in any single bird.

Strong emergence occurs when higher-level properties are genuinely novel and not fully deducible from lower-level descriptions. Consciousness is often discussed as a candidate for strong emergence because subjective experience seems difficult to derive from physical processes alone.

Emergence is central to life-first theories. If consciousness emerges from life, then we must explain what kind of emergence is involved. Is consciousness like wetness, a higher-level property of organized matter? Or is it something more radical?

D.13 Qualia

Qualia are the subjective qualities of experience: the redness of red, the painfulness of pain, the taste of sweetness, the sound of music, the feeling of warmth.

Qualia are central to the hard problem because they seem difficult to explain in purely physical or functional terms. A scientist might describe the wavelength of red light, the retinal response, and the neural processing involved in colour perception. But this does not obviously explain what red looks like from the inside.

Some philosophers argue that qualia are real and irreducible. Others argue that the concept is confused or that qualia can eventually be explained through neuroscience and cognitive science.

In this book, qualia represent the first-person dimension that any complete theory of consciousness must address.

D.14 Zombie

A philosophical zombie is a being physically and behaviourally identical to a conscious human but lacking subjective experience. It acts like a conscious person, speaks like one, and has the same physical structure, but there is nothing it is like to be that being.

David Chalmers uses the zombie argument to challenge physicalism. If zombies are conceivable, then consciousness may not be fully explained by physical facts alone.

Critics argue that zombies may be conceivable only because we do not fully understand consciousness. They may be logically incoherent if consciousness necessarily follows from the right physical organization.

Zombie arguments matter for artificial intelligence, animal consciousness, and the hard transition. A system might behave as if conscious, but how would we know whether experience is present?

D.15 The Explanatory Gap

The explanatory gap refers to the difficulty of explaining how physical processes give rise to conscious experience. Joseph Levine introduced the term in 1983 to describe the gap between physical explanation and subjective feeling.

For example, neuroscience may explain that pain involves certain neural pathways, brain regions, and physiological responses. But why should that activity feel painful? Why should it have any subjective character at all?

The explanatory gap is closely related to the hard problem, but they are not identical. The explanatory gap names the difficulty in moving from physical explanation to experience. The hard problem asks why experience exists in the first place.

This gap is one of the main reasons some philosophers turn to panpsychism, idealism, dual-aspect monism, or consciousness-first theories.

D.16 The Hard Problem

The hard problem of consciousness asks why and how physical processes give rise to subjective experience. It is “hard” because explaining behaviour, attention, memory, or neural processing does not automatically explain why there is something it is like to undergo those processes.

The hard problem is central to the life-consciousness question. If life produced consciousness, then a life-first theory must explain how biological organization crossed into experience. If consciousness is fundamental, then the hard problem may be dissolved or reframed, but new problems arise.

D.17 The Hard Transition

The hard transition is the problem of identifying when, where, and how biological information processing first became subjective experience.

It differs from the hard problem. The hard problem asks why experience exists at all. The hard transition asks when experience first appeared in evolutionary history.

The hard transition is especially important for life-first theories. If consciousness emerged from life, then there must have been a transition from non-conscious life to conscious life. The challenge is to explain what changed and why that change produced experience.

D.18 Intrinsic Nature

Intrinsic nature refers to what something is in itself, apart from its external relations, structure, or behaviour.

Russellian monism uses this concept to argue that physics describes the structure and relations of matter but not its inner nature. Consciousness may provide a clue to the intrinsic nature of physical reality.

This concept is important because it opens a possible middle path between physicalism and dualism. Consciousness may not be separate from matter, but may be the inner aspect of what physics describes externally.

D.19 Ontological Primitive

An ontological primitive is something treated as fundamental and not explained in terms of something else.

In physicalism, physical reality is often treated as the ontological primitive. In idealism, consciousness may be the ontological primitive. In some information-based theories, information or structure may be treated as primitive.

Taheri's T-Consciousness framework treats T-Consciousness as an ontological primitive: a fundamental non-material reality that is not produced by matter, energy, frequency, or neural activity.

D.20 Taheri's T-Consciousness Framework

Taheri's T-Consciousness framework is a contemporary consciousness-first framework associated with Mo-hammad Ali Taheri. It proposes that consciousness is primary, non-material, and not produced by matter or the brain.

In this framework, T-Consciousness is not understood as ordinary mental activity, neural activity, energy, frequency, or physical force. It is treated as a fundamental consciousness reality through which matter, life, and individual awareness become organized or manifest.

Philosophically, this places Taheri's framework near consciousness-first theories such as idealism, cosmopsychism, and some forms of dual-aspect thought, while remaining distinct from standard panpsychism. It is not simply the claim that every particle has its own tiny mind. Rather, it proposes a broader consciousness field or network within which life and mind are expressed.

Its relevance to this book is direct. Taheri's framework supports the view that consciousness does not emerge from life. Instead, life emerges within a prior consciousness-based order. This reframes the central question from "How did life generate consciousness?" to "How did consciousness organize into living form?"

The main philosophical challenge is testability. If T-Consciousness is non-material and not reducible to ordinary physical mechanisms, what would count as evidence for or against it? This challenge is shared with many consciousness-first theories.

D.21 Consciousness Field

A consciousness field is a proposed non-material field-like reality through which consciousness organizes, relates, or manifests matter, life, and mind.

In Taheri's framework, consciousness fields are not the same as physical fields in physics. They are not electromagnetic fields, quantum fields, frequencies, or measurable energies in the ordinary scientific sense. They are understood as non-material organizing realities.

The concept is relevant to consciousness-first models because it offers a way to describe how consciousness might be primary while still relating to material and biological organization. However, the term remains speculative unless it can be connected to clearer mechanisms, predictions, or forms of evidence.

D.22 T-Consciousness

T-Consciousness is Taheri's term for a fundamental, non-material form of consciousness that is not reducible to matter, energy, frequency, neural activity, or ordinary information.

Within Taheri's framework, T-Consciousness is treated as primary rather than emergent. It is not produced by brains. Instead, brains and living systems are understood as expressions, receivers, interfaces, or manifestations within a broader consciousness order.

In the context of this book, T-Consciousness represents a consciousness-first model. It challenges the standard life-first view and suggests that life may arise through or within consciousness rather than producing consciousness.

D.23 Consciousness-First

Consciousness-first refers to theories that treat consciousness as fundamental, primary, or prior to matter and life.

These theories include idealism, panpsychism, cosmopsychism, some forms of dual-aspect monism, process philosophy, and Taheri's T-Consciousness framework. They differ in detail, but they share the view that consciousness cannot be fully explained as a late product of non-conscious matter.

The strength of consciousness-first theories is that they address the hard problem directly. Their weakness is that they often struggle with mechanism, testability, and the relationship between consciousness and physical science.

D.24 Co-Emergence

Co-emergence is the view that life and consciousness are not entirely separate stages, but deeply linked processes.

In this book, co-emergence refers to the possibility that life and consciousness arise together through self-organization, information, meaning, agency, and embodied perspective. This does not necessarily mean that all life is fully conscious. Rather, it suggests that the roots of consciousness may be present in the same processes that make life alive.

Co-emergence offers a middle path between strict life-first emergentism and broad consciousness-first metaphysics. It treats consciousness as rooted in living organization without reducing it to mechanical function.

D.25 Moral Status

Moral status refers to whether an entity matters ethically for its own sake.

Sentient beings are usually considered to have moral status because they can experience pleasure, pain, suffering, or well-being. The more widely consciousness is distributed, the broader moral status may need to become.

The concept is important in discussions of animals, ecosystems, artificial intelligence, and simple organisms. If we cannot determine consciousness with certainty, moral status may require precaution, humility, and layered ethical frameworks.

D.26 Personhood

Personhood refers to the status of being a person, often associated with self-awareness, agency, rationality, moral responsibility, social recognition, or legal rights.

Personhood is not identical to consciousness. A newborn infant may be conscious without having full reflective personhood. An animal may be sentient without being a legal person. A river may be granted legal personhood for protection without being conscious in the human sense.

The distinction matters because moral status, legal personhood, and consciousness do not always overlap.

D.27 Reduction vs Explanation

Reduction and explanation are related but not identical. A phenomenon may be explained by lower-level processes without being eliminated.

For example, life can be explained through chemistry without saying that organisms are “nothing but” chemicals. Similarly, consciousness might be explained through biology without denying the reality of experience.

This distinction helps avoid false choices. A theory can respect neuroscience while also recognizing that subjective experience requires its own conceptual vocabulary.

D.28 First-Person Perspective

The first-person perspective is experience as lived from within. It is the perspective of “I feel,” “I see,” “I suffer,” or “I am aware.”

Consciousness is unusual because it is known most directly from the first-person perspective but studied scientifically from the third-person perspective. This creates many of the methodological challenges discussed in the book.

D.29 Third-Person Perspective

The third-person perspective is the external, observational standpoint used in science. It studies behaviour, brain activity, physiology, chemistry, and measurable processes.

Third-person methods are essential for studying life and consciousness, but they may not fully capture subjective experience. A complete approach to consciousness may need to integrate first-person, second-person, and third-person methods.

D.30 Second-Person Perspective

The second-person perspective involves direct engagement with another subject. It includes dialogue, empathy, clinical interviews, social interaction, and interpretive understanding.

Second-person methods are important in consciousness studies because they help bridge the gap between external measurement and lived experience. They are especially relevant in medicine, disability studies, psychotherapy, and developmental research.

This appendix provides a working vocabulary for the major philosophical concepts used throughout the book.

Appendix E

Consciousness Taxonomy

E.1 Purpose

This appendix provides a structured classification of the types, levels, dimensions, and theories of consciousness referenced throughout the book. The goal is not to impose a final definition of consciousness, but to give readers a map of the major ways consciousness can be classified.

Because the book moves across philosophy, biology, neuroscience, artificial intelligence, origin-of-life studies, and consciousness-first frameworks, the word “consciousness” is used in several related but distinct ways. This taxonomy helps clarify those uses.

This taxonomy is especially useful because different theories often use the same word in different ways. For example, consciousness may mean subjective experience, access to information, self-awareness, biological sentience, cosmic mind, or a fundamental non-material reality such as T-Consciousness. A clear taxonomy helps prevent these meanings from being collapsed into one another.

E.2 Types of Consciousness

E.2.1 By Philosophical Category

Table E.1: Types of consciousness by philosophical category.

Type	Definition	Key theorists / frameworks
Phenomenal consciousness	Subjective experience; what it is like to see, feel, think, or exist from the first-person perspective	Nagel, Chalmers
Access consciousness	Information available for reasoning, report, decision-making, and behavioural control	Block, Baars, Dehaene
Self-consciousness	Awareness of oneself as a subject or as being aware	Rosenthal, higher-order theorists
Reflective consciousness	The capacity to think about one's own thoughts, beliefs, identity, or experience	Higher-order theories, metacognition research
Pre-reflective consciousness	Immediate lived awareness before explicit self-reflection	Phenomenology, Merleau-Ponty
Minimal consciousness	Basic subjective awareness without complex self-reflection or language	Feinberg, Mallatt, minimal cognition debates
Proto-consciousness	A possible precursor or minimal form of experience below full consciousness	Panpsychism, basal cognition, some evolutionary theories
Sentience	The capacity for felt experience, especially pleasure, pain, suffering, or well-being	Bentham, Singer, animal ethics
Cosmic consciousness	Universe-level or reality-level consciousness from which individual minds may derive	Cosmopsychism, some contemplative traditions
T-Consciousness	Taheri's proposed fundamental non-material consciousness reality that is not produced by matter, energy, frequency, or the brain	Mohammad Ali Taheri
Consciousness field	A proposed non-material field-like organizing reality through which consciousness relates to or manifests matter, life, and mind	Taheri's T-Consciousness framework
Artificial consciousness	Possible consciousness in artificial systems such as AI, robots, or artificial life	Functionalism, AI consciousness debates

This table shows that consciousness is not a single simple category. Some uses of the term refer to lived experience, while others refer to information access, self-reflection, sentience, cosmic awareness, artificial systems, or consciousness-first metaphysics.

E.3 By Biological Level

Table E.2: Biological and artificial levels relevant to consciousness.

Level	Examples	Evidence or argument	Status
Molecular	Chemical reactions, proteins, microtubules	Speculative claims in Orch-OR and some quantum mind theories	Highly speculative
Cellular	Single cells, immune cells, developmental bioelectric networks	Sensing, memory, self-regulation, adaptive response	Cognition plausible; consciousness uncertain
Unicellular organisms	Bacteria, protists, Paramecium, Stentor	Chemotaxis, habituation-like behaviour, decision-making, environmental responsiveness	Basal cognition; consciousness debated
Multicellular non-neural life	Plants, fungi, slime molds	Signaling, learning-like behaviour, network optimization, resource allocation	Intelligence debated; consciousness uncertain
Simple nervous systems	Cnidarians, jellyfish, hydra	Nerve nets, sensory response, coordinated movement	Possible minimal sentience; uncertain
Invertebrates	Insects, crustaceans, cephalopods	Learning, flexible behaviour, nociception, problem-solving, neural integration	Increasing evidence for sentience in some groups
Vertebrates	Fish, amphibians, reptiles, birds, mammals	Central nervous systems, learning, pain behaviour, affect, perception	Strong evidence for many forms of consciousness
Primates and humans	Apes and humans	Self-recognition, language, theory of mind, first-person reports	Strongest evidence
Artificial systems	AI, robots, artificial agents	Language, learning, self-modeling, goal-directed behaviour	Unknown and debated
Artificial life	Digital organisms, synthetic cells, life-like simulations	Evolution, adaptation, self-maintenance, reproduction-like processes	Life-like status debated; consciousness uncertain

The biological taxonomy shows why the consciousness question becomes difficult at the lower levels of life. Evidence is strongest in humans and many vertebrates, increasingly plausible in some invertebrates, debated in plants and unicellular organisms, and highly speculative at the molecular level.

E.4 By Ontological Level

Table E.3: Ontological levels at which consciousness may be located.

Level	Description	Example theories
Brain-based consciousness	Consciousness exists only in organisms with suitable brains	Biological naturalism, GWT, RPT
Nervous-system consciousness	Consciousness requires nervous systems but not necessarily human-like brains	RPT, neurobiological theories, animal consciousness frameworks
Living-system consciousness	Consciousness or proto-consciousness may arise with living self-organization	FEP, autopoietic enactivism, co-emergence
Information-based consciousness	Consciousness depends on information integration or information architecture	IIT, computational theories
Function-based consciousness	Consciousness depends on functional organization, regardless of substrate	Functionalism, AST, GWT-inspired AI theories
Matter-based proto-consciousness	Matter has proto-conscious or experiential aspects	Panpsychism, Russellian monism
Cosmos-based consciousness	The universe as a whole is the primary conscious subject	Cosmopsychism
Consciousness-field ontology	Consciousness is a fundamental non-material reality expressed through consciousness fields	Taheri's T-Consciousness framework
Mind-only ontology	Consciousness or mind is the only ultimate reality	Idealism, analytic idealism

This table helps distinguish theories that locate consciousness in brains, nervous systems, life, information, function, matter, the cosmos, consciousness fields, or mind itself.

E.5 Levels and Dimensions

E.5.1 Graded Dimensions of Consciousness

Table E.4: Graded dimensions of consciousness.

Dimension	Lower end	Higher end	Relevance
Wakefulness	Deep sleep, coma, anaesthesia	Alert waking state	Measures arousal, not necessarily awareness
Awareness	No apparent awareness	Rich phenomenal experience	Core dimension of consciousness
Sentience	No evidence of feeling	Pain, pleasure, suffering, well-being	Central for moral status
Self-awareness	No self-model	Reflective awareness of oneself as a subject	Important in humans and some animals
Bodily awareness	Minimal bodily regulation	Rich embodied self-experience	Links consciousness to life and organismic regulation
Temporal depth	Immediate reaction	Memory, anticipation, planning	Important for complex cognition
Intentional complexity	Simple world-directedness	Nested beliefs, theory of mind, symbolic thought	Important for higher cognition
Integration	Fragmented processing	Unified field of experience	Important in IIT, GWT, and neurobiological theories
Reportability	No report possible	Verbal or symbolic report	Useful but not identical to consciousness
Agency	Passive response	Flexible, self-directed action	Bridges life, cognition, and consciousness
Meaning	Bare signal response	Organism-relative significance and interpretation	Central to biosemiotics and co-emergence

Consciousness is not only present or absent. It varies across dimensions such as wakefulness, awareness, selfhood, agency, integration, reportability, and meaning. This is why a graded model is often more useful than a simple yes-or-no model.

E.6 States of Consciousness

Table E.5: States of consciousness.

State	Wakefulness	Awareness	Self-awareness	Notes
Deep sleep	Low	Low or absent	None	Usually low consciousness, though some dreaming may occur in sleep states
Dreaming	Low to moderate	Moderate to high	Variable	Rich experience can occur without ordinary waking control
Lucid dreaming	Low to moderate	High	High	Dreamer recognizes the dream state
Waking state	High	High	Variable to high	Ordinary conscious life
Flow state	High	High	Often reduced	Strong task absorption with reduced reflective self-consciousness
Meditation	Variable	Often high	Variable	May involve altered attention, selfhood, or awareness
Anaesthesia	Very low or absent	Absent or severely reduced	None	Used to study loss and recovery of consciousness
Coma	Very low or absent	None or minimal	None	Severe impairment of wakefulness and awareness
Vegetative state / unresponsive wakefulness	Wakefulness cycles may remain	No clear behavioural awareness	None apparent	Raises difficult diagnostic and ethical issues
Minimally conscious state	Variable	Partial or inconsistent	Limited	Some evidence of awareness remains
Locked-in syndrome	High	High	High	Consciousness preserved but motor output severely limited
Psychedelic states	High or altered	High and altered	Variable	May involve changes in self, perception, and meaning
Dissociative states	Variable	Variable	Altered	Self-experience and embodiment may fragment
Near-death experiences	Variable / medically complex	Reported as vivid by some individuals	Often high in reports	Interpretation remains debated

States of consciousness show that wakefulness and awareness are not the same. A dreaming person may have vivid experience with reduced wakefulness, while a locked-in patient may have high awareness but little or no motor expression.

E.7 Theory Classification

E.7.1 By Ontological Commitment

Table E.6: Theories classified by ontological commitment.

Category	Theories	Core claim	Relation to life-consciousness
Physicalist	GWT, RPT, AST, emergentism, many neuroscientific theories	Consciousness arises from physical processes	Life first
Biological naturalist	Searle, neurobiological theories	Consciousness is a biological phenomenon caused by brain processes	Life first
Functionalist	Computational functionalism, some AI consciousness theories	Consciousness depends on functional organization	Potentially substrate-independent
Informational	IIT, some information-based theories	Consciousness depends on integrated information or causal structure	Co-emergence or panpsychist implications
Enactive / autopoietic	Varela, Thompson, Di Paolo	Mind arises through living self-organization and embodied sense-making	Co-emergence
Panpsychist	Strawson, Goff, some readings of Russellian monism	Consciousness or proto-consciousness is fundamental and widespread	Consciousness first or co-emergence
Cosmopsychist	Shani, Nagasawa, Goff	The universe as a whole is conscious; individuals derive from it	Consciousness first
Idealist	Berkeley, Kastrup, analytic idealism	Consciousness is the only ultimate reality	Consciousness first
T-Consciousness framework	Taheri's consciousness-field ontology	T-Consciousness is fundamental, non-material, and not produced by the brain	Consciousness first
Dualist	Substance dualism, property dualism	Mind and matter are distinct substances or properties	Separates life and consciousness
Neutral monist	Russell, James, some dual-aspect views	Mind and matter arise from a deeper neutral reality	Co-emergence
Dual-aspect monist	Spinoza, Pauli-Jung interpretations	One reality has mental and physical aspects	Co-emergence
Process philosophical	Whitehead	Reality is composed of processes or occasions of experience	Co-emergence or consciousness first
Quantum speculative	Orch-OR, Stapp, Fisher, Bohm-inspired theories	Consciousness involves quantum or deeper physical processes	Co-emergence or consciousness first

This classification shows that theories of consciousness often differ less in evidence than in starting assumptions. A physicalist, an idealist, and a T-Consciousness theorist may all discuss consciousness, but they locate it at very different levels of reality.

E.8 By Mechanism

Table E.7: Theories classified by proposed mechanism.

Mechanism type	Theories	Key mechanism	Main question
Global access	GWT, Global Neuronal Workspace	Information broadcast across a cognitive system	Does access explain experience?
Recurrent processing	RPT, Edelman's re-entry	Neural feedback loops	Is recurrence sufficient for consciousness?
Higher-order monitoring	Higher-order thought and perception theories	A mental state becomes conscious when represented by another state	Does consciousness require self-monitoring?
Attention modeling	AST	The system models its own attention	Does a model of awareness create awareness?
Information integration	IIT	Integrated causal information, Phi	Can consciousness be mathematically measured?
Predictive regulation	Predictive processing, FEP	Prediction-error minimization and active inference	Does prediction become experience?
Biological embodiment	Biological naturalism, Damasio, enactivism	Living bodies, affect, regulation, and neural processes	Does consciousness require life?
Self-production	Autopoiesis	The system continuously produces and maintains itself	Is life already proto-cognitive?
Semiotic meaning	Biosemiotics	Sign interpretation and biological meaning	Does life require meaning?
Quantum physical	Orch-OR, Stapp, Fisher	Quantum coherence, collapse, or measurement	Does quantum physics explain experience?
Consciousness-field organization	Taheri's T-Consciousness framework	Non-material consciousness fields organize or manifest matter, life, and mind	Can this be connected to testable mechanisms?
Philosophical grounding	Panpsychism, idealism, cosmopsychism	Consciousness as fundamental reality	Does this explain or relocate the mystery?

Mechanism-based classification is useful because some theories are scientific models while others are metaphysical frameworks. Global Workspace Theory, Recurrent Processing Theory, and Attention Schema Theory propose mechanisms inside cognitive systems. Idealism, panpsychism, and Taheri's T-Consciousness framework begin from broader claims about the nature of reality.

E.9 By Relation to Life

Table E.8: Theories classified by relation between life and consciousness.

View	Life first?	Consciousness first?	Core interpretation
Strict life-first	Yes	No	Consciousness is a late product of biological evolution
Neural life-first	Yes	No	Consciousness appears only with nervous systems or brains
Biological naturalism	Yes	No	Consciousness is caused by biological brain processes
Co-emergence	Partly	Partly	Life and consciousness are deeply linked through self-organization and meaning
Autopoietic-enactive view	Life begins cognition	Not necessarily	Mind is rooted in living self-production
Panpsychism	No	Yes or always present	Consciousness or proto-consciousness is fundamental in matter
Cosmopsychism	No	Yes	Cosmic consciousness precedes individual life
Idealism	No	Yes	Life appears within consciousness
Taheri's T-Consciousness framework	No	Yes	Life manifests through or within T-Consciousness and consciousness fields
Artificial consciousness functionalism	Not necessarily	No or depends	Consciousness may arise in non-living functional systems
Artificial life view	Artificial life may come first	Possibly	Consciousness may arise in life-like artificial systems

This table directly connects the taxonomy to the central question of the book. Some theories place life before consciousness. Others place consciousness before life. Co-emergence views suggest that the relationship may not be strictly sequential.

E.10 By Evidence Type

Table E.9: Evidence types used in consciousness research.

Evidence type	Most relevant to	Strength	Limitation
First-person report	Human consciousness	Direct access to reported experience	Limited to beings that can report
Behaviour	Animals, infants, AI, simple organisms	Widely applicable	Behaviour may be unconscious
Neural activity	Humans and animals with nervous systems	Strong correlation with conscious states	Correlation is not explanation
Brain stimulation / lesion evidence	Human and animal consciousness	Helps identify causal mechanisms	Mostly limited to nervous systems
Anaesthesia response	Humans, animals, some biological systems	Useful for studying loss of consciousness	Does not define consciousness by itself
Evolutionary continuity	Animal consciousness	Supports graded approaches	Does not identify exact thresholds
Basal cognition evidence	Simple organisms and plants	Shows intelligence-like life processes	Does not prove experience
Formal measures	IIT, computational theories	Provides precision	May not capture phenomenology
Artificial system behaviour	AI consciousness debates	Tests substrate independence	Intelligence may not equal sentence
Contemplative / phenomenological data	Human experience	Rich first-person detail	Hard to standardize
Ethical precaution	Animal and AI moral status	Useful under uncertainty	May overextend moral concern

Different kinds of evidence answer different questions. First-person reports are powerful but limited to beings who can report. Behaviour is widely available but may not prove experience. Neural evidence is strong in animals with nervous systems but does not apply easily to plants, microbes, or artificial systems. Formal measures offer precision but may miss lived experience.

E.11 Summary of Main Taxonomic Distinctions

Table E.10: Summary of main taxonomic distinctions.

Distinction	Meaning	Why it matters
Phenomenal vs access consciousness Cognition vs consciousness	Experience itself vs information available for use Information processing vs subjective awareness	A system may use information without experience Simple organisms may be cognitive but not conscious
Sentience vs intelligence	Capacity to feel vs capacity to solve problems	Moral status depends more on sentience than intelligence
Self-awareness vs consciousness	Reflective self-model vs basic experience	Many beings may be conscious without reflective selfhood
Life vs mind	Biological organization vs experience or cognition	Central to the book's main question
Simulation vs realization	Imitating a process vs actually instantiating it	Crucial for AI and artificial life
Weak vs strong emergence	Reducible higher-level pattern vs irreducible novelty	Central to emergentist theories
Matter-based vs consciousness-based ontology	Matter as primary vs consciousness as primary	Determines whether life or consciousness comes first
Bottom-up vs top-down explanation	Parts produce wholes vs wholes or fields constrain parts	Important in origin-of-life and consciousness-first models
Scientific vs metaphysical theory	Empirically testable model vs fundamental worldview	Important for evaluating testability
Consciousness field vs physical field	Non-material organizing reality vs measurable physical field	Important for placing T-Consciousness correctly

E.12 Taxonomic Caution

This taxonomy should be read as a map, not a final theory. Different traditions divide consciousness in different ways. Some theories treat consciousness as neural access. Others treat it as subjective experience. Others treat it as integrated information, biological meaning, cosmic awareness, or a non-material consciousness field.

The same term may therefore have different meanings depending on context. For example, “consciousness” in Global Workspace Theory does not mean exactly the same thing as “consciousness” in idealism, panpsychism, or Taheri’s T-Consciousness framework. Similarly, “information” in Shannon theory, biology, IIT, and biosemiotics does not always mean the same thing.

The purpose of this taxonomy is to make these differences visible. It helps readers see that disagreement about consciousness is often not only disagreement about evidence. It is also disagreement about classification, definition, and metaphysical starting point.

Appendix F

References

F.1 Core Texts on Consciousness

- Baars, B. J. (1988). *A Cognitive Theory of Consciousness*. Cambridge University Press.
- Block, N. (1995). On a confusion about a function of consciousness. *Behavioral and Brain Sciences*, 18(2), 227–247.
- Chalmers, D. J. (1995). Facing up to the problem of consciousness. *Journal of Consciousness Studies*, 2(3), 200–219.
- Chalmers, D. J. (1996). *The Conscious Mind: In Search of a Fundamental Theory*. Oxford University Press.
- Crick, F., & Koch, C. (1990). Towards a neurobiological theory of consciousness. *Seminars in the Neurosciences*, 2, 263–275.
- Dehaene, S. (2014). *Consciousness and the Brain: Deciphering How the Brain Codes Our Thoughts*. Viking.
- Dehaene, S., & Changeux, J.-P. (2011). Experimental and theoretical approaches to conscious processing. *Neuron*, 70(2), 200–227.
- Dennett, D. C. (1991). *Consciousness Explained*. Little, Brown and Company.
- Edelman, G. M. (2004). *Wider Than the Sky: The Phenomenal Gift of Consciousness*. Yale University Press.
- Feinberg, T. E., & Mallatt, J. M. (2016). *The Ancient Origins of Consciousness: How the Brain Created Experience*. MIT Press.
- Graziano, M. S. A. (2013). *Consciousness and the Social Brain*. Oxford University Press.
- Lamme, V. A. F. (2006). Towards a true neural stance on consciousness. *Trends in Cognitive Sciences*, 10(11), 494–501.
- Nagel, T. (1974). What is it like to be a bat? *The Philosophical Review*, 83(4), 435–450.
- Rosenthal, D. M. (2005). *Consciousness and Mind*. Oxford University Press.
- Searle, J. R. (1980). Minds, brains, and programs. *Behavioral and Brain Sciences*, 3(3), 417–424.
- Searle, J. R. (1992). *The Rediscovery of the Mind*. MIT Press.
- Tononi, G. (2004). An information integration theory of consciousness. *BMC Neuroscience*, 5, 42.
- Tononi, G. (2008). Consciousness as integrated information: A provisional manifesto. *Biological Bulletin*, 215(3), 216–242.

Tononi, G., Boly, M., Massimini, M., & Koch, C. (2016). Integrated information theory: From consciousness to its physical substrate. *Nature Reviews Neuroscience*, 17, 450–461.

F.2 Philosophy of Mind and Metaphysics

Aristotle. (1986). *De Anima*. Translated by H. Lawson-Tancred. Penguin Classics. Original work published ca. 350 BCE.

Berkeley, G. (1982). *A Treatise Concerning the Principles of Human Knowledge*. Hackett. Original work published 1710.

Chalmers, D. J. (2010). *The Character of Consciousness*. Oxford University Press.

Descartes, R. (1996). *Meditations on First Philosophy*. Translated by J. Cottingham. Cambridge University Press. Original work published 1641.

Goff, P. (2017). *Consciousness and Fundamental Reality*. Oxford University Press.

Goff, P. (2019). *Galileo's Error: Foundations for a New Science of Consciousness*. Pantheon.

Hume, D. (2000). *A Treatise of Human Nature*. Oxford University Press. Original work published 1739–1740.

James, W. (1890). *The Principles of Psychology*. Henry Holt.

Kant, I. (1998). *Critique of Pure Reason*. Translated by P. Guyer and A. W. Wood. Cambridge University Press. Original work published 1781.

Leibniz, G. W. (1991). *Monadology*. In *Discourse on Metaphysics and Other Essays*. Hackett. Original work published 1714.

Levine, J. (1983). Materialism and qualia: The explanatory gap. *Pacific Philosophical Quarterly*, 64(4), 354–361.

McGinn, C. (1989). Can we solve the mind-body problem? *Mind*, 98(391), 349–366.

Plato. (2008). *Timaeus and Critias*. Translated by D. Lee. Penguin Classics.

Russell, B. (1927). *The Analysis of Matter*. Kegan Paul.

Spinoza, B. (1994). *Ethics*. Translated by E. Curley. Penguin Classics. Original work published 1677.

Strawson, G. (2006). Realistic monism: Why physicalism entails panpsychism. *Journal of Consciousness Studies*, 13(10–11), 3–31.

Whitehead, A. N. (1929). *Process and Reality*. Macmillan.

F.3 Consciousness-First, Idealist, Panpsychist, and T-Consciousness Sources

Goff, P. (2017). *Consciousness and Fundamental Reality*. Oxford University Press.

Goff, P. (2019). *Galileo's Error: Foundations for a New Science of Consciousness*. Pantheon.

Kastrup, B. (2019). *The Idea of the World: A Multi-Disciplinary Argument for the Mental Nature of Reality*. Iff Books.

- Kastrup, B. (2021). *Science Ideated: The Fall of Matter and the Contours of the Next Mainstream Scientific Worldview*. Iff Books.
- Nagasawa, Y., & Wager, K. (2016). Panpsychism and priority cosmopsychism. In G. Brüntrup & L. Jaskolla (Eds.), *Panpsychism: Contemporary Perspectives*. Oxford University Press.
- Seager, W. (2010). Panpsychism, aggregation and combinatorial infusion. *Mind and Matter*, 8(2), 167–184.
- Shani, I. (2015). Cosmopsychism: A holistic approach to the metaphysics of experience. *Philosophical Papers*, 44(3), 389–437.
- Strawson, G. (2006). Realistic monism: Why physicalism entails panpsychism. *Journal of Consciousness Studies*, 13(10–11), 3–31.
- Taheri, M. A. (2006). *Cosmic Mysticism / Erfan-e-Kayhani (Halgheh)*. Persian edition.
- Taheri, M. A. (2007). *Human from Another Outlook*. Persian edition.
- Taheri, M. A. (2010). *The Relationship Between Ethics and Mysticism*. Persian edition.
- Taheri, M. A. (2011). *Human and Knowledge*. Persian edition.
- Taheri, M. A. (n.d.). *Halqeh Mysticism: Interuniversal Mysticism, Volume 1*. Interuniversal Press.
- Taheri, M. A. (n.d.). *Human Insight*. CreateSpace / Interuniversal Press.
- Taheri, M. A. (n.d.). *Human Worldview*. Interuniversal Press.
- Taheri, M. A. (n.d.). *Psymentology, Volume 1*.
- Taheri, M. A. (n.d.). *Nonorganic Viruses*. CreateSpace Independent Publishing Platform.
- Taheri, M. A. (n.d.). *Non-Organic Beings*. Interuniversal Press.

F.4 Origin of Life, Abiogenesis, and Self-Organization

- Ball, P. (2011). *Physics of Life: The Dawn of Quantum Biology*. Oxford University Press.
- Cronin, L., & Walker, S. I. (2016). Beyond prebiotic chemistry. *Science*, 352(6290), 1174–1175.
- De Duve, C. (1995). *Vital Dust: Life as a Cosmic Imperative*. Basic Books.
- Eigen, M., & Schuster, P. (1979). *The Hypercycle: A Principle of Natural Self-Organization*. Springer.
- Kauffman, S. A. (1993). *The Origins of Order: Self-Organization and Selection in Evolution*. Oxford University Press.
- Kauffman, S. A. (1995). *At Home in the Universe: The Search for the Laws of Self-Organization and Complexity*. Oxford University Press.
- Lane, N. (2015). *The Vital Question: Energy, Evolution, and the Origins of Complex Life*. W. W. Norton.
- Martin, W., & Russell, M. J. (2007). On the origin of biochemistry at an alkaline hydrothermal vent. *Philosophical Transactions of the Royal Society B*, 362(1486), 1887–1925.
- Miller, S. L. (1953). A production of amino acids under possible primitive Earth conditions. *Science*, 117(3046), 528–529.
- Oparin, A. I. (1938). *The Origin of Life*. Macmillan.
- Orgel, L. E. (2004). Prebiotic chemistry and the origin of the RNA world. *Critical Reviews in Biochemistry and Molecular Biology*, 39(2), 99–123.

- Prigogine, I., & Stengers, I. (1984). *Order Out of Chaos: Man's New Dialogue with Nature*. Bantam.
- Russell, M. J., & Hall, A. J. (1997). The emergence of life from iron monosulphide bubbles at a submarine hydrothermal redox and pH front. *Journal of the Geological Society*, 154(3), 377–402.
- Schrödinger, E. (1944). *What Is Life? The Physical Aspect of the Living Cell*. Cambridge University Press.
- Szostak, J. W. (2012). The eightfold path to non-enzymatic RNA replication. *Journal of Systems Chemistry*, 3, 2.
- Walker, S. I., & Davies, P. C. W. (2013). The algorithmic origins of life. *Journal of the Royal Society Interface*, 10(79), 20120869.

F.5 Autopoiesis, Enactivism, Biosemiotics, and Basal Cognition

- Barbieri, M. (2003). *The Organic Codes: An Introduction to Semantic Biology*. Cambridge University Press.
- Di Paolo, E. A. (2005). Autopoiesis, adaptivity, teleology, agency. *Phenomenology and the Cognitive Sciences*, 4, 429–452.
- Hoffmeyer, J. (1996). *Signs of Meaning in the Universe*. Indiana University Press.
- Hoffmeyer, J. (2008). *Biosemiotics: An Examination into the Signs of Life and the Life of Signs*. University of Scranton Press.
- Levin, M. (2019). The computational boundary of a “self”: Developmental bioelectricity drives multicellularity and scale-free cognition. *Frontiers in Psychology*, 10, 2688.
- Lyon, P. (2015). The cognitive cell: Bacterial behavior reconsidered. *Frontiers in Microbiology*, 6, 264.
- Maturana, H. R., & Varela, F. J. (1980). *Autopoiesis and Cognition: The Realization of the Living*. D. Reidel.
- Thompson, E. (2007). *Mind in Life: Biology, Phenomenology, and the Sciences of Mind*. Harvard University Press.
- Trewavas, A. (2014). *Plant Behaviour and Intelligence*. Oxford University Press.
- Varela, F. J. (1991). Organism: A meshwork of selfless selves. In A. Tauber (Ed.), *Organism and the Origins of Self*. Kluwer.
- Varela, F. J. (1996). Neurophenomenology: A methodological remedy for the hard problem. *Journal of Consciousness Studies*, 3(4), 330–349.
- Varela, F. J., Thompson, E., & Rosch, E. (1991). *The Embodied Mind: Cognitive Science and Human Experience*. MIT Press.

F.6 Evolution of Consciousness and Animal Minds

- Birch, J. (2024). *The Edge of Sentience: Risk and Precaution in Humans, Other Animals, and AI*. Oxford University Press.
- Darwin, C. (1859). *On the Origin of Species*. John Murray.
- Darwin, C. (1871). *The Descent of Man, and Selection in Relation to Sex*. John Murray.

- Darwin, C. (1872). *The Expression of the Emotions in Man and Animals*. John Murray.
- Feinberg, T. E., & Mallatt, J. M. (2016). *The Ancient Origins of Consciousness: How the Brain Created Experience*. MIT Press.
- Godfrey-Smith, P. (2016). *Other Minds: The Octopus, the Sea, and the Deep Origins of Consciousness*. Farrar, Straus and Giroux.
- Godfrey-Smith, P. (2020). *Metazoa: Animal Life and the Birth of the Mind*. Farrar, Straus and Giroux.
- Griffin, D. R. (1976). *The Question of Animal Awareness*. Rockefeller University Press.
- Seth, A. K. (2021). *Being You: A New Science of Consciousness*. Faber & Faber.

F.7 Predictive Processing, Free Energy, and Information-Based Frameworks

- Clark, A. (2015). *Surfing Uncertainty: Prediction, Action, and the Embodied Mind*. Oxford University Press.
- Friston, K. (2006). A free energy principle for the brain. *Journal of Physiology-Paris*, 100(1–3), 70–87.
- Friston, K. (2010). The free-energy principle: A unified brain theory? *Nature Reviews Neuroscience*, 11, 127–138.
- Hohwy, J. (2013). *The Predictive Mind*. Oxford University Press.
- Parr, T., Pezzulo, G., & Friston, K. J. (2022). *Active Inference: The Free Energy Principle in Mind, Brain, and Behavior*. MIT Press.
- Seth, A. K. (2013). Interoceptive inference, emotion, and the embodied self. *Trends in Cognitive Sciences*, 17(11), 565–573.

F.8 Quantum, Speculative, and Physical Frameworks

- Bohm, D. (1980). *Wholeness and the Implicate Order*. Routledge.
- Fisher, M. P. A. (2015). Quantum cognition: The possibility of processing with nuclear spins in the brain. *Annals of Physics*, 362, 593–602.
- Hameroff, S., & Penrose, R. (2014). Consciousness in the universe: A review of the “Orch OR” theory. *Physics of Life Reviews*, 11(1), 39–78.
- Penrose, R. (1989). *The Emperor’s New Mind*. Oxford University Press.
- Penrose, R. (1994). *Shadows of the Mind: A Search for the Missing Science of Consciousness*. Oxford University Press.
- Stapp, H. P. (2011). *Mindful Universe: Quantum Mechanics and the Participating Observer*. Springer.
- Tegmark, M. (2000). Importance of quantum decoherence in brain processes. *Physical Review E*, 61(4), 4194–4206.
- Wendt, A. (2015). *Quantum Mind and Social Science: Unifying Physical and Social Ontology*. Cambridge University Press.

Wheeler, J. A. (1990). Information, physics, quantum: The search for links. In W. H. Zurek (Ed.), *Complexity, Entropy, and the Physics of Information*. Addison-Wesley.

Wigner, E. P. (1961). Remarks on the mind-body question. In I. J. Good (Ed.), *The Scientist Speculates*. Heinemann.

F.9 Artificial Intelligence, Artificial Life, and Machine Consciousness

Bostrom, N. (2014). *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press.

Clark, A., & Chalmers, D. J. (1998). The extended mind. *Analysis*, 58(1), 7–19.

Dennett, D. C. (2017). *From Bacteria to Bach and Back: The Evolution of Minds*. W. W. Norton.

Floridi, L. (2011). *The Philosophy of Information*. Oxford University Press.

Russell, S. (2019). *Human Compatible: Artificial Intelligence and the Problem of Control*. Viking.

Turing, A. M. (1950). Computing machinery and intelligence. *Mind*, 59(236), 433–460.

F.10 Ethics, Moral Status, and Sentience

Bentham, J. (1789). *An Introduction to the Principles of Morals and Legislation*. T. Payne.

Birch, J. (2024). *The Edge of Sentience: Risk and Precaution in Humans, Other Animals, and AI*. Oxford University Press.

Donaldson, S., & Kymlicka, W. (2011). *Zoopolis: A Political Theory of Animal Rights*. Oxford University Press.

Gruen, L. (2011). *Ethics and Animals: An Introduction*. Cambridge University Press.

Regan, T. (1983). *The Case for Animal Rights*. University of California Press.

Singer, P. (2011). *The Expanding Circle: Ethics, Evolution, and Moral Progress*. Princeton University Press.

Singer, P. (1975). *Animal Liberation*. HarperCollins.

F.11 Indigenous, Relational, and Ecological Perspectives

Berkes, F. (2012). *Sacred Ecology*. Routledge.

Kimmerer, R. W. (2013). *Braiding Sweetgrass: Indigenous Wisdom, Scientific Knowledge, and the Teachings of Plants*. Milkweed Editions.

Latour, B. (2005). *Reassembling the Social: An Introduction to Actor-Network-Theory*. Oxford University Press.

Plumwood, V. (2002). *Environmental Culture: The Ecological Crisis of Reason*. Routledge.

Watts, V. (2013). Indigenous place-thought and agency amongst humans and non-humans. *Decolonization: Indigeneity, Education & Society*, 2(1), 20–34.
